Automatically Establishing the Execution Environment for User Applications

Paul Kolano

AMTI/NASA Ames Research Center

kolano@nas.nasa.gov

Overview of Presentation

- Background
- Establishing Execution Environments
- Implementation Details
- Conclusions
- Future Work

The Problem

- A major goal of grid computing is to transparently run applications across different resources
- Different resources may have different setups both within and between organizations
  - Different software installed
  - Different file system structures
  - Different default environment settings

The Problem (cont.)

- Running applications on new resources typically results in:
  - someexec: not found
  - /usr/libexec/ld-elf.so.1: Shared object "somelib.so" not found
  - Exception in thread "main" java.lang.NoClassDefFoundError: some/java/Class
  - Can't locate Some/Perl/Mod.pm in @INC
  - ImportError: No module named some.python.mod

The Problem (cont.)

- Users end up wasting time setting up the resources that were supposed to save them time

The Goal

- Automatically ensure that user applications will not encounter software dependency failures during execution

Typical Solutions

- Statically-linked executables
  - Result in
    - Overly large executables
    - Inefficient use of memory
    - Hard-coding library bugs into code
- Custom software packages
  - Require detailed knowledge of
    - Dependency analysis techniques
    - Differences in environment settings
      - Operating systems
      - Software types

Typical Solutions (cont.)

- Waste time and allocations better spent on actual work
  - Transferring unnecessarily large files
  - Manually preparing custom packages

Related Projects

- Globus Executable Management (GEM)
  - Copy appropriate executable from network repository
- UNICORE
  - Transform abstract executable names to absolute paths
- Automatic Configuration Service (F. Kon et al.)
  - Automatically install software as necessary for component-based applications using manually specified dependencies
- Installers, Package Managers, and Application Management Systems
  - Provide consistent set of software on one or more systems
- Replica Management Systems
  - Facilitate location, selection, and replication of datasets

Related Projects (cont.)

- None have automatic dependency analysis
- Most have no treatment of environment variables
- Most only support executables
- Some require significant administration
- Some can't dynamically install new software

Summary

- Existing approaches do not provide enough assistance
- This work describes a new grid service for automatically establishing execution environments

Establishing Execution Environments

- Determine software the application requires
- Provide location for software on execution host
  - Determine if software is already installed
  - Find a source for missing software
  - Copy missing software to execution host
- Set environment variables based on locations

Establishing Execution Environments

- Don't want service to transfer software itself
  - User may cancel job
  - Previous job operations may fail
- Instead, add file operations and modify execution operation environment settings

Jobs

- Set of operations arranged in some topology
- File Operations (minimum requirements)
  - Copy files
  - Create directories
- Execution Operations (minimum requirements)
  - Host to execute on
  - Path of application to execute
  - Environment mapping from variables to values

Example Job

// Example.java

```
import org.apache.commons.logging.
            impl.SimpleLog;
import my.TimeClass;

public class Example {
    public static void main(String argv[]) {
        SimpleLog log =
            new SimpleLog("output");
        TimeClass tc = new TimeClass();
        if (tc.isTimeEven()) {
            log.info("even");
        } else {
            log.info("odd");
        }
    }
}
```

// TimeClass.java

```
package my;

public class TimeClass {
    public boolean isTimeEven() {
        long time =
            System.currentTimeMillis();
        if (time % 2 == 0) return true;
        else return false;
    }
}
```


Example Job

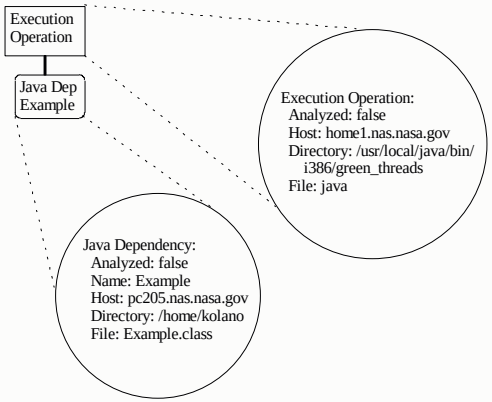

Stage 1: Dependency Analysis

- Determine application software dependencies
- Support common software types
- Concentrate initially on statically generated dependencies (i.e. do not worry about cases such as char *lib = f(); dlopen(lib))

Dependencies

- Basic information
  - Type
    - e.g. Executable
  - Name
    - e.g. w3m
  - Version range
    - e.g. [3.0, 3.1.1]
  - Feature list
    - e.g. compiled with SSL support
- Extended information
  - Source host
  - Source path
  - Target path
  - Analyzed flag

Dependencies (cont.)

- Currently supported types
  - Executable and Linking Format (ELF) objects
    - Executables
    - Shared libraries
  - Java classes
  - Perl programs
  - Python programs

Stage 1: Dependency Analysis (cont.)

- Dependency information is embedded
  - ELF executables and libraries
  - Java classes
- Dependency information must be derived
  - Perl programs
  - Python programs
  - Techniques
    - Textually traverse for relevant expressions
    - Partially evaluate using interpreter mechanisms

Stage 1: Dependency Analysis (cont.)

- Use as many existing tools as possible
  - Executables and libraries
    - ldd, elfdump
  - Java classes
    - com.sun.jini.tool.ClassDep
  - Perl programs
    - Module::ScanDeps
  - Python programs
    - modulefinder

Example Job After Dependency Analysis

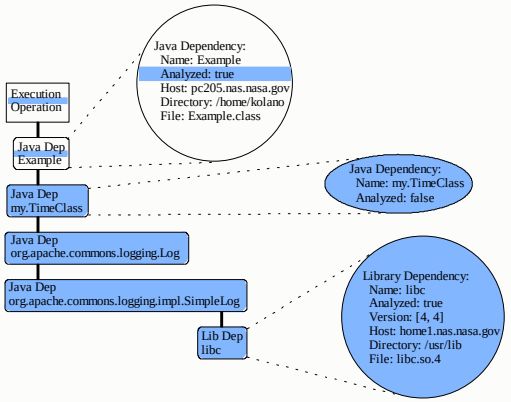

Stage 2(a): Dependency Location

- Determine existing location of software on target system
- Minimize number of files to transfer
- Searching an entire file system is impractical
- Must limit search space to specific paths
  - Which paths?
    - Add paths based on Filesystem Hierarchy Standard
    - Add paths based on user and system-default settings
  - Cannot guarantee file will be found in all cases
- Find using ls and Java, Perl, and Python interpreters

Example Job After Dependency Location

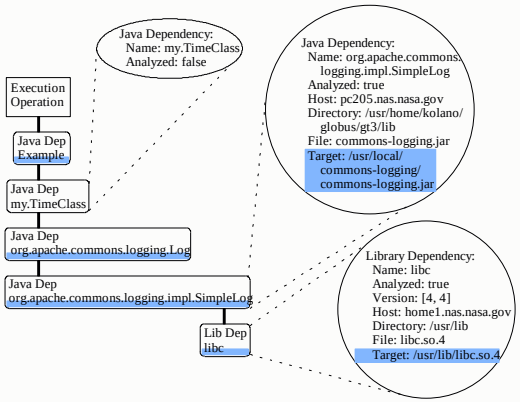

Stage 2(b): Dependency Lookup

- Find a source for missing software
- Find dependencies of missing software
- Use software catalog
  - Contains mappings from LFNs to PFNs
  - LFNs based on dependency name, type, supported operation system, and version
  - Contains dependencies of each PFN
  - Allows both centrally-managed and user-defined mappings

Example Job After Dependency Lookup

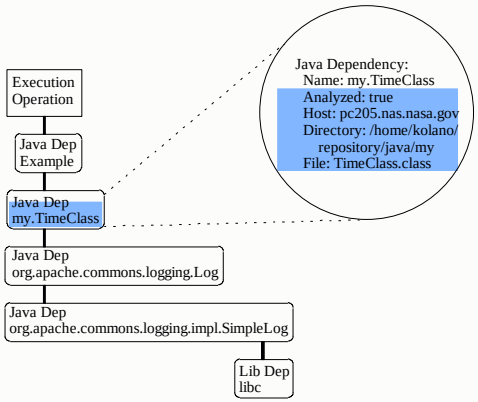

Stage 2(c): Dependency Transfer

- Copy missing software to execution host
- Must create correct directory hierarchy for Java, Perl, and Python software
  - e.g. my.TimeClass.class can only be found if it exists in some directory as .../my/TimeClass.class
- Copy at most once per job

Example Job After Dependency Transfer

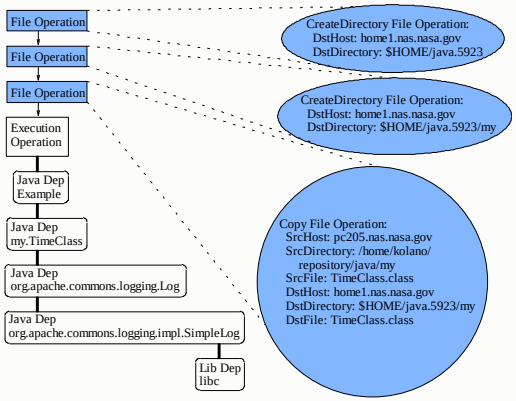

Stage 3: Variable Setup

- Application must be able to locate software
- Set environment variables to find existing and soon to be existing software
- Must consider directory hierarchy for Java, Perl, Python software
  - e.g. for my.TimeClass at /somedir/my/TimeClass.class, CLASSPATH must contain /somedir

Example Job After Variable Setup

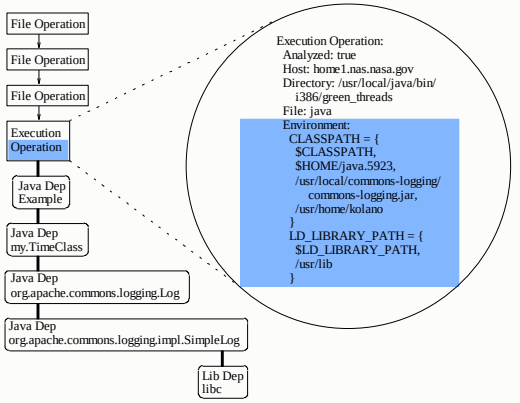

Possible Failures

- Application depends on A, but A cannot be located anywhere
- Application depends on A, which depends on B, but analysis techniques used on A are inadequate to determine B is a dependency
- Application does not depend on A, but analysis techniques report A is a dependency

Dealing with Failures

- Notify user to prevent wasted effort
- Missing software
  - Provide convenience methods to locate relevant dependencies after transformation
- False negatives and false positives
  - Cannot detect automatically
  - Currently, only Perl analysis is susceptible
  - Provide user with flexibility to compensate when necessary

Flexibility

- User has complete control of job transformation
  - Can execute stages individually
  - Can specify dependencies manually
  - Can turn analysis off for individual items
  - Can specify an exact source for software
  - Can specify an existing location on execution host
  - Can manage personal software catalog

Implementation

- Implemented in Java and Bourne shell scripts
- Runs as an OGSI-compliant grid service
- Uses OGSI GRAM service to execute analysis and location scripts
- Uses European DataGrid Local Replica Catalog as software catalog

Implementation (cont.)

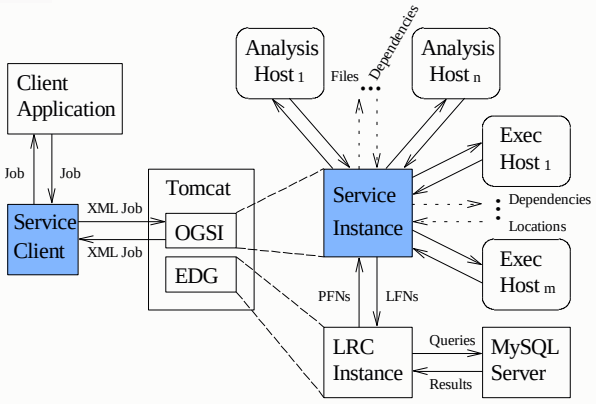

Where does this fit in?

- NASA Information Power Grid (IPG)
  - Current prototype services
    - Resource Broker
      - Select resources for jobs based on user constraints
    - Job Manager
      - Reliably execute jobs on specific resources
  - Establish environment after selection and before execution

Implementation (cont.)

- Service exists in prototype form
- All discussed functionality fully tested on FreeBSD
- Analysis and location scripts fully tested on FreeBSD, IRIX, and SunOS
- Waiting for full IPG deployment
  - GT3 stability and IRIX support
  - RB and JM are GT2 services and not yet OGSI-compliant

Conclusions

- Implemented a new OGSI-compliant service with functionality for
  - Automatically identifying application dependencies
  - Managing a flexible software catalog used as a source for key software
  - Establishing a suitable environment by transferring dependent software and setting environment variables
- Increases pool of compatible resources with little or no user intervention
- Net result is increase in user productivity

Future Work

- Software caching
- Additional dependency types
- Additional analysis capabilities
- Full IPG deployment
- Advanced software installation mechanisms
- Full version and feature support