GWD-TYPE
Category: TYPE
NAME_OF_WG_OR_RG

Author-1, Institution
Author-2, Institution
DATE
[if applicable: Revised DATE]

Resource Management on Grids

Status of This Memo

This memo provides information to the Grid community on resource management on Grids. It does not define any standards or technical recommendations. Distribution is unlimited.

Abstract

Document Abstract.

Contents

1. Introduction

Grid computing involves many forms of management: many of its technologies are, or are related to, some form of systems management. Consequently, Grid computing has many overlaps with the management of IT systems in general. Some examples of these overlaps are job management, storage management, security management, etc. As Grid technologies are increasingly adopted on IT systems in general, the overlap will become bigger and the differences will probably blur.

A second aspect of management in Grid computing is the management of the Grid itself. While it is possible to build a Grid with entities that are not manageable, that is simply not practical. For instance, any problem on the Grid would require the whole Grid to be reset, since no form of problem determination and fault recovery would be possible. Also, adaptive functionality would be impossible, since it needs both monitoring and control.

This document organizes the interfaces, services, activities, etc. that are involved in management in OGSA, including management within OGSA and the management of the OGSA infrastructure. Its targets are the CMM-WG gap analysis and the OGSA spec.

[To the OGSA-WG: this document will contain the results the gap analysis being done in the CMM-WG. However, most of its current contents (mostly an introduction), provide a management framework for OGSA, plus definitions, that can be added to the resource management sections of the OGSA spec. Also, it is hoped that its contents can help the discussions in the OGSA-WG. For instance, at times people have asked about “management”, which means many things in OGSA, or talked about “resources” and “resource managers”, which have never been clearly defined.]

2. Definitions

[To the OGSA-WG: some or all these definitions should be added to the OGSA spec glossary].

Management (in Grids or otherwise) is the process of monitoring an entity, controlling it, maintaining it in its environment, and responding appropriately to any changes of internal or external conditions.

A *manager* initiates management actions; it might be either a management console operated by a human, or a software entity that is able to monitor and control its targets without human intervention.

Manageability defines information that is useful for managing a resource or service. Manageability are those aspects of an entity that support management specifically through instrumentation of the entity to allow managers to interact with the entity. The manageability may be provided by the resource itself or by a separate means.

Manageability interfaces are sets of standardized interfaces that allow a manager to interact with an entity in order to perform common management actions on this entity such as starting it, stopping it and gathering performance data.

Manageable entities are entities that provide manageability interfaces and thus (as the name implies) can be managed. Manageable entities can be:

- physical (e.g., a node, a network switch or a disk) or logical (e.g., a process, a file system, a print job, or a service)
- discrete (e.g., a single host) or composite (e.g., a cluster)

[To OGSA: the definitions below are the important ones for OGSA, and need reviewing]

Resource is a name often used for the targets of management. There are two ways of thinking of resources:

- Grid computing tends to view a resource through its dictionary definition, i.e., “an available supply that can be used when needed”. This definition is usually applied to entities that are pooled (e.g., hosts, software licenses, IP addresses, etc.) or, similarly, entities that provide a given capacity (e.g., disks, networks, memory, etc.). In this case, the pool and the capacity are a supply, a part of which is allocated and used. However, under this definition a process, a print job, a registry service and a VO are not resources.
- Management tends to view resources as being the same as manageable entities (including entities that are not manageable entities per se, but are managed by other means, such as software licenses). This is a superset of the definition above, since it includes both pooled and manageable resources. [This definition is used by OASIS WSDM]

Since Grid management concerns both Grid computing and management, both definitions are valid, depending on the context. Unless otherwise mentioned, the latter definition of resource is used in this document. The terms resource, manageable resource and manageable entity will be used interchangeably.
[Question to OGSA: need to decide if we use both definitions in the OGSA spec as above, or choose one. Making clear that there are two nuances, and using both, might result in ambiguities but will make writing the OGSA spec a lot easier]

Resource management is a generic term for several types of management applied to the Grid definition of resources. Some examples are:

- reservation / brokering / scheduling
- aggregation (service groups, WSDM collections, VO management, etc.)
- installation / deployment / provisioning
- accounting / metering
- discovery
- problem determination, fault management

A *resource manager* is a manager that implements one or more resource management functionalities.

3. Management in OGSA

3.1 Requirements

[To the OGSA-WG: this section could be appended to section 3.5 if deemed suitable].

The main requirements for management in OGSA are:

- Scalability: management needs to scale to potentially thousands of resources. Management needs to be done in a hierarchical and/or peer-to-peer fashion to achieve this scalability, so OGSA should allow these forms of management. Hierarchical management, for instance, can be done through manageability interfaces that allow grouping resources and managing them collectively (e.g., GMA aggregators and intermediaries that implement WSDM collection interfaces), i.e.:
  - Acting as a proxy that allows a manager to perform the same action on multiple resources with a single request
  - Computing metrics for the resources (e.g., average load, average reservation rate)
  - Filtering and aggregating events
  - Polling resources for state (reserved, running, failed, idle, saturated, etc.), and providing the results when inquired, as well as sending events when the state changes

- Interoperability: management will be done through several software, hardware and service boundaries (e.g., through the boundaries between different products), so interoperability is essential to avoid “stovepipes”. Two kinds of interoperability are needed:
  - between levels: e.g., between the resources and their manager
  - on the same level: to allow the use of management-related services at the same level

  Interoperability in both cases requires that the interfaces in the boundaries (the resources and resource management functions of the system) are defined in a standard way. This applies to both Grid-specific standards and also IT management standards.
- Security: there are two security aspects in management:
  - Management of security: the management of the security infrastructure, including the management of users, VOs and access policies.
  - Secure management: using the security mechanisms on the management tasks. Management should be able to ensure its own integrity and to follow access control policies of the owners of resources and VOs.
- [Any other high-level requirements? Need to check WSDM requirements. GMA has also the following requirements: low latency, high data rate, low overhead.]

3.2 Levels

[To the OGSA-WG: this section can be added to the OGSA spec as an introduction to the resource management section (currently section 4.4), followed by description of the resource managers].

Management in OGSA consists of many interfaces, services and activities. These can be categorized in three levels, based on the levels of the OGSA service taxonomy:

- Resource level
- Platform services level
- OGSA services level

A detailed description of each level is given below. It must be noticed that the description will focus on the manageability interfaces (which correspond to the OGSA services taxonomy), not to the locus or hierarchy of implementation (which correspond to the OGSA services hierarchy).

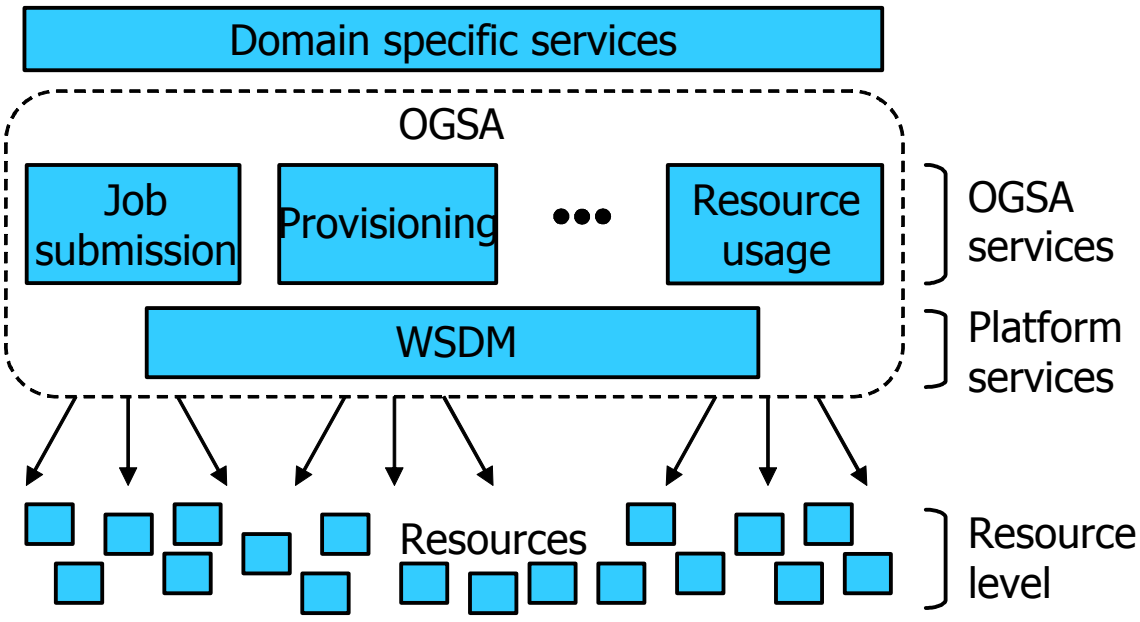

Figure 1: levels of management in the OGSA taxonomy

At the resource level, the resources are managed through their manageability interface (for discrete resources, this is usually SNMP, CIM / WBEM, JMX, or proprietary interfaces). This involves monitoring (i.e., obtaining the state of the resource, which includes events), setup and control (i.e., setting the state of the resource), and discovery.

The platform services level provides the base management behavior of resources, which forms the basis for both manageability and management on OGSA. Standardization of this base management behavior is required in order to integrate the vast number and types of resources and more limited set of resource managers introduced by multiple parties. The platform services level provide:

- The base manageability model, which (following the OGSI philosophy) represents resources as services. This allows resources in OGSA to be manipulated through the standard OGSI means for discovery, access, etc. The base manageability model includes resource identity (e.g., through GSHs), the base manageability interface, service data to represent its attributes, etc. This base manageability model allows the resources to become manageable to a minimum degree (allowing discovery, termination, introspection, monitoring, etc.). It is important to notice that the base manageability model is *not* a resource model – the resource model of the resources themselves is accessed through the base manageability model.
- Basic functionality that is common to the OGSA services, e.g.:
  - portTypes for functionalities that are common to all resources (e.g., start / stop / pause / resume)
  - Lifecycle representation and operations
  - Relationships among resources
  - Aggregation (WSDM collection interfaces)
  - Events

The functionalities above are currently being developed in the OASIS WSDM TC [mostly MUWS – management using Web services].

At the OGSA services level there are two forms of management:

- Some of the services, such as the resource managers, are by themselves are a form of management. These services (as all other services) provide their functionality (e.g., job management) through a *functional interface*.
- The services have a *manageability interface* through which the service is managed (i.e., the monitoring, introspection, termination, etc. of the service itself). This interface should extend the interface of the base manageability model, adding manageability functionality that is specific to the management of services. This interface makes the services in the OGSA services level (including resource managers) become also resources themselves (in the management sense of the word). Two interfaces are probably needed:
  - A generic interface that is common to all OGSA services (e.g., to create and destroy instances). [This is WSDM MOWS. Is this part of the platform services level?]
  - Interfaces specific to each service (e.g., monitoring of registry services, monitoring of handle resolution services, etc.).

Discovery provides a good concrete example of the differences between the resource, platform service and OGSA service levels. Discovery at the resource level it might involve scanning a network to discover the devices attached to it. Discovery at the platform services level can involve introspecting the service data of a service to find its capabilities. Discovery at the OGSA level might involve accessing one or more UDDI repositories that contain the GSHs of available resources.

The division in levels helps interoperability between levels by defining clear interfaces between them (interoperability on the same level is realized by defining the functional interfaces of the services). While it is possible to build OGSA services that bypass these levels (e.g., using a

proprietary adapter in a resource that feeds data directly to these OGSA service), that is not desirable from the point of view of interoperability (e.g., it limits the kinds of resources which the adapter will be compatible with).

4. Resource Models

[To the OGSA-WG: this is important for OGSA, but not sure if and how this section fits the OGSA spec (and if it does, where). It's still in a preliminary state.]

A Resource Model is an abstract representation of manageable IT entities, which defines their schema (conceptual hierarchy and inter-relationships) and characteristics (attributes, management operations, etc.). Examples of resource models are:

- Standard management models, which focus on system management:
  - CIM, including:
    - [Add a list of existing functionalities]
    - JSIM (Job Submission Information Model, defined by GGF's CGS-WG)
  - SNMP MIBs
  - JMX's JSR77
  - [WSDM MOWS Web service model?]
- Resource descriptions (most of them Grid-specific) which focus on resource management:
  - Reservation / brokering / scheduling
    - Unicore Resource Schema
    - Globus RSL
    - GLUE schema
    - JSDL (being defined by GGF's JSDL-WG)
  - Accounting / metering
    - Usage Record (defined by GGF's UR-WG)
  - Installation / deployment / provisioning
    - Configuration Description Language (CDL, being defined by the CDDML-WG)
    - DCML (Data Center Markup Language)
  - [Anything else?]

It must be noticed that some of the resource descriptions are not intended to be models by themselves, but they contain an implicit model (which defines, for instance, which entities exist, and what their attributes are).

[The contents of the paragraphs below are definitely a subject for discussions. First, I'm giving conclusions before offering a proof, but I think that some of you have war stories and scars that justify them. Second, even if we all agree with its contents, we need to fine-tune them to make them stronger and/or sharper. I'm accepting both opinions and text.]

Ideally, the use of a single resource model is desirable, since it makes interoperability easier to achieve when compared to mediation between models. [Insert examples]. However, usually multiple resource models will be in simultaneous use in a given grid. Thus, it is highly desirable that these resource models are coordinated to make them compatible. For instance, these multiple resource models should be a subset of a single resource model. Or, multiple resource descriptions should be "renderings" of a single resource model (with each resource description language representing this model, or a subset of it, using its own syntax, e.g., its own XML schema).

There are two areas in which there is need for coordination between resource models:

- Between the resource descriptions (to ease interoperability between OGSA services, i.e., reservation, metering, provisioning, etc.)
- Between the standard management models and the resource descriptions (to ease interoperability between resources and their resource managers)

[What does GRIP do to bridge Globus and Unicore models? Anything to say about GLUE? How far has the ontology research gone?]

5. Analysis of the OGSA Services

5.1 General analysis

The following list calls out elements of a Grid that are candidates for management, and hence need to provide manageability interfaces. The list is intended to be used to identify the types of management actions that need to be possible, and the set of common manageability interfaces that are required. Some interfaces are expected to be defined already, while others will need to be specified. The list is derived in part from Section 6 of the current draft of the GGF OGSA document, dated November 3rd, 2003.

In an OGSA Grid, all manageable resources either are, or are represented by, Grid services. By definition, any Grid service exposes some interfaces that are useful in management - e.g. its termination time and the ability to change it (possibly causing immediate termination); the handle of the factory service that created it; a means of retrieving a list of its service data elements and the ability query them, to change them, or to request notification if any of them changes.

The following items detail some specific services, and why it will be important to manage them.

[Fred's comments below between brackets. I also added letters corresponding to the three categories of services in the OGSA level: "F" when the functional interface is a management interface, and "S" when there is a manageability interface specific for the service. (I suppose that the generic manageability interface applies to all services). The "F" interfaces will be defined by other GGF WGs, but somebody (us?) has to define the "S" ones (I'd say this is a gap). Further steps: need to separate the "F" and "S" contents, and also analyze from the point of view of the platform services and models.]

- **General Grid Services.** Any Grid service will provide interfaces for at least minimal management - e.g. termination, introspection and monitoring. The OASIS WSDM TC will define some other standard manageability interfaces for Web services that should be applicable to a Grid service. We will need to determine if there are additional general interfaces that are specific to the Grid space.
- **Registry Services (S).** Registry services are likely to be deployed in every Grid. A service must be able to register itself in one or more registries so that it can be discovered, and so that its interfaces and capabilities can be queried. It is important that Registry services are available, and that they operate correctly, so managers will need to be able to monitor their operation and performance, and to create and destroy instances and copies as needed. A primary Registry service is likely to be the starting point for discovering and mapping, and hence managing, all resources in the Grid.
- **Handle Resolution Services (S).** Most Grids will need to provide access to one or more Handle Resolver services, and their performance will be critical to overall performance of the Grid. It should be possible for a manager to locate all HandleResolver services associated with the Grid, to monitor performance – e.g. request-volume and response times – and to create and destroy instances or copies as needed.
- **Factory Services (S).** It may be important to identify general Factory services as such, so that they can be managed in the same way as other key infrastructure services.
- **Virtual Organizations (F).** [OGSA 6.1] VOs can be considered as very-high-level manageable entities, and will provide significant management challenges. A manager will need to be able to discover and manage VO registries, create and destroy VOs, and manage the set of resources and users assigned to an individual VO. *There's much more we can say here.*

- **Service Configuration, Installation, Deployment & Provisioning (F).** The CDDLM working group will address how to describe configuration of services, deploy them in a Grid, and manage their deployment lifecycle (instantiate, initiate, start, stop, restart, etc.). Managers will need the ability to configure, deploy, redeploy (relocate, perhaps with a different configuration) and terminate applications and other types of services within Grids, using the interfaces defined by CDDLM. Installation and Provisioning may be separate issues.
- **Job Management (F).** Program Execution services will need to provide a way for managers to be notified as jobs are started, and either the jobs themselves or the execution services must provide an interface that allows the jobs to be managed – e.g. terminated, suspended or migrated. The Job Agreement Service [OGSA 6.20] may provide the required interfaces.
- **Choreography, Orchestration & Workflow (F).** [OGSA 6.4] Some management functions will be needed for controlling and monitoring Grid flows, but we'll need better definition in the OGSA document. This is probably low priority.
- **Transactions (?).** [OGSA 6.5] As for Choreography, Orchestration & Workflow.
- **Metering/Rating/Accounting/Billing & Payment (F).** [OGSA 6.6-9] These services all relate to measuring resource usage, and accounting and charging for it – they will not be applicable to all Grids.
  - The Metering service is effectively an infrastructure service - it must be permanently available if resource usage is to be recorded and charged for, and hence the manager must be able to monitor and control its operation as for any other critical service.
  - The Rating and Accounting services might be considered as application-level services – they are likely to be run periodically, reading and processing persistent (logged) data, and hence can be managed in the same way as any application-level service.
  - The Billing & Payment service will be a critical service for Grids that require it. This service may be internal or external, or may be an internal service that makes use of external services, such as credit card authorization services. Where needed, it will be essential that this service is operational, and a manager must be able to monitor and control it.
- **Fault Management (?).** [OGSA 6.11] A manager will need to be notified of faults, and to be able to handle them to some level. This has not yet been addressed by OGSA, and it's not clear if this would be implemented as a persistent service, or what its requirements for management might be. [OGSA should probably define the mechanisms to allow fault management (e.g., monitoring and control interfaces), but not the policies (e.g., what to do when a job crashes)]
- **Problem Determination (?).** [OGSA 6.12] A Problem Determination service, if available, is likely to be used by a manager, but may not be persistent, and its requirements for management are not clear. Not yet addressed by OGSA. [Same comment as above]
- **Logging Services (F, S).** [OGSA 6.13] Logging services are essential infrastructure services, and they must be managed accordingly. It will be necessary not only to monitor their performance, but also to deal with storage space thresholds, low-space or insufficient-space conditions, periodic purging, access control, and many other facets. Different management domains within a given Grid may have different policies for retention etc. It's likely that this will be one of the more complex management operations.
- **Messaging and Queuing (S).** [OGSA 6.14] If separate messaging and queuing services are defined, it is likely that they will become critical infrastructure services. Management requirements will include monitoring performance and managing the number of available instances and copies to handle the message volume and, if applicable, storage space.
- **Event Services (?).** [OGSA 6.15] If specialized event services are defined they will need to be managed as critical infrastructure services. [We will need common event semantics for interoperability]
- **Policy & Agreements (F, S).** [OGSA 6.16, 6.25] A Policy subsystem, when fully defined, is likely to be composed of multiple related services, including a repository. The subsystem will be a critical infrastructure component of most Grids, and the ability to monitor it and to control certain elements will be essential.

- **Data Services (F, S).** [OGSA 6.17/18] The “Base Data Services” and “Other Data Services” OGSA categories describe services that provide data representation and transformation facilities (Base Data Services), and facilities for accessing, transferring and managing replicas. In many Grids such services may be numerous and diverse; they will be fundamental to most, if not all, Grids. They will be critical infrastructure services, and their availability and performance must be monitored and managed.
- **Agreement Services (F, S).** [OGSA 6.20-22] The OGSA document lists Agreement Services for Jobs, Reservations and Data Access. All are likely to be based on the WS-Agreement specification, but each is likely to have specialized interfaces, and may require specialized management. Their correct operation and performance will be critical to a Grid, and must be monitored.
- **Queuing Service (F).** [OGSA 6.23] The OGSA document currently defines a queuing service as being a mechanism for scheduling jobs according to local policy, and it may be regarded as a part of the overall job management and execution subsystem. A manager may need to monitor the status of individual resource queues, and to be able to control them - e.g. to move jobs between queues to balance loads, to override priorities and to accommodate planned downtime.
- **Security Services (S).** Security services are not yet fully addressed in detail by OGSA, but services such as authentication and authorization services will need to be managed, and may need specialized interfaces.
- **Information and Monitoring Service (F).** There are contents there (such as persistency and archives) which are not yet covered by WSDM/CMM or OGSA.

5.2 Analysis of selected services

TBD

6. Gaps

TBD

7. TBD

- How far does the OGSA-WG (or the GGF) need to define manageability of the OGSA infrastructure? E.g.: performance monitoring of a registry.
- Virtualization? We will certainly have multiple forms of it around. It might be worth writing about it to make the term more concrete.

8. Security Considerations

As mentioned in section 3.1, security is among the main requirements on management. Security is one of the many management functionalities covered in this document.

Author Information

Editor: Fred Maciel (fred-m@crl.hitachi.co.jp)

With contributions from Jem Treadwell and Ellen Stokes, and help from Latha Srinivasan, Bryan Murray, Ravi Subramaniam, and Hiro Kishimoto.

Glossary

Recommended by not required.

Intellectual Property Statement

The GGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the GGF Secretariat.

The GGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the GGF Executive Director.

Full Copyright Notice

References

TBD