

GWD-R

Distributed Resource Management  
Application API (DRMAA) Working Group

Hrabri Rajic, Intel Americas Inc. (maintainer)  
Roger Brobst, Cadence Design Systems  
Waiman Chan, IBM  
Fritz Ferstl, Sun Microsystems  
Jeff Gardiner, John P. Robarts Res. Inst.  
Andreas Haas, Sun Microsystems  
Bill Nitzberg, Veridian Systems  
Hrabri Rajic\*, Intel Americas Inc.  
John Tollefsrud\*, Sun Microsystems  
\*co-chairs  
February, 2003

## **Distributed Resource Management Application API Specification 1.0**

### Status of this Memo

This memo is a Global Grid Forum Grid Working Draft - Recommendations (GWD-R) in-process, in general accordance with the provisions of Global Grid Forum Document GFD-C.1, the Global Grid Forum Documents and Recommendations: Process and Requirements, revised April 2002.

### Copyright Notice

Copyright © Global Grid Forum (2003). All Rights Reserved.

### **Abstract**

This document provides a description of Distributed Resource Management Application API (DRMAA), which provides a generalized API to Distributed Resource Management Systems (DRMS) to facilitate integration of application programs.

The scope of DRMAA is limited to job submission, job monitoring and control, and retrieving the finished job status. DRMAA provides application developers and DRM builders with a programming model that enables the development of distributed applications tightly coupled to an underlying DRMS. For deployers of such distributed applications, DRMAA preserves flexibility and choice in system design.

## Table of Contents

|  |    |
|--|----|
| Abstract.....                                      | 1  |
| 1. Introduction .....                              | 3  |
| 1.1 DRMAA Working Group Charter .....              | 3  |
| 1.2 DRMAA Scope .....                              | 3  |
| 1.3 Language Issues.....                           | 3  |
| 1.4 Library Issues.....                            | 3  |
| 1.5 User Program and DRMAA Interaction.....        | 4  |
| 2. API Design Issues .....                         | 5  |
| 2.1 Basic Guidelines .....                         | 5  |
| 2.2 Distributed Application Environment.....       | 5  |
| 2.2.1 File staging.....                            | 6  |
| 2.2.2 Job categories.....                          | 6  |
| 2.2.3 Native specification .....                   | 7  |
| 2.2.4 Building portals .....                       | 8  |
| 2.3 Interface Routines General Description .....   | 8  |
| 2.3.1 Init and exit routines.....                  | 8  |
| 2.3.2 Job template routines.....                   | 8  |
| 2.3.3 Job submission routines .....                | 9  |
| 2.3.4 Job monitoring and controlling routines..... | 9  |
| 2.3.5 Auxiliary routines.....                      | 10 |
| 2.3.6 DRMAA job state transition diagram.....      | 10 |
| 3. API Specification.....                          | 11 |
| 3.1 Introduction .....                             | 11 |
| 3.1.1 Error codes .....                            | 11 |
| 3.1.2 DRMAA sessions .....                         | 11 |
| 3.1.3 Run usage data.....                          | 11 |
| 3.1.4 Precedence rules.....                        | 12 |
| 3.1.5 Site specific requirements.....              | 12 |
| 3.1.6 Job valuator .....                           | 12 |
| 3.2 DRMAA API .....                                | 12 |
| 3.2.1 Initialization and exit routines.....        | 12 |
| 3.2.2 Job template routines.....                   | 13 |
| 3.2.3 Job attributes .....                         | 15 |
| 3.2.3.1 Mandatory attributes.....                  | 15 |
| 3.2.3.2 Optional attributes.....                   | 19 |
| 3.2.4 Job submission routines .....                | 21 |
| 3.2.5 Job control routines.....                    | 22 |
| 3.2.6 Auxiliary routines.....                      | 25 |
| 3.3 List of DRMAA Errors.....                      | 26 |
| 4. Security Considerations.....                    | 29 |
| Author Information .....                           | 30 |
| Intellectual Property Statement .....              | 31 |
| Full Copyright Notice .....                        | 31 |

## 1. Introduction

Distributed Resource Management Application API (DRMAA) abstracts the fundamental job interfaces of the Distributed Resource Management Systems (DRMS) and provides an API intended for distributed application developers such as Independent Software Vendors (ISVs). DRMAA provides an easy-to-use programming model, thereby encouraging adoption by both application builders and DRM system builders.

This document is DRMAA specification 1.0, the work result of the DRMAA Working Group (DRMAA-WG) in the Global Grid Forum.

### 1.1 DRMAA Working Group Charter

Develop an API specification for the submission and control of jobs to one or more Distributed Resource Management (DRM) systems.

The scope of this specification is all the high level functionality that is necessary for an application to consign a job to a DRM system including common operations on jobs like termination or suspension.

The objective is to facilitate the direct interfacing of applications to today's DRM systems by application's builders, portal builders, and Independent Software Vendors (ISVs).

### 1.2 DRMAA Scope

The scope of DRMAA is limited to job submission, job monitoring and control, and retrieving the finished job status. Resource reservation, security, etc. are topics that are beyond DRMAA scope. These topics are handled by other GGF working or research groups.

### 1.3 Language Issues

It is our position that the API should be implemented for multiple languages, C/C++ being a primary choice by virtue of its widespread use by ISVs. Additional choices are scripting languages, Perl and Python.

It is possible to design an Interface Definition Language (IDL) that will effectively resolve the issue of one interface serving multiple languages. While this is a viable approach, it was felt that it would significantly slow progress on an implementation. DRMAA interfaces are described using an IDL-like language.

Another viable approach would be to design a protocol instead of the API. It is the desire of the DRMAA-WG that the DRMAA specification be useful in the implementation of such a protocol.

### 1.4 Library Issues

A complete library implementation would support all DRMSs and have versions for either static or dynamic linkage. This is not something that will be feasible, or even desirable. A real possibility is a support for multiple, but not all DRMSs. The packaging could come as one library, where a DRMS is selected at run time by setting an environment variable for the desired DRMS, or as one DRMS link per library. The authors advocate the latter approach. In this setup the shared library is selected at run time by the end users.

It is expected that the developers will be linking the library from serial and multithreaded codes. The library should be thread safe. A thread safe DRMAA implementation allows a multi-threaded application to use DRMAA interfaces without any explicit synchronization amongst the application threads. Before a multi-threaded application can use any of DRMAA interfaces DRMAA initialization routine should be called by only one thread. The main thread could be a good choice. Similarly, DRMAA library should be disengaged by only one thread.

Debugging of distributed programs is more challenging than debugging of single machine versions. This document advocates providing production and debugging or tracing versions of the DRMAA library.

The library should support extraction of the DRMAA API version number by external programs (such as SCCS's "what" and RCS's "ident"), and allow distributed applications to obtain this information programmatically.

### **1.5 User Program and DRMAA Interaction**

All of the DRMSs are asynchronous in nature. They notify the end user of the status of a finished job via e-mail, which is insufficient to satisfy the needs of the DRMAA library users. This document proposes to deal with the asynchrony similar to Unix and Windows process interfaces, by blocking on the wait call for a specific job request or possibly for all of them in the same session. Support for reactive mode (such as used in the Globus GRAM interface) is to be addressed in future DRMAA versions, to better support the programming model used by an increasing number of applications with graphical user interfaces.

## 2. API Design Issues

Developers have been using Unix system, fork/exec, popen, and the wait interfaces for years to spawn additional processes and wait for the end of their execution to get their exit codes. Windows has equivalent utilities like CreateProcess and WaitForSingleObject. DRMAA provides its own set of interfaces that are OS neutral. It attempts to be consistent with libc interfaces.

### 2.1 Basic Guidelines

Even though the API should be self-contained, it is not always possible to consolidate all variations of end user and DRMS interactions under the API. For this reason, this document advocates that developers provide a way for the end user to specify DRMS specific options. The “job category” mechanism is intended to convey DRMS specific options in a DRMS independent fashion.

Additionally, there might be a need for DRMS specific options to be specified at run time as command line options. The end user could lose portability this way, but since most end users are only concerned with a single DRMS, this is likely to be acceptable. In contrast, DRMAA providers and ISVs are much more likely to target multiple DRMSs and therefore require a DRMS independent mechanism.

The API centers around the job\_id attribute that is passed back by the DRMS upon job submission. Job\_id is used for all the job control and monitoring purposes. An additional attribute, job\_name, which is commonly found in DRMS implementations, is part of the job description. Job\_name could be used by the developer and/or internally by the implementation to group the jobs for easier user classification and tracking. This attribute could be a key to achieve scalability for DRMAA implementations, especially since DRMS user jobs could be running concurrently with those of the other DRMS users.

There are few guidelines that were used in designing DRMAA:

- The API calling sequences should be simple and the API set small.
- The routine names should convey the semantics of the routine.
- The set should be as convenient as possible, even with the risk of being forced to emulate some functionality if missing from a DRMS.
- All job or job set manipulation is available without explicit job iterating.
- The server names are hidden, and the DRMS is a “black box”.
- The end user could specify native specification options attribute if he/she needs to interact with the DRMS beyond what is available through DRMAA job categories.
- The API should be extensible so that future implementations are backward compatible with earlier ones.
- It is expected that there may be implementation-specific details. These will need to be documented by the DRMAA library providers.

### 2.2 Distributed Application Environment

The DRMAA interfaces and programming model focus on providing sufficient mechanisms for job submission, monitoring and controlling, and obtaining the job's final status.

Ideally, DRMAA implementations and distributed applications should not be concerned with a particular DRMS environment and DRMS site policies. However, particularly for DRMAA implementations, this is not always possible to achieve and therefore solutions with minimal intrusiveness, such as "job categories" and "native specification" have been proposed. They abstract/aggregate the site-specific policies into simple strings that are interpreted by DRMAA implementations.

### **2.2.1 File staging**

DRMAA specification 1.0 does not have explicit file staging mechanisms. Setting implementation specific job template attributes could enable file staging, provided the DRMAA implementation and the DRM system supports it.

### **2.2.2 Job categories**

The DRMAA interface specification should allow ISVs to write DRM-enabled applications even though the properties of a concrete DRMS installation, in particular the configuration of the DRMS, cannot be known in advance.

Experiences made with integrations based on DRM CLI show that even when the same ISV application is run as a job with the same DRM system, the site-specific policies in effect differ widely. These policies are typically with respect to site-specific attributes such as:

- what resources are to be used by the job
- preferences where to run the job
- how the job should be scheduled relative to other jobs

For supporting the variety of policies, job specific requests expressed by DRM submit options are very common in the DRM product space.

However, for the most part, these options do not affect the job from the perspective of the ISV or from the perspective of the end user. This observation is the basis for "job categories", which insulate the ISV and end user from site-specific policies.

The job "categories concept" is the approach the DRMAA working group recommends for encapsulating site-specific details and completely hiding these details from applications making use of the DRMAA interface. The basic idea is to allow site administrators to create a job category that is suitable for an application to be dispatched by the DRMS, and the associated category name will be specified as a job submission attribute. The category name can be used by the DRMAA library to determine site-specific resource and functional requirements of jobs in the category. Such requirements need to be configurable by the site operating a DRMS and deploying an ISV application on top of it.

An example can help to illustrate this idea:

At site A, rendering application X is used in a heterogeneous clustered environment which is managed by a DRMS. Since application X is only available at a subset of these machines the administrator sets up the DRMS in a way requiring from the end-users to put a `-l X=true` into their submit command line.

At site B, the same application is used in a homogenous clustered environment with rendering application X supported at all machines managed by the DRMS. However since X jobs do compete with applications Y sharing the same resources and X applications are to be treated with higher priority than Y jobs end-users need to put a -p 1023 into their submit command line for raising the dispatch priority.

An integration based on categories will allow submitting X jobs through the DRMAA interface in compliance with the policies of both sites A and B without the need to know about these policies. The ISV does this by specifying "X" as the category used for X rendering jobs submitted through the DRMAA interface and by mentioning this in the "DRM integration" section of the X rendering software documentation.

The administrators at site A and site B read the documentation or installation instructions about the "X" DRMAA category. The documentation of their DRMS contains directions about the category support of their DRMAA interface implementation. From this documentation they learn how to configure their DRMS in a way that "-l X=true" is used for "X" jobs at site A while "-p 1023" is used at site B for those jobs.

As far as the DRMAA interface specification is concerned only a standardized mechanism for specifying the category is required. The mechanism for associating the policy related portion of the submit command line to the job is to be delivered by each DRMAA implementation. A standardization of this mechanism is beyond the DRMAA standardization effort, because it is too much related to the administrative interface and it is anticipated that for different DRMS different mechanisms will be appropriate.

### **2.2.3 Native specification**

The benefit of the categories concept from the previous section is that it provides a means for completely hiding site-specific policy details to be considered with a DRMAA job submission for a whole class of jobs. The drawback however of this concept is that it requires one job category to be maintained for each policy to be used.

To allow the DRMAA interface to also be used for the submission of jobs where job-individual policy specification is required "native specification" is supported. Native specification can be used without the requirement to maintain job categories. Instead of specifying a category name and having the DRMAA implementation associate the corresponding job submit options; the use of native specification will allow directly specifying these submit options.

An example can help to illustrate this idea:

In order to implement the example from section 2.2.2 via native specifications, the native option string "-l X=true" had to be passed directly to the DRMAA interface while "-p 1023" had to be used at site B.

As far as the DRMAA interface specification is concerned the native specification is an opaque string and interpreted by each DRMAA library. It is possible to use job categories and native specification with the same job submission for policy specification. It is assumed that in this case the DRMAA library is capable of joining the outcome of the two policy sources in a reasonable way.

#### **2.2.4 Building portals**

The nature of the DRMAA implementation, as a shared library, makes it a good candidate for inclusion in a Web Server to support a Web Portal to a DRMS.

DRMAA library could be:

- Linked by a collection of CGI scripts that are referenced by resident Web Pages.
- Linked in a Web Server as a separate module.
- Built as a Perl ( the same applies to a lesser degree to other scripting languages )
  - module that is included in mod\_perl
  - module accessed from Perl CGI scripts

The questions about maintaining state, security, and authentication and authorization, require that DRMAA implementation is viewed as just one component of a DRM Web Portal. Clearly, evaluation of the Web Portal needs and their impact on DRMAA specification is beyond the scope of the current document and DRMAA Charter.

### **2.3 Interface Routines General Description**

The routines are naturally grouped in five categories: init/exit, job template handling, job submission, job monitoring and control, and auxiliary or system routines like error message routines.

#### **2.3.1 Init and exit routines**

The calling sequence of the init routine should allow all of the considered DRM systems to be properly initialized, either by interfacing to the batch queue commands or to the DRMS API. Likewise, the exit routine should require parameters that will permit proper DRMS disengagement.

#### **2.3.2 Job template routines**

The remote jobs and their attributes are specified via the job template, which is an opaque parameter. The job attributes could be a string or a vector of string values. They are listed below.

Required job attributes are:

- Remote command to execute
- Remote command input parameters, a vector parameter
- Job state at submission
- Job environment, a vector parameter
- Job working directory
- Job category
- Native specification

- Standard input, output, and error streams
- E-mail distribution list to report the job completion and status, a vector parameter
- E-mail suppression
- Job start time
- Job name to be used for the job submission

### **2.3.3 Job submission routines**

The job submission routines come in two versions. There is one version for submitting individual jobs and one version for submitting bulk jobs.

### **2.3.4 Job monitoring and controlling routines**

The job monitoring and controlling API needs to handle:

- job releasing, stopping, resuming, and killing
- checking the exit code of the finished remote job
- checking the remote job status
- waiting for the remote job till the end of its execution
- waiting for all the jobs or a subset of the current session jobs to finish execution (this is a useful synchronization mechanism)

The Unix and Windows signals are replaced with the job control routines that have counterparts in DRM systems. The only nontraditional feature is the passing of DRMAA\_JOB\_IDS\_SESSION\_ALL string as job\_id to indicate operations on all job\_ids in the current process.

The remote job could be in following states:

- system hold
- user hold
- system and user hold simultaneously
- queued active
- system suspended
- user suspended
- system and user suspended simultaneously
- running
- finished (un)successfully


To this list we need to add a possibility of DRMAA implementation not being able to determine the status of the remote job. Note that a rejected job was not assigned a job\_id and consequently could not have a state.

### 2.3.5 Auxiliary routines

The auxiliary routines are needed to obtain a textual representation of errors and other DRMAA implementation specific information.

### 2.3.6 DRMAA job state transition diagram

The DRMAA job state transition diagram in Harel notation:



### 3. API Specification

For convenience, the API is divided in its five logical sections: init/exit, job template handling, job submission, job monitoring and control, and auxiliary routines. The interface is specified using an IDL-like language.

#### 3.1 Introduction

An IDL-like language is used here to avoid questions about allocation/de-allocation issues, the questions that are specific to an implementation language. The interface parameters could be IN, OUT, or INOUT parameters. For readers of this document who are familiar with the C programming language, they should be thought of as parameters passed by value or by reference respectively. Furthermore, the parameters could be scalar or vector values. The vector values are clearly documented.

To prevent interface name collisions all the routines have a prefix "drmaa\_".

##### 3.1.1 Error codes

All of the interfaces, except a few that return scalar values and cannot fail, return an error code on exit. Successful return is indicated by return value DRMAA\_ERRNO\_SUCCESS. All internal errors are indicated with DRMAA\_ERRNO\_INTERNAL\_ERROR error. Invalid argument is flagged as DRMAA\_ERRNO\_INVALID\_ARGUMENT. The return codes are specified and listed in section 3.3.

The error code can be provided to the drmaa\_strerror routine to retrieve a textual representation of the error. Routines may output a context specific error string which may be used in addition to the textual representation obtained from the error code. The parameter used to convey the context specific error will be ignored by the routine whenever success is returned. The length of any output context specific error string will not exceed DRMAA\_ERROR\_STRING\_BUFFER.

##### 3.1.2 DRMAA sessions

There is only one DRMAA session open at the time. Another session could be opened only after the current one is closed. Nesting of sessions is not supported. It is expected that the DRMAA library will free all the session resources, although this is not guaranteed, so old session resources are not to be used later. Job Ids are valid from one session to another. Job control routines should work correctly if a job Id came from a previous DRMAA session, provided the current DRMAA session knows how to resolve the job Id from the previous session. The burden is on the user to match previous job Id's with appropriate DRMAA sessions, i.e. DRM system servers. Re-startable applications have to make job Id's persistent in order to access the already submitted jobs.

##### 3.1.3 Run usage data

A DRMAA implementation collects remote run usage data (rusage variable) after the remote job run and job finish information (stat variable). The user can reap this data only once. The implementation is free to "garbage collect" the reaped data at a convenient time. Only the data from the current session's job Id is guaranteed to be available. Reaping data from other session job Id's is implementation specific.

### 3.1.4 Precedence rules

The attributes set by using API routines are set at the compile time. The attributes set via "job categories" are set at the installation time. The attributes set by the "native specification" are set at the run time. In principle that should determine the precedence rules, but these ideal precedence rules are not always achievable in practice due to complex interaction of attributes. Moreover, certain attributes in job categories would not be allowed to be overridden. The precedence rules are therefore implementation specific.

### 3.1.5 Site specific requirements

Job categories and native specifications are two means for describing site specific requirements. Setting job categories, if they are supported by a DRMAA implementation is implementation specific. On the other hand, setting of the native specification while straightforward in the user code could be a challenge if the user needs to provide a complex set of options. Quotation marks are especially problematic if only one variable is used for a set of native specification options.

Here are the recommendations to developers as to how to use this feature effectively:

- For each class of remote jobs give end users a chance to specify site specific environments, like a queue where to send remote jobs, architecture(s) where the remote applications are available, etc.
- Let users specify native specifications in a file, if the distributed application has several classes of jobs to submit or several DRMAA sessions.
- Applications with a graphical user interface could have a dedicated dialog for this purpose.

### 3.1.6 Job valuator

Before a submitted job enters a queue it is passed thru a valuator that determines if the job attributes as specified as valid. If yes, a job Id is returned and job is successfully queued. If not, the job is rejected and no job Id is returned and no job state is possible.

## 3.2 DRMAA API

```
/* ----- Major Assumptions/Restrictions ----- */
```

No explicit file staging.

Job Id Uniqueness -- "As unique as the underlying DRM makes them"

```
/*Global constants */
```

```
DRMAA_ERROR_STRING_BUFFER      = 1024
DRMAA_JOBNAME_BUFFER           = 1024
DRMAA_SIGNAL_BUFFER            = 32
DRMAA_TIMEOUT_WAIT_FOREVER     = -1
DRMAA_TIMEOUT_NO_WAIT          = 0
```

### 3.2.1 Initialization and exit routines

```
drmaa_init(contact, drmaa_context_error_buf)
```

IN contact /\* contact information for DRM system (string) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Initialize DRMAA API library and create a new DRMAA Session. 'Contact' is an implementation dependent string which may be used to specify which DRM system to use. This routine must be called before any other DRMAA calls, except for drmaa\_version(). If 'contact' is NULL, the default DRM system will be used. drmaa\_init() should be called by only one of the threads. The main thread could be a good choice. A call by another thread would trigger a message that nested DRMAA sessions are not allowed.

drmaa\_init routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_INVALID\_CONTACT\_STRING, DRMAA\_ERRNO\_ALREADY\_ACTIVE\_SESSION, or DRMAA\_ERRNO\_DEFAULT\_CONTACT\_STRING\_ERROR.

#### **drmaa\_exit(drmaa\_context\_error\_buf)**

INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Disengage from DRMAA library and allow the DRMAA library to perform any necessary internal clean up. This routine ends this DRMAA Session, but does not effect any jobs (e.g., queued and running jobs remain queued and running). drmaa\_exit() should be called by only one of the threads. Other thread calls to drmaa\_exit() will fail since there is no active session.

drmaa\_exit routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_DRMS\_EXIT\_ERROR or DRMAA\_ERRNO\_NO\_ACTIVE\_SESSION.

### **3.2.2 Job template routines**

#### **drmaa\_allocate\_job\_template( jt, drmaa\_context\_error\_buf )**

jt /\* job template (opaque handle) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Allocate a new job template. drmaa\_allocate\_job\_template() returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE or DRMAA\_ERRNO\_INTERNAL\_ERROR.

#### **drmaa\_delete\_job\_template(jt, drmaa\_context\_error\_buf )**

INOUT jt /\* job template (opaque handle) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Deallocate a job template. This routine has no effect on jobs. drmaa\_delete\_job\_template() returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE or DRMAA\_ERRNO\_INTERNAL\_ERROR.

**drmaa\_set\_attribute(jt, name, value, drmaa\_context\_error\_buf )**  
INOUT jt /\* job template (opaque handle) \*/  
IN name /\* attribute name (string) \*/  
IN value /\* attribute value (string) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Adds ('name', 'value') pair to list of attributes in job template 'jt'.  
Only non-vector attributes may be passed.

drmaa\_set\_attribute routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise  
DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_FORMAT,  
DRMAA\_ERRNO\_INVALID\_ARGUMENT,  
DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_VALUE, or  
DRMAA\_ERRNO\_CONFLICTING\_ATTRIBUTE\_VALUES.

**drmaa\_get\_attribute(jt, name, value, drmaa\_context\_error\_buf )**  
IN jt /\* job template (opaque handle) \*/  
IN name /\* attribute name (string) \*/  
OUT value /\* attribute value (string) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

If 'name' is an existing non-vector attribute name in the job template  
'jt', then the value of 'name' is returned; otherwise, NULL is returned.

drmaa\_get\_attribute routine returns DRMAA\_ERRNO\_SUCCESS  
on success, otherwise DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_VALUE.

**drmaa\_set\_vector\_attribute(jt, name, values, drmaa\_context\_error\_buf )**  
INOUT jt /\* job template (opaque handle) \*/  
IN name /\* attribute name (string) \*/  
IN values /\* vector of attribute value (string vector) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Adds ('name', 'values') pair to list of vector attributes in job template 'jt'.  
Only vector attributes may be passed.

drmaa\_set\_vector\_attribute routine returns DRMAA\_ERRNO\_SUCCESS  
on success, otherwise  
DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_FORMAT,  
DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_VALUE,  
DRMAA\_ERRNO\_CONFLICTING\_ATTRIBUTE\_VALUES.

**drmaa\_get\_vector\_attribute(jt, name, values, drmaa\_context\_error\_buf )**  
IN jt /\* job template (opaque handle) \*/  
IN name /\* attribute name (string) \*/  
OUT values /\* vector of attribute value (string vector) \*/

If 'name' is an existing vector attribute name in the job template 'jt',  
then the values of 'name' are returned; otherwise, NULL is returned.

drmaa\_get\_vector\_attribute routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_VALUE.

**drmaa\_get\_attribute\_names( names, drmaa\_context\_error\_buf )**

OUT names /\* vector of attribute name (string vector) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Returns the set of supported attribute names whose associated value type is String. This set will include supported DRMAA reserved attribute names and native attribute names.

**drmaa\_get\_vector\_attribute\_names( names, drmaa\_context\_error\_buf )**

OUT names /\* vector of attribute name (string vector) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Returns the set of supported attribute names whose associated value type is String Vector. This set will include supported DRMAA reserved attribute names and native attribute names.

### 3.2.3 Job attributes

#### 3.2.3.1 Mandatory attributes

The following are reserved attribute names available in all implementations of DRMAA v1.0 (and their respective meanings). Vector attributes are marked with a 'V':

**remote command to execute ( string )**

It is relative to the execution host.  
It is evaluated on the execution host.  
No binary file management is done.  
The attribute name is drmaa\_remote\_command.

**V input parameters ( vector of strings )**

These parameters are passed as arguments to the job.  
The attribute name is drmaa\_v\_argv.

**job state at submission ( string value )**

This might be useful for a rather rudimentary, but very general job dependent execution.  
The states are drmaa\_hold and drmaa\_active:  
drmaa\_active means job has been queued, and is eligible to run  
drmaa\_hold means job has been queued, but it is NOT eligible to run  
The attribute name is drmaa\_js\_state.

**V job environment ( vector of strings )**

The environment values that define the remote environment.  
Each string will comply with the format <name>=<value>.  
The values override the remote environment values if there is a collision.

If above is not possible, it is implementation dependent.  
The attribute name is `drmaa_v_env`.

#### **job working directory ( string )**

This attribute specifies the directory where the job is executed.  
If not set, it is implementation dependent.  
Evaluated relative to the execution host.

A `$drmaa_hd_ph$` placeholder at the begin denotes the remaining portion of the `directory_name` as a relative directory name resolved relative to the job users home directory at the execution host.

The `$drmaa_incr_ph$` placeholder can be used at any position within the `directory_name` of parametric job templates and will be substituted by the underlying DRM system with the parametric jobs' index.

The `directory_name` must be specified in a syntax that is common at the host where the job is executed.

If set and no placeholder is used an absolute directory specification is expected.

If set and the directory does not exist the job enters the state `DRMAA_PS_FAILED`.

The attribute name is `drmaa_wd`.

#### **job category ( string )**

An opaque string specifying how to resolve site-specific resources and/or policies.  
The attribute name is `drmaa_job_category`.

#### **native specification ( string )**

An opaque string that is passed by the end user to DRMAA to specify site-specific resources and/or policies.

The attribute name is `drmaa_native_specification`.

#### **e-mail address ( vector of strings )**

It is used to report the job completion and status.

The attribute name is `drmaa_v_email`.

#### **e-mail suppression ( string )**

It is used to block sending e-mail by default.

1 block

0 do not block.

The attribute name is `drmaa_block_email`

#### **job start time ( string )**

This attribute specifies the earliest time when the job may be eligible to be run.

This is a required attribute named `drmaa_start_time`

The value of the attribute will be of the form

`[[[CC]YY/]MM/DD] hh:mm[:ss] [{}+}UU:uu]`

where

CC is the first two digits of the year (century-1)

YY is the last two digits of the year

MM is the two digits of the month [01,12]

DD is the two digit day of the month [01,31]  
hh is the two digit hour of the day [00,23]  
mm is the two digit minute of the day [00,59]  
ss is the two digit second of the minute [00,61]  
UU is the two digit hours since (before) UTC  
uu is the two digit minutes since (before) UTC  
If the optional UTC-offset is not specified, the offset associated with the local timezone will be used.  
If the day (DD) is not specified, the current day will be used unless the specified hour:mm:ss has already elapsed, in which case the next day will be used.  
Similarly for month (MM), year (YY), and century-1 (CC).

Example:

The time: Sep 3 4:47:27 PM PDT 2002,  
could be represented as: 2002/09/03 16:47:27 -07:00

### **job name**

A job name will be comprised of alphanumeric and \_ characters.  
The drmaa-implementation will not provide the client with a job name longer than DRMAA\_JOBNAME\_BUFFER -1 (1023) characters.  
The drmaa-implementation may truncate any client-provided job name to an implementation-defined length which shall be at least 31 characters.  
The attribute name is drmaa\_job\_name

### **input stream ( string )**

Specifies the jobs' standard input.  
Unless set elsewhere, if not explicitly set in the job template, the job is started with an empty input stream.  
If set specifies the network path of the jobs input stream file of the form [hostname]:file\_path  
When the drmaa\_transfer\_files job template attribute is supported and contains the character 'i', the input file will be fetched by the underlying DRM system from the specified host or from the submit host if no hostname is specified.  
When the drmaa\_transfer\_files job template attribute is not supported or does not contain the character 'i', the input file is always expected at the host where the job is executed irrespectively of a possibly hostname specified.  
The \$drmaa\_incr\_ph\$ placeholder can be used at any position within the file\_path of parametric job templates and will be substituted by the underlying DRM system with the parametric jobs' index.  
A \$drmaa\_hd\_ph\$ placeholder at the begin of the file\_path denotes the remaining portion of the file\_path as a relative file specification resolved relative to the job users home directory at the host where the file is located.  
A \$drmaa\_wd\_ph\$ placeholder at the begin of the file\_path denotes the remaining portion of the file\_path as a relative file specification resolved relative to the jobs working directory at the host where the file is located.  
The file\_path must be specified in a syntax that is common at the host where the file is located.  
If set and the file can't be read the job enters the state DRMAA\_PS\_FAILED.  
The attribute name is drmaa\_input\_path.

### **output stream ( string )**

Specifies how to direct the jobs' standard output.

If not explicitly set in the job template, the whereabouts of the jobs output stream is not defined.

If set specifies the network path of the jobs output stream file of the form [hostname]:file\_path

When the drmaa\_transfer\_files job template attribute is supported and contains the character 'o', the output file will be transferred by the underlying DRM system to the specified host or to the submit host if no hostname is specified.

When the drmaa\_transfer\_files job template attribute is not supported or does not contain the character 'o', the output file is always kept at the host where the job is executed irrespectively of a possibly hostname specified.

The \$drmaa\_incr\_ph\$ placeholder can be used at any position within the file\_path of parametric job templates and will be substituted by the underlying DRM system with the parametric jobs' index.

A \$drmaa\_hd\_ph\$ placeholder at the begin of the file\_path denotes the remaining portion of the file\_path as a relative file specification resolved relative to the job users home directory at the host where the file is located.

A \$drmaa\_wd\_ph\$ placeholder at the begin of the file\_path denotes the remaining portion of the file\_path as a relative file specification resolved relative to the jobs working directory at the host where the file is located.

The file\_path must be specified in a syntax that is common at the host where the file is located.

If set and the file can't be written before execution the job enters the state DRMAA\_PS\_FAILED.

The attribute name is drmaa\_output\_path.

#### **error stream ( string )**

Specifies how to direct the jobs' standard error.

If not explicitly set in the job template, the whereabouts of the jobs error stream is not defined.

If set, specifies the network path of the jobs error stream file of the form [hostname]:file\_path

When the drmaa\_transfer\_files job template attribute is supported and contains the character 'e', the output file will be transferred by the underlying DRM system to the specified host or to the submit host if no hostname is specified.

When the drmaa\_transfer\_files job template attribute is not supported or does not contain the character 'e', the error file is always kept at the host where the job is executed irrespectively of a possibly hostname specified.

The \$drmaa\_incr\_ph\$ placeholder can be used at any position within the file\_path of parametric job templates and will be substituted by the underlying DRM system with the parametric jobs' index.

A \$drmaa\_hd\_ph\$ placeholder at the begin of the file\_path denotes the remaining portion of the file\_path as a relative file specification resolved relative to the job users home directory at the host where the file is located.

A \$drmaa\_wd\_ph\$ placeholder at the begin of the file\_path denotes the remaining portion of the file\_path as a relative file specification resolved relative to the jobs working directory at the host where the file is located.

The file\_path must be specified in a syntax that is common at the host where the file is located.

If set and the file can't be written before execution the job enters the state DRMAA\_PS\_FAILED.

The attribute name is drmaa\_error\_path.

#### **join files ( string )**

Specifies if the error stream should be intermixed with the output stream.

If not explicitly set in the job template the attribute defaults to 'n'.

Either 'y' or 'n' can be specified.

If 'y' is specified the underlying DRM system will ignore the value of the `drmaa_error_path` attribute and intermix the standard error stream with the standard output stream as specified with `drmaa_output_path`.  
The attribute name is `drmaa_join_files`.

### 3.2.3.2 Optional attributes

The following are reserved attribute names available which are not required to be implemented by a conforming DRMAA v1.0 implementation. For attributes that are implemented, the meanings are required to be as follows:

Note that the list of attributes which are implemented can be programmatically obtained using the `drmaa_get_attribute_names` and `drmaa_get_vector_attribute_names` routines.

#### **transfer files ( string )**

Specifies how to transfer files between hosts.  
If not explicitly set in the job template the attribute defaults to ".  
Any combination of 'e', 'i' and 'o' can be specified.  
Whether the character 'e' is specified impacts the behaviour of the `drmaa_error_path` attribute.  
Whether the character 'i' is specified impacts the behaviour of the `drmaa_input_path` attribute.  
Whether the character 'o' is specified impacts the behaviour of the `drmaa_output_path` attribute.  
The attribute name is `drmaa_transfer_files`.

#### **absolute job termination time ( string )**

This attribute specifies a deadline after which the DRMS will terminate a job.

This is a reserved attribute named `drmaa_deadline_time`

The value of the attribute will be of the form

`[[[CC]YY/]MM/]DD] hh:mm[:ss] [-|+}UU:uu]`

where

CC is the first two digits of the year (century-1)

YY is the last two digits of the year

MM is the two digits of the month [01,12]

DD is the two digit day of the month [01,31]

hh is the two digit hour of the day [00,23]

mm is the two digit minute of the day [00,59]

ss is the two digit second of the minute [00,61]

UU is the two digit hours since (before) UTC

uu is the two digit minutes since (before) UTC

If an optional portion of the time specification is omitted, then the termination time will be determined based upon the job's earliest start time.

If the day (DD) is not specified, the earliest start day for the job will be used unless the specified hour:mm:ss precedes the corresponding portion of the job start time, in which case the next day will be used.

Similarly for month (MM), year (YY), and century-1 (CC).

Example:

The time: Sep 3 4:47:27 PM PDT 2002,  
could be represented as: 2002/09/03 16:47:27 -07:00

#### **wall clock time limit ( string )**

This attribute specifies when the job's wall clock time limit has been exceeded. The DRMS will terminate a job which has exceeded its wall clock time limit. Note that the suspended time is also accumulated here.

This is a reserved attribute named `drmaa_wct_hlimit`

The value of the attribute will be of the form

`[[h:]m:]s`

where

h is one or more digits representing hours

m is one or more digits representing minutes

s is one or more digits representing seconds

Example:

To terminate a job after 2 hours and 30 minutes,  
any of the following can be passed:

2:30:0, 1:90:0, 150:0

#### **soft wall clock time limit ( string )**

This attribute specifies an estimate as to how long the job will need wall clock time to complete. Note that the suspended time is also accumulated here.

This attribute is intended to assist the scheduler.

If the time specified is insufficient, the `drmaa`-implementation may impose a scheduling penalty.

This is a reserved attribute named `drmaa_wct_slimit`

The value of the attribute will be of the form

`[[h:]m:]s`

where

h is one or more digits representing hours

m is one or more digits representing minutes

s is one or more digits representing seconds

#### **job run duration hlimit ( string )**

This attribute specifies how long the job may be in a running state before its limit has been exceeded, and therefore is terminated by the DRMS.

This is a reserved attribute named `drmaa_run_duration_hlimit`

The value of the attribute will be of the form

`[[h:]m:]s`

where

h is one or more digits representing hours

m is one or more digits representing minutes

s is one or more digits representing seconds

### **job run duration slimit ( string )**

This attribute specifies an estimate as to how long the job will need to remain in a running state to complete.

This attribute is intended to assist the scheduler. If the time specified is insufficient, the drmaa-implementation may impose a scheduling penalty.

This is a reserved attribute named `drmaa_run_duration_slimit`

The value of the attribute will be of the form

`[[h:]m:]s`

where

h is one or more digits representing hours

m is one or more digits representing minutes

s is one or more digits representing seconds

### **3.2.4 Job submission routines**

#### **drmaa\_run\_job(job\_id, jt, drmaa\_context\_error\_buf )**

OUT job\_id /\* job identifier (string) \*/  
IN jt /\* job template (opaque handle) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Submit a job with attributes defined in the job template 'jt'.

The job identifier 'job\_id' is a printable, NULL terminated string, identical to that returned by the underlying DRM system.

drmaa\_run\_job routine returns `DRMAA_ERRNO_SUCCESS` on success, otherwise `DRMAA_ERRNO_TRY_LATER`, `DRMAA_ERRNO_DENIED_BY_DRM`, `DRMAA_ERRNO_DRM_COMMUNICATION_FAILURE`, or `DRMAA_ERRNO_AUTH_FAILURE`.

#### **drmaa\_run\_bulk\_jobs(job\_ids, jt, start, end, incr, drmaa\_context\_error\_buf )**

OUT job\_ids /\* job identifiers (array of strings) \*/  
IN jt /\* job template (opaque handle) \*/  
IN start /\* beginning index ( unsigned integer?) \*/  
IN end /\* ending index ( unsigned integer?) \*/  
IN incr /\* loop increment (integer) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Submit a set of parametric jobs, dependent on the implied loop index, each with attributes defined in the job template 'jt'.

The job identifiers 'job\_ids' are all printable, NULL terminated strings, identical to those returned by the underlying DRM system. Nonnegative loop bounds are mandated to avoid file names that start with minus sign like command line options.

The special index placeholder is a DRMAA defined string

`drmaa_incr_ph /* == $incr_pl$ */`

this is used to construct parametric job templates.

For example:

```
drmaa_set_attribute(pjt, "stderr", drmaa_incr_ph + ".err" ); /*  
C++/java string syntax used */
```

drmaa\_run\_bulk\_jobs routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_TRY\_LATER, DRMAA\_ERRNO\_DENIED\_BY\_DRM, DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE, or DRMAA\_ERRNO\_AUTH\_FAILURE.

### 3.2.5 Job control routines

```
drmaa_control(job_id, action, drmaa_context_error_buf )  
IN job_id                /* job identifier (string) */  
IN action                 /* control action (const) */  
INOUT drmaa_context_error_buf /*Contains a context sensitive error upon failed return*/
```

Start, stop, restart, or kill the job identified by 'job\_id'.  
If 'job\_id' is DRMAA\_JOB\_IDS\_SESSION\_ALL, then this routine acts on all jobs \*submitted\* during this DRMAA session.

The legal values for 'action' and their meanings are:

|                          |                                  |
|--------------------------|----------------------------------|
| DRMAA_CONTROL_SUSPEND:   | stop the job,                    |
| DRMAA_CONTROL_RESUME:    | (re)start the job,               |
| DRMAA_CONTROL_HOLD:      | put the job on-hold,             |
| DRMAA_CONTROL_RELEASE:   | release the hold on the job, and |
| DRMAA_CONTROL_TERMINATE: | kill the job.                    |

This routine returns once the action has been acknowledged by the DRM system, but does not necessarily wait until the action has been completed.

drmaa\_control routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE, DRMAA\_ERRNO\_AUTH\_FAILURE, DRMAA\_ERRNO\_RESUME\_INCONSISTENT\_STATE, DRMAA\_ERRNO\_SUSPEND\_INCONSISTENT\_STATE, DRMAA\_ERRNO\_HOLD\_INCONSISTENT\_STATE, DRMAA\_ERRNO\_RELEASE\_INCONSISTENT\_STATE, DRMAA\_ERRNO\_INVALID\_JOB.

```
drmaa_synchronize(job_ids, timeout, dispose, drmaa_context_error_buf )  
IN job_ids                /* job identifiers (array of strings) */  
IN timeout                /* how long we block in this call (signed long) */  
IN dispose                /* dispose reaping information (boolean)*/  
INOUT drmaa_context_error_buf /*Contains a context sensitive error upon failed return*/
```

Wait until all jobs specified by 'job\_ids' have finished execution. If 'job\_ids' is DRMAA\_JOB\_IDS\_SESSION\_ALL, then this routine waits for all jobs \*submitted\* during this DRMAA session. To prevent blocking indefinitely in this call, the caller could use timeout specifying after how many seconds to time out in this call.

If the call exits before timeout, all the jobs have been waited on or there was an interrupt.

If the invocation exits on timeout, the return code is `DRMAA_ERRNO_EXIT_TIMEOUT`. The caller should check system time before and after this call in order to check how much time has passed.

Dispose parameter specifies how to treat reaping information:

|       |     |  |
|-------|-----|--|
| True  | = 1 | "fake reap", i.e. dispose of the rusage data |
| False | = 0 | do not reap                                  |

`drmaa_synchronize` routine returns `DRMAA_ERRNO_SUCCESS` on success, otherwise `DRMAA_ERRNO_DRM_COMMUNICATION_FAILURE`, `DRMAA_ERRNO_AUTH_FAILURE`, `DRMAA_ERRNO_EXIT_TIMEOUT`, or `DRMAA_ERRNO_INVALID_JOB`.

#### **`drmaa_wait(job_id, stat, timeout, rusage, drmaa_context_error_buf)`**

|  |  |
|--|--|
| IN <code>job_id</code>                     | /* job identifier (string) or <code>DRMAA_JOB_IDS_SESSION_ANY</code> (string) */ |
| OUT <code>job_id</code>                    | /* job identifier of ended job (string) or NULL */                               |
| OUT <code>stat</code>                      | /* status code of job (integer) */   |
| IN <code>timeout</code>                    | /* how long we block in this call (signed long) */                               |
| OUT <code>rusage</code>                    | /* resource usage (string array) */  |
| INOUT <code>drmaa_context_error_buf</code> | /*Contains a context sensitive error upon failed return*/                        |

This routine waits for a job with `job_id` to fail or finish execution. If the special string `DRMAA_JOB_IDS_SESSION_ANY` is provided as the `job_id`, this routine will wait for any job from the session. This routine is modeled on the `wait3` POSIX routine. The `timeout` value is used to specify the desired behavior when a result is not immediately available.

The value `DRMAA_TIMEOUT_WAIT_FOREVER` (-1) can be specified to wait indefinitely for a result. The value `DRMAA_TIMEOUT_NO_WAIT` (0) may be specified to return immediately if no result is available. Alternatively, a number of seconds may be specified to indicate how long to wait for a result to become available.

If the call exits before timeout, either the job has been waited on successfully or there was an interrupt.

If the invocation exits on timeout, the return code is `DRMAA_ERRNO_EXIT_TIMEOUT`. The caller should check system time before and after this call in order to check how much time has passed.

The routine reaps jobs on a successful call, so any subsequent calls to `drmaa_wait` should fail returning an error `DRMAA_ERRNO_INVALID_JOB` meaning that the job has been already reaped. This error is the same as if the job was unknown. Failing due to an elapsed timeout has an effect that it is possible to issue `drmaa_wait` multiple times for the same `job_id`. When successful, the `rusage` information will be provided as an array of strings, where each string complies with the format `<name>=<value>`.

The string portion `<value>` contains the amount of resources consumed by the job and is opaque.

`drmaa_wait` routine returns `DRMAA_ERRNO_SUCCESS` on success, otherwise

DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE,  
DRMAA\_ERRNO\_AUTH\_FAILURE,  
DRMAA\_ERRNO\_NO\_RUSAGE,  
DRMAA\_ERRNO\_EXIT\_TIMEOUT,  
DRMAA\_ERRNO\_NO\_RUSAGE, or  
DRMAA\_ERRNO\_INVALID\_JOB.

The 'stat' drmaa\_wait parameter is used in a series of functions, defined below, for providing more detailed information about job termination if available. An analogous set of macros is defined in POSIX for analyzing wait3(2) OUT parameter 'stat'. The misleading upper-case function names reminding to macros are changed to lower-case names.

**drmaa\_wifexited(OUT exited, IN stat, INOUT drmaa\_context\_error\_buf)**

Evaluates into 'exited' a non-zero value if stat was returned for a job that terminated normally. A zero value can also indicate that although the job has terminated normally an exit status is not available or that it is not known whether the job terminated normally. In both cases drmaa\_wexitstatus() will not provide exit status information. A non-zero 'exited' value indicates more detailed diagnosis can be provided by means of drmaa\_wifsignaled(), drmaa\_wtermsig() and drmaa\_wcoredump().

**drmaa\_wexitstatus(OUT exit\_status, IN stat, INOUT drmaa\_context\_error\_buf)**

If the OUT parameter 'exited' of drmaa\_wifexited() is non-zero, this function evaluates into 'exit\_code' the exit code that the job passed to \_exit() (see exit(2)) or exit(3C), or the value that the child process returned from main.

**drmaa\_wifsignaled(OUT signaled, IN stat, INOUT drmaa\_context\_error\_buf )**

Evaluates into 'signaled' a non-zero value if status was returned for a job that terminated due to the receipt of a signal. A zero value can also indicate that although the job has terminated due to the receipt of a signal the signal is not available or that it is not known whether the job terminated due to the receipt of a signal. In both cases drmaa\_wtermsig() will not provide signal information.

**drmaa\_wtermsig(OUT signal, IN stat, INOUT drmaa\_context\_error\_buf )**

If the OUT parameter 'signaled' of drmaa\_wifsignaled(stat) is non-zero, this function evaluates into signal a string representation of the signal that caused the termination of the job. For signals declared by POSIX, the symbolic names are returned (e.g., SIGABRT, SIGALRM). For signals not declared by POSIX, any other string may be returned.

**drmaa\_wcoredump(OUT core\_dumped, IN stat, INOUT drmaa\_context\_error\_buf )**

If the OUT parameter 'signaled' of drmaa\_wifsignaled(stat) is non-zero, this function evaluates into 'core\_dumped' a non-zero value if a core image of the terminated job was created.

**drmaa\_wifaborted( OUT aborted, IN stat, INOUT drmaa\_context\_error\_buf )**

Evaluates into 'aborted' a non-zero value if 'stat' was returned for a job that ended before entering the running state.

**drmaa\_job\_ps( IN job\_id, OUT remote\_ps, drmaa\_context\_error\_buf );**

IN job\_id /\* job identifier (string) \*/  
OUT remote\_ps /\* program status (constant) \*/  
INOUT drmaa\_context\_error\_buf /\*Contains a context sensitive error upon failed return\*/

Get the program status of the job identified by 'job\_id'.  
The possible values returned in 'remote\_ps' and their meanings are:

|                                     |   |
|-------------------------------------|---|
| DRMAA_PS_UNDETERMINED<br>determined | = 00H : process status cannot be determined       |
| DRMAA_PS_QUEUED_ACTIVE              | = 10H : job is queued and active                  |
| DRMAA_PS_SYSTEM_ON_HOLD             | = 11H : job is queued and in system hold          |
| DRMAA_PS_USER_ON_HOLD               | = 12H : job is queued and in user hold            |
| DRMAA_PS_USER_SYSTEM_ON_HOLD        | = 13H : job is queued and in user and system hold |
| DRMAA_PS_RUNNING                    | = 20H : job is running                            |
| DRMAA_PS_SYSTEM_SUSPENDED           | = 21H : job is system suspended                   |
| DRMAA_PS_USER_SUSPENDED             | = 22H : job is user suspended                     |
| DRMAA_PS_DONE                       | = 30H : job finished normally                     |
| DRMAA_PS_FAILED                     | = 40H : job finished, but failed.                 |

DRMAA should always get the status of job\_id from DRM system, unless the previous status has been DRMAA\_PS\_FAILED or DRMAA\_PS\_DONE and the status has been successfully cached. Terminated jobs get DRMAA\_PS\_FAILED status.

drmaa\_synchronize routine returns DRMAA\_ERRNO\_SUCCESS on success, otherwise DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE, DRMAA\_ERRNO\_AUTH\_FAILURE, or DRMAA\_ERRNO\_INVALID\_JOB.

### 3.2.6 Auxiliary routines

**error\_string drmaa\_strerror ( errno );**

IN errno /\* Erno number (integer) \*/  
RETURNS /\* Readable text version of errno (constant string) \*/

Get the error message text associated with the errno number.

**contact drmaa\_get\_contact();**

OUT contact /\* Current contact information for DRM system (string) \*/

### **drmaa\_version(major, minor)**

OUT major /\* major version number (non-negative integer) \*/  
OUT minor /\* minor version number (non-negative integer) \*/

Returns the major and minor version numbers of the DRMAA library;  
for DRMAA 1.0, 'major' is 1 and 'minor' is 0.

### **DRM\_system drmaa\_get\_DRM\_system( )**

RETURNS /\* DRM system implementation information \*/

Output (string) is implementation dependent and could contain the DRM system and the implementation vendor as its parts.

## **3.3 List of DRMAA Errors**

----- these are relevant to all sections -----

### **DRMAA\_ERRNO\_SUCCESS**

Routine returned normally with success.

### **DRMAA\_ERRNO\_INTERNAL\_ERROR**

Unexpected or internal DRMAA error like memory allocation,  
system call failure, etc.

### **DRMAA\_ERRNO\_DRM\_COMMUNICATION\_FAILURE**

Could not contact DRM system for this request.

### **DRMAA\_ERRNO\_AUTH\_FAILURE**

The specified request is not processed successfully due to  
authorization failure.

### **DRMAA\_ERRNO\_INVALID\_ARGUMENT**

The input value for an argument is invalid.

----- init and exit specific -----

### **DRMAA\_ERRNO\_INVALID\_CONTACT\_STRING**

Initialization failed due to invalid contact string.

### **DRMAA\_ERRNO\_DEFAULT\_CONTACT\_STRING\_ERROR**

DRMAA could not use the default contact string to connect to DRM  
system.

DRMAA\_ERRNO\_DRMS\_INIT\_FAILED  
Initialization failed due to failure to init DRM system.

DRMAA\_ERRNO\_ALREADY\_ACTIVE\_SESSION  
Initialization failed due to existing DRMAA session.

DRMAA\_ERRNO\_NO\_ACTIVE\_SESSION  
Exit routine failed because there is no active session.

DRMAA\_ERRNO\_DRMS\_EXIT\_ERROR  
DRM system disengagement failed.

----- job attributes specific -----  
DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_FORMAT  
The format for the job attribute value is invalid.

DRMAA\_ERRNO\_INVALID\_ATTRIBUTE\_VALUE  
The value for the job attribute is invalid.

DRMAA\_ERRNO\_CONFLICTING\_ATTRIBUTE\_VALUES  
The value of this attribute is conflicting with a previously set attributes.

----- job submission specific -----  
DRMAA\_ERRNO\_TRY\_LATER  
Could not pass job now to DRM system. A retry may succeed however (saturation).

DRMAA\_ERRNO\_DENIED\_BY\_DRM  
The DRM system rejected the job. The job will never be accepted due to DRM configuration or job template settings.

----- job control specific -----  
DRMAA\_ERRNO\_INVALID\_JOB  
The job specified by the 'jobid' does not exist.

DRMAA\_ERRNO\_RESUME\_INCONSISTENT\_STATE  
The job has not been suspended. The RESUME request will not be processed.

DRMAA\_ERRNO\_SUSPEND\_INCONSISTENT\_STATE

The job has not been running, and it cannot be suspended.

DRMAA\_ERRNO\_HOLD\_INCONSISTENT\_STATE

The job cannot be moved to a HOLD state.

DRMAA\_ERRNO\_RELEASE\_INCONSISTENT\_STATE

The job is not in a HOLD state.

DRMAA\_ERRNO\_EXIT\_TIMEOUT

We have encountered a time-out condition for `drmaa_synchronize` or `drmaa_wait`.

DRMAA\_ERRNO\_NO\_RUSAGE

This error code is returned by `drmaa_wait()` when a job has finished but no `rusage` and `stat` data could be provided.

#### **4. Security Considerations**

Security issues are not discussed in this document. The scheduling scenario described here assumes that security is handled at the point of job authorization/execution on a particular resource.

## Author Information

Roger Brobst  
rbrobst@cadence.com  
Cadence Design Systems, Inc  
555 River Oaks Parkway  
San Jose, CA 95134

Waiman Chan  
waimanc@us.ibm.com  
International Business Machines Corporation  
2455 South Road  
Poughkeepsie, New York 12601

Fritz Ferstl  
fritz.ferstl@sun.com  
Sun Microsystems GmbH  
Dr.-Leo-Ritter-Str. 7  
D-93049 Regensburg  
Germany

Jeffrey T. Gardiner  
gardiner@imaging.robarts.ca  
Robarts Research Institute  
PO Box 5015, 100 Perth Drive  
London ON, N6A 5K8  
Canada

Andreas Haas  
andreas.haas@sun.com  
Sun Microsystems GmbH  
Dr.-Leo-Ritter-Str. 7  
D-93049 Regensburg  
Germany

Bill Nitzberg  
bill@computer.org  
Veridian, PBS Products  
2672 Bayshore Pkwy, Suite 810  
Mountain View, CA 94043

Hrabri L. Rajic  
hrabri.rajic@intel.com  
Intel Americas Inc.  
1906 Fox Drive  
Champaign, IL 61820

John Tollefsrud  
j.t@sun.com  
Sun Microsystems  
200 Jefferson Drive UMPK29-302  
Menlo Park, CA 94025

## **Intellectual Property Statement**

The GGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the GGF Secretariat.

The GGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the GGF Executive Director.

## **Full Copyright Notice**

Copyright (C) Global Grid Forum (date). All Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the GGF or other organizations, except as needed for the purpose of developing Grid Recommendations in which case the procedures for copyrights defined in the GGF Document process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the GGF or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and THE GLOBAL GRID FORUM DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE."