

Network Service Interface Topology Service

Status of This Document

Group Working Draft (GWD), candidate Recommendations Proposed (R-P).

Copyright Notice

Copyright © Open Grid Forum (2008-2012). Some Rights Reserved. Distribution is unlimited.

Abstract

This document describes a normative schema which allows the description of service plane objects required for the Network Service Interface Connection Service. Additionally it describes a set of distribution mechanisms for the network topology descriptions.

Contents

1 Introduction

The NSI Connection Service requires topology descriptions to do pathfinding. In order to do that some representation of the topology is required. Once represented, some form of topology distribution is also needed. This document describes some requirements for the NSI Topology Service, suggests a short-term implementation and a strategy for better long-term support.

In the first section we describe what is necessary for the topology to support, what kind of elements should be in there. In the next section we describe the distribution requirements, some possible solutions and a recommended solution for the short-term and also for the longer term

1.1 Scope

The Network Markup Language is designed to create a functional description of multi-layer networks and multi-domain networks. An example of a multi-layered network can be a virtualised network, but also using different technologies. The multi-domain network descriptions can include aggregated or abstracted network topologies. NML can not only describe a static network topology, but also its capabilities and its configuration.

NML is aimed at logical connection-oriented network topologies. It can also be used to describe physical networks or packet-oriented networks, although the current base schema does not contain classes or properties to explicitly deal with signal degradation, or complex routing tables.

NML only attempts to describe the data plane of a computer network, not the control plane. It does contain extension mechanism to easily tie it with network provisioning standards and with network monitoring standards.

1.2 Notational Conventions

The keywords “MUST”, “MUST NOT”, “REQUIRED”, “SHALL”, “SHALL NOT”, “SHOULD”, “SHOULD NOT”, “RECOMMENDED”, “MAY”, and “OPTIONAL” are to be interpreted as described in [?].

This schema defines classes, attributes, relations, parameters and logic. Objects are instances of classes, and the type of an object is a class.

Names of classes are capitalised and written in italics (e.g. the *Node* class). Names of relations are written in camel case and in italics (e.g. the *hasNode* relation).

2 Representation of Network Topologies

In order to use the NSI, some form of topology representation is required. An introduction to this representation and the issues involved in creating network representations for the NSI is described below. A diagram that provides some generic insight into NSI Topology is provided in figure ??.

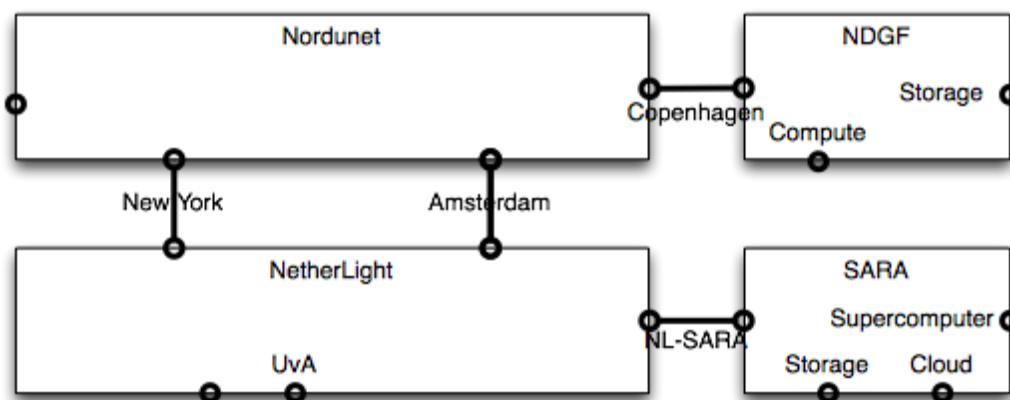


Figure 1: Abstract view of an example NSI Topology

2.1 Introduction to concept of STPs

The basic network topology for the Network Service Interface (NSI) consists of networks, points and connections. The NSI can be used to request a connection between two different points, which is then implemented using the connection(s) between those points. Since each of the points in the network topology can terminate a network service, they are called Service Termination Points (STPs). The figure above contains several of these, for example the Storage point at SARA, another example is the connection point between the network on the edge of the SARA network which connects to the Netherlight network.

2.2 Identifying STPs in a request

The STPs in the Network Topology generally have two different roles:

- Endpoints within the network – points which are of interest to users, used as source and destination points
- One Part of a Service Demarcation Point (SDP) – meaning connections to other networks, used as transit points

For the NSI network service, knowledge of the connections between networks is necessary to enable pathfinding. It is not strictly necessary to know all of the user endpoints in a network, as long as it is known in which network the endpoint is. It is assumed that once the network is reached, the endpoint within the network can be reached as well.

To allow for pathfinding it is necessary that the Topology Service can identify the Networks, and the connections between those Networks, the SDPs. It is not necessary to distribute all of the endpoints within the network, which makes the distribution process much simpler.

This means that a request to NSI must contain the source and destination network, as well as the connection points within those respective networks.

2.3 Explicit routing using STPs

In a regular request only the source and destination STPs and their networks are specified. The selected path between those STPs is left to the NSA managing the request. In NSI v2.0 it is also possible to steer the requested path into a specific direction by defining intermediate STPs that the path must touch.

Instead of a normal request with just a source and destination STP, the explicitly routed request will contain a path element which contains an indication of the path that should be taken through the NSI Network.

If the Path object contains a description of the complete path end-to-end then this is simply a question of availability. However, problems can arise if the path object only contains a single “explicit route object”(ERO) that it must touch. With the current NSI implementation and its bi-directional model, there is no way to know from which way to cross that ERO. The proposal is to use uni-directional path elements, to avoid ambiguity in the path direction (the uni-directionality of the ERO implicitly defines the direction).

2.4 Requirements for Topology Descriptions

Taking the description above into account, and with the general idea of the Network Service Interface in mind, we come to the following requirements for the network topology description:

Scalable : An NSA does not need to be aware of all STPs in other networks;

Compact : A Topology description should be able to group individual data transport capabilities in one object rather than specifying each possible VLAN for example;

Abstract : The topology description should list the connections between domains, not how these connections are implemented;

Compatible : It should be possible to relate the NSI topology to other topologies, e.g. as used for monitoring. Preferably the NSI topology should be compatible with the NML topology;

Flexible : The topology description should support future extensions, e.g. different connection types (unidirectional, bidirectional, or multipoint connections), different levels of abstraction of the network (subtopologies), or multihoming scenarios;

Unambiguous : The direction of an STP in an ERO should be unambiguous.

2.5 Topology Identifiers

To satisfy the unambiguity and flexibility requirements, we propose to describe *STPs* as two unidirectional *Ports*, since these have a single direction, there can be no misunderstanding. These unidirectional ports also easily allow point-to-multipoint requests.

To satisfy the compact requirement, we propose to allow *PortGroups* over *Ports*. A *PortGroup* groups together several *Ports* which have a single identifying attribute, for example a VLAN label.

To satisfy the scalability and flexibility requirement we propose to add the Topology identifier as an added context for an STP in a request. This makes most of the global path through the network clear, without having specific knowledge about the internal endpoints. These can be handled by the NSAs responsible for those Topologies.

To satisfy the compatibility and flexibility requirements we propose to use full URIs for each component in a request. Globally unique identifiers make it possible to have delegated subtopologies without having to rewrite identifiers for example.

This means that a connection request for NSI SHOULD use the following tuple to identify a source or destination STP:

(Topology ID, source PortGroup ID, sink PortGroup ID)

3 Topology Description Syntax

3.1 STP Identifiers

A source or destination of a connection request is identified by the Topology identifier and two unidirectional *Ports* or *PortGroups*. Each of these must be globally unique identifiers.

A recommended way of constructing such an identifier is by using the `urn:ogf:network` namespace, for example `urn:ogf:network:example.net:2012:A1`.

This identifier has three components: the prefix, `urn:ogf:network` which describes that it is a network identifier, the authoring namespace, `example.net:2012` which is the DNS name and a (at least) year to make a globally unique prefix¹, and the local component, `A1` defined by the originating network.

3.2 STP Groups

Endpoints in a network often have a technology label associated with them, for example VLANs or wavelengths. Rather than describing each of these available labels as individual STPs, we introduce the STP Group, equivalent to an NML PortGroup.

An STP with a specific label can then be selected using the query component syntax as specified in [RFC3986], so for example:

`urn:ogf:network:example.net:2012:A2?vlan=1781` is a way to phrase a request to an STP with VLAN 1781 part of the STP Group identified by `urn:ogf:network:example.net:2012:A2`.

If no specific label or attribute is given to select an STP from an STP group, the NSA for that network will select one from that STP group. The confirmation back to the requester will contain the fully specified STP selected for the request. An example for this kind of request is by specifying an STP which has VLAN labels, but not requesting a specific VLAN label. Continuing the example above, the STP `urn:ogf:network:example.net:2012:A2` has been specified to have a specific VLAN range available. A request with just that identifier as the destination will allow the pathfinder to select a VLAN on that specific endpoint, and return it to the user, using the query component.

3.3 DTOX Syntax

Version 1 of the NSI Connection Service specification left the topology definition out of scope. This has left a huge gap on the operational side, where implementers have had to cooperate to create a common file to represent the topologies of each of the domains, and how to share that data. The Distributed TOpology eXchange (DTOX) working group of GLIF jumped to the opportunity and quickly provided a topology format. This was heavily inspired by the NML work in progress, but also contained some additions specific to NSI. This has allowed us to gain some experience in required elements for a topology format, and the way it could be exchanged.

The current DTOX format contains the following elements:

¹ The date component in the identifier is optional but recommended. The DNS name is a temporary lease, which can change hands, so in order to guarantee uniqueness, the year component can be added.

<i>NSI Concept</i>	<i>Representation</i>
STP	2x nml:Port / nml:BidirectionalPort
Connected To	nml:alias
NSNetwork	nml:Topology
Has STP	nml:hasPort
Located at	nml:locatedAt
Location	nml:Location
GPS coords	nml:lat, nml:long
NSA	nsi:NSA
Network managed by NSA	nsi:managedBy
Admin Contact	nsi:adminContact
Provider endpoint URL	nsi:csProviderEndpoint
Control-plane connections	nsi:peersWith

Table 1: Relation of NSI and NML terminology

STP Service Termination Point.

connectedTo relation to form an SDP with another STP

NSNetwork Network Service Network

hasSTP to define STP containment

locatedAt to define a location of a network

Location

lat, long define GPS coordinates

NSA Network Service Agent

managing to relate it to an NSNetwork

adminContact to describe contacts for the administrator

csProviderEndpoint to define the URL at which the NSA is reachable

The above format is simple, but has proven to be very effective. The STP elements are the most important ones which provide the identifiers and connectivity information necessary to do path calculations between domains.

The NSI topology representation will make use the NML topology representation as much as possible to build on standardized work in that group. This will be extended with some NSI specific terminology as shown in Figure ??.

3.4 Topology Description Example

A simple example NSI Network topology description is provided below. This example describes only the NDGF network as depicted in Figure ???. The complete topology description for Figure ??? is available in Appendix ???.

```

ndgf:NordicDataGridFacility a nml:Topology ;
    nml:version "2011112901" ;
    nml:name "NDGF" ;
    nml:locatedAt ndgf:location ;
    nml:hasOutboundPort ndgf:dk-ndgf-nordunet ;
    nml:hasInboundPort ndgf:nordunet-dk-ndgf ;
    nml:hasOutboundPort ndgf:ndgf-storage ;
    nml:hasInboundPort ndgf:storage-ndgf ;
    nsi:managedBy ndgf:nsa .
ndgf:location a nml:Location ;
    nml:lat "55.637"^^<http://www.w3.org/2001/XMLSchema#float> ;
    nml:long "12.641"^^<http://www.w3.org/2001/XMLSchema#float> .
ndgf:nsa a nsi:NSA ;
    nsi:csProviderEndpoint "http://nsa.ndgf.org/" .
ndgf:dk-ndgf-nordunet a nml:PortGroup ;
    nmleth:vlan "1780-1783" ;
    nml:alias nordunet:dk-ndgf-nordunet .
ndgf:nordunet-dk-ndgf a nml:PortGroup ;
    nmleth:vlan "1780-1783" ;
    nml:alias nordunet:nordunet-dk-ndgf .
ndgf:ndgf-storage a nml:PortGroup ;
    nmleth:vlan "1780-1783" .
ndgf:storage-ndgf a nml:PortGroup ;
    nmleth:vlan "1780-1783" .

```

The above example provides the minimal information to expose, a Topology, a Location, an NSA, two PortGroups for the connection with NorduNet, and finally two PortGroups describing the Storage endpoint in the network, all with VLAN ranges.

The Topology element is used to hide internal connectivity, and a full-mesh is assumed. By adding more NML topology information, it is possible to include more detailed descriptions of the internal network.

The NSA provides the management information for networks, how the NSI interface can be reached, and who actually maintain the NSA.

The Location element has also proven to be quite useful in allowing us to quickly create stunning visualizations using Google Earth.

3.5 Subtopologies

A topology can be further subdivided into subtopologies if required. If for example NORDUnet decides to split their USA and Scandinavian networks, we end up with a topology such as shown in figure ??.

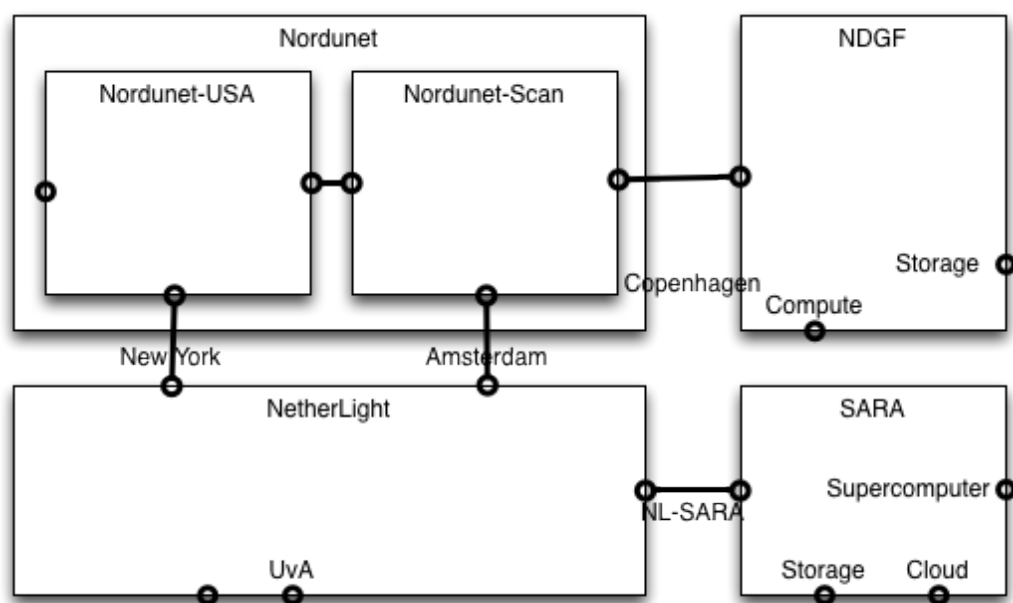


Figure 2: An example topology where Nordunet uses subtopologies

If described correctly, the other domains do not have to update their topologies or connections. NORDUnet just updates its version number, and makes a new topology file available with subtopologies. The description of the new NORDUnet topology is given at the end of Appendix A.

4 Security Considerations

5 Contributors

Jeroen J. van der Ham (Editor)

Faculty of Science, Informatics Institute, University of Amsterdam

Science Park 904, 1098 XH Amsterdam

The Netherlands

Email: vdham@uva.nl

6 Acknowledgments

The authors like to thank the NML working group members for their patience. The NML group has operated in the web of infrastructure groups and is thankful for all the input from the NM, NMC and NSI working-groups.

7 Intellectual Property Statement

The OGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the OGF Secretariat.

The OGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the OGF Executive Director.

8 Disclaimer

This document and the information contained herein is provided on an “As Is” basis and the OGF disclaims all warranties, express or implied, including but not limited to any warranty that the use of the information herein will not infringe any rights or any implied warranties of merchantability or fitness for a particular purpose.

9 Full Copyright Notice

Copyright © Open Grid Forum (2008-2012). Some Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included as references to the derived portions on all such copies and derivative works. The published OGF document from which such works are derived, however, may not be modified in any way, such as by removing the copyright notice or references to the OGF or other organizations, except as needed for the purpose of developing new or updated OGF documents in conformance with the procedures defined in the OGF Document Process, or as required to translate it into languages other than English. OGF, with the approval of its board, may remove this restriction for inclusion of OGF document content for the purpose of producing standards in cooperation with other international standards bodies.

The limited permissions granted above are perpetual and will not be revoked by the OGF or its successors or assignees.

A Example Topology Description

Below is the complete topology description of Figure ?? written in NML using the Notation3 syntax.

```
@prefix nml:      <http://schemas.ogf.org/nml/2012/10/base#> .
@prefix nmleth:   <http://schemas.ogf.org/nml/ethernet/2012/10#> .
@prefix nsi:      <http://schemas.ogf.org/nsi/topology/2012/10#> .
@prefix ndgf:     <urn:ogf:network:ndgf.org:2012:> .
@prefix nordunet: <urn:ogf:network:nordu.net:2012:> .
@prefix nl:       <urn:ogf:network:netherlight.net:2010:> .
@prefix sara:     <urn:ogf:network:sara.nl:2011:> .

ndgf:NordicDataGridFacility a nml:Topology ;
    nml:version "2011112901" ;
    nml:name "NDGF" ;
    nml:locatedAt ndgf:location ;
    nml:hasOutboundPort ndgf:dk-ndgf-nordunet ;
    nml:hasInboundPort ndgf:nordunet-dk-ndgf ;
    nml:hasOutboundPort ndgf:ndgf-storage ;
    nml:hasInboundPort ndgf:storage-ndgf ;
    nsi:managedBy ndgf:nsa .
ndgf:location a nml:Location ;
    nml:lat "55.637"^^<http://www.w3.org/2001/XMLSchema#float> ;
    nml:long "12.641"^^<http://www.w3.org/2001/XMLSchema#float> .
ndgf:nsa a nsi:NSA ;
    nsi:csProviderEndpoint "http://nsa.ndgf.org/" .
ndgf:dk-ndgf-nordunet a nml:PortGroup ;
    nmleth:vlan "1780-1783" ;
    nml:alias nordunet:dk-ndgf-nordunet .
ndgf:nordunet-dk-ndgf a nml:PortGroup ;
    nmleth:vlan "1780-1783" ;
    nml:alias nordunet:nordunet-dk-ndgf .
ndgf:ndgf-storage a nml:PortGroup ;
    nmleth:vlan "1780-1783" .
ndgf:storage-ndgf a nml:PortGroup ;
    nmleth:vlan "1780-1783" .

nordunet:Nordunet a nml:Topology ;
    nml:version "2012061801" ;
```

```
nml:name "Nordunet" ;
nml:hasOutboundPort nordunet:nordunet-dk-ndgf ;
nml:hasOutboundPort nordunet:nordunet-surfnet-NYC ;
nml:hasOutboundPort nordunet:nordunet-surfnet-AMS ;
nml:hasInboundPort nordunet:dk-ndgf-nordunet ;
nml:hasInboundPort nordunet:surfnet-NYC-nordunet ;
nml:hasInboundPort nordunet:surfnet-AMS-nordunet .
nordunet:nordunet-dk-ndgf a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias ndgf:nordunet-dk-ndgf .
nordunet:dk-ndgf-nordunet a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias ndgf:dk-ndgf-nordunet .
nordunet:nordunet-surfnet-AMS a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias nl:nordunet-surfnet-AMS .
nordunet:nordunet-surfnet-NYC a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias nl:nordunet-surfnet-NYC .
nordunet:surfnet-AMS-nordunet a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias nl:surfnet-AMS-nordunet .
nordunet:surfnet-NYC-nordunet a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias nl:surfnet-NYC-nordunet .

nl:NetherLight a nml:Topology ;
  nml:version "2011062101" ;
  nml:name "NetherLight" ;
  nml:hasOutboundPort nl:surfnet-NYC-nordunet ;
  nml:hasOutboundPort nl:surfnet-AMS-nordunet ;
  nml:hasOutboundPort nl:surfnet-SARA ;
  nml:hasInboundPort nl:nordunet-surfnet-NYC ;
  nml:hasInboundPort nl:nordunet-surfnet-AMS ;
  nml:hasInboundPort nl:SARA-surfnet .
nl:nordunet-surfnet-AMS a nml:PortGroup ;
  nml:eth:vlan "1780-1783" ;
  nml:alias nordunet:nordunet-surfnet-AMS .
nl:nordunet-surfnet-NYC a nml:PortGroup ;
```

```

        nmleth:vlan "1780-1783" ;
        nml:alias nordunet:nordunet-surfnet-NYC .
nl:SARA-surfnet a nml:PortGroup ;
        nmleth:vlan "1780-1783" ;
        nml:alias sara:SARA-surfnet .
nl:surfnet-AMS-nordunet a nml:PortGroup ;
        nmleth:vlan "1780-1783" ;
        nml:alias nordunet:surfnet-AMS-nordunet .
nl:surfnet-NYC-nordunet a nml:PortGroup ;
        nmleth:vlan "1780-1783" ;
        nml:alias nordunet:surfnet-NYC-nordunet .
nl:surfnet-SARA a nml:PortGroup ;
        nmleth:vlan "1780-1783" ;
        nml:alias sara:surfnet-SARA .

sara:SARA a nml:Topology ;
        nml:version "2010072401" ;
        nml:name "SARA" ;
        nml:hasOutboundPort sara:SARA-surfnet ;
        nml:hasInboundPort sara:surfnet-SARA .
sara:SARA-surfnet a nml:PortGroup ;
        nmleth:vlan "1780-1783" ;
        nml:alias nl:SARA-surfnet .
sara:surfnet-SARA a nml:PortGroup ;
        nmleth:vlan "1780-1783" ;
        nml:alias nl:surfnet-SARA .

```

An example of an updated topology for NorduNet with subtopologies. This new topology does not require any changes on the other network topologies.

```

nordunet:Nordunet a nml:Topology ;
        nml:version "2012080401" ;
        nml:name "Nordunet" ;
        nml:hasTopology nordunet:NordunetScandinavia ;
        nml:hasTopology nordunet:NordunetUSA .

nordunet:NordunetScandinavia a nml:Topology ;
        nml:version "2012080401" ;
        nml:name "Nordunet Scandinavia" ;

```

```
nml:hasOutboundPort nordunet:nordunet-dk-ndgf ;
nml:hasOutboundPort nordunet:nordunet-surfnet-AMS ;
nml:hasOutboundPort nordunet:Scan-to-USA-trunk ;
nml:hasInboundPort nordunet:dk-ndgf-nordunet ;
nml:hasInboundPort nordunet:surfnet-AMS-nordunet ;
nml:hasInboundPort nordunet:Scan-from-USA-trunk .

nordunet:dk-ndgf-nordunet a nml:PortGroup ;
  nml:alias ndgf:dk-ndgf-nordunet .

nordunet:nordunet-dk-ndgf a nml:PortGroup ;
  nml:alias ndgf:nordunet-dk-ndgf .

nordunet:nordunet-surfnet-AMS a nml:PortGroup ;
  nml:alias nl:nordunet-surfnet-AMS .

nordunet:Scan-from-USA-trunk a nml:PortGroup ;
  nml:alias nordunet:USA-to-Scan-trunk .

nordunet:Scan-to-USA-trunk a nml:PortGroup ;
  nml:alias nordunet:USA-from-Scan-trunk .

nordunet:surfnet-AMS-nordunet a nml:PortGroup ;
  nml:alias nl:surfnet-AMS-nordunet .

nordunet:NordunetUSA a nml:Topology ;
  nml:version "2012080401" ;
  nml:name "Nordunet USA" ;
  nml:hasOutboundPort nordunet:nordunet-surfnet-NYC ;
  nml:hasOutboundPort nordunet:USA-to-Scan-trunk ;
  nml:hasInboundPort nordunet:surfnet-NYC-nordunet ;
  nml:hasInboundPort nordunet:USA-from-Scan-trunk .

nordunet:nordunet-surfnet-NYC a nml:PortGroup ;
  nml:alias nl:nordunet-surfnet-NYC .

nordunet:surfnet-NYC-nordunet a nml:PortGroup ;
  nml:alias nl:surfnet-NYC-nordunet .

nordunet:USA-from-Scan-trunk a nml:PortGroup ;
  nml:alias nordunet:Scan-to-USA-trunk .

nordunet:USA-to-Scan-trunk a nml:PortGroup ;
  nml:alias nordunet:Scan-from-USA-trunk .
```