

Working Group

DRMAA: Distributed Resource Management Application API

WG co-chairs:

John Tollefsrud, j.t@sun.com

Hrabri Rajic, Hrabri.rajic@intel.com

www.mcs.anl.gov/~jms/ggf-sched

DRMAA WG at GGF5

Working Session

07/22/02

13:00-14:30

Working Session

07/23/02

15:00-16:30

Location:

CR Harris 1

DRMAA WG Working Session Agenda

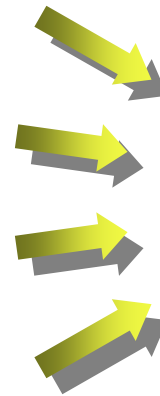
- **GGF Prelude**
 - Please sign the attendance sheet
 - IPR Rules of GGF
- **Brief DRMAA Introduction**
 - Why DRMAA?
 - The DRMAA Charter
- **DRMAA WG History**
- **DRMAA Document Discussion**
- **Next Steps**

Why DRMAA?

- **Adoption of distributed computing solutions in industry is both widespread and ‘early adopter’**
 - Commercial applications by independent software vendors (ISVs)
 - Commercial distributed resource management (DRM) systems
 - Scripted command-line integration by end users
 - Very little direct interfacing of ISV apps to DRM systems
- **Adoption is self-limiting to industries where gain exceeds the pain**
- **Fundamental shift in the adoption pattern requires shifting the DRM integration to the ISV**

Distributed Resource Management (DRM) Systems

- **Batch/job management systems**
- **Local Job schedulers**
- **Queuing systems**
- **Workload management systems**



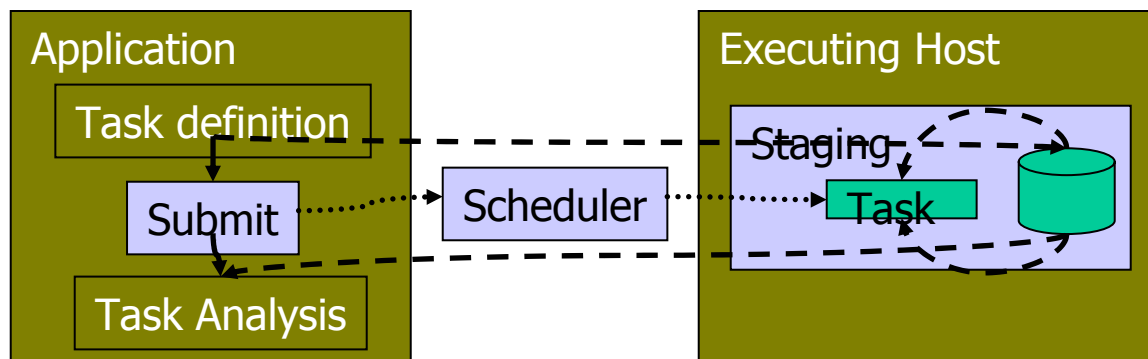
All are DRM Systems

Motivation for DRMAA

There are many DRM solutions available to end users and things keep changing

Independent Suppliers	Open Source / University	OEM Proprietary	Peer-to-Peer
Platform Computing <i>LSF</i>	Veridian <i>OpenPBS</i>	IBM <i>LoadLeveler</i>	TurboLinux <i>Enfuzion</i>
Veridian <i>PBS Pro</i>	Univ of Wisc <i>Condor</i>	Sun <i>Sun Grid Engine</i>	Entropia United Devices Parabon
Condor Inc. <i>Condor</i>	Sun <i>Grid Engine</i>		

Resource Management Systems Differ



- **Core services are fundamentally the same**
 - especially from the users perspective
- **DRM programming interfaces differ**
 - ISVs are disinclined to use

DRMAA Charter

- **Develop an API specification for the submission and control of jobs to one or more Distributed Resource Management (DRM) systems.**
- **The scope of this specification is all the high level functionality which is necessary for an application to consign a job to a DRM system including common operations on jobs like termination or suspension.**
- **The objective is to facilitate the direct interfacing of applications to today's DRM systems by application's builders, portal builders, and Independent Software Vendors (ISVs).**

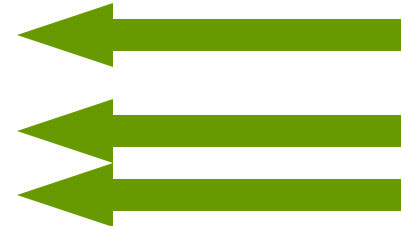
Characterizing DRMAA

- **High level attributes**
 - Application centric
 - Ease of use for end users
 - Focused on programming model
- **Benefits**
 - Faster distributed application deployment
 - Opportunity for new applications
 - Increased end user confidence
 - Improvements in Resource Management Systems
 - Distributed application portability

Scope: Run a Job API

(Steps from: "Ten Actions when SuperScheduling", GGF SchedWD 8.5, J.M. Schopf, July 2001)

- **Phase 1: Resource Discovery**
 - Step 1 Authorization Filtering
 - Step 2 Application requirement definition
 - Step 3 Minimal requirement filtering
- **Phase 2 System Selection**
 - Step 4 Gathering information (query)
 - Step 5 Select the system(s) to run on
- **Phase 3 Run job**
 - Step 6 (optional) Make an advance reservation
 - **Step 7 Submit job to resources**
 - Step 8 Preparation Tasks
 - **Step 9 Monitor progress (maybe go back to 4)**
 - **Step 10 Find out Job is done**
 - Step 11 Completion tasks



What have been the Issues?

- **Languages**

- C/C++
- Perl, Python
- Fortran, Java

- **General features**

- Asynchronous job monitoring
- Protocol based
- Scalability

- **Libraries**

- Serial / thread safe
- Tracing / diagnosis

- **Advanced features**

- Debugging support
- Data streaming
- Security

Timetable

- **API defined by mid 2002**
 - E.g., Jul'02: DRMAA v1.0 GWD submitted for review

DRMAA Activity since GGF4

- **Bi-weekly con calls**
 - Toll Free: (877)288-4427 Code: 691169
Next call: Tues, July 30, 9AM PDT (email to j.t@sun.com)
- **Two proposals merged to a single proposal**

DRMAA Proposal 01 Overview

- **Authors: Andreas Haas, Fritz Ferstl**
- **Presenter: Andreas Haas**

Interface Form

- C-API library interface - no protocol
 - Simplifies utilization by ISV's
- Shared library binding
 - Prerequisite to allow end user to select DRM technology of their choice
- Library supports only one DRM system per implementation
 - Simultaneous support of different DRM systems is beyond the scope of our project

Job Template

- Functions to create/delete job template
 - `job_template *newJobTemplate(void)`
 - `void deleteJobTemplate(job_template *jt)`
- Job template can be used multiple times for submitting jobs
- DRMAA WG will agree on a limited ***standard*** set of job attributes
- DRMAA interface to provide a hook for other, e.g. DRM-specific job attributes

Job Template, cont'd

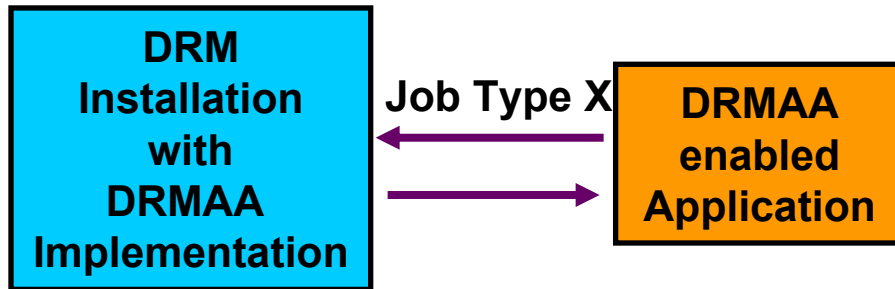
- Access functions allow setting of all job attributes that have a meaning for the DRM system
 - For DRMAA agreed attributes (real time limit, job path, job arguments, ..) via well-defined access functions, e.g.
 - int setRealTimeLimit(job_template *jt, long rtl)
 - int setJobPath(job_template *jt, char *path)
 - For non-agreed DRM specific (project, account, ...) or application specific (database affinity, ...) attributes via generic function
 - int setAttrib(job_template *jt, char *name, char *value)

Job Template, cont'd

- Job template attributes need translation into DRM system job attributes in a way specific to site/appl needs:
 - At site A jobs of type X can only be run on a subset of hosts, while at site B jobs of type X can be run on all hosts
 - At site A jobs of type Y compete with Z jobs and thus are dispatched with a high priority, while at site B jobs of type Y are less important
- Each implementation of API library is responsible for that mapping

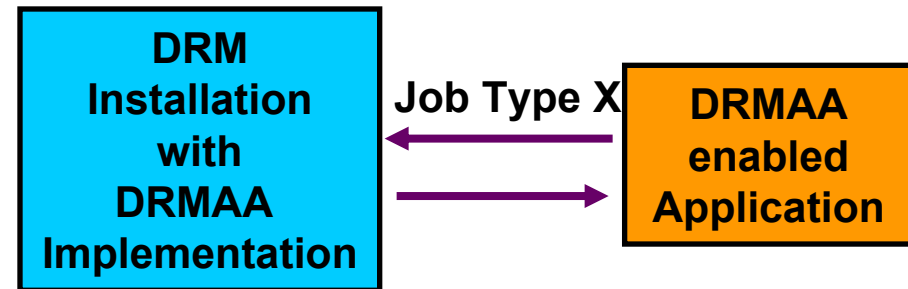
Job Categories

Site A




- Cluster consists of machines where X jobs run and others where they don't run

Site B



- X jobs run at all machines in cluster

Job Submission

- Jobs submitted to the DRM system are identified via a job identifier
- For flexibility reasons a job identifier should be of type `char *`
- Single job identifiers are returned by
 - `int runJob(job_template *jt, char *job_id)`
-  • Bulk job submissions return multiple job identifiers
 - `int runArrayJob(job_template *jt, int n, char *jid[])`

Job Monitoring, Ctrl & Termination

NEW

- Monitoring functions
 - `char *monitorJob(char *jobid);`
- Control functions
 - `int suspendJob(char *jobid);`
 - `int resumeJob(char *jobid);`
 - `int terminateJob(char *jobid);`
- Blocking and non-blocking waiting for one or more jobs to finish (like `wait4(2)`)
 - `char *waitJob(char *jobid, int *status, int options, char **rusage);`

DRMAA Proposal 02 Overview

- **Author and Presenter: Hrabri Rajic**

DRMAA Guidelines

- The API calling sequences should be simple and the API set small.
- The routine names should convey the semantic of the routine.
- The set should be as convenient as possible, even with the risk of being forced to emulate some functionality if missing from a DRMS.
- All jobs manipulation per process is available without explicit job iterating.
- The servers names are hidden, the DRMS is a black box.
- Consistent API structure
 - Err return parameter
- Data structures not exposed

DRMAA Library features

- Shared
- One or more DRMSs targeted
- Versioning
 - Backwards compatible
 - Stub API resolution for forward compatibility

API groups

- Init/exit
- Job submit
 - Individual jobs
 - One time
 - Multiple times -- templates
 - Bulk jobs
- Job monitoring and control
- Auxiliary or system routines
 - trace file specification
 - error message routines

Execution environment

- File system duality
 - Shared
 - Appls. comes to data
 - Distributed
 - Data come to appls.
 - Unique dir per job
 - Copying files?
- Environment passing
 - Use default behavior
 - Export local env.
 - Use remote env.
 - Env specified via API
- Handling of DRMS job execution options
 - Translation/consolidation
 - Allowing native

Native DRMS Options

- The end user interacts with the DRMS via `native_resource_options` parameter.
 - Simple solution
 - DRMAA implementation ignores the DRMAA DRMS implicitly used and disallowed options
 - Dist. Appls. Developers and DRMS vendors are not involved in the local environment spec.
 - The burden is on the end users to define the execution environment
 - Need to know DRMS
 - Need to know the remote application installation

DRMAA WG Agenda -Working Session-

Working Session

02/18/02

15:00-16:30

- **Detailed Review of proposals and proposed enhancements**
 - **Andreas Haas**
 - **Hrabri Rajic**
- **Review of comments received**
- **Topics - Considerations**
 - **API Guidelines**
 - **Resource Options**
 - **Job Categories/Classes/Templates**
 - **Environment Handling**
 - **User/Installer/App Developer experience**
 - **Error Handling / System diagnosis**
 - **Performance Measurement**
 - **Language Binding**

DRMAA WG Agenda -Presentation Session-

- **Why DRMAA?**
- **The DRMAA Charter**
- **DRMAA WG Activity since GGF3**
- **Proposal overviews**
 - Andreas Haas
 - Hrabri Rajic
- **Working session Agenda for GGF4**
- ➔ • **Q & A, comments**
- **Break**
- **Resume Working Session at 15:00**