



DRMAA: Distributed Resource Management Application API

Peter Tröger,
Hasso-Plattner-Institute (HPI) @ University of Potsdam

GGF 14 DRMAA session
Chicago, June 28, 2005

Agenda

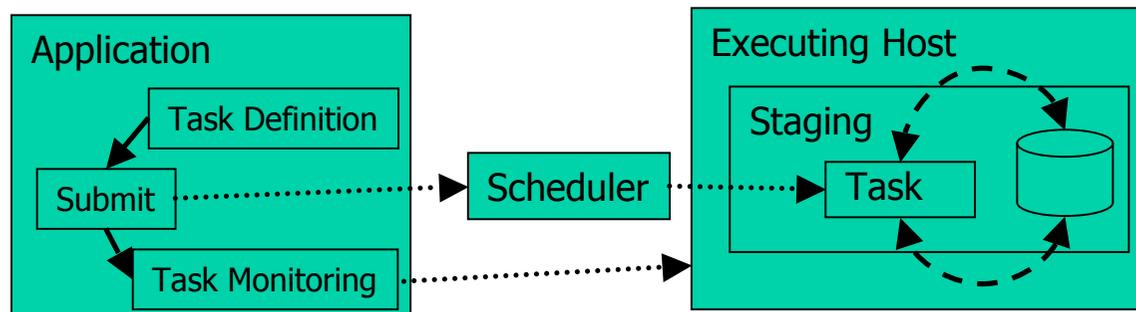
- First things first
 - GGF IP
 - Sign-up sheet
 - Note takers
- DRMAA introduction
- Status of the DRMAA implementations
- Status of the DRMAA documents
- Open floor, open issues

Scope: Run a Job API

(Steps from: "Ten Actions when SuperScheduling", GGF SchedWD 8.5, J.M. Schopf, July 2001)

- **Phase 1: Resource Discovery**
 - Step 1 Authorization Filtering
 - Step 2 Application requirement definition
 - Step 3 Minimal requirement filtering
- **Phase 2 System Selection**
 - Step 4 Gathering information (query)
 - Step 5 Select the system(s) to run on
- **Phase 3 Run job**
 - Step 6 (optional) Make an advance reservation
 - **Step 7 Submit job to resources** 
 - Step 8 Preparation Tasks
 - **Step 9 Monitor progress (maybe go back to 4)** 
 - **Step 10 Find out Job is done** 
 - Step 11 Completion tasks

Resource Management Systems Differ Across Each Component



	Interface Format	Execution Environment	Platform Mix
LSF	Has API plus Batch Utilities via "LSF Scripts"	User: Local disk exported System: Remote initialized (option)	Unix \leftrightarrow Windows
Grid Engine	GDI API Interface plus Command line interface	System: Remote initialized, with SGE local variables exported	Unix only
PBS	API (script option) Batch Utilities via "PBS Scripts"	System: Remote initialized, with PBS local variables exported	Unix \leftrightarrow Windows
DataSynapse	Proprietary API.	User: Remote initialized	Unix \leftrightarrow Windows

DRMAA Charter

- Develop an API specification for **the submission and control of jobs** to one or more Distributed Resource Management (DRM) systems.
- The scope of this specification is all the high level functionality which is necessary for an application to consign a job to a DRM system including **common operations on jobs** like termination or suspension.
- The objective is to **facilitate the direct interfacing of applications to today's DRM systems** by application's builders, portal builders, and Independent Software Vendors (ISVs).

DRMAA History

- BOF at GGF 3 in Frascati, Oct 2001
- WG status at GGF 4, Toronto, February 2002
- Participation from Altair (PBS), Sun Microsystems (SGE), Intel, IBM (LoadLeveler), University of Wisconsin (Condor), Cadence (Rocks system), Globus project
- Sideline engagement from EnFuzion, Entropia, Platform (LSF), GridIron project, United Devices
- June 2004: DRMAA 1.0 document accepted as proposed recommendation by GFSC

What have been the Issues?

- **General features**
 - Session concept
 - Asynchronous job monitoring
 - Scalability
 - Native features
- **Language bindings**
 - C/C++
 - Perl, Python
 - Fortran, Java
- **Libraries**
 - Serial / thread safe
 - Tracing / diagnosis
- **Advanced features**
 - Debugging support
 - File staging
 - Security
 - Job categories

Submit, control & monitor, and query status of jobs

DRMAA API Function Groups

- Init / Exit
- Job template handling
 - Allocation / Deletion
 - Job template parameter setter/getter routines
- Job submission
 - Individual jobs
 - One time
 - Multiple times – just re-adjust the job template (parameter sweep)
 - Bulk jobs - implicit parameterization
- Job monitoring and control
- Auxiliary or system routines
 - Error message routines
 - Informational interfaces

Job Template

- Description of all job requirements / parameters
- Mandatory and optional parameters
- Same intention as JSDL, but designed as ‘smallest common denominator’ between possible backend’s
- Functions to create/delete job templates
 - `job_template *drmaa_allocate_job_template (void)`
 - `void drmaa_delete_job_template (job_template *jt)`
- **Setter/getter job template routines**
 - `int drmaa_set_attribute(job_template *jt, char *name, char *value);`
 - `int drmaa_set_vector_attribute(job_template *jt, char *name, char **values);`
 - `char* drmaa_get_attribute(job_template *jt, char *name);`
 - `char** drmaa_get_vector_attribute(job_template *jt, char *name);`

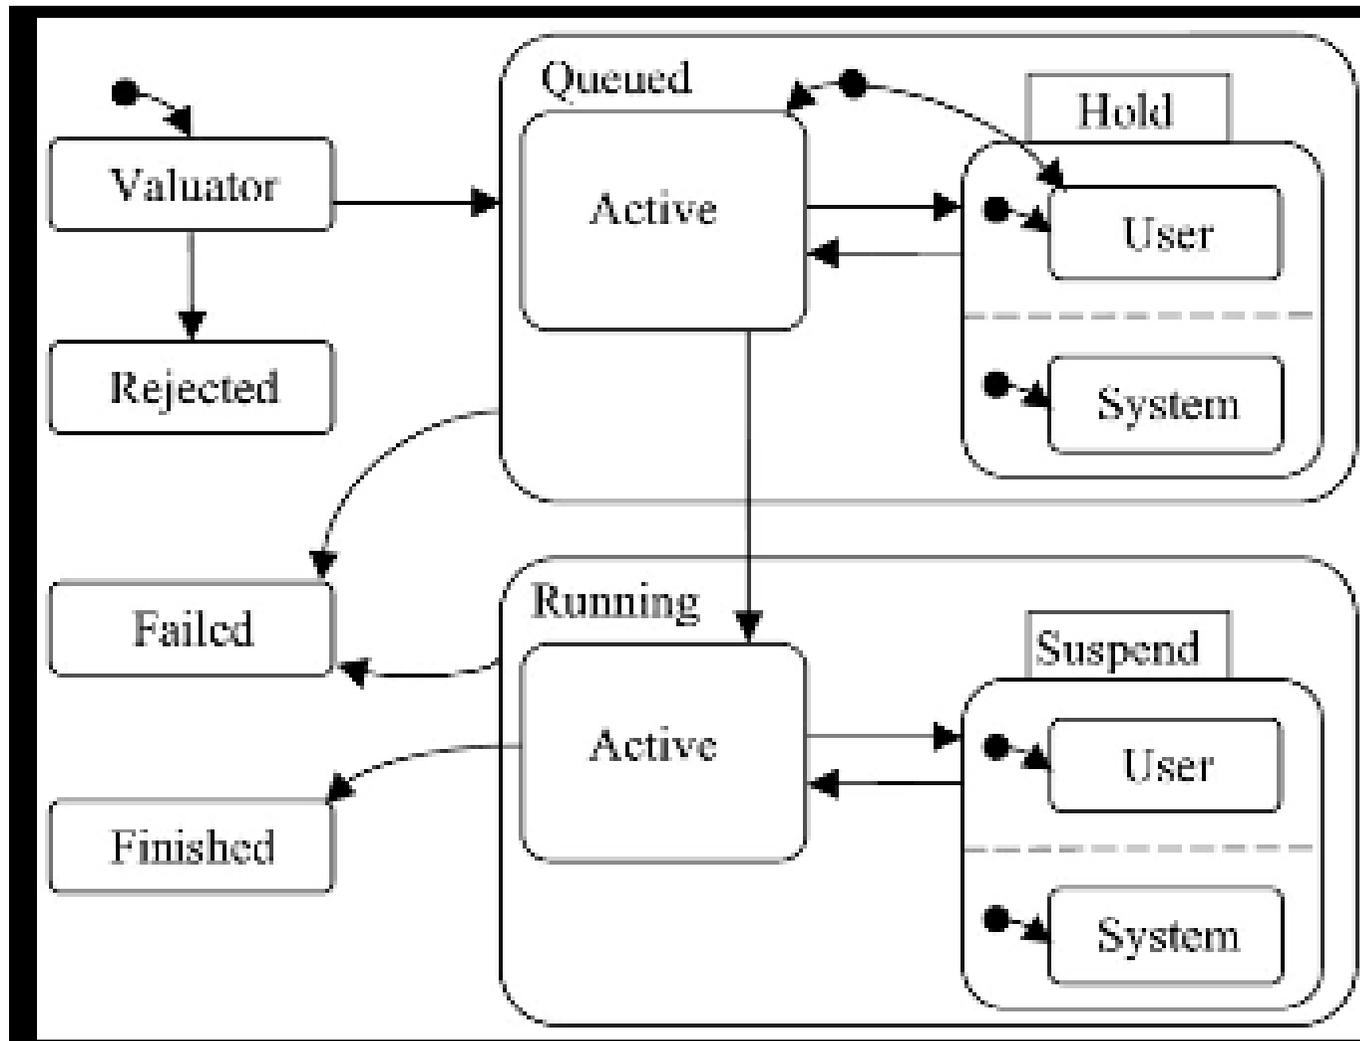
Job Submission

- Jobs submitted to the DRM system are identified via an opaque job identifier (char*)
- Single job identifiers are returned by
 - `int drmaa_run_job(job_template *jt, char *job_id)`
- Bulk job submissions return multiple job identifiers
 - `int drmaa_run_bulk_job(char **job_ids, job_template *jt, int start, int end, int incr)`

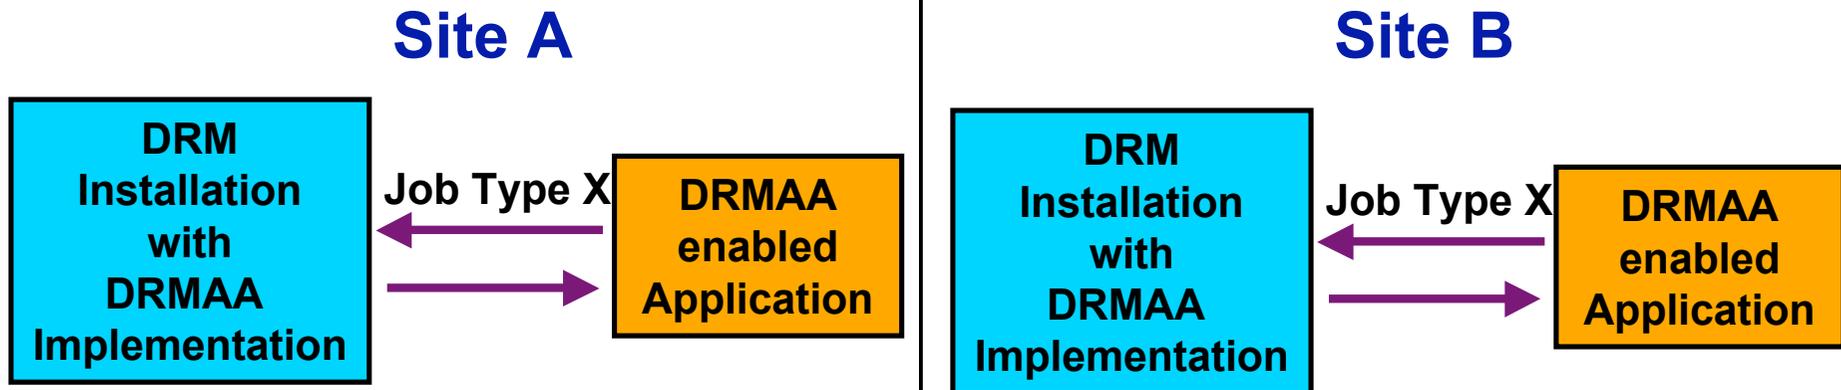
Job Monitoring, Control, and Status

- **Monitoring/Control functions**
 - `int drmaa_control(char *job_id, int action);`
 - `int drmaa_synchronize(char **job_ids);`
 - `int drmaa_job_ps(char *job_id, int *remote_ps);`
- **Blocking and non-blocking waiting for one or more jobs to finish (like POSIX wait4(2))**
 - `char *drmaa_wait(char *jobid, int *status, int timeout, char **rusage);`
 - `drmaa_wif[exited|signaled|aborted]` and friends to get more information about failed jobs

Job State Transition



Job Categories



- Cluster consists of machines where X jobs run and others where they don't run

- X jobs run at all machines in cluster

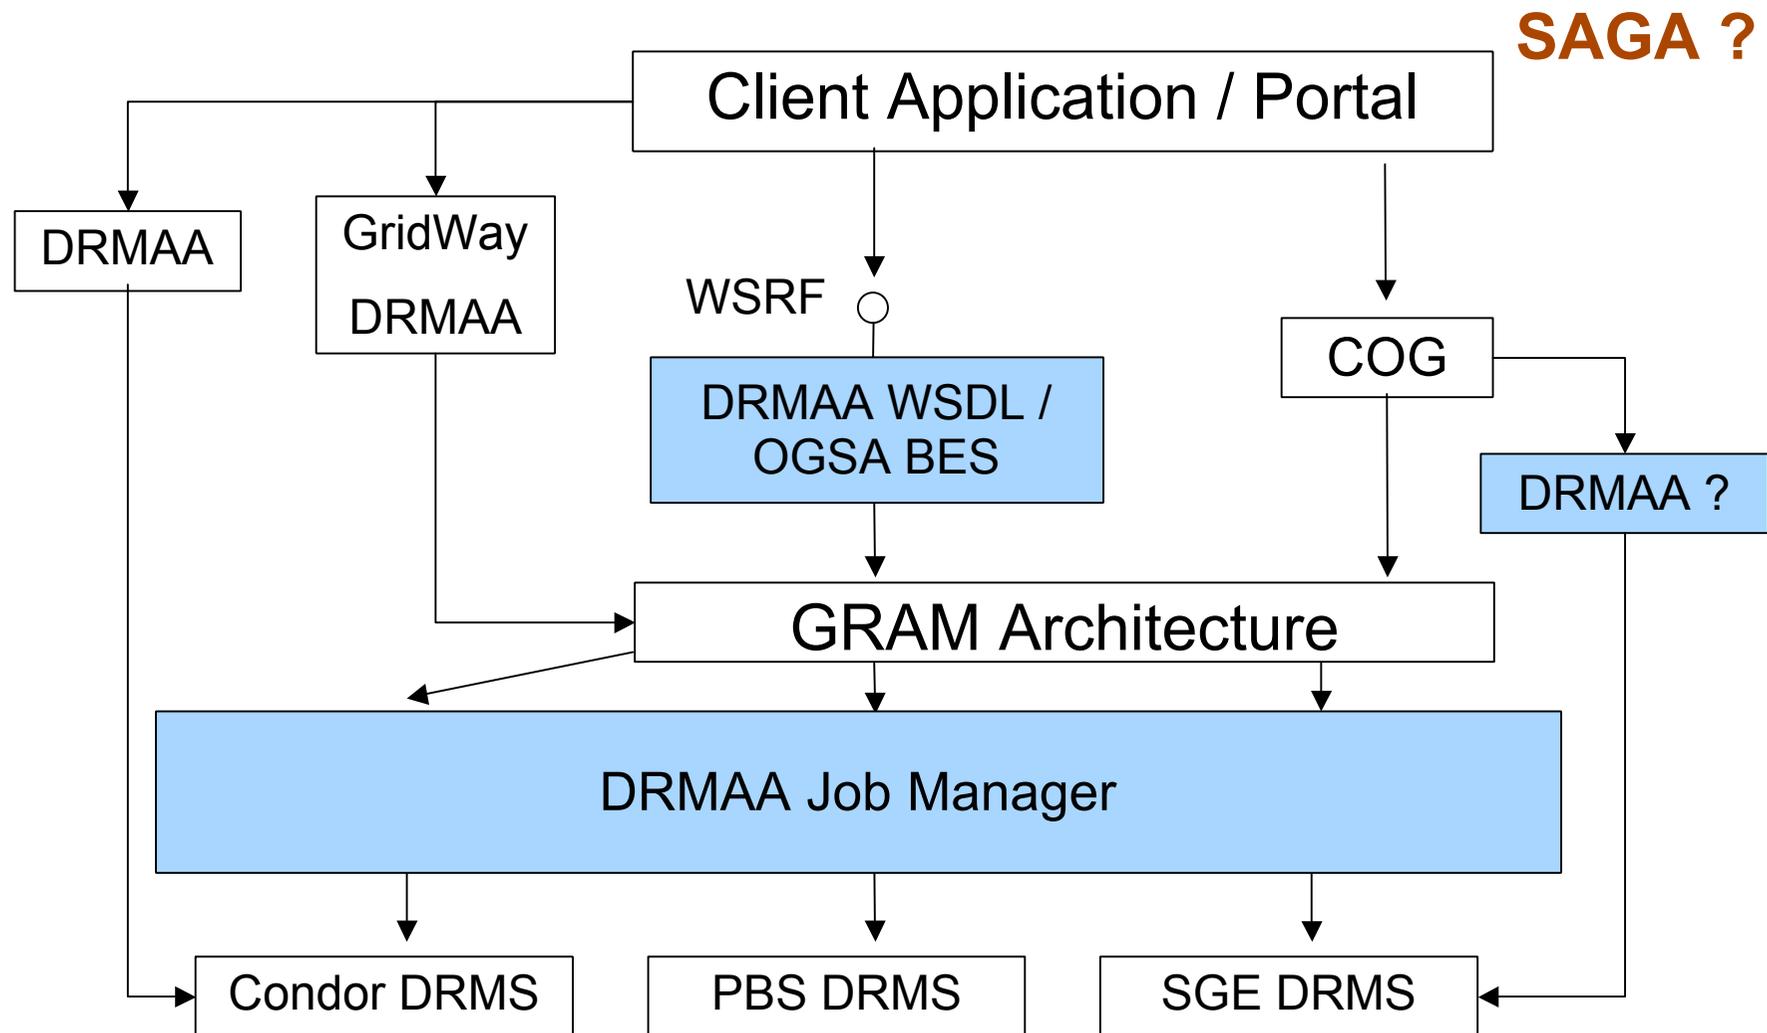
Native DRMS Options

- The end user interacts with the DRMS via `native_specification` parameter
 - Simple solution
 - DRMAA implementation ignores the DRMAA DRMS implicitly used and disallowed options
 - Dist. Appls. Developers and DRMS vendors are not involved in the local environment spec.
 - The burden is on the end users to define the execution environment
 - Need to know DRM
 - Need to know the remote application installation

DRMAA Placement

- On top of DRM systems
- On top of Globus
- Beneath GRAM
- UNICORE TSI interface to DRMSs
- CoG adapter
- On top of CoG
- Interfaced by a Portal, application, shell
- Portable command line utilities (qsub, qstat)

A World of Submission API's



DRMAA in Practice

- **Multiple implementations since 2004**
 - Product implementation in Sun Grid Engine 6
 - C- and Java-binding implementation
 - Prototype in Condor 6.7 series
 - C-binding implementation
 - CPAN Perl DRMAA module (Tim Harsch)
 - On-top-of DRMAA C-library
 - GridWay DRMAA implementation
 - Allows DRMAA on-top-of Globus
 - Prototype for Globus 3 DRMAA job manager (HPI)
 - Based on DRMAA Perl implementation
- **Tutorials, programming examples, test suites**
 - <http://gridengine.sunsource.net>
 - <http://www.dcl.hpi.uni-potsdam.de/research/drmaa>
 - GGF12 tutorial, JavaOne 05 tutorial materials

Need for Improvement !?!

- Feedback from practical usage of available implementation(s) fed's back into spec
 - Just look at the GridForge tracker and the SGE / Condor mailing lists
- Some details under-specified
 - Behavior in multi-threaded environments
- C-centric style makes it hard to develop conformant object-oriented bindings
 - Massive feedback from Java and .NET language binding specification work
- Several missing error codes
- But: Keep the API as small as it is !!!

DRMAA IDL Spec

- Started work in early 2005
- Based on Java- and .NET-binding experiences
- Specification through standardized OMG Interface Definition Language (IDL)
 - No, that does not mean CORBA ;-)
 - Example: W3C DOM specification
 - DRMAA language bindings will (and should) **not** rely on IDL language bindings from OMG
 - Complicated, weird semantics
 - Simple custom binding by specifying consistent mapping rules
 - Examples for Java in the IDL-spec
 - Usage of IDL avoids wording issues (i.e. ‚attribute‘ vs. ‚property‘)
 - Allows for true language-independent description of namespaces, enumerations, constants, and time values

DRMAA IDL Spec (contd.)

- Improved, more consistent description text for all functions
 - More details regarding advanced OO-specific features (multiple session objects, exception hierarchies)
 - Consider languages with introspection functionalities
 - Some details about RPC-DRMAA scenarios (SOAP, RMI, ...)
- More parameter placeholders (e.g. for job ID)
- A lot more possible error codes for the operations

Backward Compatibility

- C- and Java bindings in their current state can be mostly derived also from the IDL spec
 - Demand for consistent name mapping might change one or two method names in the C-binding
 - Introduction of new job state / error codes does not break existing applications
 - DRMAA has already a notion of versioning
- .NET binding will be re-designed based on the IDL spec
- No official binding documents for Perl and Python so far

DRMAA Documents

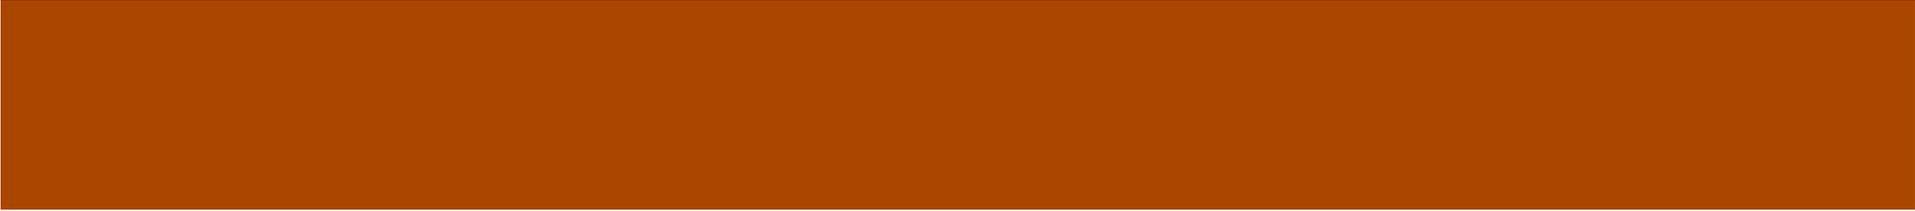
- DRMAA GFD-P-R or GFD.22 document
 - Since June 2004
- C binding v1.0
 - Ready for submission to GFCS
- C binding experimental document v0.98
- Java binding 0.6.1
 - Fairly complete
- .NET binding v0.2
 - Needs a synch with IDL-spec
- IDL document
 - v0.3 and nearly feature complete
 - Need to augment with the DRMAA GFD-P-R text
 - Will be submitted as a standalone GFD-P-R doc

Next Steps

- Putting the DRMAA-IDL spec in the GGF document chain (GGF14 version is nearly final)
- Prerequisite for some announced activities
 - .NET-binding implementation (HPI)
 - On-top-of DRMAA C-library
 - Improved Condor C library
 - Join the efforts at <http://sf.net/projects/condor-ext>
 - Python binding specification
 - Maybe more implementations
- DRMAA collaborates with SAGA, JSDL, and OGSA-BES for identifying synergy effects

Conclusion

- Please take part in the discussion
 - Bi-weekly con calls
 - Toll Free: (866)545-5198 Code: 6898552
 - Regular: (865)521-8904
 - GridForge tracker
 - E-mail: *drmaa-wg@gridforum.org*
 - Archive: Use the link at *http://drmaa.org*
- Please implement DRMAA and tell us your experience
 - It's easy, Dan did it 4 times ;-) ...



Thank you !