



GGF12: DRMAA Tutorial C and Java Language Bindings

Daniel Templeton
Senior Staff Engineer
Sun Microsystems GmbH



Agenda

- Architecture Overview
- API Overview
 - Session
 - Job Execution
 - Job Synchronization
 - Job Control & Monitoring
- Bindings
 - C Binding
 - JavaTM Language Binding

Architectural Overview

- API for job submission and control
 - Keep it simple
 - Keep the number of functions down
 - Avoid tough issues
 - Address the areas of agreement
 - Leave room for areas of disagreement

Architectural Overview

- Implementable by any* DRM in any* language
 - N1™ Grid Engine – C, Perl, & Java language bindings
 - Condor – C binding, Perl binding?
 - Globus – C binding?

* OK. Not really, but we tried to make it as universal as possible

Design & Impl Considerations

- Implementations should be modular
 - Pluggable implementations
- Implementations should be multi-threaded
 - Handle synchronization like UNIX
- Single session per instance
 - Used for tracking jobs
 - Can only init and exit
- No user handling
 - Implies a single user per instance

Applications

- Bindings may be either client- or server-side
 - N1GE bindings use the facilities available to the client-side utilities, such as qsub
 - N1GE DRMAA bindings can only be used from submit hosts
 - DRMAA uses N1GE communications mechanisms
 - Web Service binding would be server side
 - Client would only need the WSDL for the service
 - Globus binding could be either
 - Client-side library
 - Uses web services to talk to Globus
 - Server-side web service hosted within Globus

API Overview

- All APIs based on an über spec
 - Language independent binding – version 1.0
 - C language binding – version 0.9.5
 - Java language binding – version 0.4.2
 - .Net language binding – version 0.2
 - Unspecified Perl binding
- Binding specs don't all agree 100%
 - Some things don't make sense in some languages
 - Original, language independent spec has a C slant
- Attempt to agree as much as possible
 - Reconciling .Net and Java language bindings

Session Handling

- Only two things to do with a session:
 - Init
 - Creates a session for the instance
 - Takes a string which specifies to server instance
 - More info on next slide
 - Exit
 - Destroys a session
 - Session exists from init to exit
 - Used by other calls to identify jobs
 - Exit should always be called
 - Allows implementation to do necessary cleanup

Instance Handling

- Get Contact
 - Before init, returns a list of available server instances
 - Items in the list can be passed to init
 - After init, returns the bound server instance
- Get DRM System
 - Before init, returns a list of available DRMs
 - After init, returns the DRM associated with the bound server instance

Instance Handling

- Get DRMAA Implementation
 - Before init, returns a list of available DRMAA implementations
 - After init, returns the bound implementation
- Get Version
 - Returns the major and minor version for the bound implementation
 - May only be called after init

Job Description

- Encapsulated in a job template
 - Properties sheet with predefined properties
 - Functions provided for accessing properties
 - Set Attribute, Get Attribute, Set Vector Attribute, Get Vector Attribute
 - List of defined Property names
 - Java language binding and uses JavaBean™ property accessors instead
 - e.g. setWorkingDirectory(), getWorkingDirectory()
 - .Net language binding offers both
 - Open debate

Job Templates

- Job Template is independent of jobs
 - 1:n relationship
 - Static – not updated during job execution
- Allocate Job Template
 - Creates a new job template
- Delete Job Template
 - Frees an allocated job template

Job Attributes: Required

- Must be implemented by all implementations
- Remote Command to Execute
 - The executable to be run as the job
- Input Parameters
 - The args to be passed to the remote command
- Job State at Submission
 - Whether the job should be started in a suspended or runnable state
- Job Environment
 - The environment variable settings for the job

Job Attributes: Required

- Job Working Directory
 - The directory where the remote command should be executed
- Job Start Time
 - The earliest time that the job may run
 - Not a deadline
- Job Name
 - The name to be assigned to the job
- Input/Output/Error Stream
 - The input/output/error stream path

Job Attributes: Required

- Input/Output/Error Stream
 - The input/output/error stream path
- Join Files
 - Whether to attach (dup2) the error stream to the output stream

Job Attributes: Required

- Email Address
 - The address to which to send email
 - Email is only sent when the underlying DRM decides to
 - i.e. The DRM's settings control when email is sent
- Email Suppression
 - Whether to prevent the sending of email
 - Can only prevent email from being sent
 - Cannot force email to be sent if DRM settings say not to

Job Attributes: Required

- Two catch-all attributes
- Native Specification
 - Allows a user to specify settings not allowed by DRMAA
 - For example, in N1GE, -ckpt for checkpointing
- Job Category
 - Allows an administrator to assign default DRM settings to groups of jobs
 - Not true job category
 - Reference into DRM specific settings
 - “Dissolves” at submission time

Job Attributes: Optional

- Not required of every implementation
- Implementations may also implement other attributes
- Transfer Files
 - Whether to treat the input, output, and error paths as locations for file staging
- Absolute Job Termination Time
 - Deadline after which the DRM will terminate the job

Job Attributes: Optional

- Wall Clock Time Limit
 - The amount of wall clock time the job is allowed to execute before being terminated
- Soft Wall Clock Time Limit
 - The amount of wall clock time the job is expected to need to execute
- Job Run Duration Hlimit
 - The amount of CPU time the job is allowed to use before being terminated
- Job Run Duration Slimit
 - The amount of CPU time the job is expected to

Job Execution

- Two ways to run a job:
- Run Job
 - Runs a single job based on the parameters in the supplied job template
 - Returns an opaque job id
- Run Bulk Jobs
 - Runs a number of jobs all based on the same job template
 - AKA a parametric job
 - Number calculated from start, end, and increment
 - Returns a list of opaque job ids
 - More about this with the C language binding
- e.g, equivalent to “qsub [-t start-end:increment]”

Job Synchronization

- Wait
 - Waits for a single job to complete
 - Can wait for any, i.e. the next, job to complete
 - Returns information about the job's execution
 - Exit status
 - Resource usage
 - List of “name=value” strings
 - Other terminating info
 - Can only wait for a job once
 - Wait reaps the job's exit info
 - Additional waits return an error

Job Synchronization

- Synchronize
 - Waits for all jobs in a list to complete
 - Can wait for all jobs in the session to complete
 - All jobs submitted before the synchronize
 - Does not return exit information
 - May reap the exited jobs' exit info – dispose parameter
 - Can be used in conjunction with Wait to get exit info
 - e.g synchronize on the id list, then wait for each id in the list

Finish Status

- Finish status is an opaque value interpreted by:
 - If Exited
 - Whether the job exited normally
 - Exit Status
 - The jobs exit code
 - If Signaled
 - If the job exited due to a signal
 - Terminating Signal
 - The terminating signal
 - If Core Dumped
 - If the job created a core dump file
 - If Aborted
 - If the job was terminated abnormally

Job Control

- Control
 - Works on a single job id or all jobs
 - Can:
 - Terminate
 - Suspend
 - Resume
 - Hold
 - Release
 - May return before action completes

Job Monitoring

- Job PS
 - Returns the status of a particular job:
 - Queued & Active
 - Running
 - System Hold
 - User Hold
 - User & System Hold
 - System Suspended
 - User Suspended
 - User & System Suspended
 - Done
 - Failed
 - Undetermined

Bindings

- C language Binding
 - Very close to platform-independent spec
 - Should use shared libraries
 - Modular C for plugability
- Java language Binding
 - Strays from platform-independent spec a little
 - OO implementation requires different perspective
 - JavaBeans properties
 - Java language provides some things for free
 - Strings
 - Exception handling
 - Collections

The rest of the presentation is intended for programmers

C Language Binding

- 0.9.5 is the current C language binding spec
 - Likely final
- Implemented by Sun™ (N1 Grid Engine 6.0) and Condor (Condor 6.7)

Error Handling

- Every function returns an error code
 - DRMAA_ERRNO_SUCCESS – good
 - Everything else – varying degrees of bad
- Most take an error buffer parameter
 - Contains error message on error
- Actually specified in platform independent spec
 - Belongs in C binding spec
- Data is returned from functions via pointers

Buffer Lengths

- DRMAA_ERROR_STRING_BUFFER
 - Maximum error string length
- DRMAA_JOBNAME_BUFFER
 - Maximum job name length
- DRMAA_SIGNAL_BUFFER
 - Maximum signal name length
- DRMAA_ATTR_BUFFER
 - Maximum attribute name length
- DRMAA_CONTACT_BUFFER
 - Maximum contact string length
- DRMAA_DRM_SYSTEM_BUFFER
 - Maximum DRM system name length

Session Handling

- ```
int drmaa_init (
 const char *contact,
 char *error_diagnosis,
 size_t error_diag_len
);
```
- ```
int drmaa_exit (
    char *error_diagnosis,
    size_t error_diag_len
);
```
- Idiom: signal handler thread that calls drmaa_exit() on SIGINT.

Example: Session Handling

```
char error[DRMAA_ERROR_STRING_BUFFER];
int errnum = 0;

errnum = drmaa_init (argv[0], error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't init DRMAA library: %s\n", error);
    return 1;
}

/* Do Stuff */

errnum = drmaa_exit (error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't exit DRMAA library: %s\n", error);
    return 1;
}

return 0;
```

Instance Handling

- `int drmaa_get_contact (char *contact, size_t contact_len, char * error_diagnosis, size_t error_diag_len);`
- `int drmaa_get_DRM_system (char *drm_system, size_t drm_system_len, char * error_diagnosis, size_t error_diag_len);`
- `int drmaa_get_DRMAA_implementation (char *drmaa_impl, size_t drmaa_impl_len, char * error_diagnosis, size_t error_diag_len);`

Instance Handling

- int drmaa_version (
 unsigned int *major,
 unsigned int *minor,
 char * error_diagnosis,
 size_t error_diag_len
);

Example: Instance Handling

```
char contact[DRMAA_CONTACT_BUFFER];
char drm_system[DRMAA_DRM_SYSTEM_BUFFER];
char drmaa_impl[DRMAA_DRM_SYSTEM_BUFFER];
unsigned int major = 0;
unsigned int minor = 0;

drmaa_get_contact (contact, DRMAA_CONTACT_BUFFER, NULL,
                      0);
drmaa_get_DRM_system (drm_system, DRMAA_DRM_SYSTEM_BUFFER,
                        NULL, 0);
drmaa_get_DRMAA_implementation (drm_system,
                                    DRMAA_DRM_SYSTEM_BUFFER,
                                    NULL, 0);
drmaa_version (&major, &minor, NULL, 0);

printf ("Contact: %s\n", contact);
printf ("DRM System: %s\n", drm_system);
printf ("DRMAA Implementation: %s\n", drmaa_impl);
printf ("Version: %d.%d\n", major, minor);
```

Job Templates

- `int drmaa_allocate_job_template (`
`drmaa_job_template_t **jt,`
`char *error_diagnosis,`
`size_t error_diag_len`
`);`
- `int drmaa_delete_job_template (`
`drmaa_job_template_t *jt,`
`char *error_diagnosis,`
`size_t error_diag_len`
`);`

Example: Job Templates

```
char error[DRMAA_ERROR_STRING_BUFFER];
int errnum = 0;
drmaa_job_template_t *jt = NULL;

/* Init Session */

errnum = drmaa_allocate_job_template (&jt, error,
                                         DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't allocate job template: %s\n", error);
    return 1;
}

/* Run Job */

errnum = drmaa_delete_job_template (jt, error,
                                         DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't delete job template: %s\n", error);
    return 1;
}

/* Exit Session */
```

Scalar Job Attributes

- ```
int drmaa_set_attribute (
 drmaa_job_template_t *jt,
 const char *name,
 const char *value,
 char * error_diagnosis,
 size_t error_diag_len
);
```
- ```
int drmaa_get_attribute (
    drmaa_job_template_t *jt,
    const char *name,
    char *value,
    char * error_diagnosis,
    size_t error_diag_len
);
```

Vector Job Attributes

- ```
int drmaa_set_vector_attribute (
 drmaa_job_template_t *jt,
 const char *name,
 const char *value[], /* NULL-terminated */
 char * error_diagnosis,
 size_t error_diag_len
);
```
- ```
int drmaa_get_vector_attribute (
    drmaa_job_template_t *jt,
    const char *name,
    drmaa_attr_values_t **values,
    char * error_diagnosis,
    size_t error_diag_len
);
```

Attribute Values

- drmaa_attr_values_t type
 - Implementation independent vector
- int drmaa_get_next_attr_value (
 drmaa_attr_values_t *values,
 char *values,
 int value_len
);
- int drmaa_release_attr_values (
 drmaa_attr_values_t *values
);

Example: Job Attributes

```
char error[DRMAA_ERROR_STRING_BUFFER];
int errnum = 0;
drmaa_job_template_t *jt = NULL;

/* Init Session & Allocate Job Template */

errnum = drmaa_set_attribute (jt, DRMAA_REMOTE_COMMAND, "sleeper.sh",
                           error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't set remote command: %s\n", error);
    return 1;
}

errnum = drmaa_set_vector_attribute (jt, DRMAA_V_ARGV, argv, error,
                                      DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't set remote command args: %s\n", error);
    return 1;
}

/* Run Job, Delete Job Template, & Exit Session */
```

Job Execution

- ```
int drmaa_run_job (
 char *job_id,
 size_t job_id_len,
 drmaa_job_template_t *jt,
 char *error_diagnosis,
 size_t error_diag_len
);
```
- ```
int drmaa_run_bulk_jobs (
    drmaa_job_ids_t **jobids,
    drmaa_job_template_t *jt,
    int start,
    int end,
    int increment,
    char *error_diagnosis,
    size_t error_diag_len
);
```

Job IDs

- drmaa_job_ids_t type
 - Implementation independent vector
- int drmaa_get_next_job_id (
 drmaa_job_ids_t *values,
 char *value,
 int value_len
);
- int drmaa_release_job_ids (
 drmaa_job_ids_t *values
);

Example: Job Execution

```
char error[DRMAA_ERROR_STRING_BUFFER];
char jobid[DRMAA_JOBNAME_BUFFER];
int errnum = 0;
drmaa_job_template_t *jt = NULL;

/* Init Session, Allocate Job Template, & Set Attributes */

errnum = drmaa_run_job (jobid, DRMAA_JOBNAME_BUFFER, jt,
                         error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't run job: %s\n", error);
    return 1;
}
else {
    printf ("Your job has been submitted with id %s\n", jobid);
}

/* Delete Job Template, & Exit Session */
```

Example: Bulk Job Execution

```
char error[DRMAA_ERROR_STRING_BUFFER];
char jobid[DRMAA_JOBNAME_BUFFER];
int errnum = 0;
drmaa_job_template_t *jt = NULL;
drmaa_job_ids_t *ids = NULL;

/* Init Session, Allocate Job Template, & Set Attributes */

errnum = drmaa_run_bulk_jobs (&ids, jt, 1, 30, 2, error,
                           DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't run bulk jobs: %s\n", error);
    return 1;
}
else {
    while (drmaa_get_next_job_id (ids, jobid, DRMAA_JOBNAME_BUFFER) ==
           DRMAA_ERRNO_SUCCESS) {
        printf ("A job task has been submitted with id %s\n", jobid);
    }

    if (drmaa_release_job_ids (ids) != DRMAA_ERRNO_SUCCESS) ...
}

/* Delete Job Template, & Exit Session */
```

Job Synchronization

- ```
int drmaa_wait (
 const char *job_id,
 char *job_id_out,
 size_t job_id_len,
 int *stat,
 signed long timeout,
 drmaa_attr_values_t **rusage,
 char *error_diagnosis,
 size_t error_diag_len
);
```

# Bulk Job Synchronization

- ```
int drmaa_synchronize (
    const char *job_ids[],
    signed long timeout,
    int dispose,
    char *error_diagnosis,
    size_t error_diag_len
);
```
- job_ids must be NULL-terminated

Example: Job Synchronization

```
char error[DRMAA_ERROR_STRING_BUFFER];
char jobid[DRMAA_JOBNAME_BUFFER];
char jobid_out[DRMAA_JOBNAME_BUFFER];
int errnum = 0;
int status = 0;
drmaa_attr_values_t *rusage = NULL;

/* Init Session, Allocate Job Template, Set Attributes, & Run Job */

errnum = drmaa_wait (jobid, jobid_out, DRMAA_JOBNAME_BUFFER, &status,
                      DRMAA_TIMEOUT_WAIT_FOREVER, &rusage, error,
                      DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't wait for job: %s\n", error);
    return 1;
}
else {
    /* Print Job Finish Status */
}

/* Delete Job Template, & Exit Session */
```

Example: Bulk Job Synchronization

```
char error[DRMAA_ERROR_STRING_BUFFER];
char *jobids[2] = {DRMAA_JOB_IDS_SESSION_ALL, NULL}; /* Idiom */
int errnum = 0;

/* Init Session, Allocate Job Template,
   Set Attributes, & Run Jobs */

errnum = drmaa_synchronize (jobids, DRMAA_TIMEOUT_WAIT_FOREVER,
                           1, error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't wait for jobs: %s\n", error);
    return 1;
}

/* Delete Job Template, & Exit Session */
```

Example: Hybrid Job Synchronization

```
/* Init Session, Allocate Job Template, Set Attributes, & Run Jobs */
errnum = drmaa_synchronize (jobids, DRMAA_TIMEOUT_WAIT_FOREVER,
                           0, error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't wait for jobs: %s\n", error);
    return 1;
} else {
    for (count = start; count <= end; count += incr) {
        drmaa_wait (DRMAA_JOB_IDS_SESSION_ANY, jobid,
                      DRMAA_JOBNAME_BUFFER,
                      &status, DRMAA_TIMEOUT_WAIT_FOREVER, &rusage,
                      error,
                      DRMAA_ERROR_STRING_BUFFER);

        /* Print Job Finish Status */
    }
}

/* Delete Job Template, & Exit Session */
```

Finish Status

- `int drmaa_wifexited (int *exited, int status, char *error_diagnosis, size_t error_diag_len);`
- `int drmaa_wexitstatus (int *exit_status, int status, char *error_diagnosis, size_t error_diag_len);`
- `int drmaa_wifaborted (int *aborted, int status, char *error_diagnosis, size_t error_diag_len);`
- `int drmaa_wifsignaled (int *signaled, int status, char *error_diagnosis, size_t error_diag_len);`
- `int drmaa_wtermsig (char *signal, size_t signal_len, int status, char *error_diagnosis, size_t error_diag_len);`
- `int drmaa_wcoredump (int *core_dumped, int status, char *error_diagnosis, size_t error_diag_len);`

Example: Finish Status

```
int aborted = 0;

drmaa_wifaborted (&aborted, status, NULL, 0);

if (aborted == 1)
    printf ("Job never ran\n");
else {
    int exited = 0;

    drmaa_wifexited (&exited, status, NULL, 0);

    if (exited == 1) {
        int exit_status = 0;

        drmaa_wexitstatus (&exit_status, status, NULL, 0);
        printf ("Job exited with status %d\n", exit_status);
    } else {
        int signaled = 0;

        drmaa_wifsignaled (&signaled, status, NULL, 0);

        if (signaled == 1) {
            char termsig[DRMAA_SIGNAL_BUFFER+1];

            drmaa_wtermsig (termsig, DRMAA_SIGNAL_BUFFER, NULL, 0);
            printf ("Job exited due to signal: %s\n", termsig);
        }
    }
}
```

Job Control

- ```
int drmaa_control (
 const char *jobid,
 int action,
 char *error_diagnosis,
 size_t error_diag_len
);
```
- Actions
  - DRMAA\_CONTROL\_SUSPEND
  - DRMAA\_CONTROL\_RESUME
  - DRMAA\_CONTROL\_HOLD
  - DRMAA\_CONTROL\_RELEASE
  - DRMAA\_CONTROL\_TERMINATE

# Example: Job Control

```
char error[DRMAA_ERROR_STRING_BUFFER];
char jobid[DRMAA_JOBNAME_BUFFER];
int errnum = 0;

/* Init Session, Allocate Job Template,
 Set Attributes & Run Job */

errnum = drmaa_control (jobid, DRMAA_CONTROL_TERMINATE,
 error, DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
 fprintf (stderr, "Couldn't delete job: %s\n", error);
 return 1;
}

/* Delete Job Template, & Exit Session */
```

# Job Monitoring

- ```
int drmaa_job_ps (
    const char *jobid,
    int *remote_ps,
    char *error_diagnosis,
    size_t error_diag_len
);
```
- States
 - DRMAA_PS_QUEUED_ACTIVE
 - DRMAA_PS_RUNNING
 - DRMAA_PS_SYSTEM_ON_HOLD
 - DRMAA_PS_USER_ON_HOLD
 - DRMAA_PS_USER_SYSTEM_ON_HOL
D
 - DRMAA_PS_SYSTEM_SUSPENDED
 - DRMAA_PS_USER_SUSPENDED
 - DRMAA_PS_DONE
 - DRMAA_PS_FAILED
 - DRMAA_PS_UNDETERMINED

Example: Job Monitoring

```
char error[DRMAA_ERROR_STRING_BUFFER];
char jobid[DRMAA_JOBNAME_BUFFER];
int errnum = 0;
int status = 0;

/* Init Session, Allocate Job Template,
   Set Attributes & Run Job */

errnum = drmaa_job_ps (jobid, &status, error,
                      DRMAA_ERROR_STRING_BUFFER);

if (errnum != DRMAA_ERRNO_SUCCESS) {
    fprintf (stderr, "Couldn't get job status: %s\n", error);
    return 1;
}
else {
    /* Print Job Status */
}

/* Delete Job Template, & Exit Session */
```

Example: Compiling, Linking, & Running

- N1 Grid Engine 6.0

- Compiling
 - Include DRMAA header

```
cc -o app.o -I$SGE_ROOT/include app.c
```

- Linking
 - Include DRMAA library in library path
 - LD_LIBRARY_PATH, SHLIB_PATH, et al
 - Link in DRMAA library

```
cc -o app -ldrmaa app.o
```

- Running
 - Include DRMAA library in library path
 - LD_LIBRARY_PATH, SHLIB_PATH, et al

Java Language Binding

- 0.4.2 is the current spec
 - Close to final
 - May adapt to OO binding spec
- Implemented by Sun (N1 Grid Engine 6.0s1)
 - Built as layer on top of C binding
- Spec'ed using API/SPI model
 - Application Programming interface for using DRMAA
 - Service Provider Interface for building implementations
 - All implementations should be interchangeable

Object Model

- DRMAASessionFactory
 - Creates a DRMAASession
 - Enables SPI
- DRMAASession
 - Job submission, control, and monitoring
- JobTemplate
 - Describes a job for submission
- JobInfo
 - Finish information from DRMAASession.wait()
- DRMAAException
 - Parent of Exception tree

Exception Handling

- Exceptions instead of error codes
 - Exceptions map roughly to error codes
 - InvalidAttributeException
 - InconsistentStateException
- Signatures declare DRMAAException
 - JDBC model
 - Can catch more specific Exceptions if needed

Session Handling

- DRMAASessionFactory.getFactory ()
 - Static method – returns a DRMAASessionFactory
- DRMAASessionFactory.getSession ()
 - Returns a DRMAASession
 - Still only single-session
- DRMAASession.init (String contact)
 - Initializes a session
- DRMAASession.exit ()
 - Destroys a session
 - Must be called if init() is called
- Idiom: Use a shutdown hook to call exit()

Example: Session Handling

```
DRMAASessionFactory factory = DRMAASessionFactory.getFactory ();
DRMAASession = factory.getSession ();

try {
    session.init (args[0]);
}
catch (DRMAAEexception e) {
    System.err.println ("Error: " + e.getMessage ());
}

System.addShutdownHook (new Runnable () {
    public void run () {
        try {
            session.exit ();
        }
        catch (DRMAAEexception e) {
            System.err.println ("Error exiting: " + e.getMessage ());
        }
    }
});;

// Do Stuff
```

Instance Handling

- DRMAASession.getContact ()
 - Returns contact as a String
- DRMAASession.getDRMSystem ()
 - Returns DRM system as String
- DRMAASession.getDRMAAImplementation ()
 - Returns DRMAA implementation as a String
- DRMAASession.getVersion ()
 - Returns a DRMAASession.Version object
 - Contains major and minor version numbers as ints

Example: Instance Handling

```
DRMAASession session = null;
DRMAASession.Version version = null;

// Init Session

try {
    System.out.println ("Contact: " + session.getContact ());
    System.out.println ("DRM System: " + session.getDRMSystem ());
    System.out.println ("DRMAA Implementation: " +
                       session.getDRMAAImplementation ());

    version = session.getVersion ();

    System.out.println ("Version: " + version.major + "." +
                       version.minor);
}

catch (DRMAAEException e) {
    System.err.println ("Error: " + e.getMessage ());
}

// Exit Session
```

Job Templates

- DRMAASession.createJobTemplate ()
 - Creates a new JobTemplate
- JobTemplate.delete ()
 - Deletes the JobTemplate
- JobTemplate.set<Property> (<type> value)
 - Sets the property
 - JobTemplate.setRemoteCommand (“sleeper.sh”)
- JobTemplate.get<Property> ()
 - Returns the property's current value
 - String cmd = JobTemplate.getRemoteCommand ()·

Job Attributes

- boolean blockEmail
- String[] emailAddresses
- String errorPath
- String[] inputParameters
- String inputPath
- String jobCategory
- Properties jobEnvironment
- String jobName
- int jobSubmissionState
- boolean joinFiles
- String nativeSpecification

- String outputPath
- String remoteCommand
- Date startTime
- String workingDirectory

Optional

- Date deadlineTime
- long hardRunDurationLimit
- long hardWallclockTimeLimit
- long softRunDurationLimit
- long softWallclockTimeLimit
- byte transferFiles

Job Execution

- DRMAASession.runJob (JobTemplate jt)
 - Submits a job based on jt
 - Returns the job id as an opaque String
- DRMAASession.RunBulkJobs (JobTemplate jt, int start, int end, int increment)
 - Submits $(end - start)/increment$ jobs based on jt
 - Returns the job ids as a List of opaque Strings

Example: Job Execution

```
DRMAASession session = DRMAASessionFactory.getFactory ().getSession ();
JobTemplate jt = null;

session.init ("/sgc:default");
jt = session.createJobTemplate ();
jt.setRemoteCommand ("/sgc/examples/jobs/sleeper.sh");
jt.setInputParameters (new String[] {"300"});
jt.setOutputPath ("/dev/null");
jt.setJoinFiles (true);

String jobId = session.runJob (jt);
System.out.println ("Single job id is " + jobId);

List jobIds = session.runBulkJob (jt, 1, 10, 1);
Iterator i = jobIds.iterator ();

while (i.hasNext ()) {
    jobId = (String)i.next ();
    System.out.println ("Bulk job id is " + jobId);
}

jt.delete ();
session.exit ();
```

Job Synchronization

- DRMAASession.wait
(String jobId, long timeout)
 - Waits for the job with the given id to succeed or fail
 - Returns a JobInfo object for the job
- DRMAASession.synchronize
(List jobIds, long timeout, boolean dispose)
 - Wait for all jobs with ids in the list to succeed or fail
 - Must be used in conjunction with wait() to get the JobInfo objects

Example: Job Synchronization

```
DRMAASessionFactory factory = DRMAASessionFactory.getFactory ();
DRMAASession session = factory.getSession ();
JobTemplate jt = null;

session.init ("/sge:default");
jt = session.createJobTemplate ();
jt.setRemoteCommand ("/sge/examples/jobs/sleeper.sh");
jt.setInputParameters (new String[] {"300"});
jt.setOutputPath ("/dev/null");
jt.setJoinFiles (true);

String jobId = session.runJob (jt);
System.out.println ("Job " + jobId + " has been submitted");

JobInfo info = session.wait (jobId,
                           DRMAASession.TIMEOUT_WAIT_FOREVER);
System.out.println ("Job " + jobId + " has finished");

// Print JobInfo

// Delete Job Template & Exit Session
```

Example: Bulk Job Synchronization

```
DRMAASessionFactory factory = DRMAASessionFactory.getFactory ();
DRMAASession session = factory.getSession ();
JobTemplate jt = null;

session.init ("/sgc:default");
jt = session.createJobTemplate ();
jt.setRemoteCommand ("/sgc/examples/jobs/sleeper.sh");
jt.setInputParameters (new String[] {"300"});
jt.setOutputPath ("/dev/null");
jt.setJoinFiles (true);

List jobIds = session.runBulkJob (jt, 1, 10, 1);

System.out.println ("Jobs 1-10 submitted.");

session.synchronize (jobIds, DRMAASession.TIMEOUT_WAIT_FOREVER,
                     true);

System.out.println ("All jobs have completed.");

// Delete Job Template & Exit Session
```

Example: Hybrid Job Synchronization

```
DRMAASession session = null;
JobTemplate jt = null;

// Init Session, Create JobTemplate, & Set Attributes

List jobIds = session.runBulkJob (jt, 1, 10, 1);

System.out.println ("Jobs 1-10 submitted.");

session.synchronize (jobIds, 1000, false);

Iterator i = jobIds.iterator ();

while (i.hasNext ()) {
    JobInfo info = session.wait ((String)i.next (),
                                DRMAASession.TIMEOUT_NO_WAIT);

        // Print JobInfo
}

// Delete Job Template & Exit Session
```

Finish Status

- `JobInfo.hasExited ()`
 - Returns whether the job has exited – boolean
- `JobInfo.getExitStatus ()`
 - Returns the job's exit status – int
- `JobInfo.wasAborted ()`
 - Returns whether the job terminated abnormally – boolean
- `JobInfo.hasSignaled ()`
 - Returns whether the job terminated due to a signal – boolean
- `JobInfo.getTerminatingSignal ()`
 - Returns the name of the terminating signal – String
- `JobInfo.hasCoreDump ()`
 - Returns whether the job created a core dump – boolean
- `JobInfo.getResourceUsage ()`
 - Returns the job's resource usage – Map

Example: Finish Status

```
JobInfo info = null;

// Init Session, Create Template, Set Attributes, Submit Job, &
// Wait For Job

if (info.wasAborted ()) {
    if (info.hasCoreDump ()) {
        System.out.println ("Job dumped core");
    }
    else {
        System.out.println ("Job crashed");
    }
}
else if (info.hasExited ()) {
    System.out.println ("Job exited with " + info.getExitStatus ());
}
else if (info.hasSignaled ()) {
    System.out.println ("Job got signal:" + info.getTerminatingSignal ());
}
else {
    System.out.println ("Job exit with unclear conditions");
}

// Delete Job Template & Exit Session
```

Example: Resource Usage

```
JobInfo info = null;

// Init Session, Create Template, Set Attributes,
// Submit Job, & Wait For Job

Map rmap = info.getResourceUsage ();
Iterator i = rmap.keySet ().iterator ();

while (i.hasNext ()) {
    String name = (String)i.next ();
    String value = (String)rmap.get (name);

    System.out.println (name + " = " + value);
}

// Delete Job Template & Exit Session
```

Job Control

- DRMAASession.control (String jobId, int action)
 - Action may be:
 - DRMAASession.HOLD
 - DRMAASession.RELEASE
 - DRMAASession.SUSPEND
 - DRMAASession.RESUME
 - DRMAASession.TERMINATE

Example: Job Control

```
DRMAASessionFactory factory = DRMAASessionFactory.getFactory ();
DRMAASession session = factory.getSession ();
JobTemplate jt = null;

session.init ("/sge:default");
jt = session.createJobTemplate ();
jt.setRemoteCommand ("/sge/examples/jobs/sleeper.sh");
jt.setInputParameters (new String[] {"300"});
jt.setOutputPath ("/dev/null");
jt.setJoinFiles (true);

String jobId = session.runJob (jt);
System.out.println ("Job " + jobId + " has been submitted");

Thread.sleep (5000);
session.control (jobId, DRMAASession.TERMINATE);
System.out.println ("Job " + jobId + " has been deleted");

// Delete Job Template & Exit Session
```

Job Monitoring

- DRMAASession.getJobProgramStatus
(String jobId)
 - Returns a status code as an int
 - DRMAASession.QUEUED_ACTIVE
 - DRMAASession.RUNNING
 - DRMAASession.SYSTEM_ON_HOLD
 - DRMAASession.USER_ON_HOLD
 - DRMAASession.USER_SYSTEM_ON_HOLD
 - DRMAASession.SYSTEM_SUSPENDED
 - DRMAASession.USER_SUSPENDED
 - DRMAASession.DONE
 - DRMAASession.FAILED
 - DRMAASession.UNDETERMINED

Example: Job Monitoring

```
// Init Session, Create Template, Set Attributes, & Submit Job

switch (session.getJobProgramStatus (jobId)) {
    case DRMAASession.UNDERTERMINED:
        System.out.println ("undetermined"); break;
    case DRMAASession.QUEUED_ACTIVE:
        System.out.println ("queued and active"); break;
    case DRMAASession.SYSTEM_ON_HOLD:
    case DRMAASession.USER_ON_HOLD:
    case DRMAASession.USER_SYSTEM_ON_HOLD:
        System.out.println ("queued and on hold"); break;
    case DRMAASession.RUNNING:
        System.out.println ("running"); break;
    case DRMAASession.SYSTEM_SUSPENDED:
    case DRMAASession.USER_SUSPENDED:
    case DRMAASession.USER_SYSTEM_SUSPENDED:
        System.out.println ("suspended"); break;
    case DRMAASession.DONE:
        System.out.println ("done"); break;
    case DRMAASession.FAILED:
        System.out.println ("failed"); break;
}

// Delete Job Template & Exit Session
```

Example: Compiling & Running

- N1 Grid Engine 6.0s1

- Compiling

- Include drmaa.jar in the classpath

```
javac -classpath \
$SGE_ROOT/lib/drmaa.jar *.java
```

- Running

- Include libdrmaa in the library path

- LD_LIBRARY_PATH, SHLIB_PATH, et al

- Include drmaa.jar in the classpath

```
java -cp $SGE_ROOT/lib/drmaa.jar:. \
MyApp
```

Additional Information

- DRMAA Working Group
<http://www.drmaa.org>
<https://forge.gridforum.org/projects/drmaa-wg>
- N1 Grid Engine
<http://www.sun.com/software/gridware>
<http://gridengine.sunsource.net>
 - Tutorials
</project/gridengine/howto/drmaa.html>
/project/gridengine/howto/drmaa_java.html
- Condor
<http://www.cs.wisc.edu/condor>
- GGF
<http://www.ggf.org>



GGF12: DRMAA Tutorial C and Java Language Bindings

dan.templeton@sun.com

