



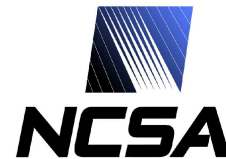
the globus alliance
www.globus.org

Basic Execution Management input from Globus Alliance

Karl Czajkowski
Univa Corporation



|epcc|



Univa





Must Support: Non-trivial Applications

- Real-time or deadline-sensitive jobs
 - ◆ Want localized, *very good* resource
- Large jobs
 - ◆ Large and/or coupled models
 - ◆ Want to *coordinate* a few good resources
- High-throughput job sets
 - ◆ Many related jobs from one user/problem
 - ◆ Many unrelated jobs from many users
 - ◆ Want scalable job control *everywhere*



Scope: BES is Step 3 Only

1. Discovery

- ◆ “What is out there? (of relevance) (to me)...”
- ◆ Finds BES providers

2. Inspection

- ◆ “How do relevant providers compare?”
- ◆ Compare policies, status, etc.

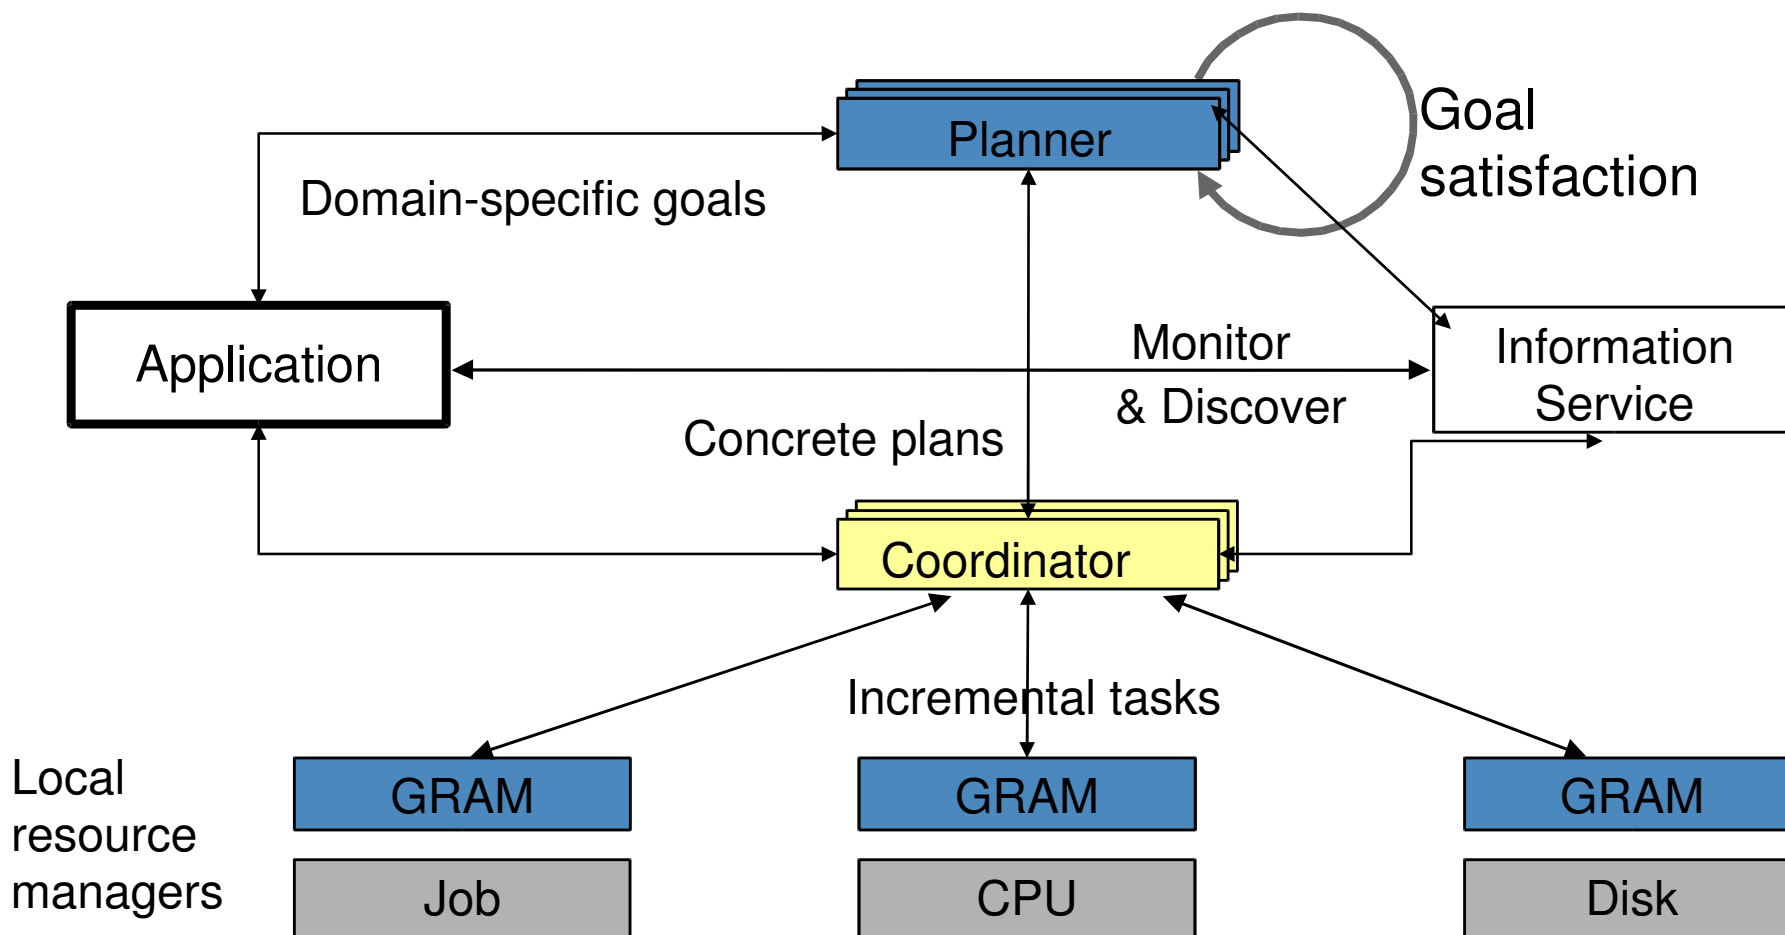
3. Agreement

- ◆ “Will/did I get what I need executed?”
- ◆ The core Execution Management problem

...Process can iterate due to adaptation



(GRAM/BES are Bottom Layer)



WS-Agreement

- Ongoing standardization effort
 - ◆ In GGF's GRAAP-WG
 - ◆ Several issues raised in public comment period
 - Need 3-6 months to address and reenter public comment?
- Generalizes GRAM ideas
 - ◆ Service-oriented architecture
 - ◆ Container/host becomes *Service Provider*
 - ◆ Tasks/jobs become *Negotiated Services*
 - ◆ Mgmt state presented as *Agreement* services
- Supports composition w/ domain terms

JSDL

- Job Submission Description Language
 - ◆ Provides an ontology for “job to be run”
- Requires external spec to put it in context
 - ◆ WS-Agreement can use JSDL as “service terms”
- A mix of abstract and concrete job bits
 - ◆ Executable+args+environment
 - Assume application already provisioned at host
 - ◆ Abstract “portable” job profiles
 - Try to cope w/ localized differences

Need to Clarify Timeline

- Theoretically: WS-Agreement + JSDL + epsilon
 - ◆ Delivery could be delayed by dependencies
 - WS-Agreement needs another revision
 - JSDL might as well, to address concrete job issues?
 - ◆ Scoping concerns
 - Is current JSDL 1.0 target complete enough? Too big?
 - No practical experience w/ WS-Agreement yet
- Conservative/fast: base on existing system(s)
 - ◆ We would help slice/dice WS-GRAM WSDL
 - ◆ Other vendors offering to help?

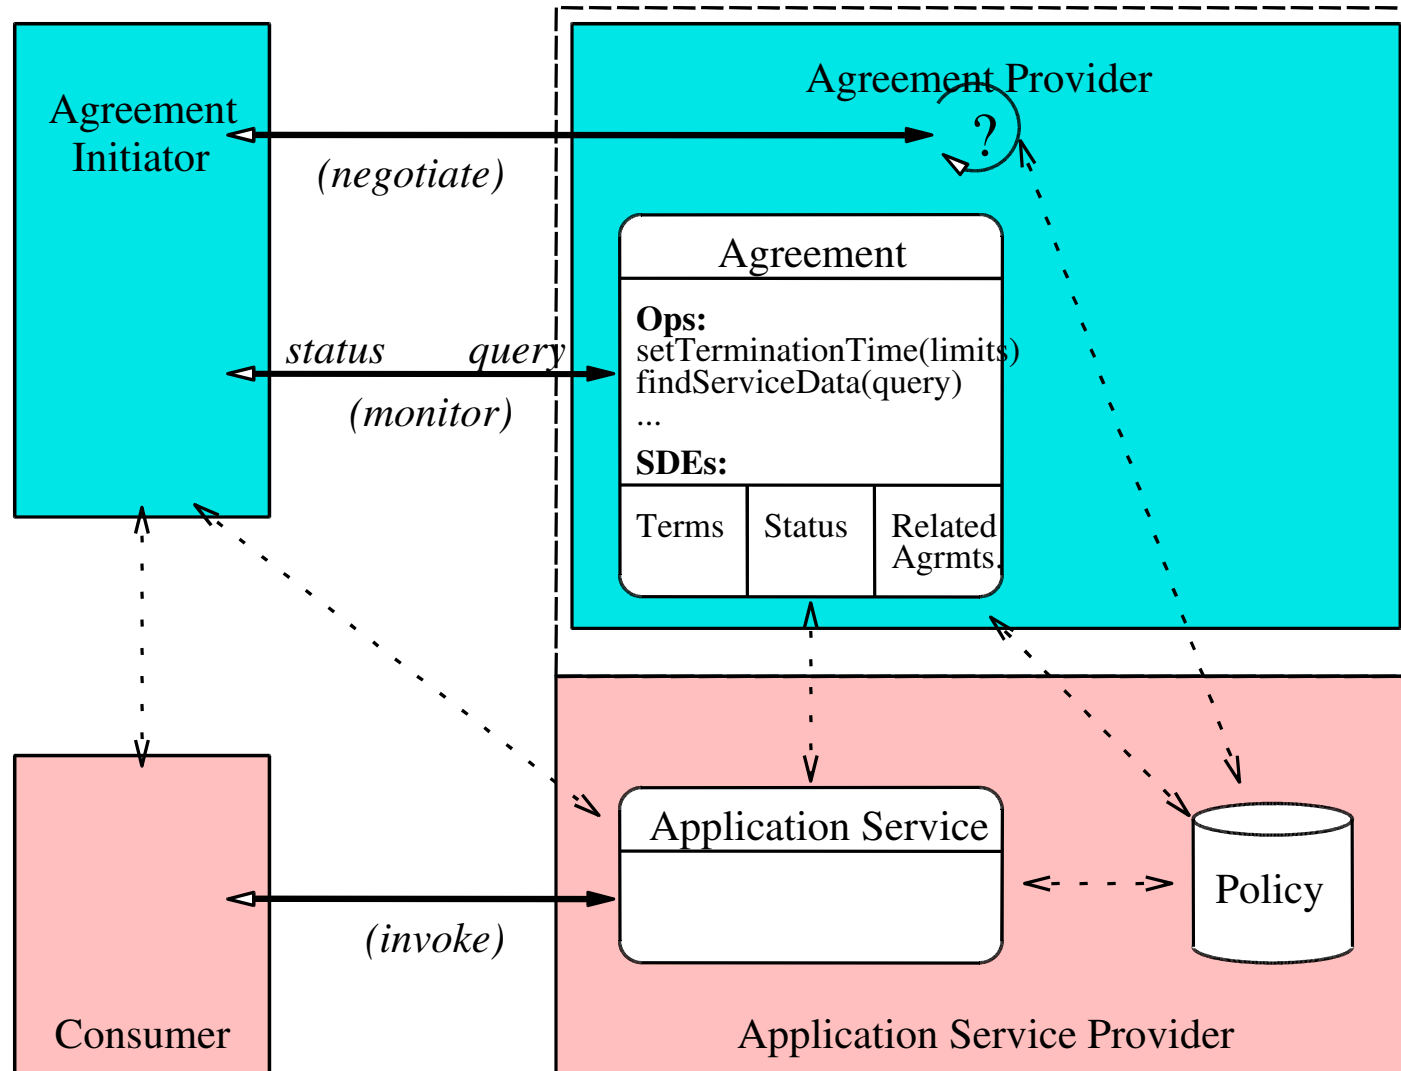
State of the Art

- Discovery is very hard and immature
 - ◆ Some viable information gathering systems
 - ◆ But information models have gaps
 - Lots of low-level “buttons and knobs” stuff, e.g. CIM
 - Some overly abstract stuff, e.g. GLUE, GRAM today
 - Complexity already a barrier to entry
 - ◆ RM policy: personalized scope/relevance
- BES needs some basic metadata/properties
 - ◆ Introspect for “dialect” support
- Assume more will be developed over time

Caution!

- Must choose scope wisely and stick to it
 - ◆ Many groups flounder on recurring debates
 - Old issues revisited for new participants
 - Old issues tripped over from new angle (and not noticed)
 - ◆ *The real work is in rendering abstract architectural consensus into concrete syntax w/o fatal warts and flaws*
- Compositionality a mixed blessing
 - ◆ Defer decisions on difficult problems
 - Can BES address file staging? Credential mgmt?
 - ◆ But, more opportunities for non-interop
 - “Power set” of different specs/extensions...

WS-Agreement Entities





Simple Negotiation

- `AgreementFactory::createAgreement()`
 - ◆ Coarse-grained
 - ◆ Conventional fault/response model
 - ◆ Batch negotiation of complex terms
 - ◆ Idiom: enables one-shot job submission
- `Agreements can be chained`
 - ◆ Establish stateful context of Agreements
 - ◆ New Agreement depends on or claims context
- `Need companion specs for advanced scenarios`



WS-Agreement is a Protocol

- WS-Agreement is a message model...
 - ◆ Not a software component
 - ...applicable to previous examples
 - ◆ Interface standard between components
 - ◆ Improve interoperability of other systems
 - ◆ To enable composition/federation
- (Possible WS-Agreement conversion examples:
- ◆ GRAM, Condor
 - ◆ Workflow, economic scheduling
 - ◆ PBS, LSF, CSF)

Specifying Terms: Who and What?

In a service provisioning **domain**

(e.g. “computational jobs”)

- A **standard** specifies **domain** terms/concepts
- A **provider** specifies its support for
 - ◆ *some or all* of the domain standard terms
 - ◆ a given term, *specifically*
 - Within *behavioral* constraints
 - Within *negotiability* constraints
 - With extra fields/sub-terms
 - Arbitrary term *properties*: e.g. *optional* or *required*
- A **client** *discovers compatible providers*



Agreement-based Jobs

- Agreement represents “queue entry”
 - ◆ Commitment with job parameters etc.
 - ◆ This is the management “proxy” for job
- Agreement Provider
 - ◆ i.e. Job scheduler/Queuing system
 - ◆ Management interface to service provider
- Service Provider
 - ◆ i.e. scheduled resource (compute nodes)
- Provided Service is the Job computation

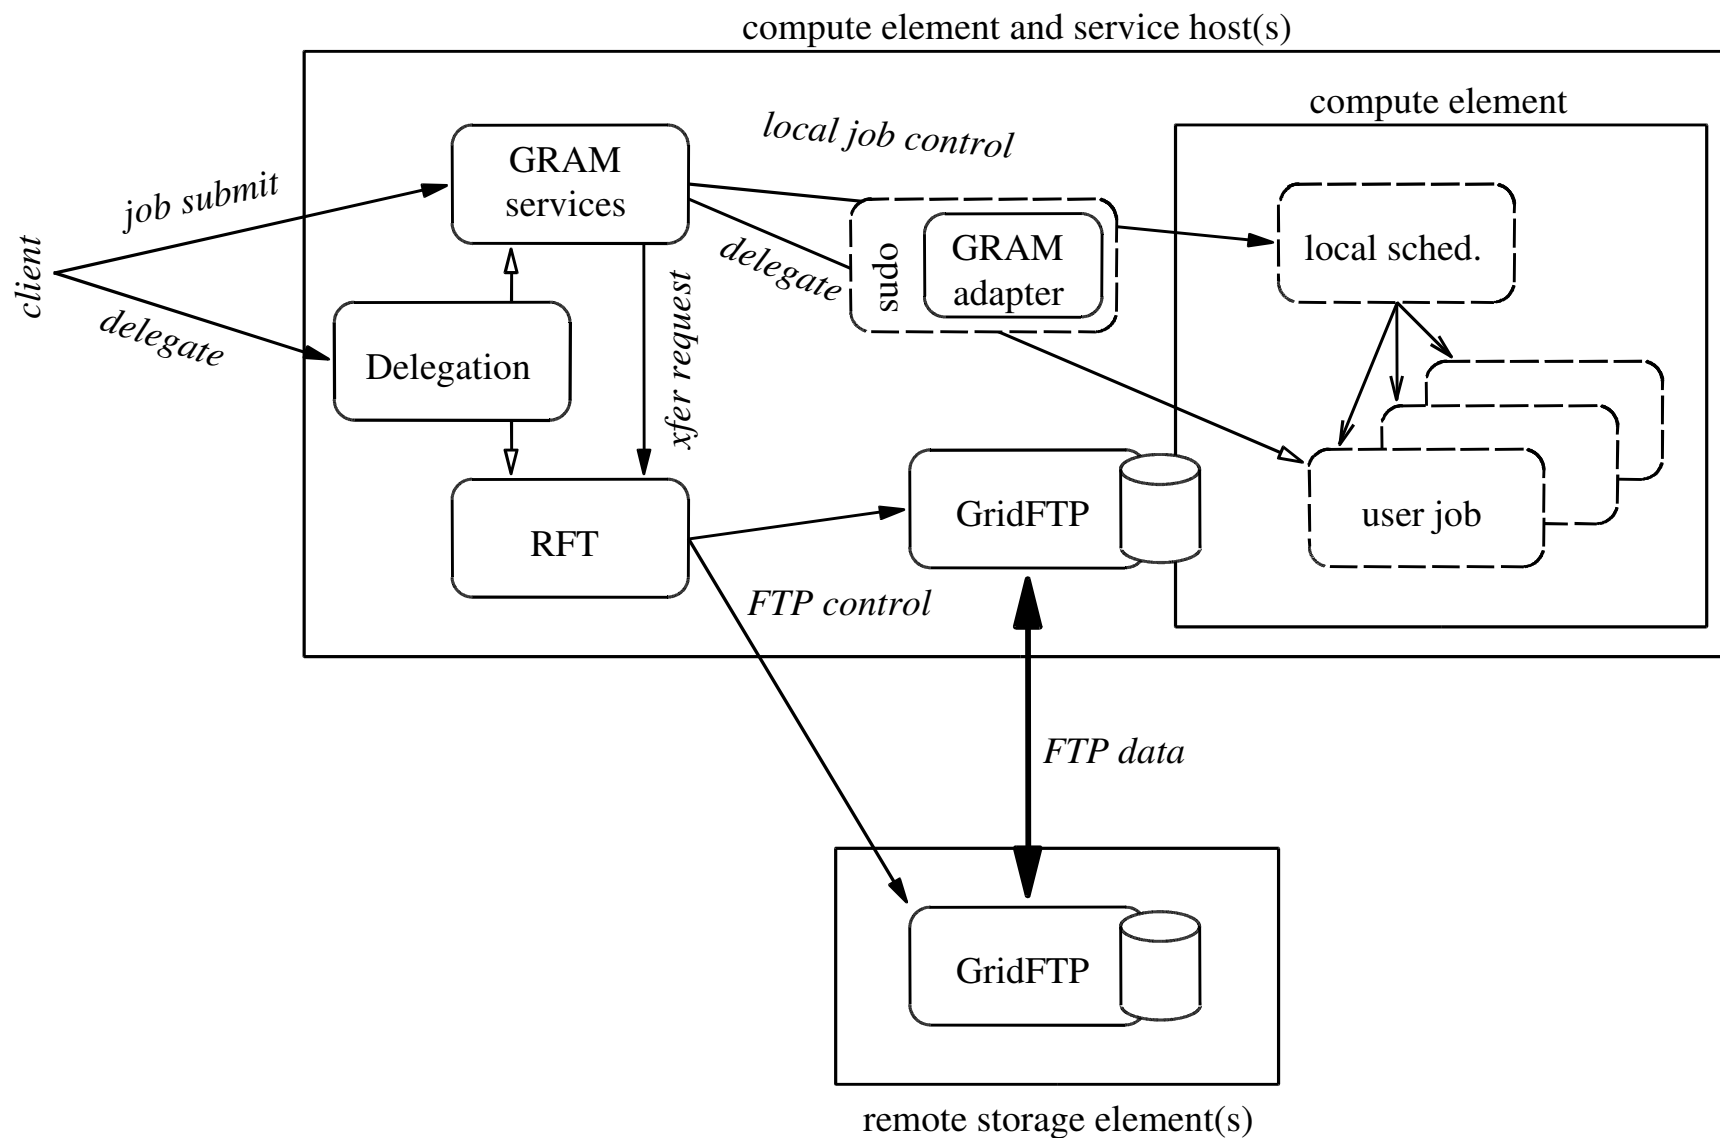


BES as WS-Agreement

- Write a BES spec
 - ◆ A profile combining WS-Agrmt + JSDL
 - ◆ Influence JSDL or spec. extensions?
 - ◆ Be conservative about generating new dialects
- Get some implementation experience
 - ◆ We know how to adapt WS-GRAM
 - ◆ Tie into existing ManagedJob resource impl.
- Put BES in comment period
 - ◆ Determine if experiences warrant revisions

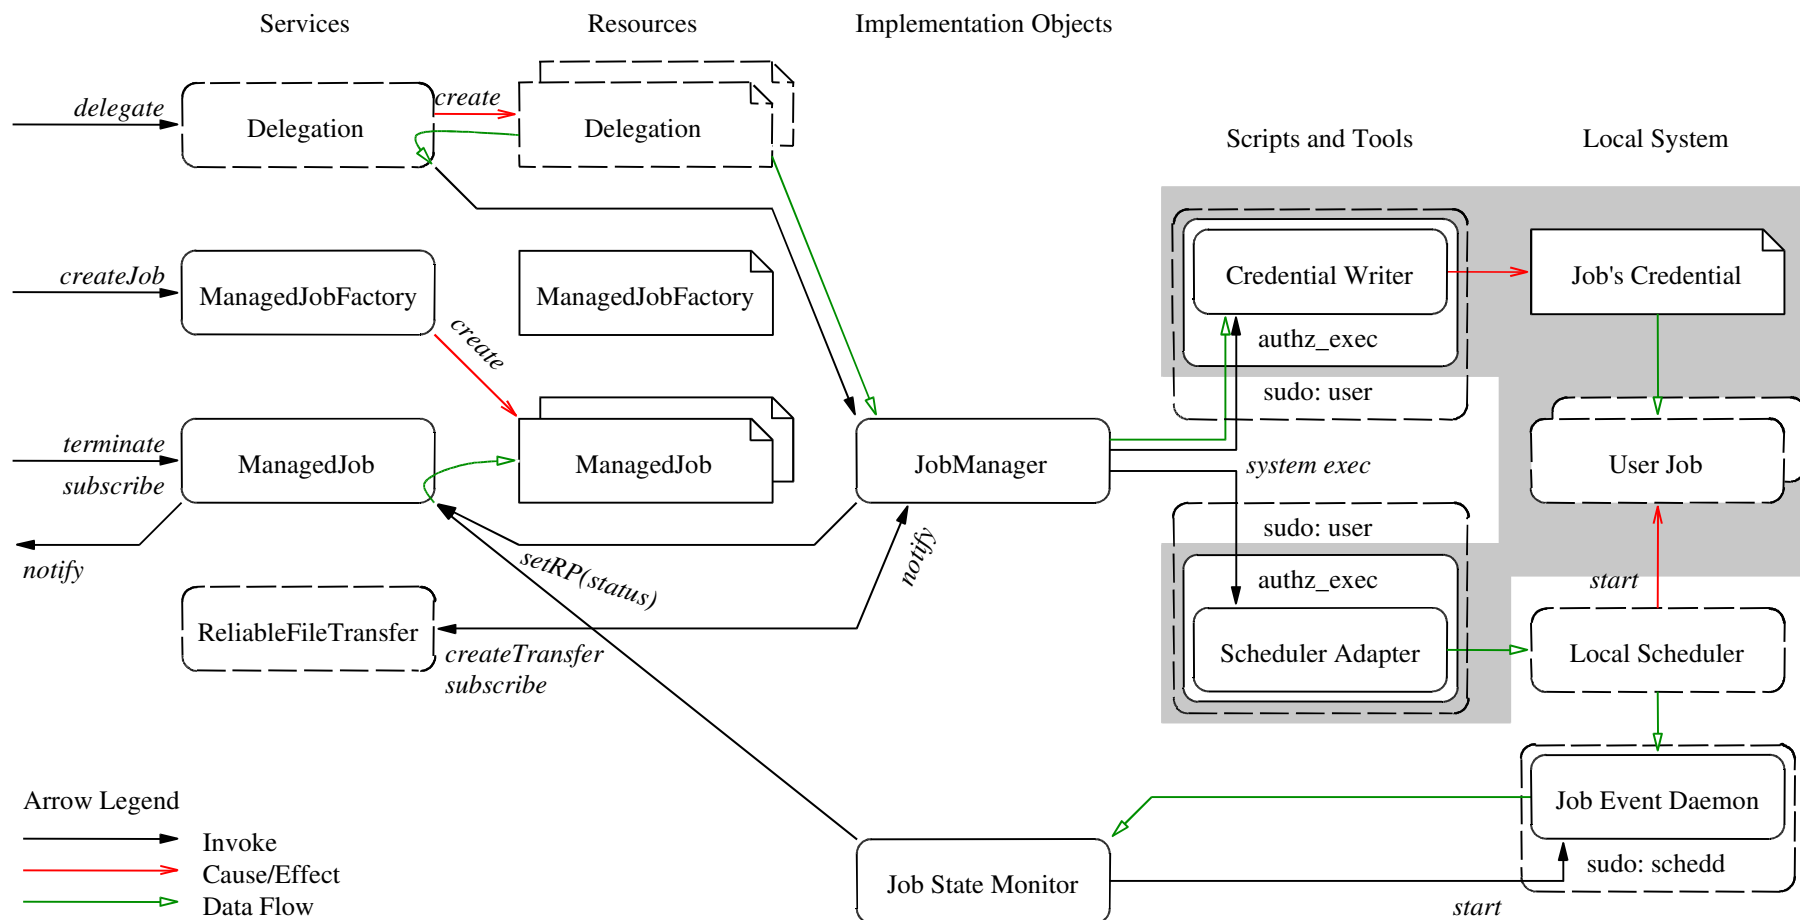


WS-GRAM Approach

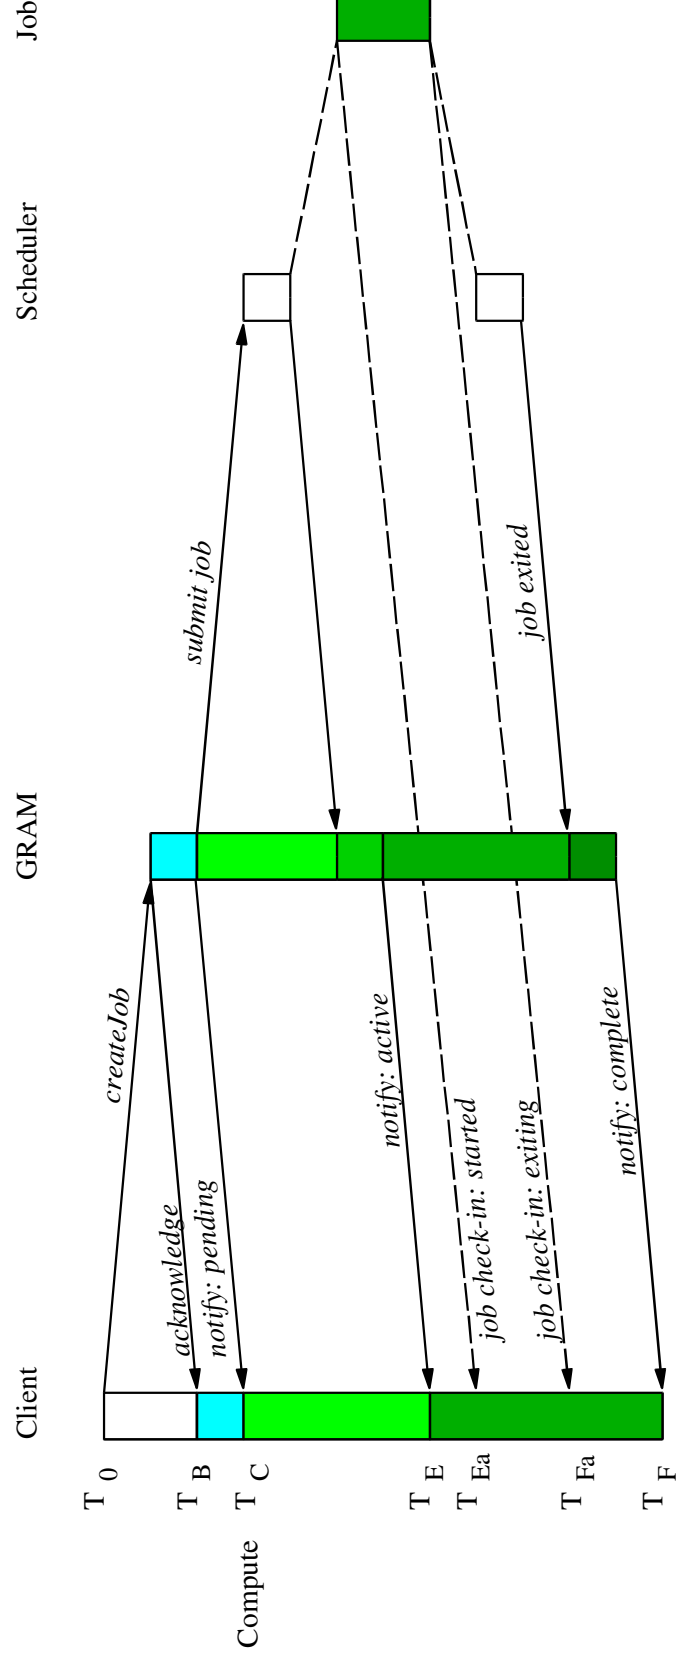




WS-GRAM Software Map



WS-GRAM Base Protocol



WS-GRAM Full Protocol

