




the globus alliance
www.globus.org

Basic Execution Management input from Globus Alliance

Karl Czajkowski
Univa Corporation





Must Support: Non-trivial Applications

- Real-time or deadline-sensitive jobs
 - ◆ Wants localized, *very good* resource
- Large jobs
 - ◆ Large and/or coupled models
 - ◆ Wants to *coordinate* a few good resources
- High-throughput job sets
 - ◆ Many related jobs from one user/problem
 - ◆ Many unrelated jobs from many users
 - ◆ Wants scalable job control *everywhere*

Distributed Execution Management

1. Discovery

- ◆ “What is out there? (of relevance) (to me)...”
- ◆ Finds BES providers

2. Inspection

- ◆ “How do relevant providers compare?”
- ◆ Compare policies, status, etc.

3. Agreement

- ◆ “Will/did I get what I need executed?”
- ◆ The core Execution Management problem

...Process can iterate due to adaptation

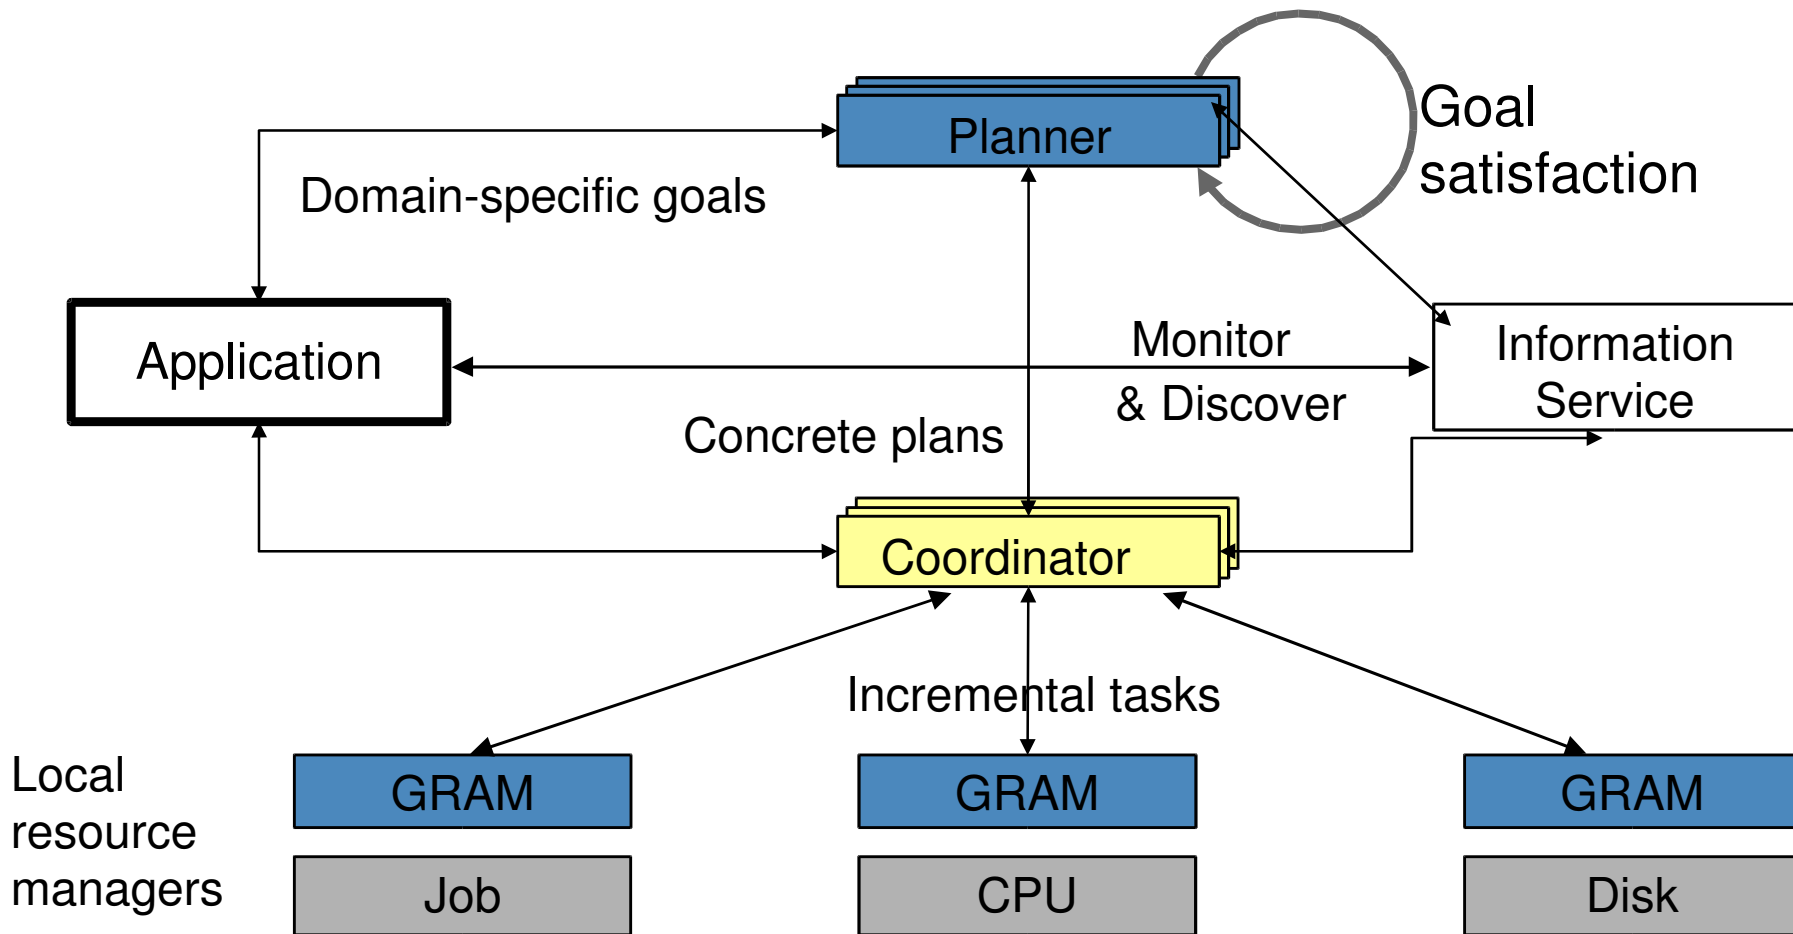


Job/Execution Duality

- GRAM supports job submission
 - ◆ A traditional “bare metal” job to run
 - ◆ With some data staging requirements
- GRAM supports execution management
 - ◆ User needs a virtual host/container
 - ◆ With some environment initialization
- Two sides of the same coin
 - ◆ All job submission IS resource virtualization
 - ◆ Some jobs more virtualizing than others!
 - Run a JVM? X Windows server? User-mode Linux?



Long-term GRAM Architecture





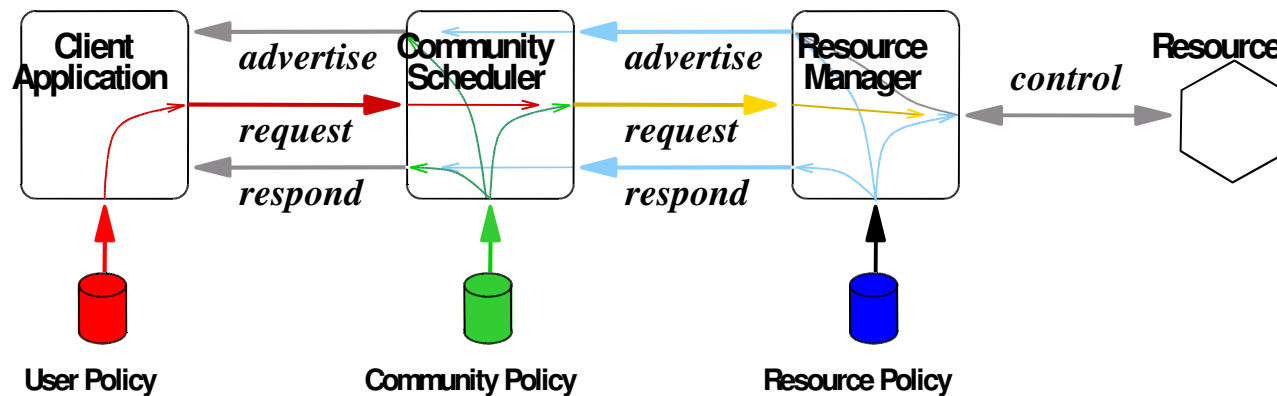
Exec. Mgmt. Mediates Conflict

- Resource Consumers/Applications Goals
 - ◆ Users: deadlines and availability goals
 - ◆ Applications: need coordinated resources
- Localized Resource Owner Goals
 - ◆ Policies distinguish users/communities
- Community Goals Emerge As:
 - ◆ “Global” optimization goals
 - ◆ Aggregate user, application and/or resource
- Reconcile demands via Agreement

An Open Negotiation Model

- Providers in a Global Context
 - ◆ Advertisement and negotiation
 - ◆ Normalized remote client interface
 - ◆ Provider maintains autonomy
- Users or Agents *Bridge* Resources
 - ◆ Drive task submission and provisioning
 - ◆ Coordinate acts across domains
- Community-based “Virtual” Providers
 - ◆ Coordination for collective interest

Intermediaries And Policy



- Resource virtualization can:
 - ◆ Abstract details of underlying resource(s)
 - ◆ Abstract cardinality of aggregates
 - ◆ Map between different resource description domains
- Policies from different domains influence agreement negotiations with intermediaries

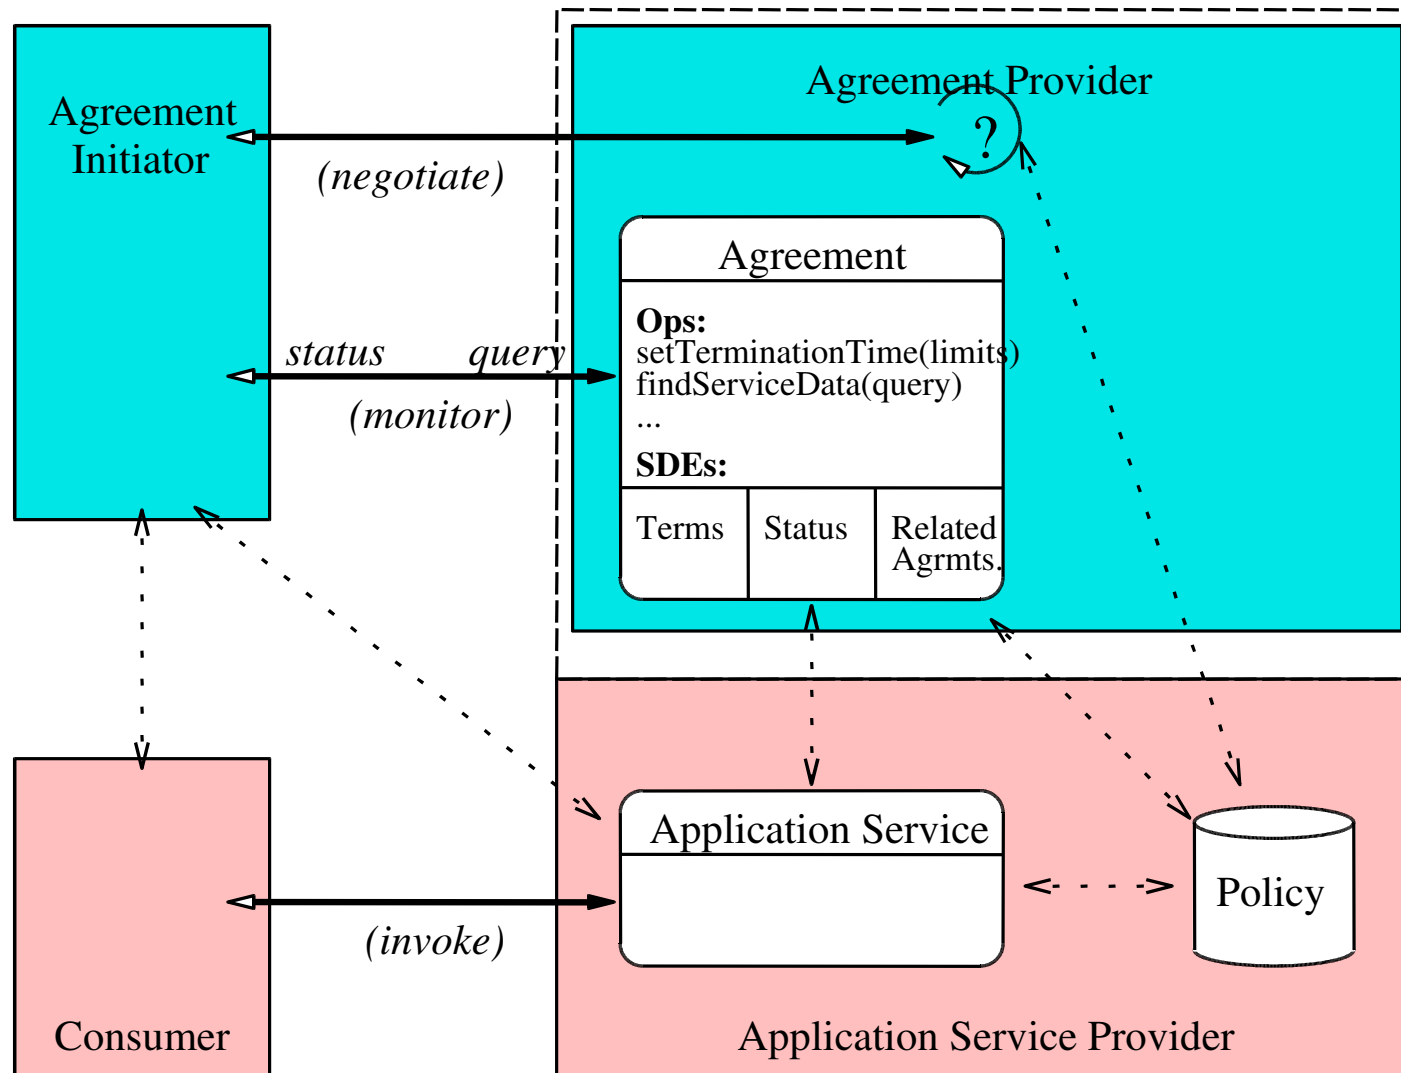
State of the Art

- Discovery is very hard and immature
 - ◆ Some viable information gathering systems
 - ◆ But information models have gaps
 - Lots of low-level “buttons and knobs” stuff, e.g. CIM
 - Some overly abstract stuff, e.g. GLUE, GRAM today
 - Complexity already a barrier to entry
 - ◆ RM policy: personalized scope/relevance
- Inspection is over-emphasized
 - ◆ Inherent race-conditions/scalability problems
- Basic allocation and “agreement” today
 - ◆ Implicit out-of-band intelligence still required

WS-Agreement

- Ongoing standardization effort
 - ◆ In GGF's GRAAP-WG
 - ◆ Several issues raised in public comment period
 - Need 3-6 months to address and reenter public comment?
- Generalizes GRAM ideas
 - ◆ Service-oriented architecture
 - ◆ Resource becomes *Service Provider*
 - ◆ Tasks become *Negotiated Services*
 - ◆ State presented as *Agreement* services
- Supports composition w/ domain terms

WS-Agreement Entities





Simple Negotiation

- `AgreementFactory::createAgreement()`
 - ◆ Coarse-grained
 - ◆ Conventional fault/response model
 - ◆ Batch negotiation of complex terms
 - ◆ Idiom: enables one-shot job submission
- `Agreements can be chained`
 - ◆ Establish stateful context of Agreements
 - ◆ New Agreement depends on/claims context
- `Need companion specs for advanced scenarios`



Agreement-based Jobs

- Agreement represents “queue entry”
 - ◆ Commitment with job parameters etc.
 - ◆ Management interface to dynamic Job(s)
- Agreement Provider
 - ◆ i.e. Job scheduler/Queuing system
 - ◆ Management interface to service provider
- Service Provider
 - ◆ i.e. scheduled resource (compute nodes)
- Provided Service is the Job computation



Advance Reservation for Jobs

- Schedule-based commitment of service
 - ◆ Requires schedule based Agreement terms
- Optional Pre-Agreement
 - ◆ Agreement to facilitate future Job Agreement
 - ◆ Characterizes virtual resource needed for Job
 - ◆ May not need full job terms
- Job Agreement almost as usual
 - ◆ May exploit Pre-Agreement, or
 - Reference existing promise of resource schedule
 - ◆ May get schedule commitment in one shot



WS-Agreement is a Protocol

- WS-Agreement is a message model...
 - ◆ Not a software component
 - ...applicable to previous examples
 - ◆ Interface standard between components
 - ◆ Improve interoperability of other systems
 - ◆ To enable composition/federation
- (Possible WS-Agreement conversion examples:
- ◆ GRAM, Condor
 - ◆ Workflow, economic scheduling
 - ◆ PBS, LSF, CSF)

Specifying Terms: Who and What?

In a service provisioning **domain**

(e.g. “computational jobs”)

- A **standard** specifies **domain** terms/concepts
- A **provider** specifies its support for
 - ◆ *some or all* of the domain standard terms
 - ◆ a given term, *specifically*
 - Within *behavioral* constraints
 - Within *negotiability* constraints
 - With extra fields/sub-terms
 - Arbitrary term *properties*: e.g. *optional* or *required*
- A **client** *discovers compatible providers*

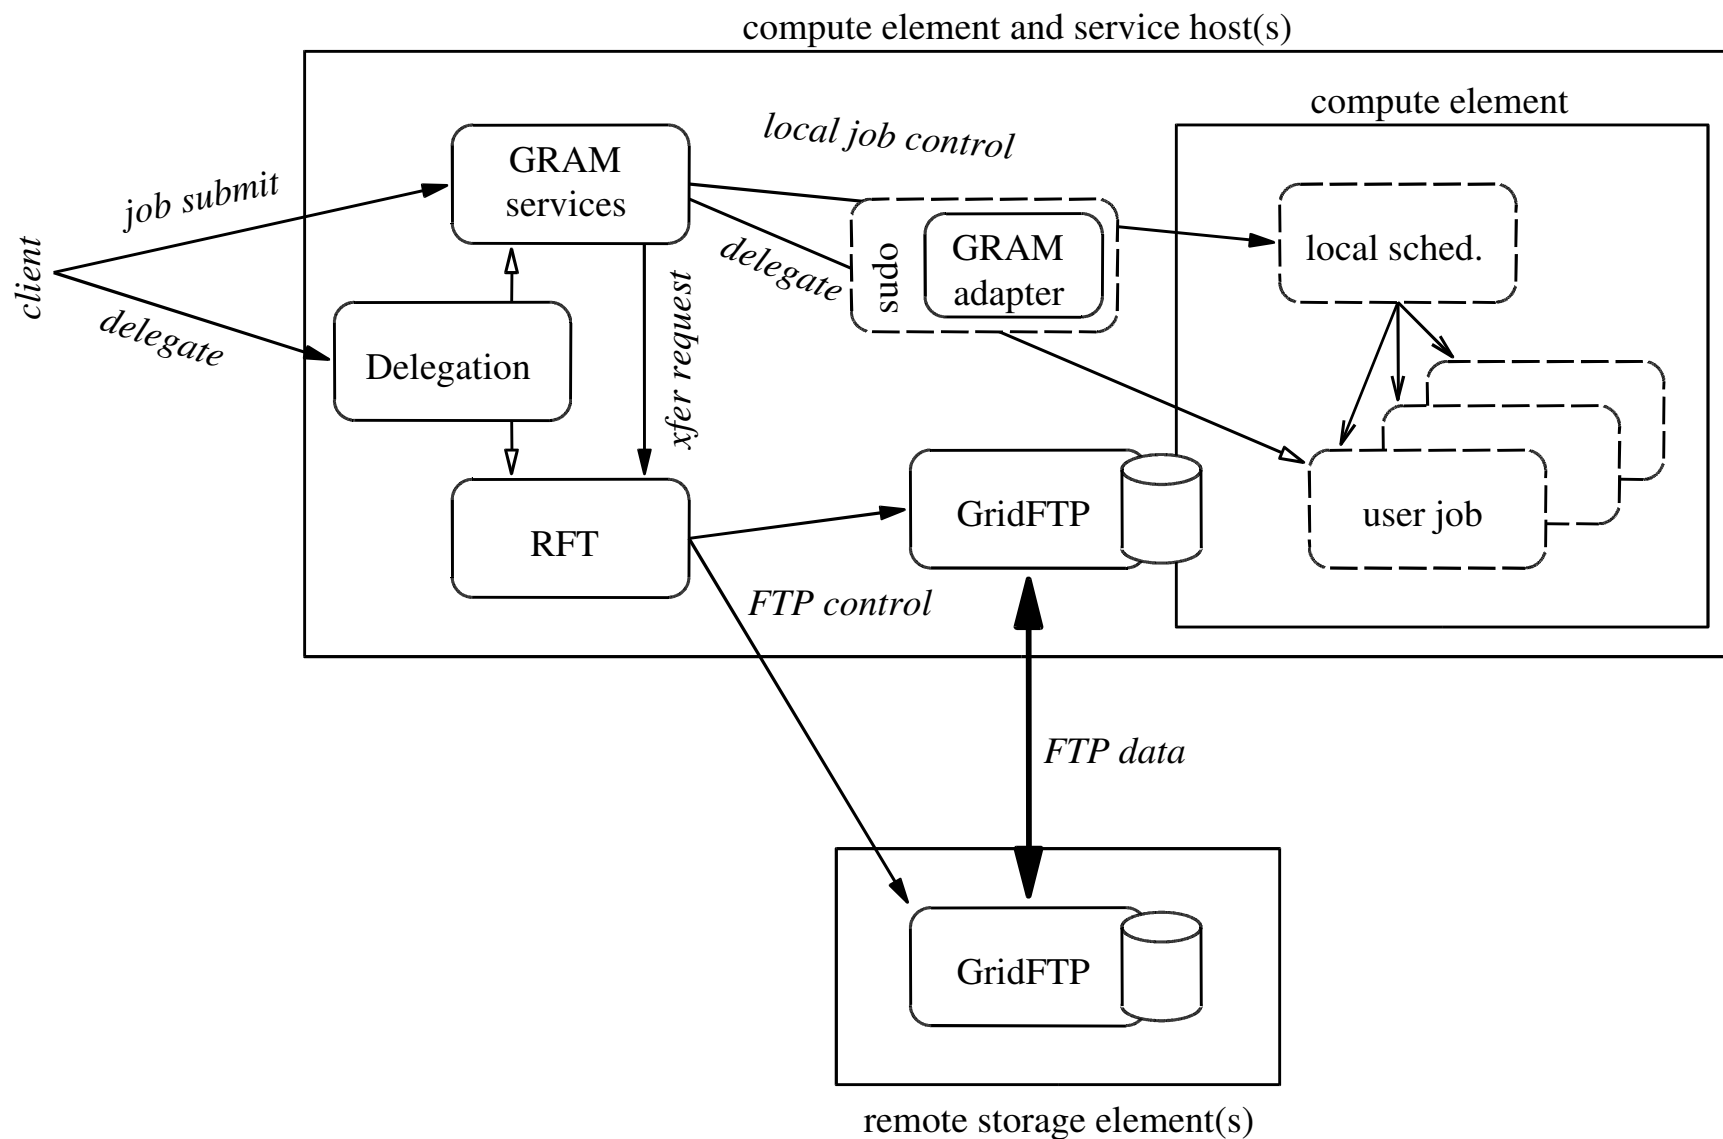


Possible GRAM Avenues

- Choose a job term language (e.g. JSDL?)
 - ◆ Don't rush for yet another job dialect
 - ◆ Define a profile for terms+mark-up?
- Provide an AgreementFactory impl.
- Provide an Agreement impl.
- Tie into existing ManagedJob resource impl.
- Support both WS-GRAM, Job Agreements
 - ◆ Ease migration

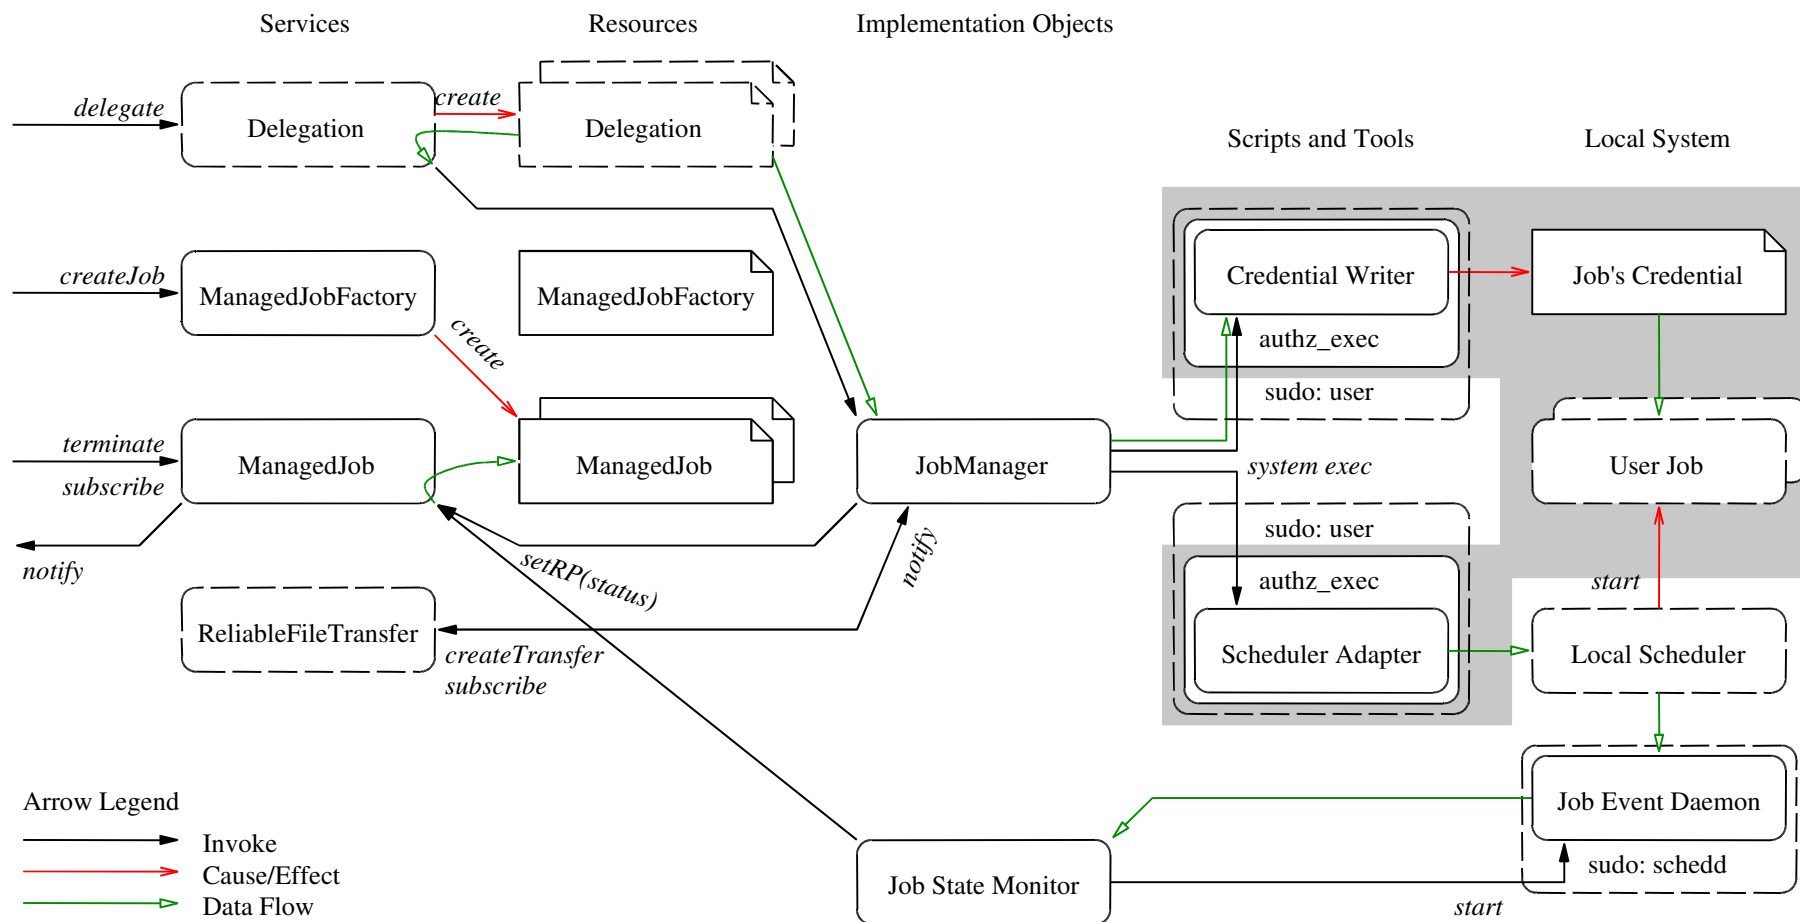


WS-GRAM Approach







WS-GRAM Software Map



WS-GRAM Base Protocol



WS-GRAM Full Protocol





Positioning for Future

- Better modularity of job manager functions
 - ◆ Track improvements in composed services
 - ◆ Experimentation w/ other job protocols
 - Share sudo/perl callouts for job control
 - Provide different “views” of job execution system
 - ◆ Reuse functions for next job standard
- Higher-level WSRF programming model
 - ◆ Return of co-allocation for MPICH
 - ◆ Unconventional job-like services next?
 - ◆ Advance reservation or co-scheduling next?