

Advance Reservation and Co-Allocation Protocol

Dean Kuo

School of Computer Science

The e-Science North West Centre

The University of Manchester

dkuo@cs.man.ac.uk

Assumptions


- Message latency is unbounded
 - Reliable messaging will only guarantee that a message will eventually be delivered in FIFO order
 - A service can not tell if another service has failed or it is a slow network or processor
- Consumer and providers can belong in different administrative domains
 - A consumer must not be able to “lock” a slot
 - Providers must have full control **AT ALL TIMES**
 - Locking introduces the risk of denial of service attacks

Cancellation

- Advance reservation needs to support cancellation
 - Providers must be able to cancel a reservation at anytime
 - Required for unscheduled downtime
 - It could cancel straight after it has agreed to a reservation request or seconds before the scheduled start time
 - Cancel if the consumer does not confirm before timeout period has expired
 - Consumer may also cancel a reservation at anytime
 - Consumer may be required to pay a cancellation fee if the reservation has been confirmed
 - Two-phase commit and Paxos commit is an agreement protocol not intended to support cancellation

Aim

- Specification of an *advance reservation protocol*
- Specify a *co-allocation protocol* using the *advance reservation protocol*
 - The following pairwise interactions are use the advance reservation protocol
 - User and co-scheduler
 - Each resource and co-scheduler
 - Support for nested configuration
 - Resource can not distinguish between a reservation directly from a user or a co-scheduler



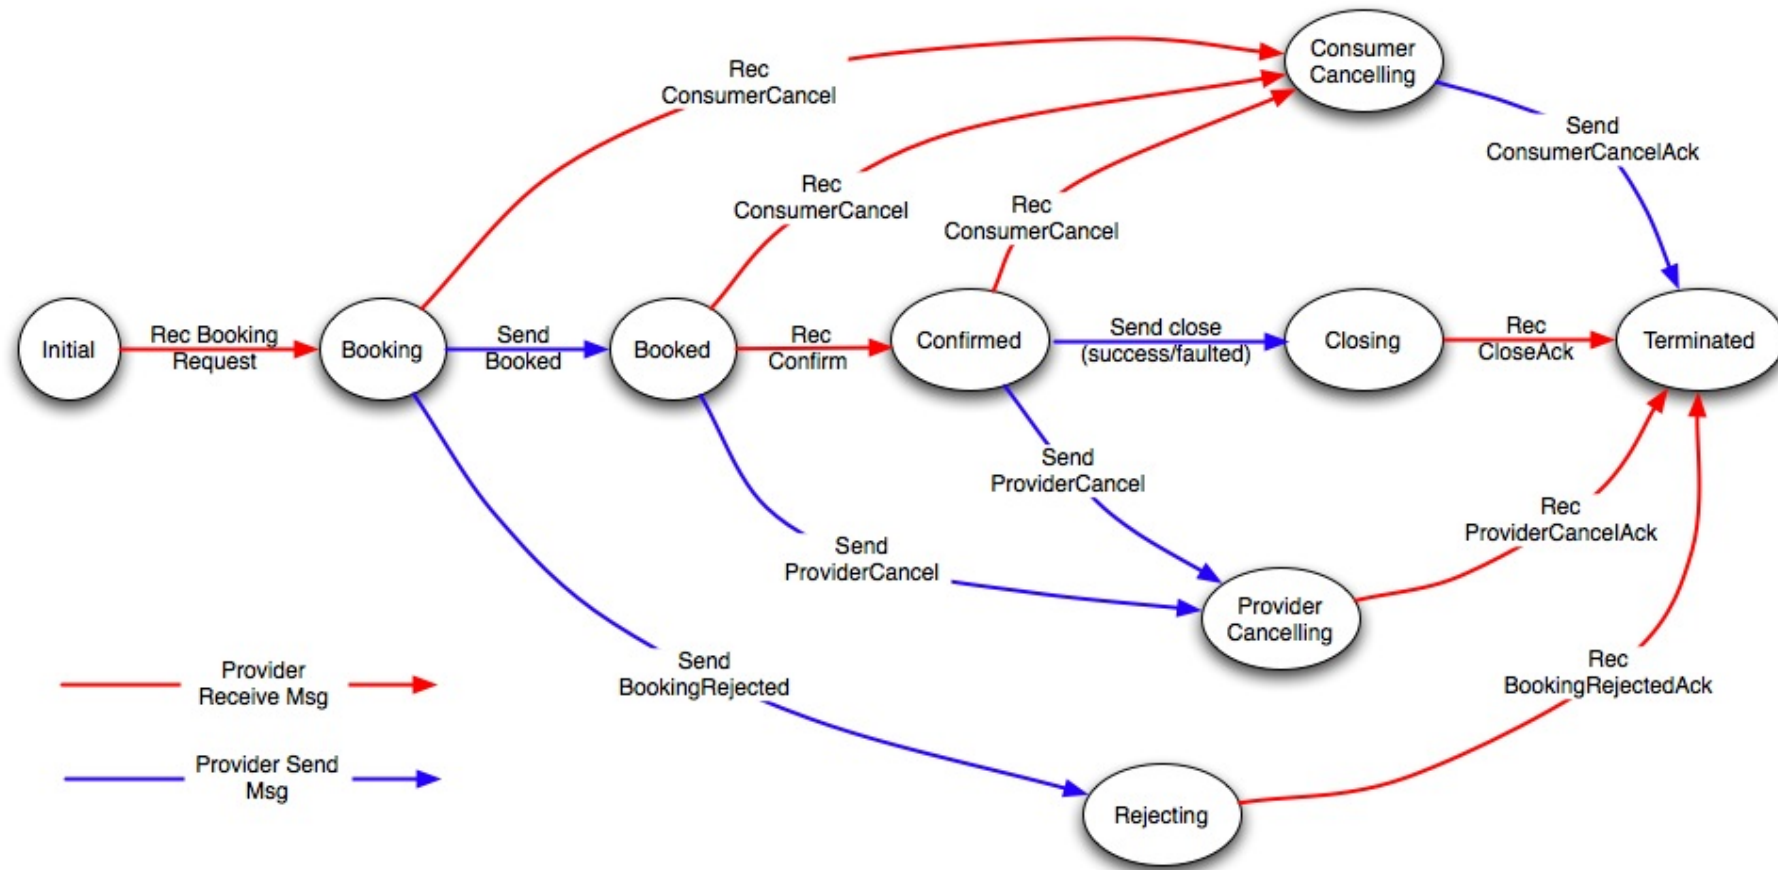
Unit of Work

- Unit of work (conversation) begins when user sends its first message
 - Terminates when the job execution completes
- The co-allocation protocol does not provide atomicity
 - It is *impossible* to support **ALL** or **NOTHING** property

Advance Reservation Protocol

- 11 Messages between consumer and provider
 - Completed execution
 - BookingReq, Booked, Confirm, Close (successful/failed), CloseAck
 - Booked message include details of pricing and cancellation policy
 - Consumer Cancellation
 - ConsumerCancel, ConsumerCancelAck
 - Provider Cancellation
 - ProviderCancel, ProviderCancelAck
 - Booking rejected
 - Reject, RejectAck

Provider Protocol



Agreed Outcomes


- Booking request rejected
- Scheduled job executed
 - Successful or faulted
- Consumer cancelled
 - Consumer cancelled before confirmation then there is no fee
 - Consumer cancelled after confirmation then fees specified in the cancellation fee applies
- Provider cancelled
- Define charging model in terms of on final agreed outcome

Race Situations

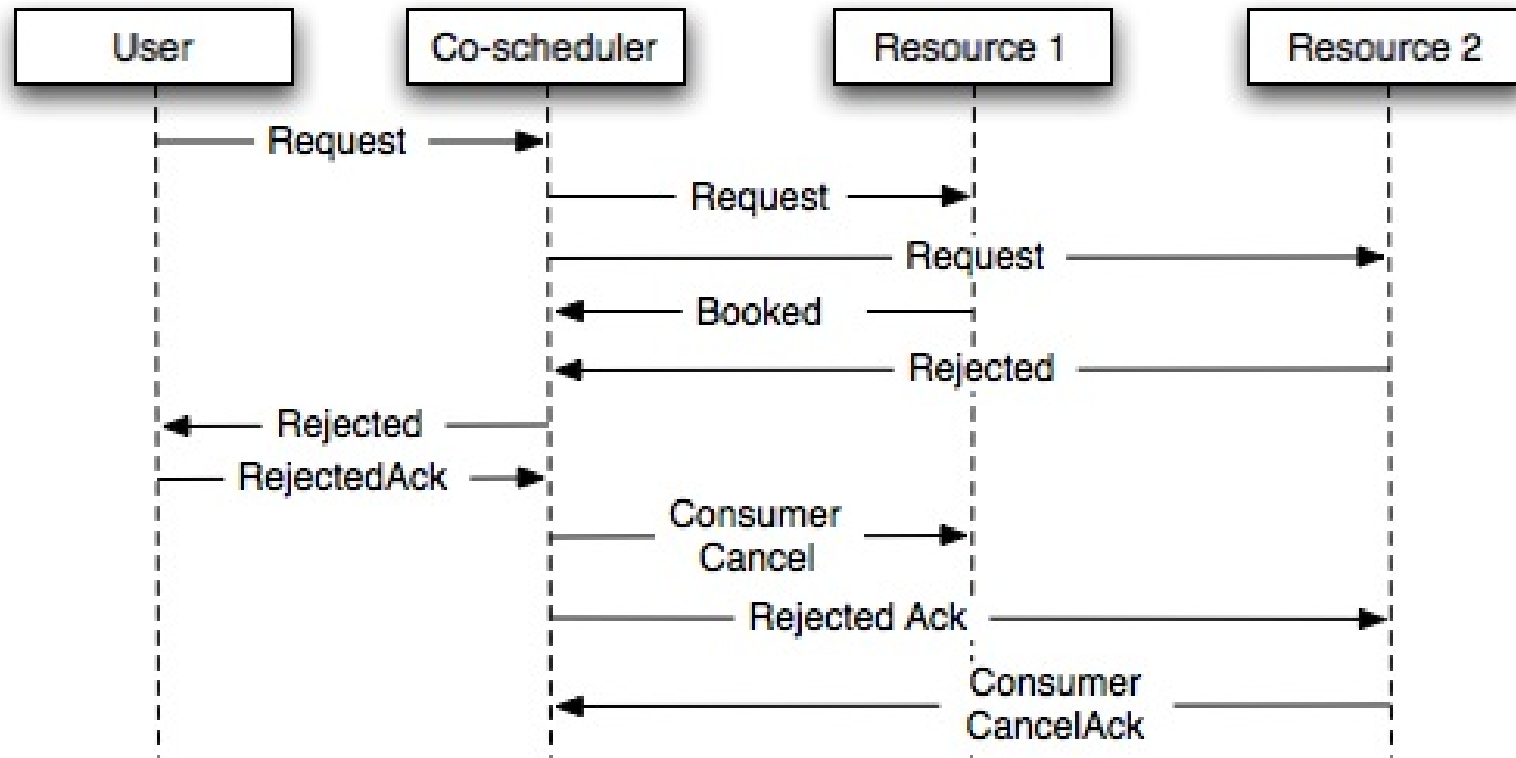
- What if consumer and provider cancel “simultaneously”?
 - Agreed outcome is “consumer cancelled”
- What if consumer cancels but when the provider receives the message, the job has already completed
 - Agreed outcome is “job completed”
- All possible race situations dealt with in the protocol

Co-Allocation

- Reservation of multiple resources
- Run the advance reservation protocol between
 - User and co-scheduler
 - Co-scheduler and Res 1
 - Co-scheduler and co-sch 2
 - Co-scheduler and Res 3
 - ...
- Main benefit
 - Simplicity and nesting



Booking Rejected



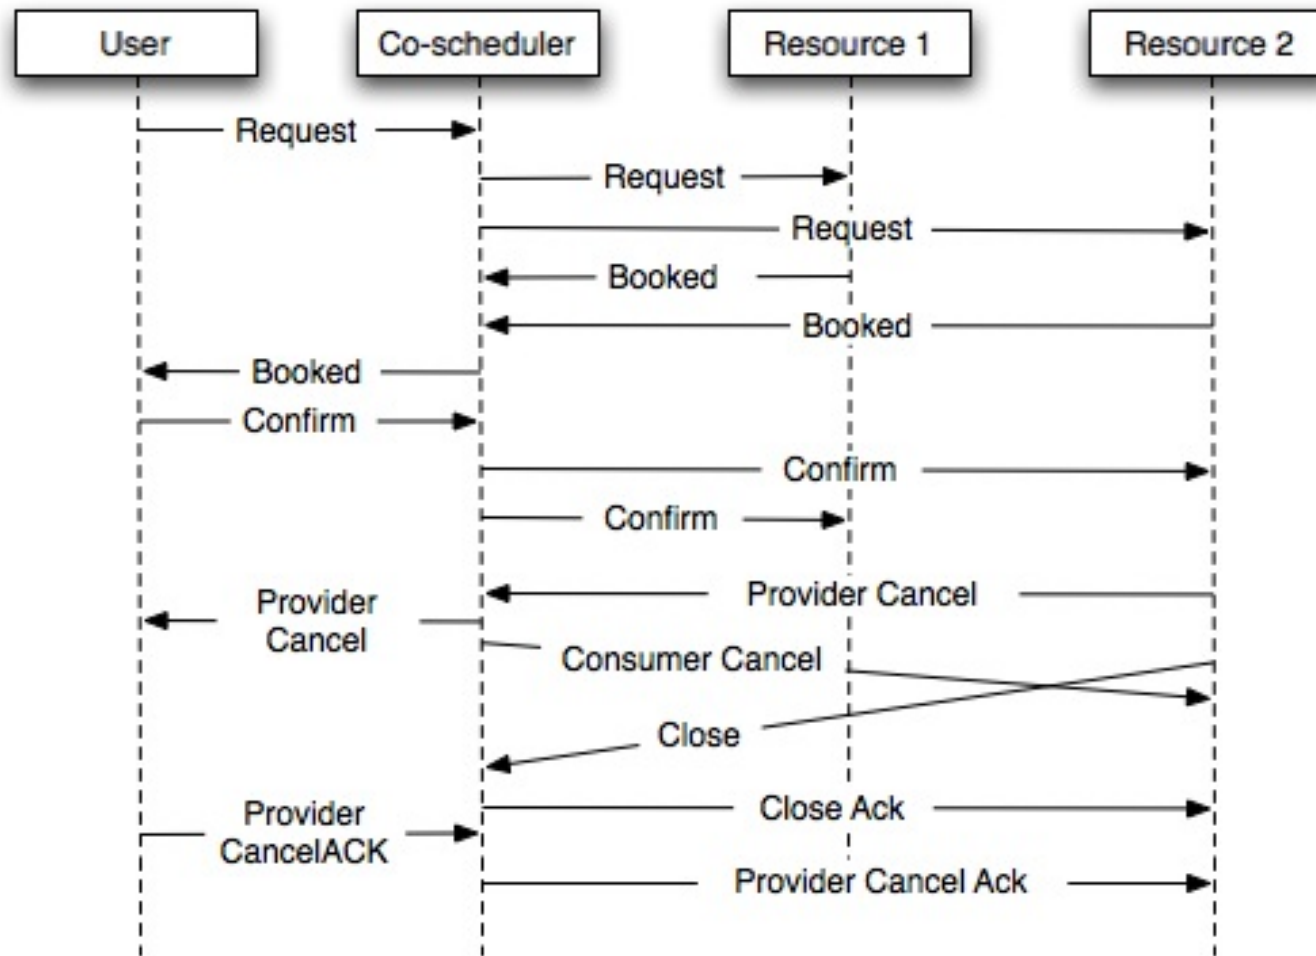
Co-Allocation Protocol - Completed

QuickTime™ and a
TIFF (LZW) decompressor
are needed to see this picture.

Co-Scheduler Cancels I

QuickTime™ and a
TIFF (LZW) decompressor
are needed to see this picture.

Co-Scheduler Cancels



Co-Allocation Protocol I

- Two phase to confirm a reservation
 - Cancellation can be initiated at anytime by any party
 - If consumer does not confirm a reservation before a timeout period then the provider cancel
 - Prevents denial of service attacks
- Atomicity is ***not*** supported
 - It is possible reservations with some resources terminated in the *completed* state while others terminate in the *provider cancelled or consumer cancelled* state
 - Co-allocation is not a consensus protocol

Co-Allocation Protocol II

- Other co-allocation strategies can also be supported besides the two phase protocol
 - Co-scheduler reserve each resource sequentially
 - Greater chance where co-scheduler needs to cancel a confirmed reservation with a resource

Correctness of Protocol

- Advance reservation protocol and co-allocation protocol have been modeled and verified using the SPIN model checker
 - Co-allocation protocol verification
 - Model contains only two resources

Charging Framework I

- Defined in terms of the agreed outcome between the user and the co-scheduler
 - Rejected
 - No fee
 - Successful completion
 - Full price as specified in the booked message
 - Faulted execution
 - No fee or whatever is defined in the policy in the booked message

Charging Framework II

- Consumer cancel
 - No fee if consumer cancels and has not sent a confirmation message
 - Fee as defined by the cancellation policy specified in the booked message
 - Co-scheduler may incur a financial loss in the co-allocation protocol
 - Less chance of incurring a loss for the two phase strategy
 - Co-scheduler can charge a premium to on all reservations to cover losses
- Provider Cancel
 - No fee or ?

Last Minute Reservations

- The need for greater fault tolerance
 - Need to know if a reservation is confirmed or not once a provider sends a “booked” message
- Can we layer or merge Paxos commit with the advance reservation and co-allocation protocol?
 - Fault tolerance + support for cancellation

Concluding Remarks

- Defined an advance reservation and co-allocation protocol based on a single protocol
 - Supports cancellation
 - A *unit of work* starts when user (consumer) initiates a reservation process and terminates when the job execution terminates
 - Includes a simple framework to support charging
 - Transaction atomicity is not possible as the protocol **MUST** be non-blocking and any party may cancel a reservation
 - Co-scheduler may incur a financial loss
 - Protocol provides a framework for explicitly defining charging models