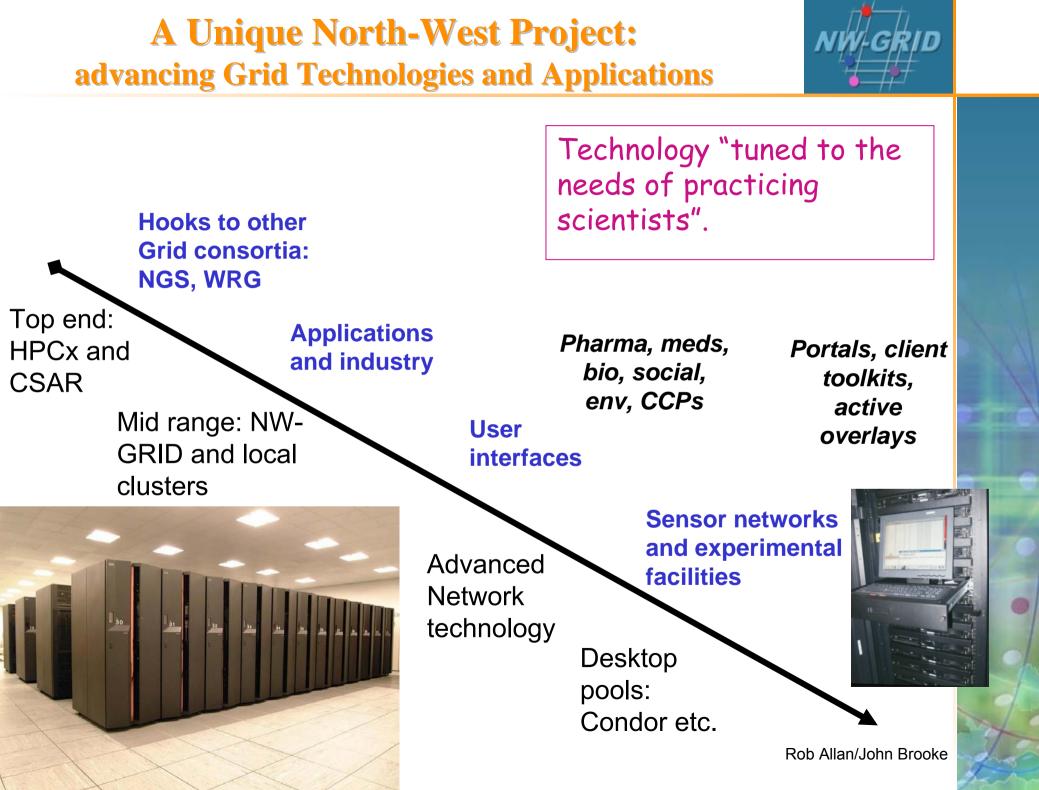


North West Grid Overview

John Brooke¹, R.J. Allan² ¹ The University of Manchester ² CCLRC Daresbury Laboratory

A world-class Grid infrastructure for the North West.

OGF20 Manchestert 7th May 2007 Rob Allany / John Brooke



The NW-GRID Project

- Funded by NWDA
- A Grid resource for the North West
- The project will be structured around three major testbed milestones:
 - Testbed 1 (2006): High-performance and parametric applications
 - Testbed 2 (2007): Safe, bookable and accountable Grid
 - Testbed 3 (2008): Collaborative real-time Grid
- Academic users
- Commercial users
- Partners

Statement

- The NW-GRID project, a collaboration between CCLRC Daresbury Laboratory and the Universities at Lancaster, Liverpool and Manchester, will establish a computational Grid comprising large-scale commodity computing systems coupled by a high-speed network. It will establish, for the region, a world-class activity in the deployment and exploitation of Grid middleware technologies (the software that glues together the various data and computing resources) and demonstrate the capabilities of the Grid in leading edge computational science and engineering applications.
- The Grid was a key component of the proposed NW science strategy and the project resonates strongly with the key elements of the NWDA's regional strategy in particular in working with targeted emerging sectors in the environment, biotechnology and pharmaceutical and complex materials areas, establishing the North-West as a global player in Grid technologies and in embedding ecompetencies across the region's business, academic and industrial base.

Key Applications and Research

NWGRID

- chemistry
- life sciences
- nano-technology
- human anatomy
- population modelling and epidemiology
- earth science
- oceanographic modelling
- environmental monitoring
- disaster response

- astronomy
- bio-informatics
- materials modelling
- social science and statistics
- management science
- complex fluids
- protein crystallography
- next generation middleware
- user interfaces
- enhancing production grids

Web site: http://www.nw-grid.ac.uk

NW-GRID Core



Four compute clusters funded by NWDA are accessible using Grid middleware and connected using a dedicated fibre network:

- Daresbury
 - dl1.nw-grid.ac.uk 96 nodes (some locally funded)
- Lancaster
 - lancs1.nw-grid.ac.uk 192 nodes (96 dedicated to local users)
- Liverpool
 - lv1.nw-grid.ac.uk 44 nodes
- Manchester
 - man2.nw-grid.ac.uk 27 nodes









OGF20 Manchestert 7th May 2007 Rob Allan/John Brooke

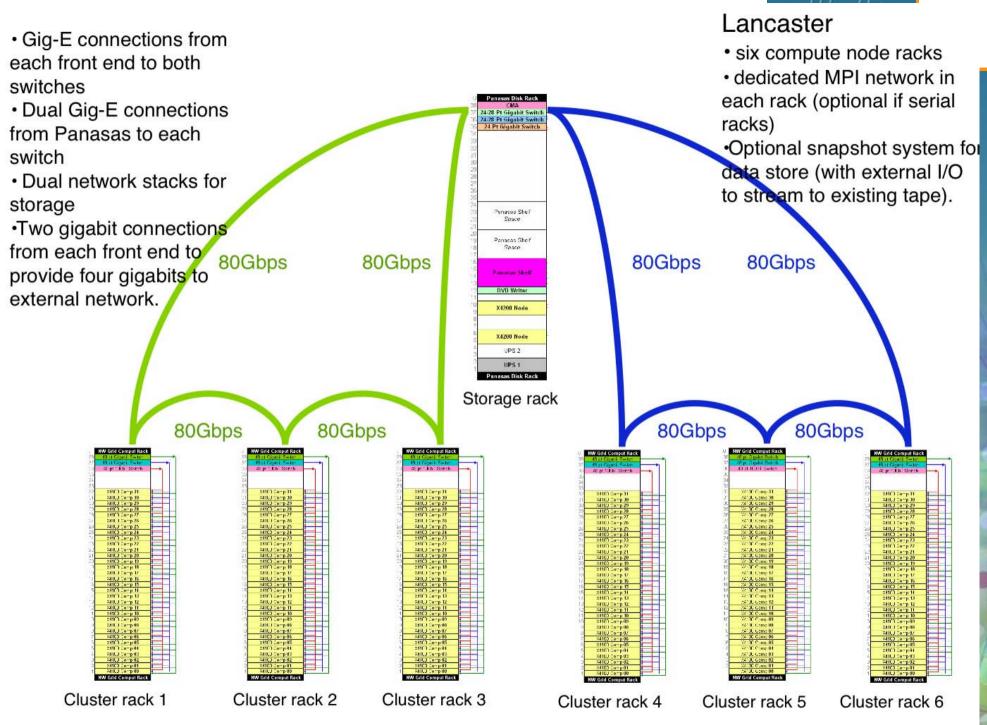
Anatomy of a Cluster

- Compute clusters from Streamline Computing
- 2x Sun x4200 head nodes
 - dual-processor single-core AMD Opteron 2.6GHz (64-bit)
 - 16GB memory
- Panasas file store
 - around 3-10TB per cluster (Manchester client only)
- Sun x4100 nodes
 - Dual-processor dual-core AMD Opteron 2.4GHz (64-bit)
 - 2-4GB memory pre core
 - 2x 73GB disks per node
 - Gbit/s ethernet between nodes
- Sun Grid Engine and Streamline cluster management
- NGS compatible Grid Middleware Globus 2.4

OGF20 Manchestert 7th May 2007

Connecting the Nodes

- within a "module" and between modules
- dual Gbit-E network
 - inter-node communications via stacked Nortel Gbit switch
 - command, control and file access via 2nd stacked Nortel Gbit switch
- Ethernet network
 - diagnostics via standard ethernet switch
- Link to Panasas if installed



OGł 7th May 2007

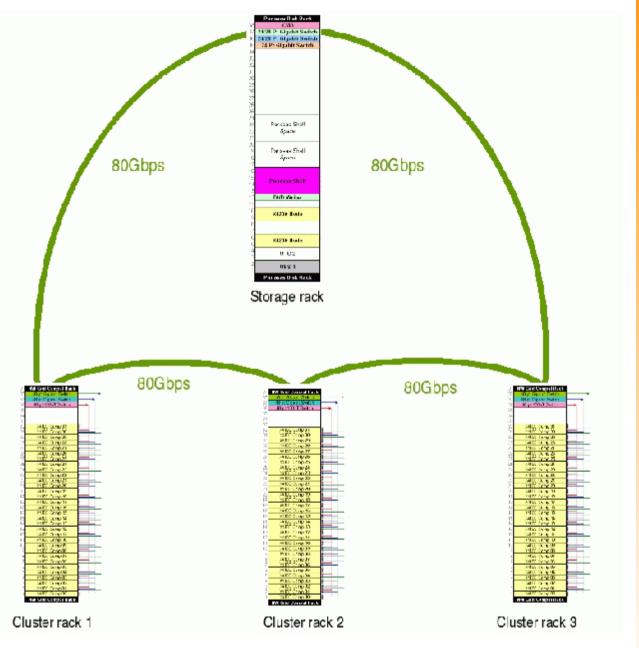
DL Cluster Detail

DL Cluster has 4 racks.

32 nodes per compute rack plus 3 switches.

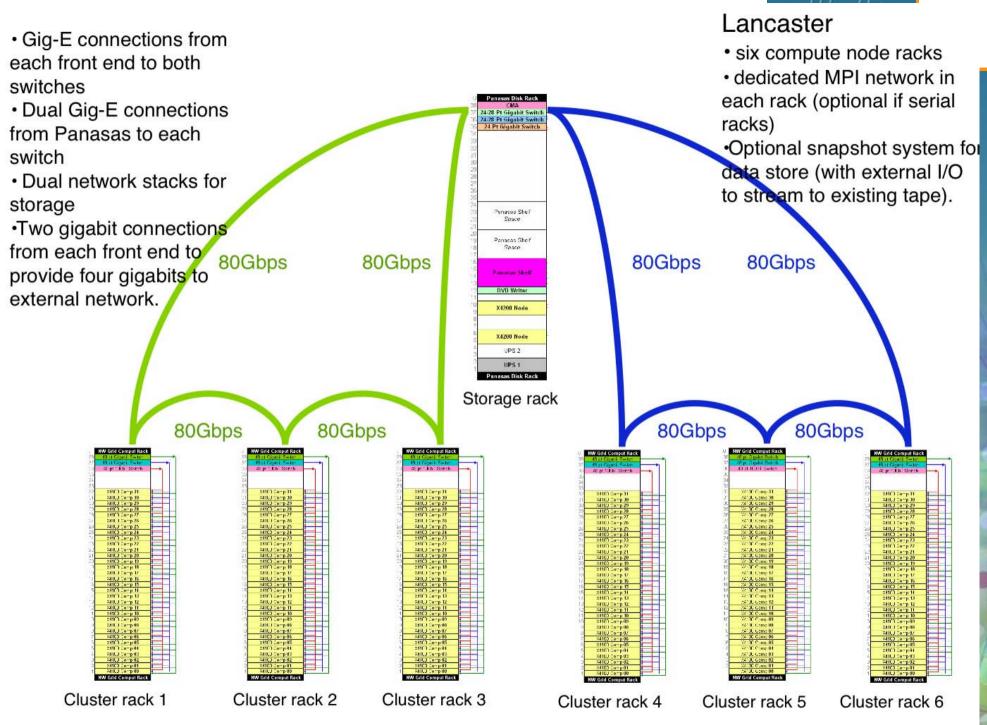
Gbit-E (80Gb/s) links between compute racks and into head node shown in green.

2x head nodes, Panasas storage and external connection in central rack.



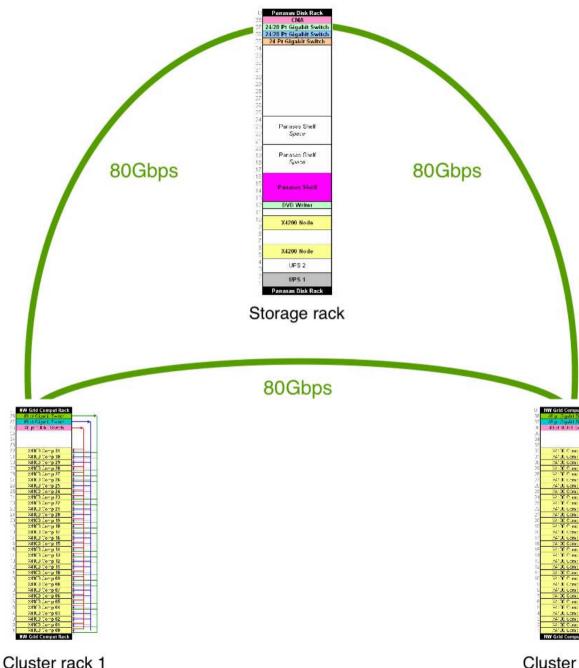
OGF20 Manchestert 7th May 2007

Rob Allan/John Brooke



OGł 7th May 2007





Liverpool

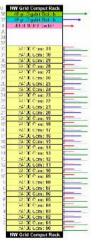
• two compute node racks One full , one part full.

dedicated MPI network per rack

 high-bandwidth storage network across all racks

 quad Gig-E connection from Panasas to storage network

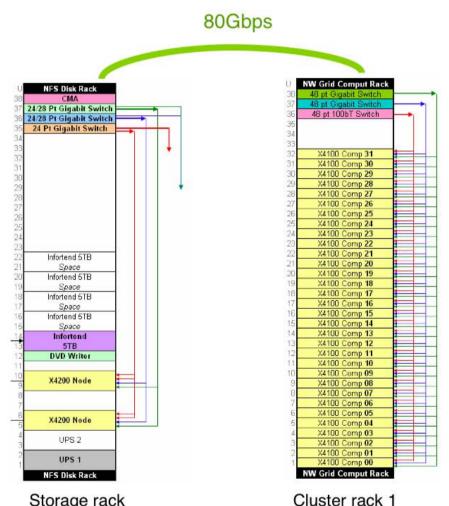
Gig-E connection from front end systems to storage network
two gigabit connections from each front end to provide four gigabits to external network.



Cluster rack 2

OGI Cluster rac 7th May 2007





Storage rack

Manchester

 one part full compute node rack with dedicated MPI network

 high-bandwidth network to storage rack

 dual Gig-E connection from NFS server to storage network

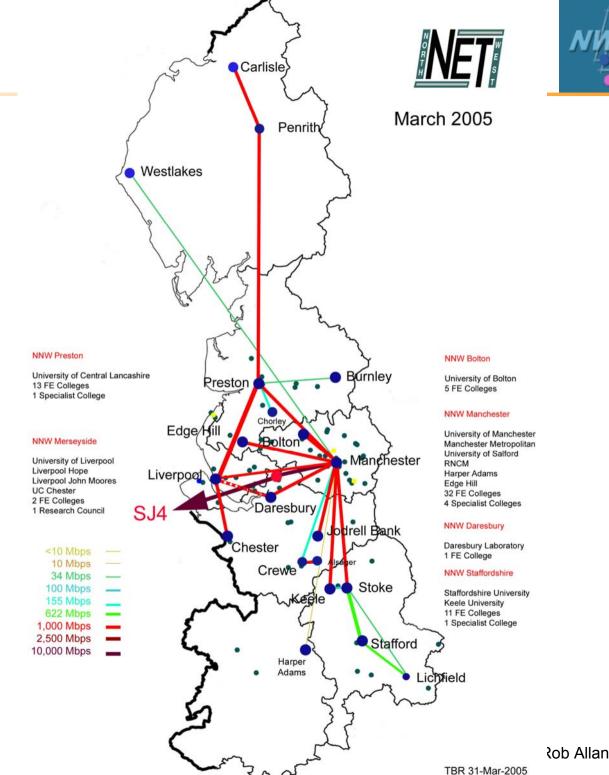
 Gig-E connection from front end systems to storage network Two gigabit connections from each front end to provide four gigabits to external network.



Connecting the Clusters

- Map of the NW on next slide
- HP Procurve/ 3Com switch is part of each control module
- External connection
 - Private Fibre network
 - SuperJANET
- Commercial access is possible by VPN and will use the fibre network





OGF20 Manchestert 7th May 2007

Rob Allan/John Brooke

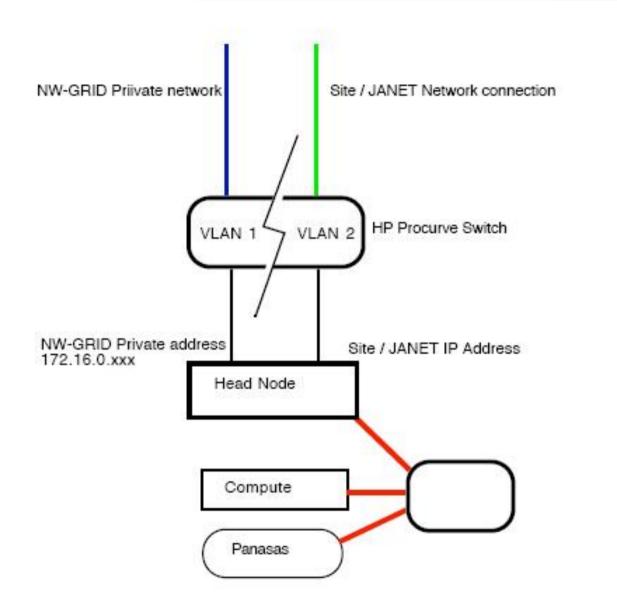
NW-GRID Fibre Network

- 1Gb/s between sites latency?
 - DL-Lancs trunk 10x 1Gb/s links
 - DL-Manchester single link
 - DL-Liverpool single link
- What can it enable us to do?
 - Scheduling, job migration, visualisation etc.
 - inter-site MPI, MPI-G
 - resource co-scheduling

•



External Connections



HARC

NWGRID

- Highly Available Resource Co-scheduler.
- John MacLaren's implementation of Paxos.
- Acceptors at different sites for fault tolerance.
- They issue prepare instructions to the RMS (e.g. SGE) then run the consensus algorithm before commit.
- Will be used in the GENIUS (Grid-Enabled Neurosurgical Imaging Using Simulation) project to link NGS, NW-GRID and TeraGrid

Other Resources

NWGRID

- Manchester see http://man4.nw-grid.ac.uk
 - man1.nw-grid.ac.uk Weyl cluster, It is a 44 node system plus login and master nodes. The address of the login node is weyl.mc.man.ac.uk and IP address is 130.88.200.240.
 Compute nodes 01 40 are dual processor AMD Opteron 248 2GHz, 78GB of disk, 2GB memory. Nodes 41 43 are fat nodes with the same spec as above except the memory is 8GB. The user, scratch and software disks are RAID, 2.5TB, NFS mounted.
 - man3.nw-grid.ac.uk Escher is a Silicon Graphics Prism (IA64) with 8 Itanium 2 processors, 4 FireGL 3 graphics pipelines running SUSE Linux Enterprise Server 9 with Silicon Graphics ProPack 4
 - man4.nw-grid.ac.uk Biruni is a Web server with Ganglia monitoring results.

Other Resources

- Daresbury
 - BlueGene L 1000 processors, for software development, available soon.
- Liverpool
 - POL Cluster 192x Xeon (32 bit), likely to become part of NW-GRID in April.
- Preston
 - 8-node SGI Altix 3700. Discussions underway with Dept. of Theoretical Astrophysics, U. Central Lancashire.



Manchester Campus Grid

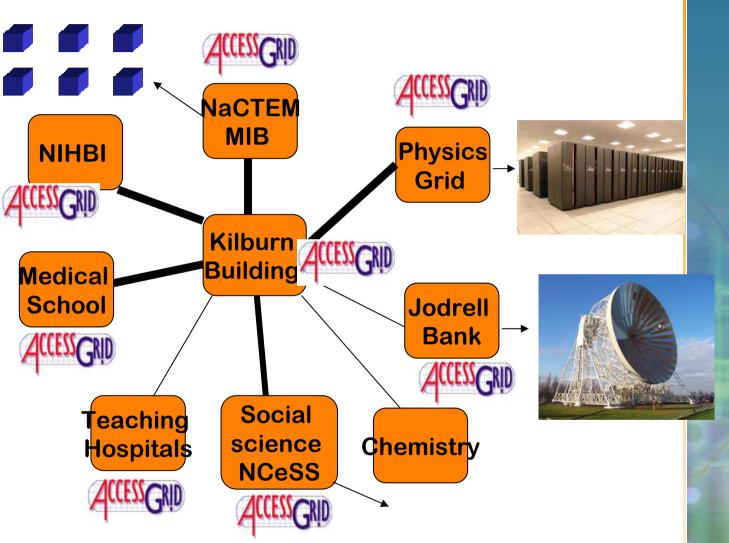
Other sites, e.g. Liverpool are also using Condor.

Pool

MIB Condor

Lancaster is an NGS partner.

HPCx at Daresbury is part of the NGS.





Bio Portal on NW-GRID

- User Interface to a number of applications
 - InterProScan
 - An integrated framework for a number of sequence searching applications
 - Exonerate
 - A generic tool for pairwise sequence comparison
- Job monitoring Interface
- Results can be requested/received via email



InterProScan Portal

on:		-	
	heck All C Clear	All	
ntScan 🔲 HMMPII	R 🗌 HMMPfan	n 🔲 HMMSmart	
		ily 🗌 SignalPHMM	1
PROTEIN 🛨 Sequ	ience in any format		
	ATIONS TO RUN C C intScan T HMMPII ileScan T ScanReg MPanther T Gene3D	ATIONS TO RUN C Check All C Clear intScan I HMMPIR HMMPfar ileScan I ScanRegExp SuperFam MPanther I Gene3D	ATIONS TO RUN C Check All C Clear All intScan F HMMPIR F HMMPfam F HMMSmart ileScan F ScanRegExp F SuperFamily F SignalPHMM

OGF20 Manchestert 7th May 2007 Rob Allan/John Brooke



Job Monitoring

2		Job Submi	ssion Portlet	
Refresh List Delete Jobs				
Job Id	Description	Resource	Status	Date Submitted
https://man2.nw-grid.ac.uk:64024/20740/1174395060/	test 2	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 12:49:08 PM
https://man2.nw-grid.ac.uk:64024/15205/1174388242/	test mail	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 10:55:30 AM
https://man2.nw-grid.ac.uk:64024/30689/1174386802/	test mail	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 10:31:29 AM
https://man2.nw-grid.ac.uk:64024/17181/1174385709/	test	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 10:13:17 AM
https://man2.nw-grid.ac.uk:64024/31605/1174384179/	test mail	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 9:47:47 AM
https://man2.nw-grid.ac.uk:64026/23425/1174383640/	test	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 9:38:47 AM
https://man2.nw-grid.ac.uk:64024/19369/1174383562/	test	man2.nw-grid.ac.uk	Job completed with message success	Mar 20, 2007 9:37:29 AM
https://man2.nw-grid.ac.uk:64028/7258/1173968912/	test9	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 2:26:38 PM
https://man2.nw-grid.ac.uk:64024/3345/1173968833/	test 8	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 2:25:20 PM
https://man2.nw-grid.ac.uk:64026/25488/1173968464/	test5	man2.nw-grid.ac.uk	Job failed with error code 131; the user proxy expired (job is still running);	Mar 15, 2007 2:19:10 PM
https://man2.nw-grid.ac.uk:64024/21594/1173968385/	test	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 2:17:51 PM
https://man2.nw-grid.ac.uk:64028/16511/1173968226/	test 3	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 2:15:13 PM
T https://man2.nw-grid.ac.uk:64026/12629/1173968144/	test 2	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 2:13:51 PM
https://man2.nw-grid.ac.uk:64024/8748/1173968066/	testcred	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 2:12:33 PM
T https://man2.nw-grid.ac.uk:64027/28417/1173960398/	testcred	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 12:04:44 PM
https://man2.nw-grid.ac.uk:64027/16713/1173959574/	test	man2.nw-grid.ac.uk	Job completed with message success	Mar 15, 2007 11:51:00 AM
F https://man2.nw-grid.ac.uk:64024/24686/1173885265/	test all	man2.nw-grid.ac.uk	Job completed with message success	Mar 14, 2007 3:12:31 PM

OGF20 Manchestert

7th May 2007

Rob Allan/John Brooke

Results Via Email

Subject: Results of job test

From: jobmailer@bio-portal.mc.manchester.ac.uk

Date: 15/03/07 11:54

To: Bharathi.Kattamuri@manchester.ac.uk

Results of job 'test' submitted to man2.nw-grid.ac.uk on 15 March 2007 11:51:00 o'clock GMT

Q9RHD9 D44DA88C544CB7C1 267 HMMPfam PF00575 S1 1 55 38.7 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 HMMPfam PF00575 S1 68 142 81.6 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 HMMPfam PF00575 S1 155 228 82.8 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 HMMSmart SM00316 S1 3 55 38.0 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 HMMSmart SM00316 S1 3 55 38.0 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 HMMSmart SM00316 S1 157 228 86.8 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 HMMPanther PTHR23270 PTHR23270 72 239 38.5 T 15-Mar-2007 Q9RHD9 D44DA88C544CB7C1 267 FPrintScan PR0681 RIBOSUMALS1 6 27 4.6E-18 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 HMMPfam PF00575 S1 155 228 82.8 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 3 55 38.0 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 70 142 81.0 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 70 142 81.0 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 157 228 86.8 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMPanther PTH23270 PTHR23270 72 239 38.5 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR0681 RIBOS/MLS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR0681 RIBOS/MLS1 85 104 4.6E-18 <t< td=""></t<>
Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 3 55 38.0 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 70 142 81.0 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 157 228 86.8 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMPanther PTHR23270 PTHR23270 72 239 38.5 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125
Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 70 142 81.0 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 157 228 86.8 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMPanther PTHR23270 PTHR23270 72 239 38.5 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR0681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR0681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR0681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR0681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 <th< td=""></th<>
Q9RHD9 D44DAE8C544CB7C1 267 HMMSmart SM00316 S1 157 228 86.8 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 HMMPanther PTHR23270 PTHR23270 72 239 38.5 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007 </td
Q9RHD9 D44DAE8C544CB7C1 267 HMMPanther PTHR23270 PTHR23270 72 239 38.5 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 6 27 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 85 104 4.6E-18 T 15-Mar-2007 Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 FPrintScan PR00681 RIBOSOMALS1 125 143 4.6E-18 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 Gene3D G3DSA:2.40.50.140 NA-bind_OB_sub 3 71 1.5E-15 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 Gene3D G3DSA:2.40.50.140 NA-bind_OB_sub 72 154 4.6E-22 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 Gene3D G3DSA:2.40.50.140 NA-bind_OB_sub 159 249 2.8E-22 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 superfamily SSF50249 Nucleic_acid_OB 3 60 1.4E-7 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 superfamily SSF50249 Nucleic_acid_OB 61 205 6.399999999999999995E-24 T
15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 ProfileScan PS50126 S1 1 55 14.869 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 ProfileScan PS50126 S1 72 142 20.809 T 15-Mar-2007
Q9RHD9 D44DAE8C544CB7C1 267 ProfileScan PS50126 S1 159 228 22.541 T 15-Mar-2007
RS16_ECOLI F94D07049A6D489D 82 Gene3D G3DSA:3.30.70.690 no description 1 79 2.9e-29 T 15-Mar-2
Y902_MYCTU CD84A335CCFFE6D7 446 Gene3D G3DSA:1.10.287.240 no description 220 292 3.3e-10 T 15-Mar-2
Y902_MYCTU CD84A335CCFFE6D7 446 Gene3D G3DSA:3.30.565.10 no description 300 443 1e-30 T 15-Mar-2
Y902_MYCTU CD84A335CCFFE6D7 446 HMMSmart SM00304 no description 170 222 3.1e-06 T 15-Mar-2007
Y902_MYCTU CD84A335CCFFE6D7 446 HMMSmart SM00388 no description 230 296 2.4e-12 T 15-Mar-2007
Y902_MYCTU CD84A335CCFFE6D7 446 HMMSmart SM00387 no description 338 446 5e-24 T 15-Mar-2007
RS16_ECOLI F94D07049A6D489D 82 ScanRegExp PS00732 RIBOSOMAL_S16 2 11 8e-5 T 15-Mar-2007
RS16_ECOLI F94D07049A6D489D 82 BlastProDom PD003791 RS16_BUCAI_P57474; 9 77 1e-19 T
15-Mar-2007

OGF20 Manchestert

7th May 2007

- NWGRID
- NW-GRID will have a dual role as a production and as an experimental Grid.
- The first phase was designed to be compatible with NGS
- The second phase will introduce experimental capability not yet available on NGS co-scheduling
- The third phase will integrate sensor nets using middleware based on overlays
- Eventually we would like to see fuller integration both with NGS and with campus Grids.

To find out more



- NW-GRID stand at UK e-Science Village here at OGF20
- Web site http://www.nw-grid.ac.uk
- ESNW representation on UK Engineering Task Force will enable software developed on NW-GRID to be evaluated for UK use.
- International collaborations, e.g. with TeraGrid