

Group: GGF SAGA-CORE-WG
Category: Recommendation
Title: SAGA API 1.0
Authors: Tom Goodale <goodale@cct.lsu.edu>
Shantenu Jha <s.jha@ucl.ac.uk>
Thilo Kielmann <kielmann@cs.vu.nl> ?????
Andre Merzky <andre@merzky.net>
John Shalf <jshalf@ncsa.uiuc.edu>
Christopher Smith <csmith@platform.com>
Date: May 01 2006
\$Revision: 1.5 \$

+-----+

TODO: check with new GFD.63

Intellectual Property Statement

=====

The GGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses. to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the GGF Secretariat.

The GGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the GGF Executive Director.

+-----+

Copyright Notice

=====

Copyright (C) Global Grid Forum (date). All Rights Reserved.
Distribution of this memo is unlimited.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be

prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the GGF or other organizations, except as needed for the purpose of developing Grid Recommendations in which case the procedures for copyrights defined in the GGF Document process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the GGF or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and THE GLOBAL GRID FORUM DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE."

+-----+

Contents	3
----------	---

Contents

Titlepage	1
Contents	3
Introduction	4
1 Introduction	4
1.1 API Scope	4
object	13
session	17
context	21
error	28
tasks	37
monitoring	48
attributes	66
namespaces	73
files	90
logicalfiles	100
job	108
stream	124
examples	137

1 Introduction

This document describes SAGA version 1.0, the *Simple API for Grid Applications*. SAGA has been defined as a high-level API that directly addresses the needs of application developers. The purpose of SAGA is twofold:

1. Provide a **simple** API that can be used with much less effort compared to the vanilla interfaces of existing grid middleware. A guiding principle for achieving this simplicity is the *80-20 rule*: serve 80 % of the use cases with 20 % of the effort. (compared to serving 100 % of all possible requirements)
2. Provide a standardized, common interfaces across various grid middleware systems and their versions.

1.1 API Scope

It had been decided to orient SAGA's functionality on actual application needs. For this purpose, the SAGA group has collected an as broad as possible set of use cases. The received use cases have been published as GFD.xx [?]. From these use cases, the requirements on a SAGA API have been derived. The requirements analysis has been published as GFD.yy [?].

In detail, the SAGA-RG received about 12-15 use cases (with some overlap) in the latter half of 2004, in response to an open call for use cases. A list of the Use Cases can be found on the Wiki and the mailing list archives:

Wiki: <http://wiki.cct.lsu.edu/saga/space/Use+Cases>

email archive: http://www.gridforum.org/mail_archive/saga-rg/threads.html

In addition, several use cases were received by direct communication with other GGF groups: nine use cases originate from the GridRPC-WG. They have varying degrees of detail and completeness, and are available as *gridrpc_use_cases.zip* at the Wiki.

We have also looked at the OGSA use cases, published as GFD.29 [?]. The most significant OGSA use cases in GFD.29 were determined to be the two scientific grid use cases; Severe Storm Modeling and National Fusion Collaboratory. The OGSA use case document describes each use case at a rather high level, and unfortunately does not include API information.

There are also several groups in GGF that have worked on high level interfaces and API's; relevant API's and frameworks that have helped guiding the scope and designing the SAGA API can be found at: <http://wiki.cct.lsu.edu/saga/space/Related+Grid+APIs>.

In addition to these groups within the GGF, there are several projects and groups that have worked on APIs and frameworks that are similar in spirit to a SAGA API. These in turn have helped motivate a SAGA API. See the latter half of <http://wiki.cct.lsu.edu/saga/space/Related+Grid+APIs> for a partial listing. The work of all these groups has been considered while defining the SAGA API, and close collaboration with these groups ensures good mapping consistency between the APIs. Simply put, SAGA covers all those efforts, simplifies them, and provides a consistent look and feel. The SAGA Requirements Document [?] discusses this topic in more detail.

Based upon the use cases received and the discussions of the SAGA design team, and the group as a whole, it was felt that a two phase approach should be taken for developing the API.

1. The packages for phase (or version) one would be for the most significant and mature areas like file transfer, streams, job submission. They can be understood as those which are required the most, if not all use cases, which are well understood, and for which prototypes exist (at least partially). Along with this first version of areas, the API also needs to define necessary auxiliary APIs, such as as session and error handling etc.

This document specifies exactly the API for these areas, being referred to as SAGA Version 1.0.

2. The packages for the second phase will be those which are to be found in some use cases only, like computational steering, or for an advert service.

Based upon the above reasoning, areas of functionality that were included in SAGA Version 1.0 are the following. They are also referred to as SAGA's *packages*.

- jobs
- files (and logical files)
- streams
- auxiliary API's for
 - session handle and security context
 - errors
 - asynchronous method calls (tasks)
 - attributes

Possible areas of functionality (packages) to be included in future SAGA Versions are:

- steering and monitoring
- possibly combining logical/physical files (read on logical files)
- advert service (see GAT [?]: persistent information service)
- GridRPC [?]
- GridCPR [?]
- Task dependencies (simple work flows and batches)
- extensions to existing classes

The versions as described above do not imply a hierarchy of API interfaces: all packages are motivated by their use cases, there is no split into 'lower level' and 'higher level' packages. The only exception is the group of auxiliary API's.

Dependencies between packages have been kept to a minimal level, to allow each package to be used independently of any other; this also may allow partially conformant API implementations (see below).

Style and Design Issues:

An Object Oriented (OO) approach was adopted, as it is easier to produce a procedural API from an OO API than the converse, and one of the goals of SAGA is to provide APIs which are as natural as possible in each implementation language. Advanced OO features such as polymorphism were avoided, both for simplicity and also to avoid complications when mapping to a procedural language.

The design team chose SIDL (Scientific Interface Definition Language; <http://www.llnl.gov/CASC/components/babel.html>) to specify the API. This provides a language neutral representation of the API, but with well-defined syntax and clear mapping to implementation languages.

The need for asynchronous calls was explicitly stated by the user community, as reasonable synchronous behaviour cannot always be expected in Grids. The team discussed the merits of

callback and polling mechanism and agreed that a non-blocking polling mechanism would be used initially. The SAGA task interface provides a mechanism to associate an asynchronous call with each blocking API call, and a polling mechanism to determine the state of the task; in the future a callback mechanism may augment this.

Object Oriented versus Procedural, Language issues:

The abstract SAGA API specification is object oriented, and specified in SIDL. Normative bindings for specific languages, both object oriented and procedural, will be defined in additional documents.

In several places, flags are denoted as bitfields (specifically, integer enums which can be combined by logical AND and OR), this is for notational convenience, and a language binding should use the most natural mechanism available.

The document contains several examples illustrating the use of the APIs, and these have naturally been shown using specific languages, such as C++. These examples should not be taken as normative, but merely as illustrative of the use of the API. When normative language bindings are available these examples may be revised to reflect these bindings. In order to give a feeling of the Look-and-Feel in other languages, Appendix A lists some of the examples in different languages; again, Appendix A is illustrative, not normative.

SIDL

As stated above, SIDL was chosen as specification language for this API specification. However, the document stays not true to the SIDL language in several places. This section gives a very short introduction to SIDL, and also lists a number of 'Notes to Implementors' on how to interpret this specification.

SIDL from the Babel project is similar to COM and CORBA IDL, but has an emphasis on scientific computing, with support of multi dimensional arrays etc. Although the SAGA spec does not use these features extensively, the multi language scope of Babel for mappings from SIDL to programming languages appealed to the authors of this specification.

The key SIDL concepts used in this document are

- package: specifies a name space (see note below)
- interface: a set of methods
- class: stateful object and set of methods
- method: a service a caller can invoke on a object
- type: constraint to value of method parameters

SIDL supports single inheritance of classes, and multiple inheritance of interfaces. Method definitions have signatures, which define what parameters are accepted on method invocation. These parameters can be

- in: passed by value
- out: passed by reference, assumed NULL
- inout: passed by reference

A implementation can destroy and re-allocate parameters which are passed by reference, no assumptions in that respect should be made to the implementation.

Notes to implementors:

SIDL has the notion of packages, which are equivalent to Java packages or C++ name spaces. Packages are used in this specification, for the purpose of cross referencing different API sections. The packages are not supposed to show up in the implementations class names or name spaces, apart from the top level 'saga' name space.

SIDL does also have the notion of 'versions', which are actually required on packages. We do not use versions in this specification, as the specification itself is versioned.

SIDL defines a string to be a char*. We feel however that strings have more powerful and native expressions in some languages (such as C++, Perl and Java), and use string for these native types. a char* is expressed in this document as array<byte,1> .

This specification defines all method calls as VOID (or rather does not specify a return type for any method call at all). Instead we define out parameters. Out parameters in SIDL are passed by reference. However, for this specification we expect language bindings to use one output parameter as return value to function calls where appropriate.

We are using output parameter for the following reason: most calls in the SAGA specification can be use asynchronously. As such, they return a task object, and have additional output

parameters as described in the spec. Only the synchronous versions do not return tasks, and can hence use normal output parameters.

+-----+

Compliant Implementations:

=====

A implementation of the SAGA API is "SAGA compliant" if it implements all objects and methods of the SAGA API specification. However, the implementation MAY throw a NOT_IMPLEMENTED exception for calls; however, that exception is to be used only in necessary cases, for example if a underlying Grid Middleware does not provide some capability, and if that capability can also not be emulated. The implementation MUST carefully document and motivate the use of the NOT_IMPLEMENTED exception.

The semantics for all methods is explicitly described for all methods, and must be followed by compliant implementations, unless explicitly stated otherwise.

Also, a compliant implementation MUST follow the SAGA API specification both syntactically and semantically. The consistency model supported by the implementation MUST be documented. The thread safety or unsafety of the implementation MUST be documented. The underlying middleware security model MUST be documented, as MUST its mapping to the SAGA::Context object.

A implementation is "partial SAGA compliant" if it follows the conformance guidelines above, but implements only some packages (some objects are not implemented).

All other implementations of the API are not "SAGA compliant".

Please note that the current specification does not as of yet define 'subsystems' sufficiently well -- that will be fixed in the final specification.

Notes to implementors:

Early versus late bindings:

An implementation may choose to use late binding to middleware. That means that the middleware binding might

change between subsequent SAGA calls. For example, a file open might be performed via the HTTP binding, but a subsequent read() on that file might be performed with GridFTP.

Late bindings has some advantages in terms of flexibility and error recovery. However, it implies a certain amount of object state to be kept on client side, which might have semantic consequences. For example, a file write might fail on HTTP for some reasons, but might succeed via GridFTP. The situation might be inversed for file reads. In order to allow alternating access via both protocols, the file pointer information (e.g. the file object state) MUST be hold on client side.

It is left to the later experience documents about the SAGA API implementations to discuss potential problems arising from early/late binding implementations.

+-----+

Security considerations:

=====

As the API is supposed to get implemented on very different types of Grid (and non-Grid) middleware, it does not specify a single security model, but rather provides hooks to interface to various security models - see the documentation of SAGA.Context for details. A SAGA implementation is considered secure if it fully supports/implements the security model of the middleware layer(s) it depends upon, and does not provide any (intentional and unintentional) means to bypass that security model, and does not weaken the security model policies.

+-----+

Relation to OGSA:

=====

Although we feel that widely off target, the SAGA API specification effort has often been compared to, and seen in rivalry to the OGSA standardization effort. That is NOT correct. Reasons are the following:

- OGSA strives to define interfaces on Service and middleare level,

SAGA aims on application level.

- OGSA strives to be complete, and to fully cover any potential Grid Service in its architectural frame.

SAGA is by definition incomplete (80:20 rule), and aims for coverage of the mostly used grid functionalities on application level, with NO ambition to be complete in any sense.

- OGSA is an Architecture (or a framework for an architecture).

SAGA is an API.

- OGSA cannot sensibly interface to SAGA.

SAGA implementations can interface to (a subset of) OGSA compliant services (and in fact usually will do so).

- The OGSA spec aims at middleware developers.

The SAGA spec aims at application developers.

For these and more reasons we think that SAGA and OGSA are complementary, but by NO means competitive. The only communality we are aware of is the broadness of both approaches: both OGSA and SAGA strive to cover more than one specific area of, well, middleware and application, respectively.

+-----+

The 'URL Problem':

=====

The end user might expect the SAGA API, as a high level and simple API, to handle protocol specific issues transparently. In particular, she might expect that SAGA gracefully and intelligently handles an URL such as `http://host.net/tmp/file` even if HTTP as protocol is, in fact, not available at `host.net`, but for example the FTP protocol is,

However, that innocent looking problem has farreaching consequences, and in fact is, to our best knowledge, unsolved. Consider the following server setup on `host.net`:

FTP Server: server root: /var/ftp/pub/
HTTP Server: server root: /var/http/htdocs/

The entity described by the two URLs

`http://host.net/tmp/file`
`ftp://host.net/tmp/file`

does hence refer to distinct files on host.net! Even worse:
it might be impossible to access the HTTP file space via the
FTP service, and vice versa.

Similar consideration hold for absolute file names, and for
file names relative to the users home dir: consider

`http://host.net/~{ }merzky/tmp/file`

That URL point on my linux box to

`file:///home/merzky/public/_html/tmp/file`

and not, as could have been expected, to

`file:///home/merzky/tmp/file`

Hence, a reliable translation of URLs between different
protocols (schemes) is only possible, if the exact server
setup of all affected protocol serving services is known.
That knowledge is often not available.

Further, even if a correct translation of protocols and hence
URLs succeeds, there is no guarantee that the referred file is
actually available via that protocol -- that again depends on
the service configuration.

SAGA 'solution' to the 'URL Problem':

- 1) A SAGA compliant implementation MAY be able to
transparently translate URLs, but is not required to do so.
Further, that behaviour CAN vary during the runtime of the
program.
- 2) The SAGA API specification allows the use of the
placeholder 'any' (as in `any://host.net/tmp/file`). An SAGA
compliant implementation MAY be able to choose a suitable
protocol.
- 3) Abstract name spaces, such as the name space used by
Replica Systems, or by Grid File Systems, efficiently and
transparently hides that problem from the end user. SAGA

encourages implementations to use such name spaces.

- 4) A URL which cannot be handled for the stated reasons MUST cause the exception `Incorrect_URL` to be thrown. Note that this holds only for those cases, where a given URL cannot be handled *\emph{as such}*, e.g. because the protocol is unsupported, any:// cannot be handled, or a necessary URL translation failed. Any other error related to the URL (e.g. file at service is not available) MUST be indicated by the error codes/exceptions as state at the method specifications in this document.

+-----+

Additional Notes

=====

- For files, flags are used to specify if a open is truncating, creating, and/or appending to an existing entity. For jobs, and in particular for file staging, the LSF scheme is used (e.g. "url >> local_file" for appending a remote file to a local one after staging). We are aware of that seeming inconsistency. However, we think that a forcefull unification of both schemes would be more awkward to use, and at the same time less useful.
- About consistency we had a lengthy discussion, with the agreement that the consistency model is to be defined and documented by the implementation. The API spec itself does not assume any specific consistency model, as we feel that (a) POSIX consistency is not achievable within reasonable effort/performance, (b) if the user assumes the worst (no consistency), he will still be able to make good use of the API, and (c) reality will be somewhere in the middle.

+-----+

```

+-----+
      #####
      #      # #####      # #####      #####      #####
      #      # #      #      # #      #      #      #
      #      # #####      # #####      #      #
      #      # #      #      # #      #      #      #
      #      # #      # #      # #      #      #      #
      #####      #####      #####      #####      #####      #
+-----+

```

Summary:

=====

The basic SAGA object interface provides methods which are essential for all SAGA objects. For now, it provides a unique ID which helps to maintain list of SAGA objects in the application, and inspection to allow to test for the objects type and the attached session.

```

+-----+

```

Specification:

=====

```

package saga.object
{
    enum object_type
    {
        Unknown          = -1,
        Exception         =  1,
        Session           =  2,
        Context           =  3,
        NSEntry           =  4,
        NSDirectory       =  5,
        File              =  6,
        Directory         =  7,
        LogicalFile       =  8,
        LogicalDirectory  =  9,
        JobServer         = 10,
        Job               = 11,
        JobDescription    = 12,
        Session           = 13,
        StreamServer      = 14,
        Stream            = 15,
    }
}

```

```

interface object
{
    get_id      (out int      id      );
    get_type    (out object_type type  );
    get_session (out session   session);
}
}

```

+-----+

Details:

=====

```

class object:

- get_id:
  Purpose: query the object ID
  Format:  get_id      (out int id);
  Inputs:  none
  Outputs: id          id of object

- get_type:
  Purpose: query the object type
  Format:  get_type    (out object_type type);
  Inputs:  none
  Outputs: type        type of object

- get_session:
  Purpose: query the objects session
  Format:  get_session  (out session s);
  Inputs:  none
  Outputs: s           session of object
  Notes   - if no specific session was attached to the
            object on creation time, the default SAGA
            session is returned.
            - some objects don't have sessions attached,
              such as job_description. The session
              returned is NULL then (that includes the
              session object itself).

```

+-----+

Examples:

=====

```

// c++ example

// have 2 objects, streams and files, and do:
// - read 100 bytes

```

```
// - skip 100 bytes
// - read 100 bytes

char out1[100];
char out2[100];
char out_last[100];

// create map
std::map <saga::task, saga::object> tmap;

// create objects, and map
saga::file f (url[1]);
saga::stream s (url[2]);

s.connect ();

// create tasks for reading first 100 bytes, and map
saga::task t1 = f.read <saga::task> (100, buf1, &out1); tmap[t1] = f;
saga::task t2 = s.read <saga::task> (100, buf2, &out2); tmap[t2] = s;

// put in same container...
saga::task_container tc;

tc.add (t1);
tc.add (t2);

// ... and wait who gets done first
while ( saga::task t = tc.wait () )
{
    // depending on type, skip 100 byte then create a
    // new task for the next read, and re-add to the tc

    if ( tmap[t].get_type () == saga::object::File )
    {
        saga::file (tmap[t]).seek (100, SEEK_SET);
        tc.add (saga::file (tmap[t]).read <saga::task> (100, &out_last))
    }
    else
    if ( tmap[t].get_type () == saga::object::Stream )
    {
        saga::stream (tmap[t]).read (100, NULL); // ignore result
        tc.add (saga::stream (tmap[t]).read <saga::task> (100, &out_last))
    }
    else
    {
        throw saga::exception ("Something went terribly wrong");
    }

    // tc is filled again, we run forever, read/seeking from
    // whoever we find after the wait.
}
```

+-----+

Notes:

=====

Really useful will the base saga object once we add
serialization - it will then be possible to serialize objects
(e.g. tasks!), and to resume operations on them at a later
point.

+-----+


```

-----+
#####
#   # #####   ####   ####   #   ####   #   #
#   #   #   #   #   #   #   #   #   #   #
#####   #####   #####   #####   #   #   #   #   #
#   #   #   #   #   #   #   #   #   #   #   #
#   #   #   #   #   #   #   #   #   #   #   #
#####   #####   #####   #####   #   ####   #   #
-----+

```

Summary:

=====

The session object provides the functionality of a session handle, which isolates independent sets of SAGA objects from each other. Sessions also support the management of security information (see `saga::context`).

+-----+

Specification:

=====

```

package saga.session
{
  class session
  {
    CONSTRUCTOR      (out session      obj);
    DESTRUCTOR       (in  session      obj);

    add_context      (in  context      context);
    remove_context   (in  context      context);
    list_contexts    (out array<context,1> contexts);
  }
}

```

+-----+

Details:

=====

```

class session:
-----

```

Multiple sessions can co-exist. A single session can be shared between threads.

A context (which encapsulates security information in SAGA) can be attached to a session. A SAGA implementation MAY allow to attach more than one context to a single session. However, a single context instance can get attached only once to a specific session instance.

A session can be used as first parameter to all SAGA object instantiation calls. SAGA objects created from other SAGA objects inherit its session. Only some objects do not need a session handle on creation time, and can hence be shared between sessions. That includes:

- context
- job_description
- metric
- utility classes

If the session handle is omitted as first parameter, a default session handle is used, with default security context(s) attached. Example:

```
// create a file object in a specific session:
saga::file f (session, url);

// create a file object in the default session:
saga::file f (url);
```

Any SAGA operation CAN throw a `IncorrectSession` exception if involves two different session handles.

- CONSTRUCTOR
 - Purpose: create the object
 - Format: CONSTRUCTOR (out session obj)
 - Inputs: none
 - Outputs: obj: the newly created object
- DESTRUCTOR
 - Purpose: destroy the object
 - Format: DESTRUCTOR (in session obj)
 - Inputs: obj: the object to destroy
 - Outputs: none
- add_context
 - Purpose: attach a security context to a session handle
 - Format: add_context (in context context);
 - Inputs: context Security context to add
 - Outputs: none
 - Throws: BadParameter
 - AlreadyExists

```

- remove_context
  Purpose: detach a security context from a session handle
  Format:  add_context          (in context context);
  Inputs:  context              Security context to remove
  Outputs: none
  Throws:  BadParameter
           DoesNotExist

- list_contexts
  Purpose: retrieve all contexts attached to a session
  Format:  list_contexts        (out array<context>
                                contexts);

  Inputs:  none
  Outputs: contexts             list of contexts of this
                                session

  Note:    - a empty list is returned if no context is
            attached, yet.
            - contexts may get added to a session by default.
            hence the returned list MAY be non empty even
            if no add_context was ever called before.

```

+-----+

Examples:

=====

```

// c++ example
saga::session s;
saga::context c (saga::context::X509);

s.add_context  (c);

saga::directory d ("gsiftp://remote.net/tmp/", s);
saga::file      f = dir.open ("data.txt");

// file has same session attached as dir

```

+-----+

Notes:

=====

Most libraries use session handles to distinguish scope (security, settings, lifetime) of objects etc.

GAT used a context object, which is a session handle with attached information (security context, preferences) and some methods (get 'self', init environment, current state, ...).

Proposal: re-use what we have in GAT: falls back to the well known paradigm of a session handle), plus give potential for a few more features (see above).

+-----+

```

+-----+
#####
#   #   #####   #   #   #####   #####   #   #   #####
#       #   #   ##   #   #   #       #   #   #
#       #   #   #   #   #   #   #   #####   ##   #
#       #   #   #   #   #   #   #   #       ##   #
#   #   #   #   #   #   ##   #   #       #   #   #
#####   #####   #   #   #   #####   #   #   #
+-----+

```

Summary:

=====

The context provides the functionality of a security information container. A context is created, and attached to a session handle. As such it is available to all objects instanciated in that session. Multiple contexts can co-exist on one handle. A single context can be shared between threads and sessions. SAGA Objects created from other SAGA Objects inherit its session and also its context(s).

A implementation CAN implement various types of Security contexts. or just one type. The type of context to be created is specified by a enum which is the only argument to the context constructor. The default type is unknown. Other methods than get_type MUST NOT be called on a context with type Unknown.

Every context has a specific set of attributes which can be set/get via the SAGA attribute interface. Exactly what attributes a context offers depends on its type. A context MUST issue an error if attributes not corresponding to its type are set or requested.

For incoming interactions (streams, monitoring, steering), read only contexts are used to inform the application about the requestor identity. To support that, a number of specific getter methods are available.

```

+-----+

```

Specification:

=====

```

package saga.context
{
    enum context_type
    {

```

```

    Unknown          = -1,
    X509              = 1, // Globus
    MyProxy           = 2, // X509 extended
    KeyStore          = 3  // Unicore
    SSH               = 4, // SSH
    Kerberos          = 5, // Kerberos
    UserPass          = 6, // Unix default
}

class context : implements-all saga::object
               implements-all saga::attribute
               implements-all saga::monitorable
{
    CONSTRUCTOR (in context_type type,
                 in session      session);
                 out context      context);
    DESTRUCTOR  (in context      context);

    get_type    (out context_type type);
}

```

+-----+

Details:

=====

class context:

- CONSTRUCTOR:

Purpose: create a security context

Format: CONSTRUCTOR (in context_type type,
in session session,
out context context);

Inputs: type type of context

Outputs: context the newly created context

Throws: BadParameter

Notes: - if NO session handle is defined for at context
creation, the resulting context is NOT bound
to any session, any must be added to a session
in order to get used.

- BadParameter is thrown if a context type is
not supported (NOT NotImplemented).

- DESTRUCTOR:

Purpose: destroy a security context

Format: DESTRUCTOR (in context context);

Inputs: context the context to destroy

Outputs: none

```

- get_type:
  Purpose: query the context type
  Format:  get_type           (out context_type type);
  Inputs:  none
  Outputs: type               type of context

```

+-----+

Examples:

=====

```

// c++ example
saga::session s;
saga::context c (saga::context::X509);

s.add_context (c);

saga::directory d ("gsiftp://remote.net/tmp/", s);
saga::file      f = dir.open ("data.txt");

// file has same session attached as dir

```

+-----+

Notes:

=====

- Following attributes MUST be supported by the corresponding context types:

Unknown:

No attributes supported

X509:

Set/Get:

```

X509_Proxy      (/tmp/x509...)
X509_CertDir    (/etc/grid-security/certificates/)

```

Get:

```

X509_RemoteDN
X509_RemoteHost
X509_RemotePort

```

MyProxy

Set/Get:

```

MyProxy_UserName (anonymous)
MyProxy_Password (anon)

```

Get:

```

MyProxy_RemoteUserName

```

```

MyProxy_RemoteHost
MyProxy_RemotePort

```

```

KeyStore
Set/Get:
    KeyStore_Location    ($HOME/.keystore)
    KeyStore_Password    (anon)
Get:
    KeyStore_RemoteUserName
    KeyStore_RemoteHost
    KeyStore_RemotePort

```

```

SSH
Set/Get:
    SSH_PrivKey          ($HOME/.ssh/id_dsa)
    SSH_PubKey           ($HOME/.ssh/id_dsa.pub)
Get:
    SSH_RemoteUserName
    SSH_RemoteHost
    SSH_RemotePort

```

```

Kerberos
Set/Get:
    Kerberos_Ticket      (/tmp/kticket...) ?
Get:
    Kerberos_RemoteUserName
    Kerberos_RemotePort
    Kerberos_RemotePort

```

```

UserPass
Set/Get:
    UserPass_UserName    (anonymous)
    UserPass_Password    (anon)
Get:
    UserPass_RemoteUserName
    UserPass_RemoteHost
    UserPass_RemotePort

```

- Other types MAY be specified by a SAGA implementation.
- Default values can be specified by a SAGA implementation.
- Should we also specify the default values? Mostly simple I guess. But then the defaults may differ per platform and installation, so leaving that to the implementation gives more flexibility...

+-----+

Examples:

=====

```
// c++ example
saga::context c_1 (saga::context::SSH); // default attribs apply
saga::context c_2 (saga::context::FTP);

c_2.set_attribute ("UserName", "myself");
c_2.set_attribute ("Password", "secret");

saga::session s;
s.add_context (c_1);
s.add_context (c_2);

saga::file f ("any://remote.net/tmp/data.txt", s);

// file can be accessed now via ssh or ftp
```

+-----+

Notes:

=====

For encapsulating security information, a security context is created and associated with a context (aka session handle). The security context can hold information about X509 certificates, private/public keys, username/password, kerberos tickets etc., and provides these information to the SAGA implementation as needed.

A SAGA implementation MAY be able to attach more than one security context to one context.

Moved from Stream:

We need to do something with a SAGA context and security contexts.

context: This is an opaque datastructure that is used throughout the SAGA APIs. It hides key state information such as the security context and other shared data. It is passed in explicitly in order to support thread safety.

interface security_info:

security_info encapsulates information about the host or

authenticated user on the other end of a stream/socket connection.

The information encapsulated by this object can be used to make authorization/access-control decisions based on the identity of the remote user or host.

The security_info is an opaque structure that can be interrogated (via a different api) to determine the identity of the connected host. This information is essential for supporting Authorization and access control mechanisms. convenience functions that encode some of the most commonly required information used to make authorization decisions. Additional information that can be used to make authorization decisions or provide other identifying features for the remotely connected host or user can be interrogated using the SAGA "parameters" API that the security_info object implements. These parameters are always interrogated as string-based Key-value pairs.

- get_remote_user_name

Purpose: Gets the username associated with the remotely connected entity (if available).

Format: get_remote_user_name (out string username);

Inputs: none

Outputs: username: username assoc with remote connection

Notes: - returns NULL string if UserName not available.

- get_remote_dn

Purpose: Gets the distinguished name associated with the remotely connected entity (if available).

Format: get_remote_dn (out string dn);

Inputs: none

Outputs: dn Distinguished Name assoc with remote connection

Notes: - returns NULL string if that information is not available.

- get_remote_host

Purpose: Gets the hostname of the other side of connected stream (if available).

Format: get_remote_host (out string hostname);

Inputs: none

Outputs: hostname: hostname assoc with remote connection

Notes: - returns NULL string if that information is not available.

- get_remote_port

Purpose: Gets the portnumber of the other side of

connected stream (if available).
Format: get_remote_port (out int port);
Inputs: none
Outputs: port: portnumber assoc with remote
connection
Notes: - returns '0' if that information is not
available.

+-----+

```

+-----+
#####
#      #####      #####      #####      #####
#      #      #      #      #      #      #      #
#####      #      #      #      #      #      #
#      #####      #####      #      #      #####
#      #      #      #      #      #      #      #
#####      #      #      #      #      #####      #      #
+-----+

```

Summary:

=====

Each SAGA API call has an associated list of exceptions it can throw. These exceptions implement the `saga::exception` interface.

All objects in SAGA implement `error_handler`, which allows a user of the API to query for the latest error associated with a saga object. In languages with exception facilities, such as Java, C++ and Perl, the language binding may allow exceptions to be thrown instead. Bindings for languages without exception handling capabilities will provide a mechanism to examine nested exceptions.

The details of the error handling mechanisms will be defined in the respective language bindings.

```

+-----+

Specification:
=====

```

```

package saga.error
{
    class exception
    {
        CONSTRUCTOR      (in Object      object,
                          in string      message);

        DESTRUCTOR      (void);

        what              (out string      message);
        get_message      (out string      message);
        get_object       (out Object      object);
    }

    interface error_handler
    {

```

```

    get_error          (out exception    error);
    get_error_message  (out string      message);
    has_error_type     (out boolean     state);
  }
}

```

+-----+

Details:

=====

SAGA provides a set of well defined error states (exceptions) which MUST be supported by the implementations. As to whether these error states are non-critical, critical or fatal depends on (A) the specific implementation (one implementation might be able to recover from an error while another implementation might not), (B) the specific application use case (for example the error 'file does not exist' is fatal or not depends if the application really needs information from that file).

The exceptions available in SAGA are listed below, with a number of explicit examples on when that exception should be thrown. These examples are not normative, but illustrative. The spec defines the set of allowed exceptions for each method explicitly - that set is normative.

The exceptions below are sorted, most specific ones first, least specific ones last. On any error condition, the most specific exception possible MUST be thrown.

The SAGA spec defines what exceptions can be thrown by what method. However, depending on the implementation, other exceptions can be thrown as well. For example, an implementation might have authorization on specific attribute settings, and could throw a `AuthorizationFailed` exception on attempts to write that attribute, even if that is not specified in the SAGA spec. New SAGA exception types SHOULD NOT be defined by the implementation.

Depending on the language bindings, the listed exceptions are derived from the base SAGA exception types, or are error codes with that specific name etc. For details, see the language bindings.

NotImplemented:

If a method is specified in the SAGA API, but cannot be provided by a specific SAGA implementation, this exception MUST be thrown. See also the notes about

compliant implementations in the instruction.

IncorrectSession:

A method was invoked which effects two object instances which belong to different SAGA sessions. Currently, the SAGA API does not provide any method which could potentially have colliding sessions - that exception is defined for future SAGA extensions, e.g. work flows.

AuthenticationFailed:

An operation failed because none of the available contexts of the used session could be used for successful Authentication.

Example:

- a remote host did not accept a X509 certificate because the respective CA is unknown there.

AuthorizationFailed:

An operation failed because none of the available contexts of the used session could be used for successful Authorization. Authentication succeeded.

Example:

- although a certificate was valid on a remote GridFTP server, the ID could not be mapped to a valid local user ID.

PermissionDenied:

A operation failed because the identity used for the operation did not have sufficient permissions to perform the operation successfully. Authentication and authorization have been successful.

Example:

- although a user could login to a remote host via GridFTP and could be mapped to a local user, the write on /etc/passwd failed.

BadParameter:

This exception indicates that at least one of the parameters of the method call is ill-formed, invalid, out of bound or otherwise not usable. The error message **MUST** give specific information on what parameter caused that exception, and why.

Examples:

- a specified context type is not supported by the implementation
- a file name specified is invalid, e.g. too long, or contains characters which are not allowed
- an ivec for scattered read/write is invalid, e.g. has offsets which are out of bound, or non-allocated buffers
- a buffer to be written and the specified lengths are incompatible
- an enum specified is not known
- flags specified are incompatible (ReadOnly | WriteOnly)

IncorrectState:

This exception indicates that the object a method was called on is in a state where that method cannot possibly succeed. A change of state might allow the method to succeed with the same set of parameters.

Examples:

- calling read on a stream which is not connected
- calling write on a file which is opened read only
- calling run on a task which was canceled
- calling resume on a job which is not suspended

AlreadyExists:

This exception indicates that an operation cannot succeed because an entity to be created or registered already exists, or is already registered and cannot be overwritten. Explicit flags on the method invocation may allow the operation to succeed, if for example they indicate that overwrite is allowed.

Examples:

- a file to be created already exists
- a target for a file move already exists

- a name to be added to a logical file is already known
- a metric to be added to a object has the same name as an existing metric on that object
- a context to be added to a session was added earlier

DoesNotExist:

This exception indicates that an operation cannot succeed because a required entity is missing. Explicit flags on the method invocation may allow the operation to succeed, if they for example indicate that create is allowed.

Examples:

- a file to be moved does not exist
- a directory to be listed does not exist
- a name to be replicated is not in a replica set
- a name to be deleted is not in a replica set
- a metric asked for is not known to the object
- a context asked for is not known to the session
- a task asked for is not in a task_container
- a attribute asked for is not supported
- a job asked for by id is not known by the backend

ReadOnly:

A attribute or metric was attempted to be changed but is read-only, e.g. is provided only for informational purposes. That exception does NOT apply for files or streams which are in incorrect state (i.e. not readable or writable) - that would cause a `IncorrectState` exception.

Examples:

- attempt to change an attribute which is read only
- attempt to change or update a metric which is read only

ReadError:

This exception indicates that a read operation on a file, directory or stream failed, although the object in question has been in the correct state (i.e. readable). On `NonBlocking` objects, reads might frequently fail but might succeed in a later call (`EAGAIN`) - in such cases this exception MUST NOT be thrown, as that situation does not indicate an error.

Examples:

- a read on a file failed because the file was opened read-only
- a read on a stream failed because no data are available

WriteError:

This exception indicates that a write operation on a file, directory or stream failed, although the object in question has been in the correct state (i.e. writable). On NonBlocking objects, writes might frequently fail but might succeed in a later call (EAGAIN) - in such cases this exception MUST NOT be thrown, as that situation does not indicate an error.

Timeout:

This exception indicates that a remote operation did not complete successfully because the network communication or the remote service timed out. That exception MUST NOT be thrown if a timed wait times out - that is indicated by the waits return value, and does not pose an error condition. The time waited before a Timeout is indicated depends on the implementation and on the backend, and SHOULD be documented where it is implementation specific.

Examples:

- a remote file authorization request timed out
- a remote data base access timed out
- a host name resolution timed out
- a started file transfer stalled and timed out

NoSuccess:

This exception indicates that an operation failed semantically, e.g. the operation was not successfully performed. This exception is the least specific exception defined in SAGA, and CAN be used for all error conditions which do not indicate a more specific exception.

Examples:

- a once open file is not available right now
- a backend did not answer a request about job state
- a file copy was interrupted mid-stream

```
class exception:
```

```
-----
```

This is the basic exception interface for all exceptions thrown by a SAGA object implement.

Note that `saga::exception` does NOT implement the `saga::object` interface.

- CONSTRUCTOR

Purpose: create the exception

Format: CONSTRUCTOR (in object object,
out exception e);

Inputs: object: the object associated with the exception.

Outputs: e: the newly created exception

- DESTRUCTOR

Purpose: destroy the exception

Format: DESTRUCTOR (in exception e);

Inputs: e the exception to destroy

Outputs: none

- what

what is an alias for `get_message`.

- get_message

Purpose: gets the message associated with an exception

Format: `get_message` (out string message);

Inputs: none

Outputs: message the error message

- get_object

Purpose: gets the SAGA object associated with exception

Format: `get_message` (out object o);

Inputs: none

Outputs: o: the object associated with the exception

```
+-----+
```

Examples:

```
=====
```

```
// c++ example
int main ()
{
    try
    {
        saga::file f ("file://localhost/etc/passwd");
        f.copy (/home/user/passwd.bak");
    }
    catch ( const saga::exception::PermissionDenied & e )
    {
        std::cerr << "SAGA error: No Permissions!" << std::endl;
    }
    catch ( const saga::exception & e )
    {
        std::cerr << "SAGA error: " << e.what () << std::endl;
    }

    return (0);
}
```

+-----+

Notes:

=====

There was discussion of using tagging interfaces, such as

```
class LibraryException extends exception {
}

class LibraryFatalException implements-all LibraryException {
}

class LibraryRecoverableException
                                implements-all LibraryException {
}

class BackendException extends exception {
}

class BackendFatalException implements-all BackendException {
}

class BackendRecoverableException
                                implements-all BackendException {
}
```

In languages which allow exceptions, to allow a user of the API to additionally catch exceptions based upon the types, however such a tagging interface approach does not enforce that an exception falls into one of the types.

+-----+

```

+-----+
          #####
          #      ##      #####      #      #
          #      #      #      #      #      #
          #      #      #      #####      #####
          #      #####      #      #      #
          #      #      #      #      #      #
          #      #      #      #####      #      #
+-----+

```

Summary:

=====

Operations performed in widely distributed environments may take a long time to complete, and thus it is desirable to have the ability to perform operations in an asynchronous manner. There are many possible ways in which an asynchronous API may be developed --- the notes for this API contain several possibilities.

The main requirements the SAGA design team faced were ease of implementation in different languages, the ability to be implemented in a single-threaded environment, generality and ease of use.

This document defines an API and a pattern which associates a 'task' with each outstanding asynchronous operation. Each task represents an asynchronous version of one SAGA API method, and may have no one-to-one correspondence with any external process, such as a job.

```

+-----+

```

Specification:

=====

```

package saga.task
{
    enum state
    {
        Unknown      = -1,
        New           =  1,
        Running       =  2,
        Done          =  3,
        Failed        =  4
    }
}

```

```

class task : implements-all saga::object
               implements-all saga::monitorable
{
    // no constructor
    DESSTRUCTOR    (in Task          obj);

    run            (void);
    cancel         (void);
    wait           (in double         timeout,
                   out boolean        finished);

    get_state      (out State          state);
    get_error      (out Exception      e);

    rethrow        (void);

    // Metric:
    //    name: state
    //    desc: "fires if on task state change, and
    //           has the value of the task state enum."
    //    mode: Read
    //    Unit: 1
    //    Type: Enum
    //    Value: "Unknown"
}

class task_container : implements-all saga::object
                      implements-all saga::monitorable
{
    CONSTRUCTOR    (out task_container obj);
    DESSTRUCTOR    (in task_container obj);

    add            (in Task          task,
                   out int           cookie);
    remove         (in int           cookie);

    run            (void);
    cancel         (void);
    wait           (in double         timeout,
                   int bool          all,
                   out Task          finished);

    get_tasks      (out array<Task,1> tasks);
    get_states     (out array<State,1> states);

    // Metric:
    //    name: state
    //    desc: "fires if on changes of any task in container,"

```

```

    //          and has the value of that tasks cookie."
    //      mode:  Read
    //      Unit:  1
    //      Type:  int
    //      Value: "0"
}
}

```

Details:

=====

// FIXME

Each object in the SAGA API defines a `create_task_factory` method, which creates a corresponding factory object implementing the same set of methods as the original object, but returning a `saga::task` object.

E.g. the `saga::file` class has a corresponding `saga::directory_task_factory` class, objects of which are instantiated by invoking `directory.create_task_factory`. This `directory_task_factory` object has the same methods as those of the `directory` object; invoking any of these methods creates a task object representing an asynchronous call.

Input and Output arguments of API calls MUST not be accessed or changed until the asynchronous task has completed; i.e. until 'wait' has been invoked on the task object and returned that the task state is Done or Failed. Output values are only defined if the task is in Done state.

enum state:

A task can be in one of several possible states:

New

The task has been created but not yet started.
Tasks start in this state, it is initial

Running

The `run()` method has been invoked on the task.

Failed

The asynchronous operation has unsuccessfully finished, or has been cancelled.
This state is final.

Done

The asynchronous operation has successfully finished.

This state is final.

class Task:

Objects of this class represent asynchronous API calls. They are only created by invoking a method on a saga object which returns a task object (with saga::ASync or saga::Task).

- DESTRUCTOR

Purpose: destroy the object

Format: DESTRUCTOR (in Task obj)

Inputs: obj: the object to destroy

Outputs: none

- run

Purpose: Start the asynchronous operation.

Format: run ();

Inputs: none

Outputs: none

Throws: IncorrectState

- wait

Purpose: Wait for the task to finish.

Format: wait (in double timeout,
out boolean done);

Inputs: timeout: number of seconds to wait

Outputs: done: indicating if the task is
done running

Throws: IncorrectState

Notes:

- timeout < 0.0 wait forever
- timeout = 0.0 return immediately
- timeout > 0.0 wait for this number of seconds
- a task must be in New, Running, Done or Failed state for wait to succeed. Otherwise an exception is thrown.
- if the task is Running or New, and on timeout of wait is still Running or New, false is returned
- if the task is Running or New, and on timeout of wait is not Running or New anymore, true is returned. True hence indicates that the task reached a final state.

- cancel

Purpose: Cancel the asynchronous operation.

Format: cancel ();

Inputs: none

Outputs: none

Notes: - cancel can return immediately, even if the

operation is neither done nor definitely cancelled. The state remains 'Running' until the cancel operation succeeded - the state then changes to 'Failed'.

- Information relying on the task state are undefined after calling cancel.

- getState

Purpose: Get the state of the task.

Format: getState (out State state);

Inputs: none

Outputs: state: state of the task.

- getError

Purpose: Get the error of a failed task.

Format: getState (out exception e);

Inputs: none

Outputs: e: exception of task.

- rethrow

Purpose: re-throw any error a failed task caught.

Format: throw ();

Inputs: none

Outputs: none

class task_container:

When there are many asynchronous tasks it would be inefficient to invoke the wait() method on each one sequentially. The task_container class provides a mechanism to wait and other operations for a set of tasks.

- add

Purpose: Add a Task to a task_container.

Format: add (in Task task);

Inputs: task: task to add to the task_container

Outputs:

- remove

Purpose: Remove a Task from a task_container.

Format: remove (in Task task);

Inputs: task: task to remove from the task_container

Outputs: none

Throws: DoesNotExist

- run

Purpose: Start all asynchronous operations in the container.

```

Format:  run ();
Inputs:  none
Outputs: none
Throws:  IncorrectState

```

- wait

```

Purpose: Wait for one or more of the tasks to finish.
Format:  wait          (in  double timeout,
                        in  bool   all
                        out Task   done);

Inputs:  timeout:      number of seconds to wait
          all:          if true wait for all tasks
Outputs: done:          finished task
Throws:  IncorrectState
Notes:   - < 0.0 wait forever
          - = 0.0 return immediately
          - > 0.0 wait for this number of seconds
          - for the exception condition, see the wait
            description in the task class.
          - if 'all' is true, the wait call returns only
            if all tasks in the container are finished,
            or on timeout, whatever occurs first.
            The output task is then any of the finished
            tasks.

```

- cancel

```

Purpose: Cancel all the asynchronous operations in the
         container.
Format:  cancel ();
Inputs:  none
Outputs: none
Notes:   - see semantics of task cancel.

```

- get_states

```

Purpose: Get the states of all tasks in the task
         task_container.
Format:  getStates      (out array<State,1> states);
Outputs: States:        array of States for
                        tasks in task_container

```

- get_tasks

```

Purpose: Get the tasks in the task task_container.
Format:  get_tasks      (out array<Task,1> tasks);
Outputs: tasks:         array of Tasks in
                        task_container

```

+-----+

Examples:

```
=====
```

```
// c++ example, partly pseudocode
saga::directory dir;
saga::job      job;

...

/* Create Tasks */
saga::task t1 = dir.ls      <saga::task> (result);
saga::task t2 = dir.copy    <saga::task> (source,target);
saga::task t3 = dir.move    <saga::task> (source,target);
saga::task t4 = job.checkpoint <saga::task> ();
saga::task t5 = job.signal   <saga::task> (SIG_USR);

// Start Tasks
t1.run ();
t2.run ();
t3.run ();
t4.run ();
t5.run ();

// put all tasks into container
saga::task_container tc;

tc.add (t1);
tc.add (t2);
tc.add (t3);
tc.add (t4);
tc.add (t5);

// take one out again
tc.remove (t5);

// wait for all tasks in container to finish
tc.wait ();

// wait for the last task
t5.wait ();
```

```
+-----+
```

Notes:

```
=====
```

Error Handling:

```
=====
```

```
{
    task.run ();
    task.wait ();
```

```

    if ( task.get_state = saga::task::Failed )
    {
        try {
            task.rethrow ();
        }
        catch ( saga::exception e )
        {
            std::cout << "task failed: " << e.what () << std::endl;
        }
    }
}

```

Task models:

=====

We had six different task models, as shown in example form below. Model (E) has no compile-time sanity checking. Model (F) allows only one asynchronous operation per object. Once these models were eliminated, the choice between the remaining four was a matter of aesthetics as they all have equivalent functionality.

The task container could have more methods to ease retrieval and manipulation of tasks. E.g. the ability to label tasks and retrieve by label.

```

-----
Directory dir ("foo://bar/baz")
Job          job = ...

```

Model A)

In this model there is a Task class associated with each API class, which is created by a create_task method. Once a Task object has been created the asynchronous operation is invoked on it to associate an operation with the Task.

```

/* Create Tasks */
dir_task dt1 = dir.create_task ();
dir_task dt2 = dir.create_task ();
dir_task dt3 = dir.create_task ();
job_task jt1 = job.create_task ();
job_task jt2 = job.create_task ();

/* Invoke operations on Task Objects */
dt1.ls      ();
dt2.copy    (source,target);
dt3.move    (source,target);

```

```
jt1.checkpoint ();
jt2.signal      (USR);
```

```
/* Start Tasks */
```

```
dt1.run ();
dt2.run ();
dt3.run ();
jt1.run ();
jt2.run ();
```

```
-----
```

Model B)

In this model there is a `task_factory` class associated with each API class, which is created by a `create_task_factory` method. Once a `task_factory` object has been created the asynchronous operation is invoked on it to create a Task object.

```
/* Create Task factories */
```

```
dir_task_factory dtf = dir.create_task_factory ();
job_task_factory jtf = job.create_task_factory ();
```

```
/* Create Tasks */
```

```
task t1 = dtf.ls      ();
task t2 = dtf.copy    (source,target);
task t3 = dtf.move    (source,target);
task t4 = jtf.checkpoint ();
task t5 = jtf.signal  (USR);
```

```
/* Start Tasks */
```

```
t1.run ();
t2.run ();
t3.run ();
t4.run ();
t5.run ();
```

```
saga::directory dir (url);
saga::directory::task_factory dft = dir.get_task_factory ();
```

```
std::list <std::string> listing;
```

```
        dir.list (listing);
saga::task t = dtf.list (listing);
```

```
t.run ();
```

```
-----
```

Model C)

In this model there is an object as an attribute on each API object. Invoking an operation on this object creates a Task.

```

/* Create Tasks */
task t1 = dir.task.ls          ();
task t2 = dir.task.copy       (source,target);
task t3 = dir.task.move       (source,target);
task t4 = job.task.checkpoint ();
task t5 = job.task.signal     (USR);

/* Start Tasks */
t1.run ();
t2.run ();
t3.run ();
t4.run ();
t5.run ();

```

Model D)

In this model there is an equivalent for each API call which creates an asynchronous task.

```

/* Create Tasks */
task t1 = dir.task_ls          ();
task t2 = dir.task_copy       (source,target);
task t3 = dir.task_move       (source,target);
task t4 = job.task_checkpoint ();
task t5 = job.task_signal     (USR);

/* Start Tasks */
t1.run ();
t2.run ();
t3.run ();
t4.run ();
t5.run ();

```

Model E)

In this model, there is a `get_task` method associated with each API object, which creates a Task given a string argument defining the operation.

```

/* Create Tasks */
task t1 = dir.get_task ("ls");
task t2 = dir.get_task ("copy",source,target);
task t3 = dir.get_task ("move",source,target);
task t4 = job.get_task ("checkpoint");
task t5 = job.get_task ("signal",USR);

/* Start Tasks */
t1.run ();

```

```
t2.run ();
t3.run ();
t4.run ();
t5.run ();
```

Model F)

In this model, there is an asynchronous version of each API call, and each API class has a 'wait' method. As there is no Task object, only one asynchronous operation may be outstanding on any object.

```
dir.async_ls      ();
job.async_checkpoint ();
job.wait          ();
dir.wait          ();
dir.async_copy    (source,target);
dir.wait          ();
dir.async_move    (source,target);
job.async_signal  (USR);
job.wait          ();
dir.wait          ();
```

+-----+

```

+-----+
#      #
##    ##   ##   #   #   #   #####   ##   #####   #   #   #   ##
#   #   #   #   #   #   #   #   #   #   #   #   #   #   #   #
#   #   #   #   #   #   #   #   #   #   #   #   #   #   #   #
#       #   #   #   #   #   #   #   #   #   #   #   #   #   #   #
#       #   #   #   #   #   #   #   #   #   #   #   #   #   #   #
#       #   #   #   #   #   #   #   #   #   #   #   #   #   #
#       #   #   #   #   #   #   #   #   #   #   #   #   #   #

      ##
      #   #
      ##
      ###
      #   #   #
      #       #
      ###   #

#####
#       #   #####   #####   #####   #   #   #   #####
#       #       #       #       #   #   #   ##   #   #   #
#####   #   #####   #####   #       #   #   #   #   #
      #   #       #       #       #####   #   #   #   #   #   #
#       #   #       #       #       #       #   #   #   #   #
#####   #   #####   #####   #       #   #   #   #   #
+-----+

```

Summary:

=====

The ability to query Grid entities about state is requested in several SAGA use cases. Also, the SAGA Task model incorporates a certain amount of task monitoring.

This package definition approaches the problem space of monitoring to unify the various usage patterns (see details and examples), and to transparently incorporate SAGA task monitoring. The paradigm is realised by introducing monitorable SAGA objects, which expose metrics to the application, which represent values to be monitored.

A closely related topic is Computational Steering, which is (for our purposes) not seen independently from Monitoring: in the SAGA approach, the Steering mechanisms extend the Monitoring mechanisms by the ability to push values back to the monitored entity, i.e. to introduce writable monitorables.

+-----+

Specification:

=====

```
package saga.monotoring
{
    // a metric represents an entity / value to be monitored.
    class metric : implements-all saga::object
                  implements-all saga::attribute
    {
        CONSTRUCTOR      (in  string      name,
                          in  string      desc,
                          in  string      mode,
                          in  string      unit,
                          in  string      type,
                          in  string      value,
                          out metric      metric);

        DESTRUCTOR       (in  metric      metric);

        // manage callbacks on the metric
        // add a callback, which gets active whenever
        // the metric changes (fires)
        add_callback      (in  call_back    * cb,
                          out int          cookie);

        // remove the callback
        remove_callback   (in  int          cookie,
                          out call_back    * cb);

        // push a new value to the consumers
        fire              (void);

        // Attributes:
        //   name:  name
        //   desc:  name of metric
        //   mode:  ReadOnly
        //   type:  String
        //   value: naming conventions as described below apply
        //
        //   name:  description
        //   desc:  description of metric
        //   mode:  ReadOnly
        //   type:  String
        //
        //   name:  mode
        //   desc:  access mode of metric
        //   mode:  ReadOnly
    }
}
```

```

//  type: String
//  value: 'ReadOnly' or 'ReadWrite'
//
//  name: unit
//  desc: unit of metric
//  mode: ReadOnly
//  type: String
//
//  name: type
//  desc: value type of metric
//  mode: ReadOnly
//  type: String
//  value: 'Int', 'Float', 'Bool' or 'String'
//
//  name: value
//  desc: value of metric
//  mode: depending on the mode attribute above
//  type: String
//  value: formating restrictions described below apply
//
}

// callbacks are used for asynchroneous notification of
// metric changes (== events)
interface call_back
{
    callback          (in  metric          metric,
                      out bool            keep);
}

// SAGA objects which provide metrics and steering via
// metrics, implement the monitorable interface
interface monitorable
{
    // introspection
    list_metrics      (out array<string,1> names);

    // get an availble metric for monitoring/steering
    get_metric        (in  string          name,
                      out metric          metric);

    // add an application hook for monitoring/steering
    add_metric        (in  metric          metric);

    // removes an application hook for monitoring/steering
    remove_metric     (in  string          name);

    // add a callback, which gets active whenever
    // the respective metric changes (fires)
    add_callback      (in  string          name,

```

```

        in call_back    * cb,
        out int         cookie);

    // remove the callback
    remove_callback      (in int         cookie,
                          out call_back  * cb);

    // steering: push the updated metric value (fire)
    // on a ReadWrite metric
    fire_metric          (in string      name);
}
}

```

+-----+

Details:

=====

```

class metric:
-----

```

The fundamental object introduced in this package is a metric. A metric represents an observable, which can be readable, writeable, or read/writable. The availability of a readable observable corresponds to monitoring; of a writable observable corresponds to steering. A metric is 'Final' when its values cannot change anymore, ever (i.e. progress is '100%', job state is 'Done' etc).

The approach is severely limited by the use of saga attributes for the description of a metric, as these are only defined in terms of string typed keys and values. An extension of the attribute definition by typed values will greatly improve the usability of this package, but will also challenge its semantic simplicity.

The metric MUST provide access to following attributes (examples given):

```

name:      short human readable name.
           - ex:  file.copy.progress

desc:      extensive human readable description
           - ex:  "This metric gives the state of
                  an ongoing file transfer as
                  percent completed."

mode:      "Read", "Write", "ReadWrite" or "Final"
           - ex:  "ReadWrite"

```

```

unit:      Unit of values
           - ex: "percent (%)"
           - ex: "Unit"

type:      "String", "Int", "Float" etc.
           - ex:  "Float"

value:     value of the metric
           - ex:  "20.5"

```

The name of the metric must be unique, as it is used in most methods to identify the metric of interest. The use of a dot-delimited name space for metrics as in the example above is encouraged, as it greatly benefits the interactive handling of metrics. The first element of the name space SHOULD be the SAGA class the metric belongs to, the second element SHOULD describe the operation the metric describes (if applicable, otherwise leave out), the third element SHOULD indicate the description of the metric (e.g. 'state' or 'progress' or 'temperature'). Illustrative examples for metric names are:

```

- file.copy.progress
- file.move.progress
- file.progress      // applies to any operation on file
- job.state
- job.temperature    // a custom observable on a job

```

All attributes, apart from 'value', are ReadOnly, and are initialized in the metric constructor. If any other attribute than 'value' changes, a new metric needs to be created.

The following attribute values MUST be interpreted case insensitive: name, type, mode, unit.

Metric definitions in the SAGA specification

The SAGA specification defines a number of metrics which MUST or CAN be supported, for various SAGA objects. An example such a definition is (from the SAGA stream object):

```

class stream ...
{
    ...

    // Metrics:
    //   name:  read
    //   desc:  fires if a stream gets readable

```

```
// mode: Read
// unit: 1
// type: Bool
// value: True
}
```

These specifications are NORMMATIVE, even if described as comments in the SIDL specification! The specified metrics MUST be supported by an implementation, unless noted otherwise, e.g.:

```
// note: MAY be supported
```

Implementations MAY add custom metrics, which SHOULD be documented if possible. However, metrics CAN also be added at runtime - that is, for example, required for computational steering of custom applications.

A metric can 'appear' and 'go away' during the lifetime of an object (again, computational steering provides the obvious use case for this). Any operation on a metric which got removed ('dead metric') MUST throw an DoesNotExist exception. However, existing class instances of a dead metric MUST stay valid, and expose the same life time as any other 'life metric'. Attributes of a dead metric MUST be readable for the lifetime of the object. The Mode attribute of such an instance MUST be changed to Final by the implementation, and no other changes are allowed after that change.

- CONSTRUCTOR

Purpose: create the object

```
Format: CONSTRUCTOR      (in string  name
                           in string  desc,
                           in string  mode,
                           in string  unit,
                           in string  type,
                           in string  value,
                           out metric  obj);
```

```
Inputs: name:      name of metric
        desc:      description of metric
        mode:      mode of metric
        unit:      unit of metric value
        type:      type of metric
        value:     initial value of metric
```

```
Outputs: obj:      the newly created object
```

```
Notes:  - a metric is NOT attached to a session, but
         can be used for monitoring objects from
         different sessions.
         - the string arguments given are used to
```

```

        initialise the attributes of the metric, which
        are subsequently ReadOnly (see description
        above).

- DESTRUCTOR
  Purpose: destroy the object
  Format:  DESTRUCTOR          (in  metric obj)
  Inputs:  obj:                the object to destroy
  Outputs: none

// manage callbacks on the metric
- add_callback
  Purpose: add asynchron notifier callback to watch metric
           changes
  Format:  add_callback        (in  call_back    * cb,
                               out int          cookie);
  Inputs:  cb:                callback class instance
  Outputs: cookie:            handle for this callback,
                               to be used for removal

  Throws: IncorrectState
  Notes:  - IncorrectState is thrown if the metric is Final
           - the 'callback' method on cb will be invoked on
             any change of the metric value
           - if the 'callback' method returns true, the
             callback is kept registered; if it returns
             false, the callback is called only once, and
             is un-registered after completion. Note that,
             due to threading, multiple cb instances can be
             active at the same time, so returning false can
             lead to race conditions.
           - the cb is passed by reference, not by value, so
             that no copy of private cb data is implied.
             The user has to ensure data consistency if a cb
             instance is used multiple times.

- remove_callback
  Purpose: remove a callback from a metric
           changes
  Format:  remove_callback     (in  int          cookie,
                               out call_back    * cb);
  Inputs:  cookie:            handle identifying the cb to
                               be removed
  Outputs: cb:                the removed cb reference
  Notes:  - if the callback was removed earlier, or
           was unregistered by returning false, no
           exception is thrown. However, the returned cb
           reference is NULL.
           - the removal only affects the cb reference
             identified by 'cookie', even if the same
             reference was registered multiple times.
           - no assumption is made about threading: a

```

returned cb can very well be active at the time of removal, e.g. can in a nother thread execute the callback method. The user has to ensure the correct shutdown.

- fire

Purpose: push a new metric value to the backend

Format: fire (void);

Throws: IncorrectState

Notes:

- IncorrectState is thrown if the metric is not Writable - so also if the metric was Writable, but is meanwhile flagged Final! To catch race condition triggered exceptions, each file should be try'ed/catched.
- it is not necessary to change the value of a metric in order to fire it.
- 'set_attribute ("value", "...") on a metric does NOT imply a fire. Hence the value can be changed multiple times, but unless fire() is explicitly called, no consumer will notice.
- any callback registered on the metric in the firing application is invoked on fire().

-

interface callback:

- callback

Purpose: asynchronous handler for metric changes

Format: callback (in metric metric,
out bool keep);

Inputs: metric: the metric causing the
callback invocation

Outputs: keep: indocates if cb stays
registered

Notes:

- see notes to the metric class above
- if 'keep' is returned as true, the callback stays registered, and will be invoked again on the next metric update.
- if 'keep' is returned as false, the callback gets unregistered, and will not be invoked again on metric updates, unless it gets re-added by the user.
- if 'keep' is returned as false, all active callback are finished normally, even if they run longer than the instance returning 'false'. However, it is guaranteed that no new activation occurs after one instance returns 'false'.

- 'metric' is the metric the callback is invoked on - that means that this metric recently changed its value. Note that this change is semantically defined by the metric: the string of the 'value' attribute of the metric might have the same value in two subsequent invocations of the callback.
- a callback can be added to a metric multiple times.
- a callback can be added to multiple metrics at the same time.

interface monitorable:

The monitorable interface is implemented by those SAGA objects which can be monitored, i.e. which have one or more associated metrics. The interface allows introspection of these metrics, and allows to add notification callbacks to these metrics.

Several methods on this interface reflect similar methods on the metric class - the additional string argument 'name' identifies the metric these methods act upon. The semantics of these calls are identical to the specification above.

The method add_metric() allows to implement steerable applications. Metrics added with this method SHOULD be available on the saga job object representing the application.

```
// introspection
- list_metrics
  Purpose: list all metrics associated with the object
  Format:  list_metrics      (out array<string,1> names);
  Outputs: names:           array of names identifying
                           the metrics associated with
                           the object instance

  Notes:  - several SAGA objects are required to expose
           certain metrics (e.g. 'task.state'). However,
           in general that assumption cannot be made, as
           implementations might be unable to provide
           metrics.
           - no order is implied on the returned array
           - the returned array is guaranteed to have no
             double entries

- get_metric
  Purpose: returns a metric instance, identified by name
```


-
- Format: `get_metric` (in string name,
out metric metric);
- Inputs: name: name of metric to be returned
- Outputs: metric: metric instance identified by name
- Throws: `DoesNotExist`
- Notes: - multiple calls of this method return multiple identical instances of the metric.
- `add_metric`
- Purpose: add a metric instance to the application instance
- Format: `add_metric` (in metric metric);
- Inputs: metric: metric to be added
- Throws: `DoesAlreadyExist`
- Notes: - a metric is uniquely identified by its name attribute - no two metrics with the same name can be added.
- any callbacks already registered on the metric stay registered (state of metric is not changed)
- `remove_metric`
- Purpose: remove a metric instance
- Format: `remove_metric` (in string name);
- Inputs: name: identifies metric to be removed
- Throws: `BadParameter`
- Notes: - only previously added metrics can be removed; default (saga defined or implementation specific) metrics cannot be removed, attempts to do so raise a `BadParameter` exception.
- `add_callback`
- Purpose: add a callback to the specified metric
- Format: `add_callback` (in string name,
in call_back * cb,
out int cookie);
- Inputs: name: identifies metric to which cb is to be added
- cb: reference of callback class instance to be registered
- Outputs: cookie: handle to be used for removal of the callback
- Throws: `DoesNotExist` - metric is unknown
- Notes: - notes to the `add_callback` method of the metric class apply
- `remove_callback`
- Purpose: remove a callback from the specified metric
- Format: `remove_callback` (in string name,

```

        in int    cookie);
Inputs:  name:      identifies metric for which cb
           is to be removed
        cookie:     identifies the cb to be
           removed
Throws:  DoesNotExist - metric is unknown
Notes:   - notes to the remove_callback method of the metric
           class apply

```

```

- fire_metric
Purpose: push a new metric value to the backend
Format:  fire_metric (int string name);
Inputs:  name:      identifies metric to be fired
Throws:  DoesNotExist
        IncorrectState
Notes:   - notes to the fire method of the metric
           class apply

```

+-----+

Examples:

=====

callback example: trace all metric changes:

```

// callback definition
class trace_cb : public saga::callback
{
public:
    bool callback (saga::metric m)
    {
        std::cout << "metric " << m.get_attribute ("name")
                    << " fired." << std::endl;
        return true; // stay registered
    }
}

// the application
int main ()
{
    ...

    // if the callback defined above is added to all known
    // metrics of all saga objects, a continuous trace of state
    // changes of these saga objects will be written to stdout
    trace_cb cb;

    saga::job j = [...]; // details see description of saga::job

```

```

        j.add_metric ("task.state", &cb);

        ...
    }

```

monitoring example: monitor a write task

```

// c++ example for task state monitoring
class write_metric_cb : public saga::callback
{
private:
    saga::task t_;

public:
    write_metric_cb (const saga::task & t) { t_ = t; }

    bool callback (saga::metric & m)
    {
        std::cout << "bytes written: "
                    << m.get_attribute ("value")
                    << std::endl;

        std::cout << "task state:  "
                    << t_.t_state ()
                    << std::endl;

        return (false); // keep callback registered
    }
};

int main (int argc, char** argv)
{
    ssize_t    len = 0;
    std::string str ("Hello SAGA\n");
    std::string url (argv[1]);

    saga::file  f (url);
    saga::task  t = f.write <saga::task> (str, &len);

    // assume that file has a 'progress' metric indicating
    // the number of bytes already written.  In general,
    // the list of metric names has to be searched for an
    // interesting metric, unless it is a default metric as
    // specified in the SAGA spec.

    // add the callback
    write_metric_callback cb (t);
    f.add_callback ("progress", &cb);
}

```

```

    // wait until task is done, and give cb chance to get
    // called a couple of times
    t.wait ();
}

```

steering example: steer a remote job

```

// example for computaional steering (metric is writable).
class observer_cb : public saga::metric::callback
{
private:
    saga::task t;

public:
    bool callback (saga::metric & m)
    {
        int val = atoi ( m.get_attribute ("value") );

        std::cout << val << " is the new value." << std::endl;

        return (false); // keep callback registered
    }
};

// the steering appliciation
int main (int argc, char** argv)
{
    saga::job_service js;

    saga::job j = js.run ("remote.host.net",
                          "my_remote_application");

    // Assume that job has a 'param_1' metric representing
    // a integer parameter for the remote application.
    // First add observer metric - that causes the
    // asynchronous printout of any changes to the value
    // of the 'param_1' metric
    observer_cb cb;
    j.add_callback ("param_1", &cb);

    // then get metric for active steering
    saga::metric m = j.get_metric ("param_1");

    for ( int i = 0; i < 10; i++ )
    {
        // if param_1 is ReadOnly, set_value would throw
        // 'IncorrectState'
        m.set_attribute ("value", std::string (i));
    }
}

```

```

    // push the pending change out to the receiver
    m.fire ();

    // callback should get called NOW + 2*latency
    // That means fire REQUESTS the value change, but only
    // the remote job can CHANGE the value - that change
    // needs then reporting back to us.

    // give steered application some time to react
    sleep (1);
}
}

```

steering example: BE a steerable job

the example shows a job which

- creates a metric to expose a Float steerable parameter
- on each change of that parameter computes a new isosurface

```

// callback - on any change of the metric value, e.g. due to
// steering from a remote GUI application, a new iso surface
// is computed

```

```

class my_cb : public saga::callback
{
public:
    // the callback gets called on any
    bool callback (saga::metric m)
    {
        // get the new iso-value
        float iso = atof (m.get_attribute ("value"));

        // compute an isosurface with that iso-value
        compute_iso (iso);

        // keep this callback alive, and get called again on
        // the next metric event.
        return (false);
    }
}

```

```

int main ()
{
    // create a metric for the iso-value of an isosurfacer
    saga::metric m ("application.isosurfacer.isovalue",
                    "iso-value of the isosurfacer",
                    "ReadWrite",
                    "");
}

```

```

        "Float",
        "1.0");

    // add the callback which reacts on changes of the
    // metric's value (cookie is ignored)
    my_cb cb;
    m.add_callback (&cb);

    // get job handle for myself
    saga::job_server js;
    saga::job self = js.get_self ();

    // add metric to job
    self.add_metric (m);

    // now others can 'see' the metric, e.g. via
    // job.list_metrics ();

    // now, the callback could also have been added with:
    // self.add_metric ("application.isosurfacers.isovalue", cb);

    // compute isosurfaces for the next 10 minutes -
    // the real work is done in the callback
    sleep (600);

    // on object (self) destruction, metrics and callback
    // objects are destroyed as well
    return (0);
}

```

monitoring example: callback for stream connects

```

// my callback container class
class my_cb : public saga::callback
{
    privat:
        // we keep a stream server and a single client stream
        saga::stream_server ss_;
        saga::stream        s_;

    public:
        // constructor initialises these (note that the
        // client stream should be not connected at this
        // point
        my_cb (saga::stream_server ss,
              saga::stream        s )
        {
            ss_ = ss;

```

```
s_ = s;
}

~my_cb (void) { }

// the callback gets called on any
bool callback (saga::metric m)
{
    // the stream server got an event triggered, and
    // should be able to create a client socket now.
    s_ = ss_.wait ();

    if ( s_.state == saga::stream::open )
    {
        // have a client stream, we are done
        // don't call this cb again!
        return (true);
    }

    // no valid client stream obtained: keep this
    // callback alive, and get called again on the
    // next metric event.
    return (false);
}
}

int main ()
{
    // create a stream server, and an not connected
    // stream
    saga::stream_server ss;
    saga::stream        s;

    // give both to our callback class, and register that
    // callback with the 'client_connect' metric of the
    // server. That causes the callback to be invoked on
    // every change of that metric, i.e. on every event
    // that changes that metric, i.e. on every client
    // connect attempt.
    my_cb cb (ss, s);
    ss.add_callback ("client_connect", &cb);

    // now we serve incoming clients forever
    while ( true )
    {
        // check if a new client is connected
        // the stream state would then be Open
        if ( s.state == saga::stream::Open )
        {
            // a client got conncted!
            // handle open socket

```

```

    s.write ("You say hello, I say good bye!\r\n", 32);

    // and close stream
    s.close ();

    // the stream is not Open anymore. We re-add the
    // callback, and hence wait for the next client
    // to connect.
    ss.add_callback ("client_connect", &cb);
}
else
{
    // no client yet, idle, or do something useful
    sleep (1);
}
}

// we should never get here
return (-1);
}

```

+-----+

Notes:

=====

- possible deviation: allow only one CB per metric:
no add/remove, but set/reset CB
- other useful attributes might be:
 - update-mode (discreet, cont, static)
This attribute would describe how often an attached callback gets activated.
discreet: on changes of the metric value
 example: job state (value is always defined, but changes infrequently)
 - cont: as fast as possible
 example: time (value always defined, and changes continuously)
 - static: once
 example: complete (value changes once)
- value-range (1-10, [1,2,3], ...)

That attribute obviously depends on type and unit of the metric, and is therefore not useful unless types attribute values or typed metric values are introduced in SAGA.

In particular computational steering would benefit from value ranges, as that allows client side range checks, and the creation of sensible client side user interfaces.

+-----+

```

+-----+
#
# # ##### ##### ##### # ##### # # ##### #####
# # # # # # # # # # # # #
# # # # # # # ##### # # # #####
##### # # ##### # # # # #
# # # # # # # # # # # #
# # # # # # # ##### ##### # #####
+-----+

```

Summary:

=====

There are various places in the SAGA APIs where attributes need to be associated with objects, for instance in the job descriptions. This API section provides a common interface for storing and retrieving attributes.

Objects implementing this interface can maintain a set of attributes. These attributes can be considered as a set of key-value pairs attached to the object. The key-value pairs are string based for now, but might cover other types in later versions of this API specification.

The SAGA spec defines attributes which **MUST** be supported by the various SAGA objects, and their default values, and also defines those which **CAN** be supported. An implementation **CAN** issue a warning if a specified attribute cannot be supported.

Implementations are strongly discouraged to include other, implementation specific attributes in SAGA, as that would bind applications to that specific implementation, and limit portability, which is a declared goal of the SAGA approach.

```

+-----+

```

Specification:

=====

```

package saga.attribute
{
    interface attribute
    {
        // setter / getters
        set_attribute      (in string      key,
                           in string      value);
        get_attribute      (in string      key,
                           out string     value);
    }
}

```

```

    set_vector_attribute    (in  string      key,
                           in  array<string,1> values);
    get_vector_attribute    (in  string      key,
                           out array<string,1> values);
    remove_attribute       (in  string      key);

    // inspection methods
    list_attributes        (out array<string,1> keys);
    equals                 (in  string      key,
                           in  string      val,
                           out bool        test);
    has_attribute          (in  string      key,
                           out bool        test);
    has_vector_attribute   (in  string      key,
                           out bool        test);
    is_readonly            (in  string      key,
                           out bool        test);
    is_writable             (in  string      key,
                           out bool        test);
}
}

```

Details:

=====

The attribute interface in SAGA provides a uniform paradigm to set and query parameters and properties of SAGA objects. Although the attribute interface is generic by design (i.e. it allows arbitrary keys and values to be used), its use in SAGA is limited to a finite and well defined set of keys, unless otherwise specified.

In several languages, attributes can much more elegantly expressed by native means - e.g. by using hash tables in Perl. Bindings for such languages MAY allow to use a native interface additionally to the one described here.

Many SAGA objects have very frequently use attributes. To simplify usage of these objects, setter and getter methods MAY be defined by the various language bindings, additionally to the interface described below. For attributes of native non string types, these setter/getters MAY be typed.

In order to limit semantic and syntactic ambiguities (e.g. due to spelling deviations), language bindings SHOULD define known attribute keys as constants, such as (in C):

```
#define SAGA_BUFFERSIZE "BufferSize"
```

```
...
stream.set_attribute (SAGA_BUFFERSIZE, "1024");
```

The distinction between scalar and vector attributes is somewhat artificial, and is supposed to help those languages where that nature of attributes cannot be handled transparently, e.g. by overloading. Bindings for languages such as Python, Perl and C++ CAN hide that distinction as long as both access types are supported.

To simplify handling of scalar/vector attributes, vector attributes can be specified as comma delimited strings.

```
val 1: home, sweet home
val 2: Global GF
val 3: SAGA
string: "home\, sweet home, Global GF, SAGA"
```

That format is returned if scalar getters are used for vector attributes, and can be used for scalar setters for vector attributes. Vector setters/getters handle scalar attributes as vectors of length one.

The order of the elements of vector attributes is ignored, and CAN be changed by the SAGA implementation. The equals method does also not rely on ordering (i.e. '"one" "two"' equals '"two" "one"').

Values of boolean type attributes MUST be expressed as one of the following: '0', '1', 'true', 'false', 'TRUE', 'FALSE'.

Values of floating point type attributes are expressed as they would result in a printf of the format "%lld", as defined by POSIX.

Values of integer type attributes are expressed as they would result in a printf of the format "%Lf", as defined by POSIX.

Values of string type attributes are expressed as-is, however, comma, backslashes and leading spaces are escaped by a backslash, as described above.

Some of the means above are aimed at the prevention of abuse of the attribute interface for implementation specific extensions and semantic overloading. However, we think that the attribute interface is flexible enough to accommodate small changes to the SAGA API in future versions of this specification with minor or no code changes, as long as the attribute interface as described below is used - both SAGA implementors and users are hence encouraged to prefer that interface to the convenience declarations described above.

```
interface attribute:
```

```
-----
```

```
- set_attribute
```

```
Purpose: set an attribute to a value.
```

```
Format: set_attribute      (in string key,  
                           in string value);
```

```
Inputs: key:               attribute key  
        value:             value to set the attribute  
                           to,
```

```
Outputs: none
```

```
Throws: ReadOnly
```

```
Notes:  - a value of NULL means to remove the attribute  
        - a empty string means to set an empty value  
          (the attribute remains present).
```

```
- get_attribute
```

```
Purpose: get an attributes value
```

```
Format: get_attribute      (in string key,  
                           out string value);
```

```
Inputs: key:               attribute key  
Outputs: value:            value of the attribute
```

```
Throws: DoesNotExist
```

```
- set_vector_attribute
```

```
Purpose: set an attribute to an array of values.
```

```
Format: set_vector_attribute (in string      key,  
                              in array<string,1> values);
```

```
Inputs: key:               attribute key  
        values:            array of values for the  
                           attribute
```

```
Outputs: none
```

```
Throws: ReadOnly
```

```
- get_vector_attribute
```

```
Purpose: get the array of values associated with an  
       attribute
```

```
Format: get_vector_attribute (in string      key,  
                              out array<string,1> values);
```

```
Inputs: key:               attribute key  
Outputs: values:           array of values of the  
                           attribute.
```

```
Throws: DoesNotExist
```

```
- remove_attribute
```

-
- Purpose: removes an attribute completely.
 Format: `remove_attribute` (in string key)
 Inputs: key: attribute key
 Outputs: none
 Throws: `DoesNotExist`
 `ReadOnly`
 Notes: - Any following query or set operation will not find the attribute present anymore.
- `list_attributes`
- Purpose: Get the list of attribute keys.
 Format: `list_attributes` (out array<string,1> keys);
 Inputs: none
 Outputs: keys: existing attribute keys
- `equals`
- Purpose:
 Format: `equals` (in string key,
 in string val,
 out bool test);
 Inputs: key: attribute key
 val: val to compare against
 Outputs: test bool indicating success
 Throws: `DoesNotExist`
 Notes: - This method returns TRUE if the attribute identified by key has the value identified by val.
 - For vector attributes, the value has to be specified as comma delimited concatenated string of the vector elements (order of the elements is ignored).
- `has_attribute`
- Purpose:
 Format: `has_attribute` (in string key,
 out bool test);
 Inputs: key: attribute key
 Outputs: test bool indicating success
 Notes: - This method returns TRUE if the attribute identified by key exists and is a scalar attribute.
 - It does NOT throw a `DoesNotExist` exception.
- `has_vector_attribute`
- Purpose:
 Format: `has_vector_attribute` (in string key,
 out bool test);

```

Inputs:  key:          attribute key
Outputs: test          bool indicating success
Throws:
Notes:   - This method returns TRUE if the attribute
          identified by key exists and is a vector
          attribute.
          - It does NOT throw a DoesNotExist exception.

```

```

- is_readonly
Purpose:
Format:  is_readonly      (in  string key,
                           out bool  test);

Inputs:  key:          attribute key
Outputs: test          bool indicating success
Throws:  DoesNotExist
Notes:   - This method returns TRUE if the attribute
          identified by the key exists, and can be read
          by get_attribute or get_vector attribute, but
          cannot be changed by set_attribute or
          set_vector_attribute.

```

```

- is_writable
Purpose:
Format:  is_writable      (in  string key,
                           out bool  test);

Inputs:  key:          attribute key
Outputs: test          bool indicating success
Throws:  DoesNotExist
Notes:   - This method returns TRUE if the attribute
          identified by the key exists, and can be
          changed by set_attribute or
          set_vector_attribute.

```

+-----+

Examples:

=====

```

// c++ example:
job_definition d;

std::list <char*> env;
env.push_back ("a = b");
env.push_back ("c = d");

d.set_attribute      ("JobCmd", "/bin/ls");
d.set_vector_attribute ("JobEnv", env);

```

+-----+

Notes:

=====

A find on attributes (both keys and values) should be considered, as that is needed and defined on other classes (saga::locical_file, saga::advert_service) anyway.

+-----+


```

+-----+
#      #
##    #    ##    #    #    #####
# #    #    #    ##    ##    #
#    #    #    #    #    ##    #    #####
#    #    #    #####    #    #    #
#    ##    #    #    #    #    #
#    #    #    #    #    #    #    #####

#####
#    #    #####    ##    #####    #####
#    #    #    #    #    #    #    #
#####    #    #    #    #    #    #####
#    #####    #####    #    #
#    #    #    #    #    #    #    #
#####    #    #    #    #    #####    #####

+-----+

```

Summary:

=====

This file describes interfaces which operate on arbitrary hierarchical namespaces, such as those used in physical, virtual and logical file systems, and information systems.

Several SAGA packages share the notion of namespaces and operations on these namespaces. In order to increase consistency in the API, those packages should share the same API paradigm.

The API is inspired by the POSIX standard, which defines tools and calls to handle the name space of physical files (directories). The methods listed for the interfaces have POSIX like syntax and semantics.

While POSIX has an iterative interface to directory listing (i.e., `opendir`, `telldir`, `seekdir`, `readdir`), the interface included here deviates significantly from the POSIX version; this interface has fewer calls, with a different syntax, but identical semantics.

Please note that 'stat' like API calls are not covered here - they are rather meaningless on a namespace per se, but belong to the specific implementations, e.g. physical files, which implement the namespace interfaces.

```

+-----+

```

Specification:

=====

```

package saga.name_space
{
    enum flags
    {
        Unknown          = -1,
        None              = 0,
        Overwrite         = 1,
        Recursive         = 2,
        FollowSymbolic    = 4,
        Create            = 8,
        Excl              = 16,
        Lock              = 32,
        CreateParents     = 64,
        DeReference       = 128
    }

    class ns_entry : implements_all saga::object
                    implements_all saga::monitorable
    {
        CONSTRUCTOR      (in  string          url,
                        in  array<flags,1>    flags,
                        in  session           session );
                        out ns_entry          obj      );

        DESTRUCTOR      (in  ns_entry         obj      );

        // basic properties
        get_url          (out string          url      );
        get_name         (out string          name     );
        get_cwd          (out string          cwd      );

        // navigation/query methods
        read_link        (out string          link     );
        is_dir           (in  int             flags = None,
                        out boolean          test     );
        is_entry         (in  int             flags = None,
                        out boolean          test     );
        is_link          (in  int             flags = None,
                        out boolean          test     );

        // management methods
        copy             (in  string          target,
                        in  array<flags,1>    flags   );
        link             (in  string          target,
                        in  array<flags,1>    flags   );
        move             (in  string          target,
                        in  array<flags,1>    flags   );
    }
}

```

```

    remove      (void                );
    close       (void                );

    // helper methods
    convert      (in  string          template,
                  out string          new_url);
}

class ns_directory : implements-all saga::object
                    implements-all saga::monitorable
                    extends          saga::ns_entry
{
    CONSTRUCTOR    (in  string          url,
                    in  array<flags,1> flags,
                    in  session        session );
                    out ns_directory    obj    );
    DESTRUCTOR     (in  ns_directory    obj    );

    // navigation/query methods
    change_dir     (in  string          dir    );
    list           (in  string          name,
                    out array<string,1> names );
    find           (in  string          pattern,
                    in  array<flags,1> flags,
                    out array<string,1> names );
    read_link      (in  string          name,
                    out string          link   );
    exists         (in  string          name,
                    out boolean         exists );
    is_dir         (in  string          name,
                    in  int             flags = None,
                    out boolean         test  );
    is_entry       (in  string          name,
                    in  int             flags = None,
                    out boolean         test  );
    is_link        (in  string          name,
                    in  int             flags = None,
                    out boolean         test  );

    // manage entries by number
    get_num_entries (out int            num    );
    get_entry       (in  int entry,
                    out string          name   );

    // management methods
    copy           (in  string          source,
                    in  string          target,
                    in  array<flags,1> flags );
    link           (in  string          source,
                    in  string          target,

```

```

        in array<flags,1>    flags );
    move      (in string      source,
              in string      target,
              in array<flags,1> flags );
    remove    (in string      target,
              in array<flags,1> flags );
    touch     (in string      target );
    make_dir  (in string      target,
              in array<flags,1> flags );

    // factory methods
    open      (in string      name,
              in array<flags,1> flags,
              out ns_entry    entry );
    open_dir  (in string      name,
              in array<flags,1> flags,
              out ns_directory dir  );
}
}

```

+-----+

Details:

=====

Definitions:

pathname:

A 'Pathname' as accepted by this specification MUST follow the specification of pathnames as described in section 1.1.3 "Pathnames" of the Document "Namespace Service" of the Grid File System Working Group (GFS-WG) in GGF [1].

directory:

A 'Directory' represent what [1] defines as 'Virtual Directories'.

directory_entry:

A 'directory_entry' or 'Entry' represent what [1] defines as 'Junction'. Note that any type of junction defined there could be used.

Wildcards:

The API supports wildcards where appropriate. Available wildcard patterns are:

- * : matches any string
- ? : matches a single character
- [abc] : matches any of a set of characters
- [a-z] : matches any of a range of characters
- [!abc] : matches none of a range of characters
- [!a-z] : matches none of a range of characters
- {a,bc} : matches any of a set of strings

See the POSIX standard for more details. In the API, wildcards are allowed in all strings where they can be used in the respective shell commands, such as:

```
copy *.txt dir
move *.txt dir
link *.txt dir
ls *.txt
remove *.txt
```

Users are rarely aware that wildcards can be used in unorthodox places, such as:

```
move *.txt dir*
move *
```

The result of such operations is dependend on the order the wildcard expansion is performed, e.g. if 'dir*' expands to 'dir_1 dir_2', all txt files and dir_1 will end up in dir_2.

SAGA implementation MUST support wildcards for all strings where that ambiguity cannot arise, (source for move etc), and MAY support wildcards at all respective string parameters which would accepts wildcards in the respective POSIX shell calls.

For the method calls on ns_entry, NO wildcards are allowed. The methods read_link, exists, is_dir, is_entry, is_link, open and open_dir MUST NOT support wild cards. Flags MUST be applied to all elements of a wildcard expansion, even if that raises an exception for any reasons.

interface ns_directory

ns_directory defines two sets of methods: one set to navigate in the namespace hierarchy (e.g. cd, ls, find, ...), and one set to handle entities in the namespace (e.g. copy(), move(), open(), ...)

Methods for creating / destroying the ns_directory object

- CONSTRUCTOR

Purpose: create the object

Format: CONSTRUCTOR (in string url,
in array<flags,1> flags,
in Session session,
out ns_directory obj)

Inputs: url: initial working dir
flags: open mode
session: session handle for
object creation

Outputs: obj: the newly created object

Notes: - the session handle defaults to the SAGA
default session handle if not explicitly
specified
- the default flag set is (None)

- DESTRUCTOR

Purpose: destroy the object

Format: DESTRUCTOR (in ns_directory obj)

Inputs: obj: the object to destroy

Outputs: none

Methods for navigation in the namespace hierarchy:

- change_dir

Purpose: change the working directory

Format: change_dir (in string dir);

Inputs: dir: directory to change to

Outputs: none

Throws: BadParameter
DoesNotExist
IncorrectState

Notes: - dir can be relative to pwd
- similar to the 'cd' command in Unix shells,
as defined by POSIX

- list

Purpose: list entries in this directory

Format: list (in string name,
out array<string,1> names);

Inputs: name: name to list

Outputs: names: array of names existing in
the directory

Throws: BadParameter
DoesNotExist
IncorrectState

-
- Notes:
- similar to 'ls' in Unix shells, as defined by POSIX
 - name can be relative to pwd
 - if name is NULL, the current directory contents is listed.
- find
- Purpose: find entries in this directory and below
- Format: find (in string pattern,
in array<flags,1> flags,
out array<string,1> names);
- Inputs: pattern: pattern for names of
entries to be found
- flags: flags defining the operation
modus
- Outputs: names: array of names matching the
pattern
- Throws: BadParameter
IncorrectState
- Notes:
- similar to 'find' in Unix shells, as defined by POSIX, but limited to the -name option.
 - pattern MUST not include a path element
 - the find operates recursively below the current working directory if the Recursive flag is specified (default)
 - find does not follow symbolically linked directories, unless the FollowSymbolic flag is specified
 - find does list symbolically linked entries with matching name
 - the pattern follows the standard unix shell wildcard specification, as defined by POSIX
 - the matching entries returned are complete relative (to cwd) path names.
 - default flags set is 'Recursive' (1)
- read_link
- Purpose: returns the name of the link target
- Format: read_link (in string name,
out string link);
- Inputs: name: name to be resolved
- Outputs: link: resolved name
- Throws: BadParameter
DoesNotExist
IncorrectState
- Notes:
- link may be relative or absolute depending on underlying implementation. However, the returned name MUST be sufficient to access the target entry
 - resolves one link level only
 - name can be relative to pwd

-
- inspired by 'ls -L' command in Unix shells, as defined by POSIX
 - exists
 - Purpose: returns true if entry exists, false otherwise
 - Format: exists (in string name, out boolean exists);
 - Inputs: name: name to be tested for existence
 - Outputs: exists: boolean indicating existence of name
 - Throws: BadParameter
IncorrectState
 - Notes:
 - name can be relative to pwd
 - as in 'test -e' in Unix shells, as defined by POSIX
 - is_dir
 - Purpose: tests name for being a directory
 - Format: is_dir (in string name, in int flags, out boolean test);
 - Inputs: name: name to be tested
flags: flags for operation
 - Outputs: test: boolean indicating if name is a directory
 - Throws: BadParameter
DoesNotExist
IncorrectState
 - Notes:
 - returns true if entry is a directory, false otherwise
 - name can be relative to pwd
 - flag can be set to DeReference, default is None
 - as in 'test -d' in Unix shells, as defined by POSIX
 - is_entry
 - Purpose: tests name for being a ns_entry
 - Format: is_entry (in string name, in int flags, out boolean test);
 - Inputs: name: name to be tested
flags: flags for operation
 - Outputs: test: boolean indicating if name is a ns_entry
 - Throws: BadParameter
DoesNotExist
IncorrectState
 - Notes:
 - returns true if the entry is a entry, false otherwise (although a ns_dir IS_A ns_entry,

```
        false is returned on a test on a ns_dir)
    - name can be relative to pwd
    - flag can be set to DeReference, default is
      None
    - as in 'test -f' in Unix shells, as defined
      by POSIX

- is_link
  Purpose: tests name for being a link
  Format:  is_link          (in string name,
                           in int   flags,
                           out boolean test);

  Inputs:  name:           name to be tested
           flags:          flags for operation
  Outputs: test:           boolean indicating if name
                           is a link

  Throws:  BadParameter
           DoesNotExist
           IncorrectState

  Notes:   - returns true if the entry is a link, false
            otherwise
            - name can be relative to pwd
            - flag can be set to DeReference, default is
              None
            - as in 'test -l' Unix shells, as defined
              by POSIX

- get_num_entries
  Purpose: gives the number of entries in the directory
  Format:  get_num_entries  (out int   num);
  Inputs:  none
  Outputs: num:             number of entries in the
                           directory

  Throws:  IncorrectState

  Notes:   - can be used for iteration through large
            directories (see get_entry)
            - at the time of using the result of this call,
              the actual number of entries may already have
              changed (no locking is implied)

- get_entry
  Purpose: gives the name of an entry in the directory
           based upon the enumeration defined by
           get_num_entries
  Format:  get_entry        (in int   entry,
                           out string name);
  Inputs:  entry:           index of entry to get
  Outputs: name:            name of entry at index
  Notes:   - '0' is the first entry
            - there is no sort order implied by the
              enumeration, however an underlying
```

```

    implementation MAY choose to sort the entries
- subsequent calls to get_entry and/or
  get_num_entries may return inconsistent data
  if there is no locking or state tracking in
  the underlying implementation
- can be used for iteration through large
  directories

```

Methods for operation on namespace entities:

- copy

Purpose: copy the entry to another part of the namespace

```

Format: copy          (in string      source,
                      in string      target,
                      in array<flags,1> flags);

```

```

Inputs: source:      name to copy
       target:      name to copy to
       flags:       flags defining the operation
                   modus

```

Outputs: none

```

Throws: BadParameter
       DoesNotExist
       IncorrectState

```

```

Notes: - if the target is a directory the source entry
       is copied into the directory
       - it is an error if the source is a directory
         and the 'Recursive' flag is not set.
       - if the target already exists, it will be
         overwritten if the 'Overwrite' flag is set,
         otherwise it is an error
       - default flags set is empty (0)
       - both source and target can be relative to pwd
       - similar to the 'cp' command in Unix shells,
         as defined by POSIX

```

- link

Purpose: create a symbolic link from the source entry to the target entry so that any reference to the target refers to the source entry

```

Format: link          (in string      source,
                      in string      target,
                      in array<flags,1> flags);

```

```

Inputs: source:      name to link
       target:      name to link to
       flags:       flags defining the operation
                   modus

```

Outputs: none

```

Throws: BadParameter
       DoesNotExist
       IncorrectState

```

```

Notes: - if the target is a directory the source entry

```

-
- is linked into the directory.
 - if the target already exists, it will be overwritten if the 'Overwrite' flag is set, otherwise it is an error
 - default flag set is empty (0)
 - both source and target can be relative to pwd
 - similar to the 'ln -s' command in Unix shells, as defined by POSIX
 - move
 - Purpose: rename source to target, or move source to target if target is an directory.
 - Format: move (in string source, in string target, in array<flags,1> flags);
 - Inputs: source: name to move
target: name to move to
flags: flags defining the operation
modus
 - Outputs: none
 - Throws: BadParameter
DoesNotExist
IncorrectState
 - Notes:
 - if the target already exists, it will be overwritten if the 'Overwrite' flag is set, otherwise it is an error
 - moving '.' is not allowed, and throws a BadParameter exception
 - default flag set is empty (0)
 - both source and target can be relative to pwd
 - similar to the 'mv' command in Unix shells, as defined by POSIX
 - remove
 - Purpose: removes the entry
 - Format: remove (in string target, in array<flags,1> flags);
 - Inputs: target: entry to be removed
 - Outputs: none
 - Throws: BadParameter
DoesNotExist
IncorrectState
 - Notes:
 - if the entry is a directory the 'Recursive' flag MUST be set or an exception will be raised
 - default flag set is empty (0)
 - removing '.' is not allowed, and throws a BadParameter exception
 - target can be relative to pwd
 - similar to the 'rm' command in unix shells, as defined by POSIX

- close
 - Purpose: closes the object
 - Format: close (void);
 - Inputs: none
 - Outputs: none
 - Throws: IncorrectState
 - Notes:
 - IncorrectState is thrown if the object was closed before
 - any subsequent method call on the object MUST also raise IncorrectState (apart from DESTRUCTOR)

- touch
 - Purpose: creates a new entry, or updates access time
 - Format: touch (in string target);
 - Inputs: target: target to create
 - Outputs: none
 - Throws: BadParameter
IncorrectState
 - Notes:
 - target can be relative to pwd
 - similar to the 'touch' (1) call, as defined by POSIX
 - the target is created as ns_entry if it does not exist
 - the last access time of the target MUST be updated if the target exists
 - if the target string is NULL, the cwd is touched

- make_dir
 - Purpose: creates a new directory
 - Format: make_dir (in string target,
in array<flags,1> flags);
 - Inputs: target: directory to create
 - Outputs: none
 - Throws: BadParameter
AlreadyExists
IncorrectState
 - Notes:
 - if the parent directory or directories do not exist, 'CreateParents' flag MUST be set or an exception will be raised. If set, the parent directories are created as well
 - an exception MUST be raised if the directory already exists
 - default flag set is empty (0)
 - target can be relative to pwd
 - similar to the 'mkdir' (2) call, as defined by POSIX

- open_dir

-
- Purpose: creates a new ns_directory instance
- Format: open_dir (in string name,
in array<flags,1> flags,
out ns_directory dir);
- Inputs: name: directory to open
flags: flags defining the operation
modus
- Outputs: dir: opened directory instance
- Throws: BadParameter
DoesNotExist
IncorrectState
- Notes: - the cwd of the new dir object instance is set to <name>
- target can be relative to pwd
- similar to the 'opendir' (3) call in Unix, as defined by POSIX
- open
- Purpose: creates a new ns_entry instance
- Format: open (in string name,
in array<flags,1> flags,
out ns_entry entry);
- Inputs: name: entry
flags: flags defining the operation
modus
- Outputs: entry: opened entry instance
- Throws: BadParameter
DoesNotExist
IncorrectState
- Notes: - target can be relative to pwd
- similar to the 'open' (2) call in Unix, as defined by POSIX
- convert
- Purpose: converts url into a different name space
- Format: convert (in string template,
out string new_url);
- Inputs: template: template for new url
- Outputs: new_url: translated url
- Notes: - the method translates a url into a different URL which MUST point at the same endpoint
- the form of the new URL is determined by the given template
- the template gives the part of the new url which is known.
- the template has to present a complete URL _beginning_, e.g. "http://" or "http://host.net:123/", but not "host.net".
- if convert cannot determine a url in the name

space given by the template which points to the same endpoint, an empty string is returned.

+-----+

Examples:

=====

The interfaces are not implemented directly - for more examples, check out the physical and logical file specifications.

Example: provide recursive directory listing for a given directory

Note: - check for '.' and '..' recursion are left as an exercise to the reader...
 - string operations and printf statements are obviously simplified...

```
-----
// c++ example
std::string indent (int indent)
{
    std::string s = " ";

    for (int i = 0; i < indent; i++, s += " ");

    return (s);
}

void list_dir (std::string & url,
               int          indent = 0)
{
    try
    {
        // create directory and iterate over entries
        saga::ns_dir dir (url);

        printf ("\n%s ---> %s\n", indent (indent), url);

        for ( int i = 0; i < dir.get_num_entries (); i++ )
        {
            char    type = '?';
            string  info = "";

            // get name of next entry
            string  name = dir.get_entry (i);

            // get type and other infos
            if ( dir.is_link (name) )
```

```

    {
        if (dir.exists(dir.read_link (name))){info="---> ";}
        else                               {info="-|-> ";}
        info += dir.read_link (name);
        type  = 'l';
    }
    else if (dir.is_entry(name)){ type = 'f';           }
    else if (dir.is_dir  (name)){ type = 'd'; info = "/";}

    printf ("%s > %3d - %s - %s%s\n", indent (indent), i+1,
                                           type, name, info);

    // recursion on directories
    if ( dir.is_dir (name) )
    {
        list_dir (name, indent++);
    }
}

printf ("\n%s <--- %s\n", indent (indent), url);
}

// catch all errors - see elsewhere for better examples
// of error handling in SAGA
catch ( const saga::exception & e )
{
    std::cerr << "Oops! SAGA error: "
               << e.what () << std::endl;
}

return;
}

```

+-----+

References:

=====

- [1] http://forge.gridforum.org/projects/gfs-wg/document/\RNS-Proposed_Final_Draft-v1.10/en/10

+-----+

Notes, Issues and Known Limitations:

=====

A useful extension to the presented interface is a find like

method. However, the flags and options to find (1), are manyfold, and it currently it is unclear how a good mapping to an `_simple_ SAGA API` call might look like.

An Directory can be seen as a container of `directory_entries`, which can be Files, Links, Directories etc. That notion is not reflected in this version of the interface, since no call is taking such entities as argument, or is returning such entities. However, a later version of this interface may introduce this distinction if necessary - it needs then to be reflected in all classes implementing this interface.

In the current version, it is not possible to (e.g.) copy Files w/o creating a directory first. That seems in particular cumbersome if the source and target namespace of the file copy are different. However, we think that the presented approach is more coherent than the alternatives.

Similarly, 'stat' like calls seem (semantically) to specific to the specific namespace incarnation to get included in this rather generic specification.

The notion of security, permissions, ACLs, ownership etc. is missing from this version of the spec, but is crucial to it's usability and acceptance. It will get added as soon as there is an agreement on security in the SAGA API in general.

The URL Problem (see intro):

In order to settle the issue for the SAGA-CORE-API (NOT forever), we propose the following solution:

- A URL utility class seems useful, but would increase the syntactic load on the current set of methods (either allow all methods with strings `_and_ urls`, or require transformation from strings to urls `_always_` - both seems suboptimal).
- hence NO URL utility class is mandated.
- a `NSEntry` utility call is introduced:

```
convert (in  string target,
        out string new_url);
```

Usage examples:

```
saga::file f ("ftp://ftp.net.org/pub/data/file.txt");
std::string n1 = f.convert ("ftp://ftp.net.org:123/");
std::string n2 = f.convert ("ftp://alias.net.org/data/");
```

```
std::string n3 = f.convert ("http://www.net.org");
```

- that call tries to translate URLs, and if unsuccessful returns an empty string.

I know that is 'somewhat' of a compromise. A full fledged URL class certainly would be useful, and Java for example already has it. I would expect that most SAGA implementations will provide one anyway. But as we could not reach consensus on that, the topic seems to need more discussion, and more time which I wthink we don't have for stabilizing the API.

+-----+

```

+-----+
          #####
          #      #      #      #####
          #      #      #      #
          #####    #      #      #####
          #      #      #      #
          #      #      #      #
          #      #      #####    #####
+-----+

```

Summary:

=====

The ability to access files regardless of their location is central to many of the SAGA use cases (see below). This interface addresses the most common operations detailed in the use cases.

The interfaces are syntactically and semantically POSIX oriented, but also borrow some ideas from the GridFTP specification, which is nowadays widely used for remote data access.

Please note that the interactions with files as opaque entities (as entries in file name spaces) are covered by the `name_space` package. The interfaces presented here supplement the namespace package with operations for the reading and writing of files.

```

+-----+

```

Specification:

=====

```

package saga.file
{
    enum flags
    {
        Unknown      = -1, // same as in name_space::flags
        None          = 0, // same as in name_space::flags
        Overwrite     = 1, // same as in name_space::flags
        Recursive     = 2, // same as in name_space::flags
        FollowSymbolic = 4, // same as in name_space::flags
        Create        = 8, // same as in name_space::flags
        Excl          = 16, // same as in name_space::flags
    }
}

```

```

    Lock          = 32, // same as in name_space::flags
    CreateParents = 64, // same as in name_space::flags
    DeReference   = 128, // same as in name_space::flags
    Truncate      = 256,
    Append        = 512,
    Read          = 1024,
    Write         = 2048,
    ReadWrite     = 4096,
    Binary        = 8192
}

enum seek_mode
{
    Unknown    = -1,
    Start      = 1,
    Current    = 2,
    End        = 3
}

struct ivec
{
    long      offset; // position of data to be r/w
    long      length; // number of bytes to be r/w
    array<byte,1> buffer; // data to be r/w
    long      result; // number of bytes r/w
}

class directory : implements-all saga::object
                  implements-all saga::monitable
                  extends         saga::ns_directory
{
    CONSTRUCTOR (in string      url,
                  in session    session,
                  in int        flags = None,
                  out directory dir   );
    DESTRUCTOR  (in directory   dir   );

    get_size    (in string      name,
                  out long      size   );
    is_file     (in string      name,
                  in int        flags = None,
                  out boolean    test   );

    open_dir    (in string      name,
                  in int        flags = None,
                  out directory dir   );

    open        (in string      name,
                  in int        flags = Read,

```

```

        out file          file    );
    }

class file : implements-all saga::object
            implements-all saga::monitorable
            extends          saga::ns_entry
{
    CONSTRUCTOR (in    string          url,
                 in    session        session,
                 in    int             flags = Read,
                 out    file           file    );
    DESTRUCTOR  (in    file           file    );

    read        (in    long            len_in,
                 inout array<byte,1>    buffer,
                 out    long            len_out );
    write       (in    long            len_in,
                 in    array<byte,1>    buffer,
                 out    long            len_out );
    seek        (in    long            offset,
                 in    seek_mode        whence,
                 out    long            position );

    read_v      (inout array<ivec>      ivec    );
    write_v     (inout array<ivec>      ivec    );

    read_p      (in    string          pattern,
                 inout array<byte,1>    buffer,
                 out    long            len_out );
    write_p     (in    string          pattern,
                 in    array<byte,1>    buffer,
                 out    long            len_out );

    modes_e     (out    array<string,1> emodes );
    read_e      (in    string          emode,
                 in    string          spec,
                 inout array<byte,1>    buffer,
                 out    long            len_out );
    write_e     (in    string          emode,
                 in    string          spec,
                 in    array<byte,1>    buffer,
                 out    long            len_out );

    // Metric:
    //     name:  read
    //     desc:  fires if a file gets readable
    //     mode:  Read
    //     unit:  1
    //     type:  Bool

```

```

//      value: True
//
//      name:  write
//      desc:  fires if a file gets writable
//      mode:  Read
//      unit:  1
//      type:  Bool
//      value: "True"
//
//      name:  error
//      desc:  fires if a file gets writable
//      mode:  Read
//      unit:  1
//      type:  Bool
//      value: "True"
}
}

```

+-----+

Details:

=====

The current description covers the ubiquitous open/close/read/write/seek pattern, which is present in the vast majority of remote file access providers.

class directory

This class represents a directory containing physical files.

- CONSTRUCTOR

Purpose: open the directory

Format: CONSTRUCTOR (in string url,
in int flags,
in session session,
out directory dir)

Inputs: url: location of directory
flags: mode for opening
session: session to associate the
object with

Outputs: dir: the newly created object

Throws: BadParameter
DoesNotExist

Notes: - the session handle defaults to the SAGA
default session handle if not explicitly
specified
- the default flag set is (None)

```
- DESTRUCTOR
Purpose: destroy the directory object
Format:  DESTRUCTOR          (in  directory dir)
Inputs:  dir:                 the object to destroy
Outputs: none
```

Methods giving information about files:

```
- get_size
Purpose: returns the number of bytes in the file
Format:  get_size            (in  string  name,
                             in  int     flags,
                             out long    size);

Inputs:  name:               name of file to inspect
Outputs: size:               number of bytes in the file
Throws:  BadParameter
         DoesNotExist

Notes:   - as 'st_size' field in the Unix call 'stat',
          as defined by POSIX

- is_file
Purpose: alias for is_entry in ns_directory
```

Factory like methods for creating objects:

```
- open_dir
Purpose: creates a directory object
Format:  open_dir            (in  string  name,
                             in  int     flags,
                             out directory dir)

Inputs:  name:               name of directory to open
         flags:               flags definition operation
                             modus

Outputs: dir:                 opened directory instance
Throws:  BadParameter
         DoesNotExist

Notes:   - creates a new directory instance
          - currently there are no supported flags (FIXME)
          - default flag set is empty (None)
          - similar to opendir (3), as defined by POSIX

- open
Purpose: creates a new file instance
Format:  open                (in  string  name,
                             in  int     flags = Read,
                             out file    file);
```

Inputs: name: file to be opened
 flags: flags definition operation
 modus

Outputs: file: opened file instance

Throws: BadParameter
 DoesNotExist

Notes: - if the file does not exist, it is created if the 'Create' flag is given, otherwise it is an error
 - it is an error if the file exists and both the 'Create' and the 'Excl' flag are given. Otherwise the 'Excl' flag is ignored
 - the file is truncated to length 0 on the open operation if the 'Trunc' flag is given
 - the file is opened in append mode if the 'Append' flag is given (a seek (0, End) is performed after the open)
 - the file is locked on open if the 'Lock' flag is given. If the file is already in a locked state, the open will fail and a descriptive error will be issued. If a file is opened in locked mode, any other open on that file MUST fail, with no respect to the given flags. Note that a file can be opened in normal mode, and then in locked mode, w/o an error getting raised. The lock will get removed on destruction of the file object (that is on close). If an implementation does not support locking, an descriptive error MUST get issued if the 'Lock' flag is given.
 - default flag set is (Read).
 - similar to the open (2) call in Unix, as specified by POSIX

class file:

This class represents an open file descriptor for read/write operations on a physical file. Its concept is similar to the file descriptor returned by the open (2) call in Unix.

- CONSTRUCTOR

Purpose: create the obj

Format: CONSTRUCTOR (in string url,
 in int flags,
 in session session,
 out file obj)

Inputs: url: location of file
 flags: mode for opening
 session: session to associate the

```

                                object with
Outputs: obj:                  the newly created object
Throws:  BadParameter
        DoesNotExist
Notes:   - the session handle defaults to the SAGA
          default session handle if not explicitly
          specified
          - the default flag set is (Read)

- DESTRUCTOR
  Purpose: destroy the object
  Format:  DESTRUCTOR          (in file          obj)
  Inputs:  obj:                the object to destroy
  Outputs: none

- read
  Purpose: reads data from an open file
  Format:  read                (in long   len_in,
                              in string buffer,
                              out long   len_out);
  Inputs:  len_in:             length of data section to
                              be read
          buffer:              buffer to read into
  Outputs: len_out:            length of data read into
                              buffer
  Throws:  BadParameter
          ReadError
  Notes:   - reads up to len_in bytes from the file into
            the buffer.
            - the actually number of bytes read into buffer
              is returned in len_out. It is not an error
              to read less bytes than requested, or in fact
              zero bytes, eg. at the end of the file.
            - the file pointer is positioned at the end of
              the byte area read during this call.
            - the given buffer must be large enough to
              store up to len_in bytes
            - similar to the read (2) call in Unix, as
              specified by POSIX

- write
  Purpose: write data into an open file
  Format:  write                (in long   len_in,
                              in string buffer,
                              out long   len_out);
  Inputs:  len_in:             number of bytes to be
                              written
          buffer:              bytes to be written
  Outputs: len_out:            number of bytes written
  Throws:  BadParameter
          WriteError

```


- Notes:
- writes up to len_in bytes from buffer into the file at the current file position.
 - the file pointer is positioned at the end of the byte area written during this call.
 - similar to the write (2) call in Unix, as specified by POSIX
- seek
- Purpose: reposition the file pointer
- Format: seek (in long offset, in seek_mode whence, out long position);
- Inputs: offset: offset in bytes to move pointer
whence: offset is relative to 'whence'
- Outputs: position: position of pointer after seek
- Throws: BadParameter
WriteError
- Notes:
- seek repositions the file pointer for subsequent read, write and seek calls.
 - initially (after open), the file pointer is positioned at the beginning of the file, unless the 'Append' flag was given - then the initial position is the end of the file.
 - the repositioning is done relative to the position given in 'Whence', so relative to the Begin or End of the file, or to the current position.
 - the file pointer can be positioned after the end of the file w/o extending it. Reads behind EOF return Zeros.
 - similar to the lseek (2) call in Unix, as specified by POSIX.
- read_v
- Purpose: gather/scatter read
- Format: read_v (inout array<ivec> ivec);
- Inputs/
- Outputs: ivec: array of ivec structs defining start (offset) and length (length) of each individual read, and buffer to read into, and integer to store result into.
- Throws: BadParameter
ReadError
- Notes:
- the behaviour of each individual read is as in the normal read method.
 - the lengths returned also correspond to those

```

        of the normal read method.
    - an exception is thrown if any of the
      individual reads detects a condition which
      would raise an exception for the normal
      read method.
    - similar to the readv (2) call in Unix, as
      specified by POSIX

- write_v
  Purpose: gather/scatter write
  Format:  write_v          (inout array<ivec>  ivec);
  Inputs/
  Outputs: ivec:           array of ivec structs
                           defining start (offset) and
                           length (length) of each
                           individual write, and
                           buffers containing the data
                           to write

  Throws:  BadParameter
           WriteError

  Notes:   - the behaviour of each individual write is as
            in the normal write method.
            - the lengths returned also correspond to those
              of the normal write method.
            - an exception is thrown if any of the
              individual writes detects a condition which
              would raise an exception for the normal write
              method.
            - similar to the writev (2) call in Unix, as
              specified by POSIX

```

+-----+

Examples:

=====

Example: open a file. if its size is > 10, then read the first 10 bytes into a string, and print it.

```

-----
// c++ example
void head (const char* url)
{
    try {
        // get type and other infos
        saga::file my_file (url);

        off_t size = my_file.get_size ();

```

```

    if ( size > 10 )
    {
        char    buffer[11];
        long    buflen;

        my_file.read (10, buffer, &buflen);

        if ( buflen == 10 )
        {
            printf ("head: '%s'\n", buffer);
        }
    }

    // catch any possible error - see elsewhere for better
    // examples of error handling in SAGA
    catch ( const saga::exception & e )
    {
        std::cerr << "Oops! SAGA error: " + e.what () + std::endl;
    }

    return;
}
-----

```

+-----+

Notes:

=====

A 'stat' like method is not yet specified; the form of such an interface needs further consideration. However, the 'get_size' method provides the most frequent and well defined file size for now (the call may be deprecated when a stat specification is available).

future API version may have something like:

```

stat          (in string          name,
               out struct         statinfo );

```

+-----+

```

+-----+
#
#      ####  ####  #      ####  ##  #
#      #  #  #  #  #  #  #  #  #  #  #
#      #  #  #  #  #  #  #  #  #  #
#      #  #  #  ###  #  #  #  #####  #
#      #  #  #  #  #  #  #  #  #  #
#####  #####  #####  #  #####  #  #####

          #####
          #      #  #      #####
          #      #  #      #
          #####  #  #      #####
          #      #  #      #
          #      #  #      #
          #      #  #####  #####

+-----+

```

Summary:

=====

There are a number of replica catalogue systems implemented or in development. This API is the intersection of features common to these implementations. (TODO: enumerate these systems.)

Please note that the interactions with logical files as opaque entities (as entries in logical file name spaces) are covered by the name_space package. The interfaces presented here supplement the namespace package with operations for operating on entries in replica catalogues.

```

+-----+

```

Specification:

=====

```

package saga.logical_file
{
    enum open_dir_flags
    {
        /* Placeholder */
    }

    enum open_flags
    {

```

```

    Unknown    = -1,
    None       = 0,
    Create     = 1,
    Excl       = 2,
    Lock       = 4,
    Truncate   = 8,
    Append     = 16,
    Read       = 32,
    Write      = 64,
    ReadWrite  = 128
}

class logical_directory : implements-all saga::object
                        implements-all saga::attribute
                        implements-all saga::monitorable
                        extends          saga::ns_directory
{
    CONSTRUCTOR      (in  string          url,
                     in  array<open_dir_flags,1> flags,
                     in  session          session,
                     out logical_directory dir);
    DESTRUCTOR       (in  logical_directory dir);

    // add for inspection
    is_file          (in  string          name,
                     out boolean         test);

    // open methods
    open_dir         (in  string          name,
                     in  array<open_dir_flags,1> flags,
                     out logical_directory dir);

    open             (in  string          name,
                     in  array<open_flags,1> flags,
                     out logical_file     file);
}

class logical_file : implements-all saga::object
                   implements-all saga::attribute
                   implements-all saga::monitorable
                   extends          saga::ns_entry
{
    CONSTRUCTOR      (in  string          url,
                     in  array<open_flags,1> flags,
                     in  session          session,
                     out logical_file     file);
    DESTRUCTOR       (in  logical_file     file);
}

```

```

    add_location      (in string          name);
    remove_location  (in string          name);
    update_location   (in string          name_old,
                      in string          name_new);
    list_locations    (out array<string,1> names);
    replicate         (in string          name);
}
}

```

+-----+

Details:

=====

class logical_directory

This class represents a container for logical files in a logical file catalog. It allows traversal of the catalogs name space, and the manipulation and creation (open) of logical files in that name space.

Factory like methods for creating objects (see note in saga::name_space specification)

- CONSTRUCTOR

Purpose: create the object

Format: CONSTRUCTOR (in string url,
in array<open_dir_flags,1>
flags,
in session session,
out logical_directory
obj)

Inputs: url: location of directory
flags: mode for opening
session: session to associate with
the object

Outputs: obj: the newly created object

Throws: BadParameter
DoesNotExist

Notes: - the session handle defaults to the SAGA
default session handle if not explicitly
specified
- currently there are no supported flags
- the default flag set is (None)
- similar to opendir (3), as defined by POSIX

- DESTRUCTOR

Purpose: destroy the object

Format: DESTRUCTOR (in logical_directory obj)
 Inputs: obj: the object to destroy
 Outputs: none

- is_file
 Purpose: alias for is_entry in ns_directory
- open_dir
 Purpose: creates a new logical_directory instance
 Format: open_dir (in string name,
 in array<open_dir_flags,1>
 flags,
 out logical_directory dir);
 Inputs: name: name of directory to open
 flags: flags definition operation
 modus
 Outputs: dir: opened directory instance
 Throws: BadParameter
 DoesNotExist
 Notes: - notes to logical_directory constructor apply
- open
 Purpose: creates a new logical_file instance
 Format: open (in string name,
 in array<open_flags,1>
 flags,
 out logical_file file);
 Inputs: name: file to be opened
 flags: flags definition operation
 modus
 Outputs: file: opened file instance
 Throws: BadParameter
 DoesNotExist
 Notes: - notes to logical_file constructor apply

class logical_file:

This class provides means to handle the contents of Logical Files.

That contents consists of strings representing locations of physical files associated with the logical file. In general, these locations could be logical files as well. In fact, they are usually handled as opaque strings, and no assumption about validity or the nature of the target of the location is made. Exception: see the replicate and add_location method description.

- CONSTRUCTOR

Purpose: create the object

Format: CONSTRUCTOR (in string url,
in array<open_dir_flags,1>
flags,
in session session,
out logical_file obj)

Inputs: url: location of directory
flags: mode for opening
session: session to associate with
the object

Outputs: obj: the newly created object

Throws: BadParameter
DoesNotExist

Notes:

- the session handle defaults to the SAGA default session handle if not explicitly specified
- if the file does not exist, it is created if the 'Create' flag is given, otherwise it is an error
- it is an error if the file exists and both the 'Create' and the 'Excl' flag are given. Otherwise the 'Excl' flag is ignored
- the file is truncated to length 0 on the open operation if the 'Trunc' flag is given. For logical files that means: no physical file location is associated with the logical file.
- the file is opened in append mode if the 'Append' flag is given. For logical files that means: newly added physical file locations are appended to the set of known locations.
- the file is locked on open if the 'Lock' flag is given. If the file is already in a locked state, the open will fail and a descriptive error will be issued. If a file is opened in locked mode, any other open on that file MUST fail, with no respect to the given flags. Note that a file can be opened in normal mode, and then in locked mode, w/o an error getting raised. The lock will get removed on destruction of the file object (that is on close). If an implementation does not support locking, an descriptive error MUST get issued if the 'Lock' flag is given.
- default flag set is (Append)
Note that the default is different to the class saga::file.
- similar to the open (2) call in Unix, as specified by POSIX

- DESTRUCTOR
 - Purpose: destroy the object
 - Format: DESTRUCTOR (in logical_file obj)
 - Inputs: obj: the object to destroy
 - Outputs: none

- add_location
 - Purpose: add a name to the location set
 - Format: add_location (in string name);
 - Inputs: name: location to add to set
 - Outputs: none
 - Throws: BadParameter
AlreadyExists
 - Notes:
 - this methods adds a given string to the set of locations associated with the logical file.
 - if the location is already in the set, no error is issued.
 - the implementation may choose to interpret the locations associated with the logical file instance. It may return an error indicating an invalid location if it is unable or unwilling to handle that specific location.
 - the documentation MUST specify how valid location are constructed.

- remove_location
 - Purpose: remove a name from location set
 - Format: remove_location (in string name);
 - Inputs: name: location to remove from set
 - Outputs: none
 - Throws: BadParameter
DoesNotExist
 - Notes:
 - this method removes a given string from the set of locations associated with the logical file.
 - if the location is not in the set of locations, an error is issued.
 - if the set of locations is empty after that operation, the logical file object is still a valid object (see replicate methods description).

- update_location
 - Purpose: change a name in location set
 - Format: update_location (in string name_old,

```

        in string name_new);
Inputs:  name_old      location to be updated
        name_new      update for location
Outputs: none
Throws:  BadParameter
        DoesNotExist
Notes:   - this method removes a given string from the
          set of locations associated with the logical
          file, and adds a new string.
          - if the old location is not in the set of
            locations, an error is issued.

- list_locations
Purpose: list the locations in the location set
Format:  list_locations (out array<string,1> names);
Inputs:  none
Outputs: names:         array of locations in set
Notes:   - this method returns an array of strings
          containing the complete set of locations
          associated with the logical file.
          - an empty array returned is not an error - see
            description to the remove_location method.

- replicate
Purpose: replicate a file from any of the known
        locations to a new location, and add the new
        location to the location set on success.
Format:  replicate      (in string name);
Inputs:  name:          location to replicate to
Outputs: none
Throws:  BadParameter
Notes:   - the method requests a two step operation:
          1) copy an entity at any of the locations
             associated with the logical file to the
             given string, which represents a new
             location.
          2) perform an add_location for the given
             string.
          - the method is not required to be atomic, but:
          - the method is required to be either
            successfull in both steps, or to issue an
            error indicating if both methods failed, or if
            only one of the methods succeeded (leaving the
            system in an inconsistent state).
          - a replicate call on an instance with empty
            location set results in an error.
          - this methods requires the implementation of
            the class to interpret the locations
            associated with the logical file instance. If
            that is impossible, an error indicating an

```

invalid location must be issued.

+-----+

Examples:

=====

```
// c++ example
int main ()
{
    saga::logical_file lf ("lfn://remote.catalog.net/tmp/file1");

    lf.replicate ("gsiftp://localhost.net/tmp/file.rep");

    saga::file f ("gsiftp://localhost.net/tmp/file.rep");
    std::cout << "size of local replica: "
                << f.get_size ()
                << std::endl;
}
```

+-----+

Notes:

=====

It is recommended to interpret the locations associated with logical files with valid locations for `saga::file`, and to have the implementation using `saga::file`. That helps to program coherently with the `saga::name_space`, `saga::file` and `saga::logical_file` packages.

`logical_file` and `logical_directory` should implement the `saga::attribute` interface.

`logical_directory` should implement a `find` method searching on `saga::attribute`.

+-----+

```

+-----+
          #####
            #   ####   #####
            #  #     #  #   #
            #  #     #  #####
#         #  #     #  #   #
#         #  #     #  #   #
#####   #####   #####
+-----+

```

Summary:

=====

Many of the use cases provided to the SAGA-RG had either explicit or implied requirements for submitting jobs to grid resources, and for monitoring and controlling these submitted jobs.

This API provides an interface for submitting jobs to a grid resource, either in batch mode, or in an interactive mode. It also provides APIs for controlling these submitted jobs (e.g. to terminate, suspend, or signal a running job), and APIs for retrieving status information for both running and completed jobs.

The goals of this API are to provide enough functionality to satisfy the requirements of grid developers according to the "80-20" rule. This API is also intended to incorporate the work of the DRMAA-WG, and to extend the API based on the experience of implementing DRMAA. Much of this specification was taken directly from DRMAA, with many of the differences arising from an attempt to make the job API consistent with the overall SAGA API model. Note [1].

```

+-----+

Specification:
=====

```

```

package saga.resource
{
    // enum job_state = saga::state (from task package)

    class job_definition : implements-all saga::object
                        implements-all saga::attribute
    {
        CONSTRUCTOR (out job_definition obj);
        DESTRUCTOR  (in  job_definition obj);
    }
}

```

```

}

class job : implements-all saga::object
    implements-all saga::attribute
    implements-all saga::monitorable
    extends          saga::task
{
    /* no CONSTRUCTOR */
    DESTRUCTOR        (in  job          job);

    // job inspection
    get_job_id        (out string        job_id);
    get_state          (out state         state);
    get_state_detail   (out string        detail);
    get_job_definition (out job_definition job_def);
    get_stdin          (out opaque        stdin);
    get_stdout         (out opaque        stdout);
    get_stderr         (out opaque        stderr);

    // job management
    suspend            (void);
    resume             (void);
    checkpoint         (void);
    migrate            (in job_definition job_def);
    signal             (in int            signum);
    terminate          (void);
}

class job_service : implements-all saga::object
                    implements-all saga::monitorable
{
    CONSTRUCTOR        (in  session      session,
                        out job_service  service);
    DESTRUCTOR         (in  job_service  service);

    create_job         (in  string        rm,
                        in  job_definition job_def,
                        out job           job);
    run_job            (in  string        rm,
                        in  string        commandline,
                        out opaque        stdin,
                        out opaque        stdout,
                        out opaque        stderr,
                        out job           job);
    list               (out array<string,1> job_ids);
    get_job            (in  string        job_id,
                        out job           job);
    get_self           (out job           job);
}
}

```

+-----+

Details:

=====

```
class job_definition:
```

```
-----
```

This object encapsulates all the attributes which define a job to be run. It has no methods of its own, but implements the 'Attribute' interface in order to provide access to the job properties. The only required attribute in order to perform a valid job submission is the 'JobCmd'. Given the 'JobCmd', a job can be instantiated in many existing back end systems without any further specification.

There should be much overlap between the attributes defined within SAGA and within the JSDL specification. This list, however, will not be complete in cases where the JSDL was deemed more complicated than was required for a simple API (e.g. the notion of JSDL Profiles), or where an attribute was needed to interact with a scheduler, which was not within the stated scope of the JSDL working group (e.g. 'Queue', which is considered a "site" attribute, and thus not relevant to the pure description of a job).

At the end of the description of an attribute there is a bit in parentheses that indicates whether a particular attribute is supported within a particular system. Tokens include DRMAA, JSDL, LSF, OpenPBS, PBSPro, SGE and Condor, and are intended to be extended by members of the working group.

The attributes encapsulated within this class are:

'JobCmd'

- The command to execute. This is the only required attribute. Can be a full pathname, or a pathname relative to the 'JobCwd' as evaluated on the execution host. String. (DRMAA, JSDL, LSF)

'JobArgs'

- Positional parameters for the command. Vector of strings. (DRMAA, JSDL, LSF)

'JobState'

- The job state at submission. jobs can be submitted into a suspend or hold state such that they need manual resume before being considered for scheduling. Valid values are "Hold", "Suspend". If not specified, the job will enter the default "Pending" state. Type String. (DRMAA, LSF)

'JobEnv'

- The set of environment variables which will be exported to the environment of the started job. The string format is "name=value". Vector of strings. (DRMAA, JSDL)

'JobCwd'

- The working directory for the job. If this is a relative path, it will be treated as relative to the users home directory on the system where the job runs. String. (DRMAA, JSDL)

'JobInteractive'

- Run the job in interactive mode. This means that stdio streams will stay connected to the submitter after job submission, and during job execution. The stdio streams are retrieved by calling the getXStream methods of the jobs class. Boolean. (LSF)

'JobStdin'

- The pathname of the standard input file. If this is a relative pathname, it will be treated as relative to the users home directory on the system where the job runs. String. (DRMAA, JSDL, LSF)

'JobStdout'

- The pathname of the standard output file. If this is a relative pathname, it will be treated as relative to the users home directory on the system where the job runs. String. (DRMAA, JSDL, LSF)

'JobStderr'

- The pathname of the standard error file. If this is a relative pathname, it will be treated as relative to the users home directory on the system where the job runs. String. (DRMAA, JSDL, LSF)

'JobContact'

- A set of endpoints describing where to report job completion status, as well as other resource manager defined state transitions. The format of the string will be that of a URI (e.g. fax:+123456789, sms:+123456789, mailto:csmith@platform.com). Vector of strings. (DRMAA (email addresses), LSF (email addresses))

'JobNotification'

- A flag which indicates whether to send notifications to endpoints listed in 'JobContact'. Mostly used to shut off notifications if they are on by default. Boolean. (DRMAA, LSF)

'JobName'

- The job name to be attached to the job submission. String. (DRMAA, LSF)

'JobNative'

- The native specification as described in the DRMAA specification. Note [3]. This value is passed as is to the backend without any meaning or semantics within the SAGA API. String. (DRMAA)

'FileTransfer'

- A list of file transfer directives which can be used to transfer files to the execution host of the job before the job is run, and to transfer files from the execution host of the job when the job completes. Vector of strings. (DRMAA (limited), JSDL (much enhanced), LSF)

The syntax of a file transfer directive is modeled on the LSF syntax, and has the general syntax:

```
"local_file operator remote_file"
```

Both the local_file and the remote_file can be URLs. If they are not URLs, but full or relative pathnames, then the local_file is relative to the host where the submission is executed, and the remote_file is evaluated on the execution host of the job.

The operator is one of the following four:

- '>' - copies the local file to the remote file before the job starts. Overwrites the remote file if it exists.
- '>>' - copies the local file to the remote file before the job starts. Appends to the remote file if it exists.
- '<' - copies the remote file to the local file after the job finishes. Overwrites the local file if it exists.
- '<<' - copies the remote file to the local file after the job finishes. Appends to the local file if it exists.

'JobStartTime'

- The time after which a job is considered for scheduling. Could be viewed as a desired job start time, but that is up to the resource manager. Date/time. (DRMAA, LSF)

'Deadline'

- Specifies a hard deadline after which the resource manager should terminate the job. Date/time. (DRMAA, LSF)

'WallclockHardLimit'

- Specifies a hard limit on the amount of wall clock time in seconds that a job may consume, after which the resource manager should terminate the job. Integer. (DRMAA, JSDL, LSF)

'WallclockSoftLimit'

- Provides an estimate of the amount of wall clock time in seconds which a job will require. This attribute is intended to provide hints to the scheduler. If this time limit is reached, the action taken is specific to the resource manager and its scheduling policies. Integer. (DRMAA, LSF)

'Cputime'

- Estimated job runtime in CPU seconds. The CPU time is aggregated across all processes/threads of the job. Integer. (LSF)

'NumCpus'

- The total number of cpus requested for this job. How the cpus are allocated is determined by the policy of the resource manager, and can possibly be affected by the 'Native' attribute if the resource manager supports it. Integer. (JSDL, LSF)

'Memory'

- Estimated maximum amount of memory that the job requires in Megabytes. The memory usage of the job is aggregated across all processes of the job. Float. (JSDL, LSF)

'ProcessorType'

- Select compatible processor for job submission. The list of allowed values is taken from the JSDL specification `jsdl:ProcessorArchitectureEnumeration`. Note [4]. String. (JSDL)

'OperatingSystem'

- Select compatible operating system for job submission. The list of allowed values is taken from the JSDL specifications `jsdl:OperatingSystemTypeEnumeration`. Note [4]. String. (JSDL)

'HostList'

- A list of host names, or host group names, which can be considered by the resource manager as candidate hosts for the job. Whether or not the job actually ends up running on one of the hosts in the list, is solely at the discretion of the resource manager. Vector of strings. (JSDL, LSF)

'Queue'

- The name of a queue to place the job into. While SAGA itself does not define the semantics of "queue", many back end systems can make use of this attribute. String. (LSF)

- CONSTRUCTOR

Purpose: create the object

Format: CONSTRUCTOR (out job_definition obj)

Inputs: none

Outputs: obj: the newly created object

Notes: - a job_definition is NOT associated to a session, but can be used for job services from different sessions.

- DESTRUCTOR

Purpose: destroy the object

Format: DESTRUCTOR (in job_definition obj)

Inputs: obj: the object to destroy

Outputs: none

class job:

The job provides the manageability interface to a job submitted to a resource manager. There are two general types of methods: those for retrieving job state and information, and those for manipulating the submitted job. The methods intended to manipulate jobs cannot make any guarantees about how the resource manager will effect an action to be taken. Please see note [5].

job implements the 'Attribute' interface, and understands the following attribute names. If not noted otherwise, none of these attributes is available before the job is running, and none is guaranteed to have a value if the job is running or after the job finishes. Also, the attributes can change their values during the lifetime of the job.

'ExecutionHosts'

- The list of host names or IP addresses which were allocated to run this job. Vector of strings.

'Created'

- The time stamp of the job creation in the resource manager (i.e. the submission time). Date/time.

'Started'

-
- The time stamp indicating when the job started running.
Date/time.
 - 'Finished'
 - The time stamp indicating when the job completed.
Date/time.
 - 'Cputime'
 - The number of cpu seconds consumed by the job. The value is aggregated across all processes/threads of the job.
Integer.
 - 'MemoryUse'
 - The current aggregate memory usage in megabytes of the processes of this job, or the memory high water mark when the job is complete. Integer.
 - 'VmemoryUse'
 - The current aggregate virtual memory usage in megabytes of the processes of this job, or the virtual memory high water mark when the job is complete. Integer.
 - 'JobCwd'
 - The current working directory of the job on the execution host(s). This can be used to determine the location of files that are staged using relative file paths. String.
 - 'ExitCode'
 - The process exit code as collected by the wait(2) series of system calls. The exit code is collected from the process which was started from the 'JobCmd' attribute of the jobDefinition object. Is only available in final states, but not guaranteed. Integer.
 - 'Signaled'
 - Indicates whether the job exited due to receipt of a signal. Is only available in final states, but not guaranteed. Boolean.
 - 'Termsig'
 - The signal number which caused the job to exit. Is only available in final states, but not guaranteed. Integer.
 - CONSTRUCTOR
This class has no constructor, and can only be obtained by calling `job_server.create_job()` or `job_server.run_job()`.
 - DESTRUCTOR
Purpose: destroy the object

-
- Format: DESTRUCTOR (in job obj)
 Inputs: obj: the object to destroy
 Outputs: none
- get_job_id
 Purpose: Get the resource managers representation of the job identifier.
 Format: get_job_id (out string job_id);
 Inputs: none
 Outputs: job_id: job identifier string
- get_state_detail
 Purpose: Retrieve details of the current job state.
 Format: get_state_detail (out string detail);
 Inputs: none
 Outputs: deatail: details about the current job state
- Notes: - The SAGA job state model is a very simplified model. The job state details obtained with this calls returns information from any more explicit state model of the back end.
 - BES compliant states SHOULD be returned
 - the format of the state details SHOULD be: "<model>:<state>", with valid models being: "BES", "DRMAA", "GRAM", or implementation specified. The state should be spelled like "UpperCase" (example: "BES:StagingIn").
- get_job_definition
 Purpose: Retrieve the job_definition which was used to submit this job instance.
 Format: get_job_definition (out job_definition job_def);
 Inputs: none
 Outputs: job_def: a job_definition object
- Notes: - There are cases when the job_definition is not available, and thus this object will be null. These include cases when the job might not have been submitted through SAGA, and get_job() was used to retrieve a job, or this state information has been lost (e.g. the client application restarts and the particular SAGA implementation did not persist the information).
- get_stdin
 Purpose: retrieve input stream for a job.
 Format: get_stdin (out opaque stdin)
 Inputs: none

Outputs: stdin: standard input stream for
 the job

Notes: - If the job was submitted as interactive (the
 'JobInteractive' attribute was set at job
 submission), this method retrieves the standard
 input stream for the job. The type of the
 stream is indicated in SIDL as opaque, since
 this type will be rendered differently based
 on the language bindings, and will be made
 concrete in another specification document
 which describes language bindings.

- get_stdout

Purpose: retrieve output stream of job

Format: get_stdout (out opaque stdout)

Inputs: none

Outputs: stdout: standard output stream for
 the job

Notes: - If the job was submitted as interactive (the
 'JobInteractive' attribute was set at job
 submission), this method retrieves the
 standard output stream for the job.

- get_stderr

Purpose: retrieve error stream of job

Format: get_stderr (out opaque stderr)

Inputs: none

Outputs: stderr: standard error stream for
 the job

Notes: - If the job was submitted as interactive (the
 'JobInteractive' attribute was set at job
 submission), this method retrieves the
 standard error stream for the job.

- suspend

Purpose: Ask the resource manager to perform a suspend
 operation on the running job.

Format: suspend ();

Inputs: none

Outputs: none

Notes: - The semantics of suspend, and the action taken
 to suspend a job is resource manager specific.

- resume

Purpose: Ask the resource manager to perform a resume
 operation on the running job.

Format: resume ();

-
- Inputs: none
 Outputs: none
 Notes: - The semantics of resume, and the action taken to resume a job is resource manager specific.
- checkpoint
 Purpose: Ask th resource manager to initiate a checkpoint operation on a running job.
 Format: checkpoint ();
 Inputs: none
 Outputs: none
 Notes: - The semantics of checkpoint, and the actions taken to initiate a checkpoint, are resource manager specific.
- migrate
 Purpose: Ask the resource manager to migrate a running job to another host.
 Format: migrate (in job_definition job_def);
 Inputs: job_def: new job parameters to apply when the job is migrated
 Outputs: none
 Notes: - The call may also be used to change some parameters of a non-finished job (e.g. change runtime limit estimates, etc). The action of migration might change the job identifier within the resource manager.
 - job_def might indicate new resource requirements, for example.
- terminate
 Purpose: Ask th resource manager to terminate a dispatched job
 Format: terminate ();
 Inputs: none
 Outputs: none
 Notes: - the job can be in in Running or Suspended state
 - the semantics of terminate, or the action taken, is specific to the resource manager.
- signal
 Purpose: Ask th resource manager to deliver an arbitrary signal to a dispatched job.
 Format: signal (in int signum);
 Inputs: signum: signal number to be delivered
 Outputs: none

Notes: - The semantics of signal, or the action taken, is specific to the resource manager. There is no guarantee that the signal number specified is valid for the operating system on the execution host where the job is running.

```
class job_service:
```

```
-----
```

The job_service provides an class for job creation and discovery.

- CONSTRUCTOR

Purpose: create the object

Format: CONSTRUCTOR (in session session,
out job_service obj)

Inputs: session: session to associate with
the object

Outputs: obj: the newly created object

Notes: - the session handle defaults to the SAGA
default session handle if not explicitly
specified

- DESTRUCTOR

Purpose: destroy the object

Format: DESTRUCTOR (in job_service obj)

Inputs: obj: the object to destroy

Outputs: none

- create_job

Purpose: create a job instance

Format: create_job (in string rm,
in job_definition job_def,
out job job);

Inputs: rm: rm name or IP address of
the resource manager which
will accept and run the job
job_def: description of job to be
submitted

Outputs: job: a job object representing
the submitted job instance

Throws: IncorrectParameter

Notes: - the returned job is in the New state
- calling run() on the job will submit it to
the resource, and advance its State.

- run_job

Purpose: Run a command synchronously.

Format: run_job (in string rm,
 in string commandline,
 out opaque stdin,
 out opaque stdout,
 out opaque stderr,
 out job job);

Inputs: rm: rm name or IP address of
 the resource manager which
 will accept and run the job

 commandline: the command and arguments
 to be run

Outputs: stdin: IO handle for the running
 jobs standard input stream

 stdout: IO handle for the running
 jobs standard output

 stderr: IO handle for the running
 jobs standard error

 job: a job object representing
 the submitted job instance

Notes: - This is a convenience routine built on the
 create_job method, and is intended to simplify
 the steps of creating a job_definition,
 creating and running the job, and then
 querying the stdio streams.

- list

Purpose: Get a list o jobs which are currently known by
 the resource manager.

Format: list (out array<string,1> job_ids);

Inputs: none

Outputs: job_ids: an array of job identifiers

Notes: - The semantics of which jobs are viewable by the
 calling user context, or how long a resource
 manager keeps job information are implementation
 dependent.

- get_job

Purpose: Given a job identifier, this method returns a
 job object representing this job.

Format: get_job (in string job_id,
 out job job)

Inputs: job_id: job identifier as returned
 by the resource manager

Outputs: job: a job object representing
 the job identified by
 job_id


```

- get_self
  Purpose: This method returns a job object representing
           _this_ job, i.e. the calling application.
  Format:  get_self          (out job    job)
  Inputs:  none
  Outputs: job:              a job object representing
                             _this_ job.

```

```

job_id:
-----

```

The job ID is treaded as opaque string in the SAGA API.
 However, for reasons of compatibility and potential extended
 use of the job id information, the job id SHOULD be
 implemented as:

```
"[backend url]-[native id]"
```

For example, a job submitted via ssh, and having the unix pid
 1234, should get the job id:

```
"[ssh://remote.host.net:22/]-[1234]"
```

```
+-----+
```

```

Examples:
=====

```

Example : simple job submission and polling for finish.

```

// -----
// c++ example
std::list<char*>    transfers;
saga::job_definition jobdef;

transfers.push_back ("infile > infile");
transfers.push_back ("ftp://host.net/path/out << outfile");

jobdef.set_attribute      ("JobCmd", "myjob.sh");
jobdef.set_attribute      ("NumCpus", "16");
jobdef.set_vector_attribute ("FileTransfer", transfers);

saga::job_service myjs;
saga::job          myjob = myjs.create_job ("remote.host.net", jobdef);

myjob.run ();

while ( 1 )
{

```

```

saga::state state = myjob.get_state ();

if ( saga::job::Running == state.saga () )
{
    std::list <char*> hostlist = myjob.get_attribute
                                ('ExecutionHosts');

    // print hostlist
}
else if ( saga::job::Done == state.saga () )
{
    print "Job completed successfully.";
    exit;
}
else
{
    // saga state can be Unknown or Failed

    char* exitcode = myjob.get_attribute ('ExitCode');

    std::cout << "Job failed with " << exitcode << std::endl;
    exit (exitcode);
}

sleep (1); // idle
}
// -----

```

+-----+

Notes:

=====

References:

- [1] We expect that SAGA-API implementations may be implemented using DRMAA or may produce JSDL documents to be passed to underlying scheduling systems.
- [2] The job_state enumerated type encapsulates the possible states of a job. They are copied from the BES WG as of February 2006.
<https://forge.gridforum.org/projects/ogsa-bes-wg/document/ogsa-bes-draft-v16/en/1>
- [3] <http://www.ggf.org/documents/GWD-R/GFD-R.022.pdf>
- [4] <https://forge.gridforum.org/projects/jsdl-wg>
- [5] The API implementation is designed to be agnostic of the back end implementation, such that any back end could be

implemented to perform an action. For example, the checkpoint routine might cause an application level checkpoint, or might use the services of GridCPR.

- [6] In attributes that take paths and pathnames, there was some discussion as to whether we should require the implementation of placeholders which could represent things like 'home directory', and that are not known until the job is bound to an execution host.
- [7] There is discussion as to which interfaces might be missing. One possibility was a job history retrieval interface could be necessary. This could be used to map state transitions of a job throughout its lifetime.
- [8] The DRMAA 'job category' attribute was left out of the strawman API. During the discussions of this attribute within the design team meetings, it was deemed to simplify the API at the expense of the implementor of the back end system. Thus, it was left out pending discussion.

+-----+

```

+-----+
#####
#   # ##### ##### #####   ##   #   #
#           #   #   #   #   #   #   #   #
#####   #   #   #   #####   #   #   #   #
      #   #   #####   #   #####   #   #
#   #   #   #   #   #   #   #   #   #   #
#####   #   #   #   #####   #   #   #   #
+-----+

```

Summary:

=====

A number of use cases involved launching of remotely located components in order to create distributed applications. The use cases require simple remote socket connections to be established between these components and their control interfaces.

The target of this streams API is to establish the simplest possible authenticated socket connections with hooks to support authorization and encryption schemes. The API is

- 1) Not performance oriented: If you need performance, then it is better to program directly to the APIs of existing performance oriented protocols like GridFTP or XIO.
- 2) Focused on TCP/IP socket connections. There has been no attempt to generalize this to arbitrary streaming interfaces (although it does not prevent such things from being supported).
- 3) Does not attempt to create a programming paradigm that diverges very far from baseline BSD sockets, Winsock, or Java Sockets interfaces.

This API greatly reduces the complexity of establishing authenticated socket connections in order to communicate with remotely located components. It, however, provides very limited functionality suitable for applications that do not have too sophisticated requirements (as per 80-20 rule). As applications become more sophisticated, they can graduate to more sophisticated native APIs in order to support those needs.

```

+-----+

```

Specification:

=====

```

package saga.stream
{
    enum stream_state
    {
        Unknown      = -1,
        Error         =  1,
        Open          =  2,
        Dropped       =  3,
        NotConnected  =  4
    }

    enum activity_type
    {
        Unknown      = -1,
        Read         =  1,
        Write        =  2,
        Exception     =  4,
        Any          =  7,
    }

    class stream : implements-all saga::object
                  implements-all saga::attribute
                  implements-all saga::monitorable
    {
        CONSTRUCTOR (in    string      url,
                    in    session     session,
                    out   stream      obj);

        DESTRUCTOR  (in    stream      obj);

        get_url     (out   string      url);
        get_context (out   context     info);

        connect     (void);
        state       (out   stream_state state);
        wait        (in    activity_type what,
                    in    double      timeout,
                    out   array<activity_type,1>
                        activity);

        read        (inout array<byte,1> buffer,
                    in    long        buffer_size,
                    out   long        bytes_read);
        write       (in    array<byte,1> buffer,
                    in    long        size,
                    out   long        bytes_written);
        close       (void);
    }
}

```

```
// Attributes:
//   name: bufsize
//   desc: determines the size of the send buffer, in bytes
//   mode: ReadWrite
//   type: Int
//   value: <system dependend>
//
//   name: timeout
//   desc: determines the amount of idle time
//         before dropping the line, in seconds
//   mode: ReadWrite
//   type: Int
//   value: <system dependend>
//
//   name: blocking
//   desc: determines if read/writes are blocking
//         or not
//   mode: ReadWrite
//   type: Bool
//   value: True
//
//   name: compression
//   desc: determines if data are compressed
//         before/after transfer
//   mode: ReadWrite
//   type: Bool
//   value: <system dependend>
//
//   name: nodelay
//   desc: determines if packets are sent
//         immediatley, i.e. w/o delay
//   mode: ReadWrite
//   type: Bool
//   value: True
//
//   name: reliable
//   desc: determines if all sent data MUST arrive
//   mode: ReadWrite
//   type: Bool
//   value: True

// Metrics:
//   name: read
//   desc: fires if a stream gets readable
//   mode: Read
//   unit: 1
//   type: Bool
//   value: True
//
//   name: write
```

```
// desc: fires if a stream gets writable
// mode: Read
// unit: 1
// type: Bool
// value: True
//
// name: exception
// desc: fires if a stream has an error condition
// mode: Read
// unit: 1
// type: Bool
// value: True
//
// name: any_event
// desc: fires if the stream gets readable,
//       writable, or has an error condition
// mode: Read
// unit: 1
// type: Bool
// value: True
//
// name: dropped
// desc: fires if the stream gets dropped by the
//       remote party
// mode: Read
// unit: 1
// type: Bool
// value: True
//
// name: state
// desc: fires if the state of the stream changes,
//       and has the value of the stream state
//       enum
// mode: Read
// unit: 1
// type: enum
// value: Unknown
}
```

```
class stream_server : implements-all saga::object
                    implements-all saga::attributes
                    implements-all saga::monitorable
{
    CONSTRUCTOR      (in    string          url,
                     in    session         session,
                     out   stream_server   obj);
    DESTRUCTOR       (in    stream_server   obj);

    get_url          (out   string          url);
}
```

```

        wait          (in   double   timeout,
                       out   stream   stream);

    // Metrics:
    //   name:  client_connect
    //   desc:  fires if a client connects
    //   mode:  Read
    //   unit:  1
    //   type:  Bool
    //   value: True
}
}

```

+-----+

Details:

=====

class stream:

This is the object that encapsulates all client stream objects.

- CONSTRUCTOR

Purpose: Constructor, initializes a client client stream, for later connection to an server.

Format: CONSTRUCTOR (in string url,
in session session,
out stream stream);

Inputs: url: server location in URL
syntax
ctx: SAGA context used for
stream setup

Outputs: stream: new, unconnected stream
instance

Notes:

- the session handle defaults to the SAGA default session handle if not explicitly specified
- returned stream is NULL on error.
- server location and possibly protocol is described by the input URL.
- a saga::context is necessary to authenticate the socket.
- The socket is only connected after the connect method is called in order to support two-phase connections that appear in some authentication schemes. The state of the socket upon construction is therefore NotConnected. Once the connect() method is successfully called, the state will change to Open.

- DESTRUCTOR

Purpose: Destructor, closes any active connection and deallocates any memory consumed by the stream data structures.

Format: DESTRUCTOR (in stream stream)

Inputs: stream: stream data structure that needs to be closed and deallocated.

Outputs: none

Notes: - Because the data structures might consume some memory space internally, even closed, dropped, or failed sockets must be deallocated using the `destroystream` method.

- close

Purpose: closes an active connection

Format: close (void)

Inputs: none

Outputs: none

Throws: `IncorrectState`

Notes: - `IncorrectState` is thrown when the stream was closed previously

- connect

Purpose: Establishes a connection to the target defined during the construction of the stream.

Format: connect ();

Inputs: none

Outputs: none

Throws: `IncorrectState`

Notes: - on success, the streams state is changed to `Open`

- read

Purpose: Read a raw buffer from socket.

Format: read (inout string buffer, in long size, out long nbytes);

Inputs: buffer: Empty buffer passed in to get filled
size: Maximum number of bytes that can be copied in to the buffer.

Outputs: nbytes: number of bytes read, if successful. (0 is also valid)

Throws: `IncorrectState`

Notes: - This call is blocking. Use `wait` or `poll` methods to implement non-blocking reads.

- write
 - Purpose: Write a raw buffer to socket.
 - Format: write (in string buffer,
in long size,
out long nbytes);
 - Inputs: buffer: raw array containing data
that will be sent out via
socket
 - size: number of bytes of data in
the buffer
 - Outputs: nbytes: bytes written if successful
 - Throws: IncorrectState
 - Notes: - This call is blocking. Use wait method to
implement polling for non-blocking writes.

- state
 - Purpose: Check on the state of an active connection.
 - Format: state (out stream_state state);
 - Inputs: none
 - Outputs: state: state of stream
 - Notes: - the only valid states for a stream are:
 - Error:
 - The socket has entered a non-fatal error state. If the state is fatal, then the state will be Dropped. The reason for the error must be queried through a separate interface (not yet defined).
 - NotConnected:
 - This the state for a newly created socket where the connect method has not been invoked.
 - Open:
 - This is the state for an active/connected socket.
 - Dropped:
 - This is the state for a socket where the remote side of the socket connection has been lost or some other error has broken the connection. A socket will enter the dropped state if authentication fails for example. The actual reason for the drop must be queried through the error handling interface.
 - this method is only returning the *state* of the stream and not the reason it entered that state.
 - more states can be added as required

- wait
 - Purpose: Allows the stream to be interrogated to find out if it is ready for reading/writing, or if it has entered an error state.

Format: wait (in activity_type what,
in double timeout,
out <array,1>activity_type
cause);

Inputs: what: parameter list of activity
types to wait for
timeout: number of seconds to wait

Outputs: cause: activity type causing the
call to return

Throws: IncorrectState

Notes: - wait will only check on the conditions specified
in the 'what' parameter list (a bitmask in some
language bindings). The options are
Read:
The socket has pending data available for
reading.
Write:
The socket is available for writing.
Exception:
If the socket has entered an error state or
the remote host has dropped the connection.
Any:
This is shorthand for any of the above
- the call returns enum describing the
availability of the socket (eg. readable,
writable, or exception) masked against the
input 'what' enum list.
- the call is blocking if the timeout is any
positive value. It blocks forever (no timeout)
if the timeout value is < 0.0. The wait method
can be used for polling if the timeout is set to
zero. The wait method will only check for the
activity_type that is specified in the call (and
ignore all other issues).

- get_context
Purpose: Gets a security context object from an OPEN
(connected)

Format: get_context (out context context);

Inputs: none

Outputs: context: a context object.

Throws: IncorrectState

Notes: - throws IncorrectState exception or <null>
context object returned if the security
info is inapplicable (non-authenticated sockets)
- the context returned contains the security
information from the REMOTE party, and can be used
for authorization.
- It is assumed that the context is authenticated.

```

- get_url
  Purpose: get URL used for creating the string
  Format:  get_url          (out string url);
  Inputs:  none
  Outputs: url:             string containing the URL
                           of the connection.

  Thorws:  nothing
  Notes:   - this is the URL which can be passed to
            stream constructor to create another
            connection to the same stream_server.

```

```
class stream_server:
```

The stream_server object establishes a listening/server object that waits for client connections. It can *only* be used as a factory for Server sockets. It doesn't do any read/write I/O.

```

- CONSTRUCTOR
  Purpose: Constructor, to create a new stream_server object
  Format:  CONSTRUCTOR      (in string      url,
                           in session      session,
                           out stream_server stream);

  Inputs:  url:             channel name or url,
                           defines the source side
                           binding for the stream
                           (eg. the port number for
                           the service)

           session:         session to be used for
                           object creation

  Outputs: stream:          new stream_server object
  Notes:   - the session handle defaults to the SAGA
            default session handle if not explicitly
            specified
            - returns NULL stream_server object on error
            - the context is primarily used to hide the
              security information necessary to establish
              authenticated connections.

- DESTRUCTOR
  Purpose: Destructor for stream_server object.
  Format:  DESTRUCTOR       (in stream_server stream)
  Inputs:  stream:          stream_server object to be
                           destroyed

  Outputs: none
  Notes:   - the call cleans up any memory used by the
            stream_server object in addition to closing the

```

service port.

- wait
 - Purpose: wait for incoming client connections
 - Format: wait (in double timeout,
out stream client);
 - Inputs: timeout: number of seconds to wait
for client
 - Outputs: client: new Connected stream object
 - Throws: IncorrectState
 - Notes:
 - supports either blocking or polling for new client connections.
 - if successful, it returns a new stream object that is connected to the client.
 - unlike new client streams, the new connection is return in the Connected state.
 - returns NULL or equivalent if it times out.
 - returns NULL or equivalent if connection setup failed
 - timeout < 0.0 wait forever
 - timeout > 0.0 wait this number of seconds
 - timeout = 0.0 poll and return immediately.

- get_url
 - Purpose: get URL to be used to connect to serverstream
 - Format: get_url (out string url);
 - Inputs: none
 - Outputs: url: string containing the URL
of the connection.
 - Thorws: nothing
 - Notes:
 - this is the URL which can be passed to stream constructor to create a connection to this stream_server.

+-----+

Examples:

=====

Sample SSL/Secure Client:

Opens a stream connection using native security: context is passed in implicitly via a global SAGA context (GSI or SSL security)

```
// C++/JAVA Style
int recvlen;
saga::stream s ("localhost:5000");
```

```

s.connect ();
s.write  ("Hello World!", 12);

// blocking read, read up to 128 bytes
recvlen = s.read (buffer, 128);

/* C Style */
int recvlen;

SAGA_stream = SAGA_Stream_open ("localhost:5000");

SAGA_Stream_connect (s);
SAGA_Stream_write  (s, "Hello World!", 12);

/* blocking read, read up to 128 bytes */
recvlen = SAGA_Stream_read (s, buffer, 128);

c Fortran Style */
INTEGER  err,SAGAStrRead,SAGAStrWrite,err
INTEGER*8 SAGAStrOpen,streamhandle
CHARACTER buffer(128)
SAGAStrOpen("localhost:5000",streamhandle)
call SAGAStrConnect(streamhandle)
err = SAGAStrWrite(streamhandle,"localhost:5000",12)
err = SAGAStrRead(streamhandle,buffer,128)

```

Sample Secure Server:

Once a connection is made, the server can use information about the authenticated client to make an authorization decision

```

// C++/JAVA Style
saga::stream_server server ("tcp://localhost/5000");
saga::stream         client;
int                  done = 0;

// now wait for a connection (normally in a loop)
do {
    string value;
    // wait forever for connection
    client = server.wait ();

    // get the distinguished name (DN)
    saga::context = client.get_context ();
}

```

```

    // check if context type is X509, and if DN is the
    // authorized one
    if ( saga::context::X509 == context.type () &&
        ! strcmp (context.get_attribute ("DN"), auth_dn) )
    {
        done = 1; // allowed
    }
    else
    {
        SAGA::stream_close (client); // not allowed
    }
} while ( ! done );

// start activity on client socket...
```

Example for async stream server

```

// c++ example
class my_cb : public saga::callback
{
    privat:
        saga::stream_server ss;
        saga::stream        s;

    public:

        my_cb (saga::stream_server ss_,
               saga::stream        s)
        {
            ss = ss_;
            s   = s_;
        }

        ~my_cb (void) { }

        void callback (saga::monitorable mt,
                       saga::metric      m,
                       int                 c)
        {
            s = ss_.wait ();
            mt.remove_callback (c); // want to be called only once
        }
}

int main ()
{
    saga::stream_server ss;
    saga::stream        s;
    my_cb cb (ss, s);
```

```
    ss.add_callback ("client_connect", cb);

    while ( true )
    {
        if ( s.state != saga::stream::Open )
        {
            // no client, yet
            sleep (1);
        }
        else
        {
            // handle open socket
            s.write ("Hello Client\r\n", 14);
            s.close ();

            // restart listening
            ss.add_callback ("client_connect", cb);
        }
    }

    return (-1); // unreachable
}
```

+-----+

Notes:
=====

+-----+


```

+-----+
#
# # ##### ##### ##### # # ##### # # # # # #
# # # # # # ## # # # # # # # #
# # # # # # ##### # # # # # # ## # #
##### ##### ##### # # # # # # ## #####
# # # # # # # ## # # # # # # # #
# # # # # ##### # # ##### # # # # # #

+-----+

#####
# # # ## # # ##### # ##### #####
# # # # # ## ## # # # #
#### ## # # # ## # # # # #####
# ## ##### # # ##### # # #
# # # # # # # # # # # #
##### # # # # # # ##### ##### #####

+-----+

```

Introduction

=====

This appendix shows a couple of API examples in different languages. As stated in the global introduction, these examples are not supposed to be normative -- language bindings are outside the scope of this document. This appendix is rather supposed to illustrate how the authors imagine the use of the API in various languages.

We hope that the examples illustrate, that the API stays SIMPLE in various language incarnations, as was the major design intent for the `_S_AGA` API.

```

+-----+

```

Example 1: Files:

=====

open a file. if its size is > 10, then read the first 10 bytes into a string, print it, end return it.

Example 1a: C++

```

// c++ example
void head (const char* url)

```

```

{
    try {
        // get type and other infos
        saga::file my_file (url);

        off_t size = my_file.get_size ();

        if ( size > 10 )
        {
            char    buffer[11];
            long    bufflen;

            my_file.read (10, buffer, &bufflen);

            if ( bufflen == 10 )
            {
                printf ("head: '%s'\n", buffer);
            }
        }
    }

    // catch any possible error - see elsewhere for better
    // examples of error handling in SAGA
    catch ( const saga::exception & e )
    {
        std::cerr << "Oops! SAGA error: " + e.what () + std::endl;
    }

    return;
}

```

Example 1b: C

```

char* head (const char* url)
{
    SAGA_File my_file = SAGA_File_create (url);

    if ( NULL == my_file )
    {
        fprintf (stderr, "Could not create SAGA_File for %s: %s\n",
                  url, SAGA_Session_get_error (theSession));
        return (NULL);
    }

    off_t size = SAGA_File_get_size (my_file);

    if ( size < 0 )
    {
        fprintf (stderr, "Could not determine file size for %s: %s\n",
                  url, SAGA_Session_get_error (theSession));
    }
}

```

```

    return (NULL);
}
else if ( size > 10 )
{
    char    buffer[11];
    size_t  buflen;

    ssize_t ret = SAGA_File_read (my_file, 10, buffer, &buflen);

    if ( ret < 0 )
    {
        fprintf (stderr, "Could not read file %s: %s\n",
                  url, SAGA_Session_get_error (theSession));
        return (NULL);
    }

    if ( buflen == 10 )
    {
        buffer [11] = '\0';
        printf ("head: '%s'\n", buffer);
        return (buffer);
    }
    else
    {
        fprintf (stderr, "head: read is short! %d\n", buflen);
        return (NULL);
    }
}

fprintf (stdout, "head: file is too small! %d\n", size);

return (NULL);
}

```

Example 1c: Java

```

import saga*;

class MyClass
{
    // open a file. if its size is > 10, then read the first
    // 10 bytes into a string, print it, end return it.
    string head (URI uri)
    {
        try
        {
            saga::file f (uri);

            if ( 10 <= f.get_size () )

```

```

    {
        FileInputStream in (uri);
        byte[]          buffer = new byte[10];
        int              res    = in.read (buffer);

        if ( 10 == res )
        {
            System.out.println ("head: " + buffer);
        }
        else
        {
            System.err.println ("head: read is short! " + res);
        }

        return new String (buffer);
    }
    else
    {
        System.out.println ("file is too small: " + size);
    }
}

// catch any possible error - see elsewhere for better
// examples of error handling in SAGA
catch (...)
{
    System.out.println ("Oops!");
}

return null;
}
}

```

 Example 1d: Perl ('normal' error handling)

```

sub head ($)
{
    my $url      = shift;
    my $my_file = new saga::file ($url)
        or die ("can't create file for $url: $!\n");

    my $size     = $my_file->get_size ();

    if ( $size > 10 )
    {
        my $buffer = $my_file->read (10)
            or die ("can't read from file $url: $!\n");
    }
}

```

```
    if ( length ($buffer == 10 ) )
    {
        print "head: '$buffer'\n";
        return ($buffer);
    }
    else
    {
        printf "head: short read! %s\n", saga::get_error ();
    }
}
else
{
    print "file $url is too short: $size\n";
}

return (undef);
}
```

Example 1e: Perl (exceptions)

```
sub head ($$)
{
    my $session = shift;
    my $url      = shift;

    eval {
        my $my_file = new saga::file (session, url);

        my $size    = my_file->get_size ();

        if ( size > 10 )
        {
            my $buffer = my_file->read (10);

            if ( length ($buffer == 10 ) )
            {
                print "head: '$buffer'\n";
                return ($buffer);
            }
            else
            {
                printf "head: short read! %s\n", saga::get_error ();
            }
        }
        else
        {
            print "file $url is too short: $size\n";
        }
    }
}
```

```
if ( $@ =~ /saga/i )
{
    print "caught saga error: $@\n" if $@;
}

return (undef);
}
```

Example 1f: Fortran

TBD

Example 1f: Python

TBD

+-----+