TeraGrid

Charlie Catlett, Director
Pete Beckman, Chief Architect
University of Chicago & Argonne National Laboratory
June 2003
GGF-11
Honolulu, Hawaii

The TeraGrid

Distributed resources & expertise can be leveraged to accelerate scientific discovery

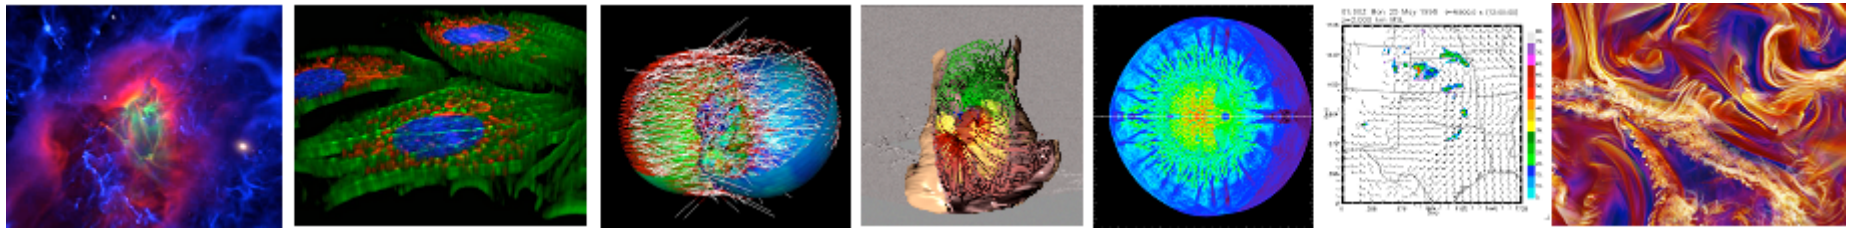

Action:

Interconnect large-scale shared scientific databases, computing systems, instruments, and facilities to improve scientific productivity by *removing barriers to collaboration and use of distributed resources*

Funding: US National Science Foundation, industry partners, states
Partners: UC/ANL, Caltech, NCSA, PSC, SDSC, TACC, ORNL, Purdue, Indiana

What we are doing

- Creating a Grid infrastructure that focuses first and foremost on ease of use for users.
  - Common TeraGrid Services & Software (CTSS)
- Creating a software and services infrastructure that is robust and persistent.
  - Inca Verification and Validation System
- Doing this in a way that is easily reproducible and/or extensible
- Not in today's talk… Science Gateways
  - TeraGrid as a wholesale service provider behind discipline-oriented portals, for instance.

Timeline and Program Overview

- Distributed TeraScale Facility (DTF; $50M, 3 yr)
  - IA-64 homogeneous systems at CIT, ANL, NCSA, SDSC
  - Use Grid technologies to embed Grid resources in existing large-scale centers
  - Offer NSF PACI user community a migration path from client-server supercomputing to Grid computing.
- Extensible TeraScale Facility (ETF; $35M, 1 yr)
  - Add Power4 (SDSC), Alpha (PSC) systems
  - Create extensible backbone network (hubs in LA, Chicago)
- TeraGrid Extension Program (TEP; $10M, 1 yr)
  - Add ORNL, TACC, Indiana, Purdue
  - Add network hub in Atlanta
- TeraGrid Operation, Mgmt & Evolution ($150M, 5 yrs)
  - Target date for Production: October 1, 2004 (now in pre-production)
  - TeraGrid OM&E Program currently under review

General Overview

- TeraGrid System Management & Integration Group (~40 staff)
- 9 Resource Providers (~80 staff at 9 sites)
  - Resources & Services Described via Grid "Service Level Agreements"
  - Resource: Computers, networks, instruments, databases/collections, storage systems…
- TeraGrid Architecture defines what it means to "Join TeraGrid"

How To Add a New Site: An Objective, Open Process

- Identify Resource
- Participate in Allocations Process (for compute resources)
  - Accounts, accounting processes
- Support CTSS user environment
  - Accept authentication and operation of shared coordinated software
- Participate in Operations Process
  - Define Contacts
  - Accept trouble tickets
- Join Security Infrastructure
  - Sign Security Memorandum, participate in reporting, vulnerability & risk analysis, etc.

Example use Scenario (one of many)

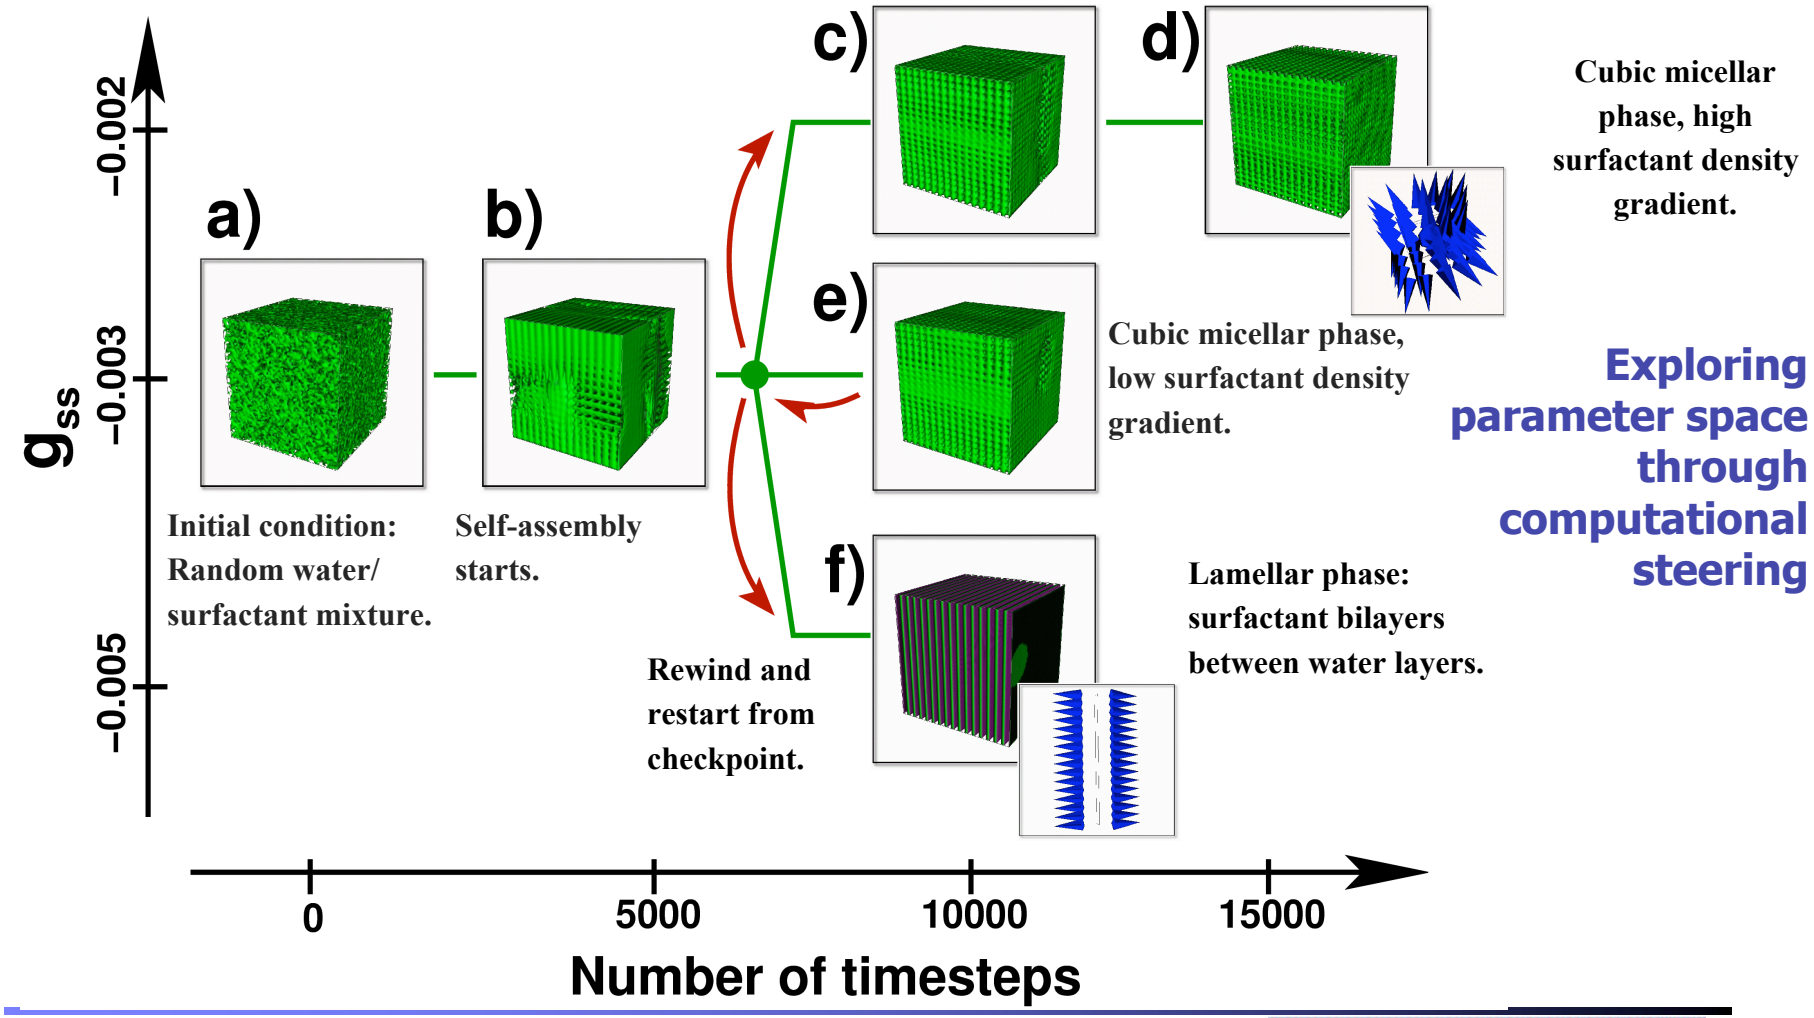

Exploring parameter space through computational steering

ETF Hardware Deployment

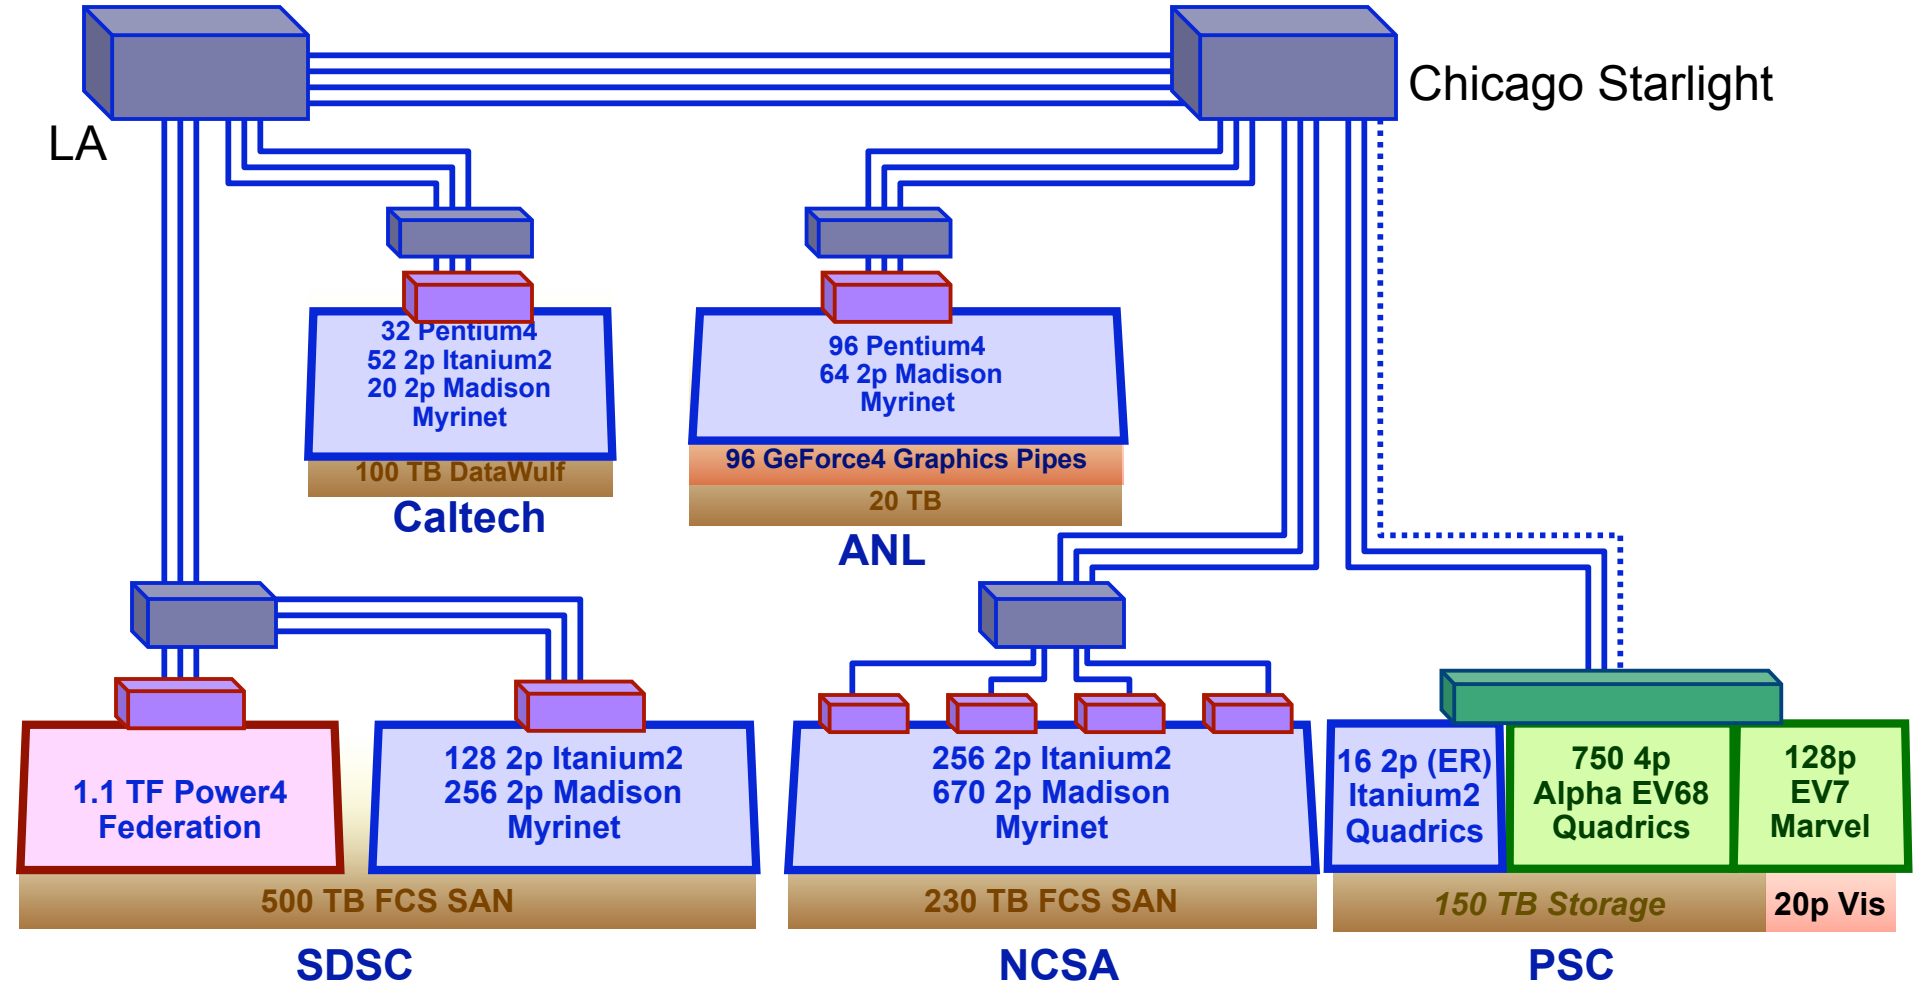

TeraGrid Network (June 2004)

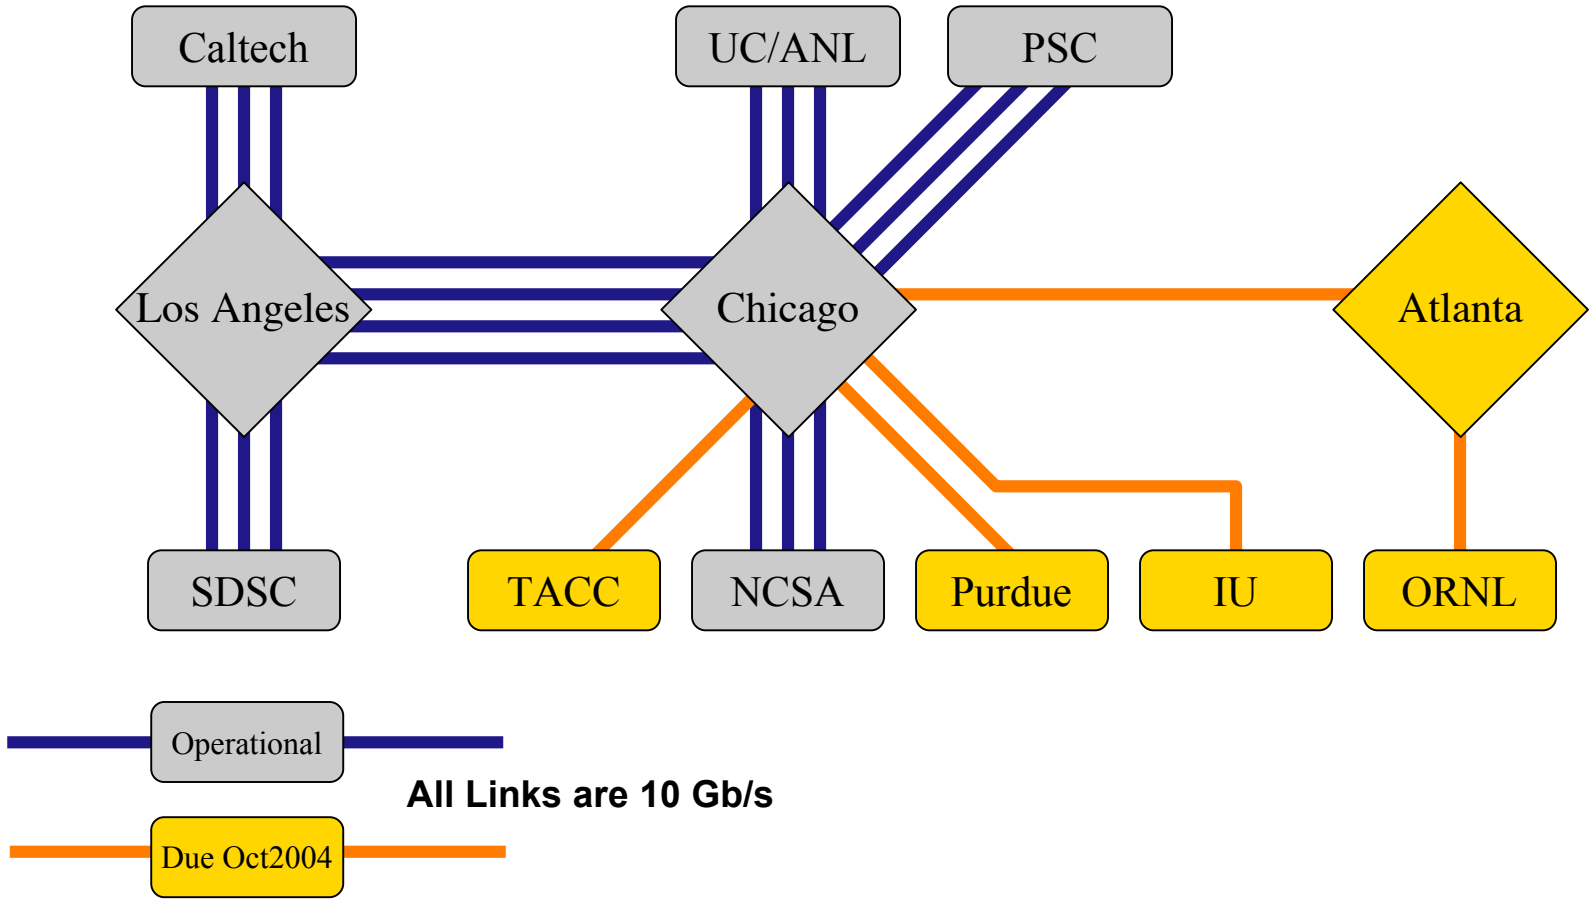

User-Centric Guiding Principles:

- The larger the pool of distributed, networked, ***unified*** resources, the greater the benefit
- Promoting Adoption:

Usability	**Stability**	**Capability**
It Must Be Easy	**It Must Be Ready**	**It Must Be Better**

CIO Magazine: "Timing is Everything. Seizing the perfect moment to present a new technology to your company can make or break a strategic plan."

Unified Policies and Common Resource Currency (TROO)

The Rule Of One

- One help desk
- One NRAC submission
- One account req. form
- One accounting currency
- One set of user policies
- One documentation set

- Result:
  - Improved Usability
  - Very attractive target for adoption
  - Unified networked resources more valuable to community (Metcalf)

Commitment To Adoption

- Learn Once, Run Anywhere (LORA)
  - A *single* training course and a *single* set of manuals preserve user investment
  - "TeraGrid Roaming" supported:
    - Develop app locally, run on any TG resource
- Result:
  - SC03 TeraGrid User Tutorial was over subscribed! Plans for road show under way
  - Smiling users
    - Improved usability & stability of TeraGrid architecture

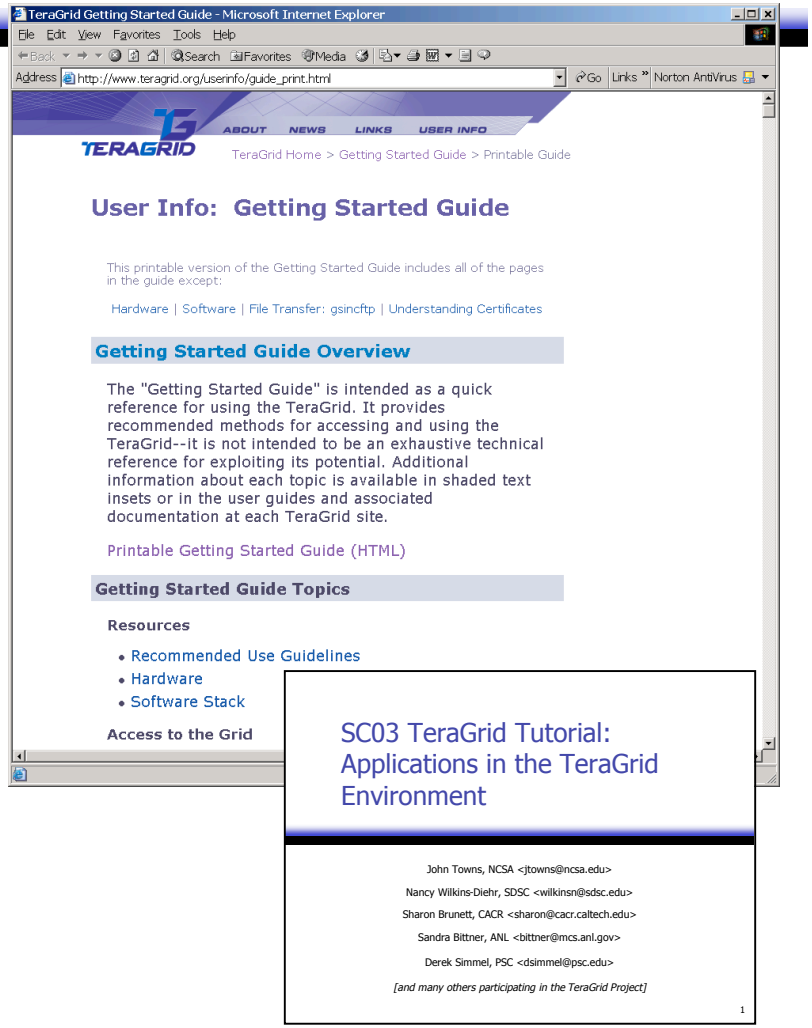

TeraGrid Architecture

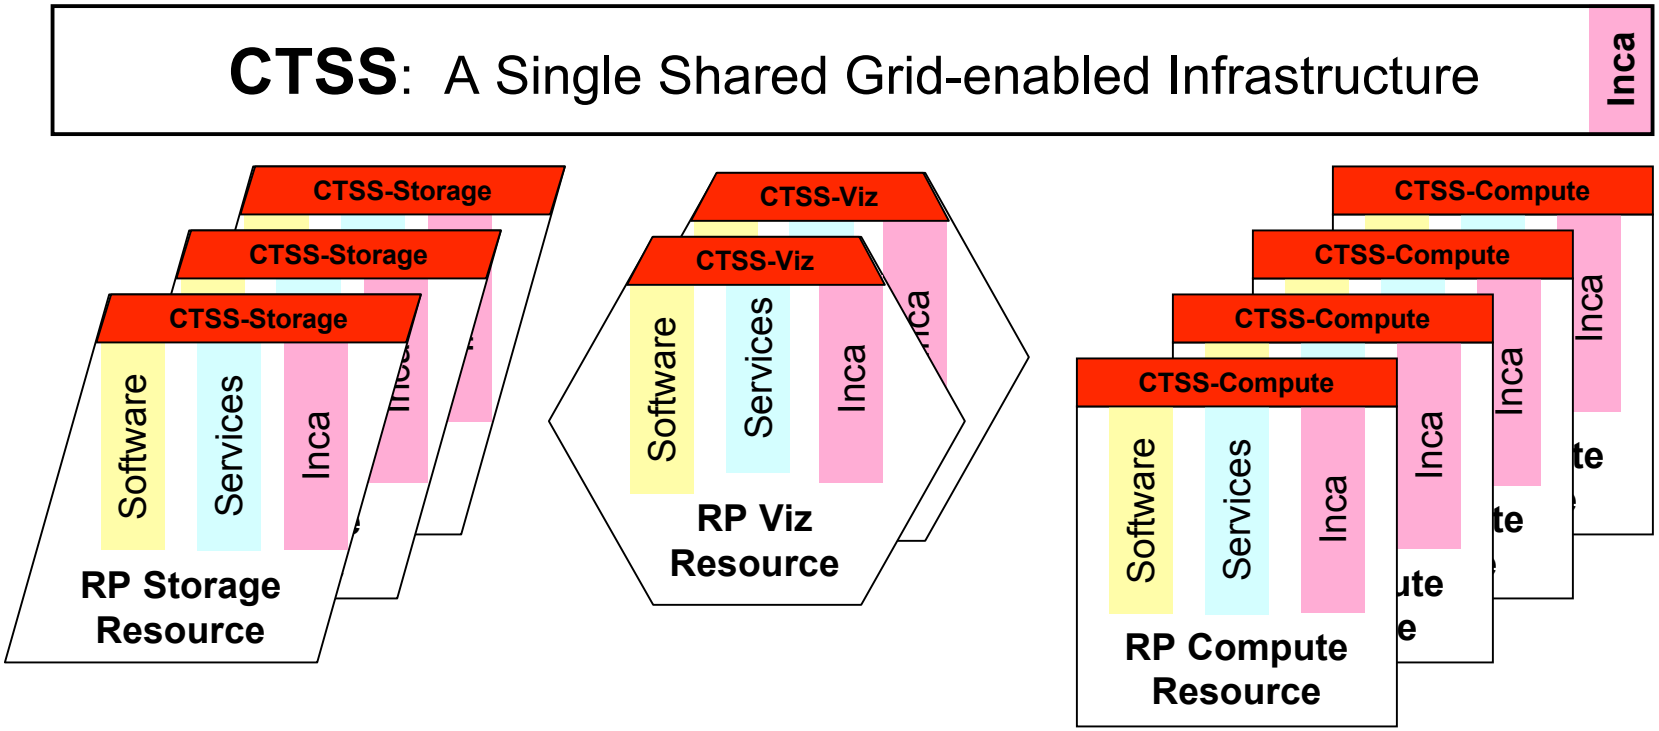

Common TeraGrid Software & Services (CTSS)

A Grid Hosting Environment

The core infrastructure for Virt Orgs to build Grid-based projects

Special Capabilities

Experimental math libraries

Unique storage system

Large shared memory arch….

Example:

Web Hosting Env.

PHP, Perl, Python scripting.

MySQL, FrontPage

100 POP accts, 100MB disk

SMTP, IMAP & Webmail

US$49 per year

TeraGrid Hosting Env.

Single Contact: help@teragrid.org

Unified Ops center

Certified Software Stack

MPICH, Globus

GridFTP, BLAS, Linpack,

Atlas, SoftEnv, gsi-ssh

$TG_SCRATCH, …

$100 Million

Classic Unix-like Environment

- /bin/sh, bin/cp, /bin/ls, Unix file system & tools, dev tools (make, compilers) etc

Common TeraGrid Software Stack

- CTSS Provides a single unified set of interoperable components and services that define the TeraGrid's Grid Hosting Enviroment and enable "TeraGrid Roaming"

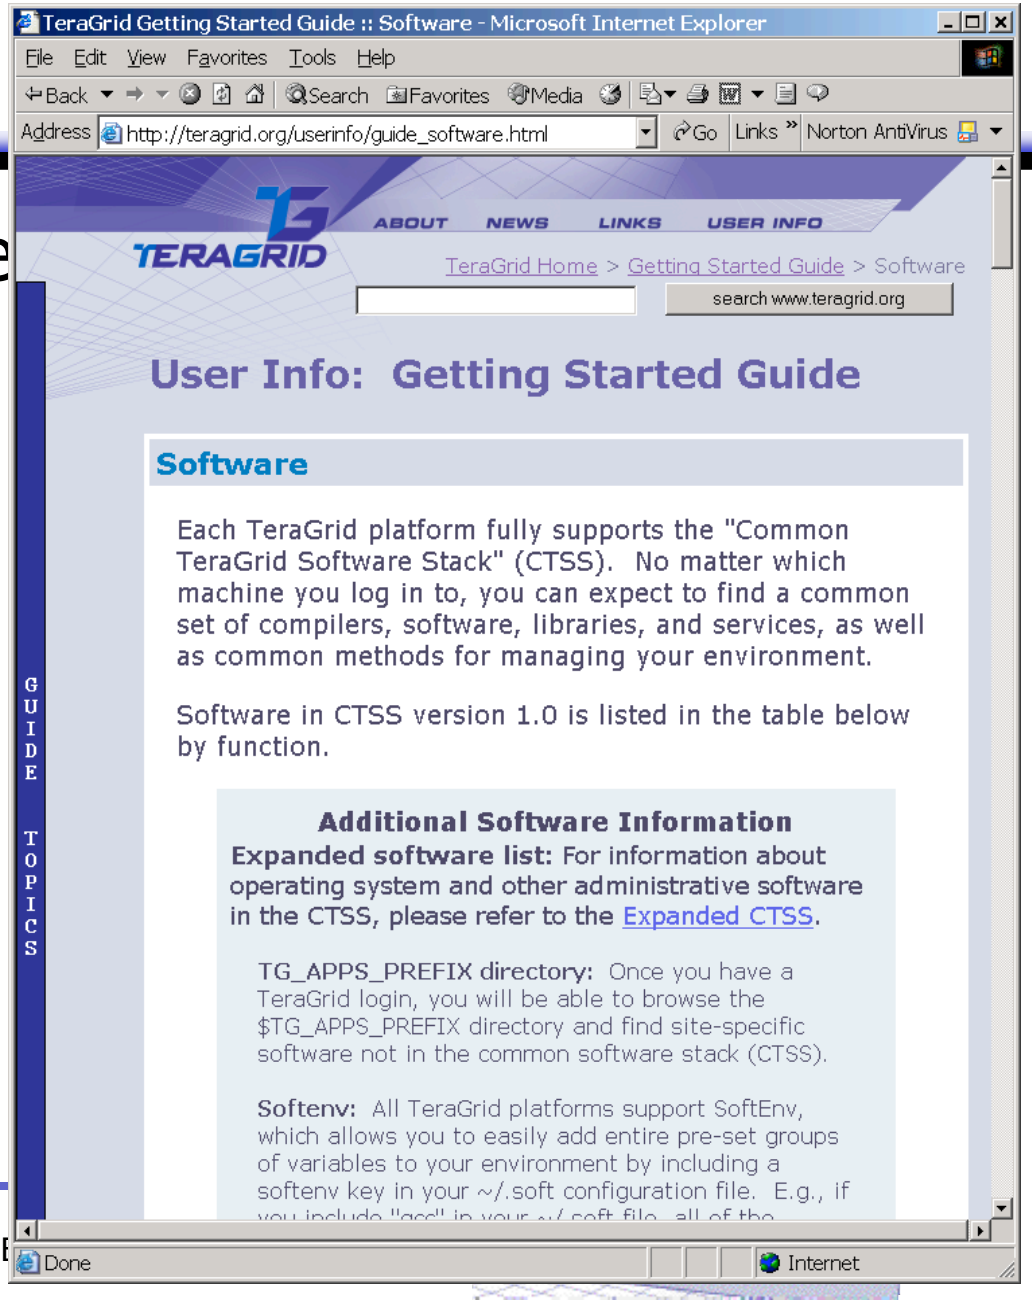

Core CTSS Components Across All Platforms: Linux, AIX, Tru64

- atlas
- blas
- condor-g
- gcc
- globus-2.2.4-gcc
- globus-2.4.3-gcc
- gsi-ncftp
- gsi-openssh
- gx-map
- hdf4
- hdf5
- Java_COG
- mpich-g2-gcc
- mpich-p4-gcc
- myproxy
- openssh
- openssl
- petsc-gcc
- python
- softenv
- srb-client
- tcl
- vmi-crm

Additional component sets, by category, include:

Intel compilers, IA64 (Myrinet, BIOS, etc), Linux kernels & patches

High Quality Production System: Integrated Verification & Validation

"And then one day the grid went down and never came back up."

New Yorker Magazine, 2003

V & V Concepts for Production-Quality Grid Systems

- Formal specifications of supported environment and "operational" must exist
- A set of independent tests should check for compliance and correctness
- A human should not be "in the loop"
- Testing should include performance, and be archived over time for trend analysis
- Scalability: Adding new sites should require minimal resources

Inca: A Test Harness Framework for Builders

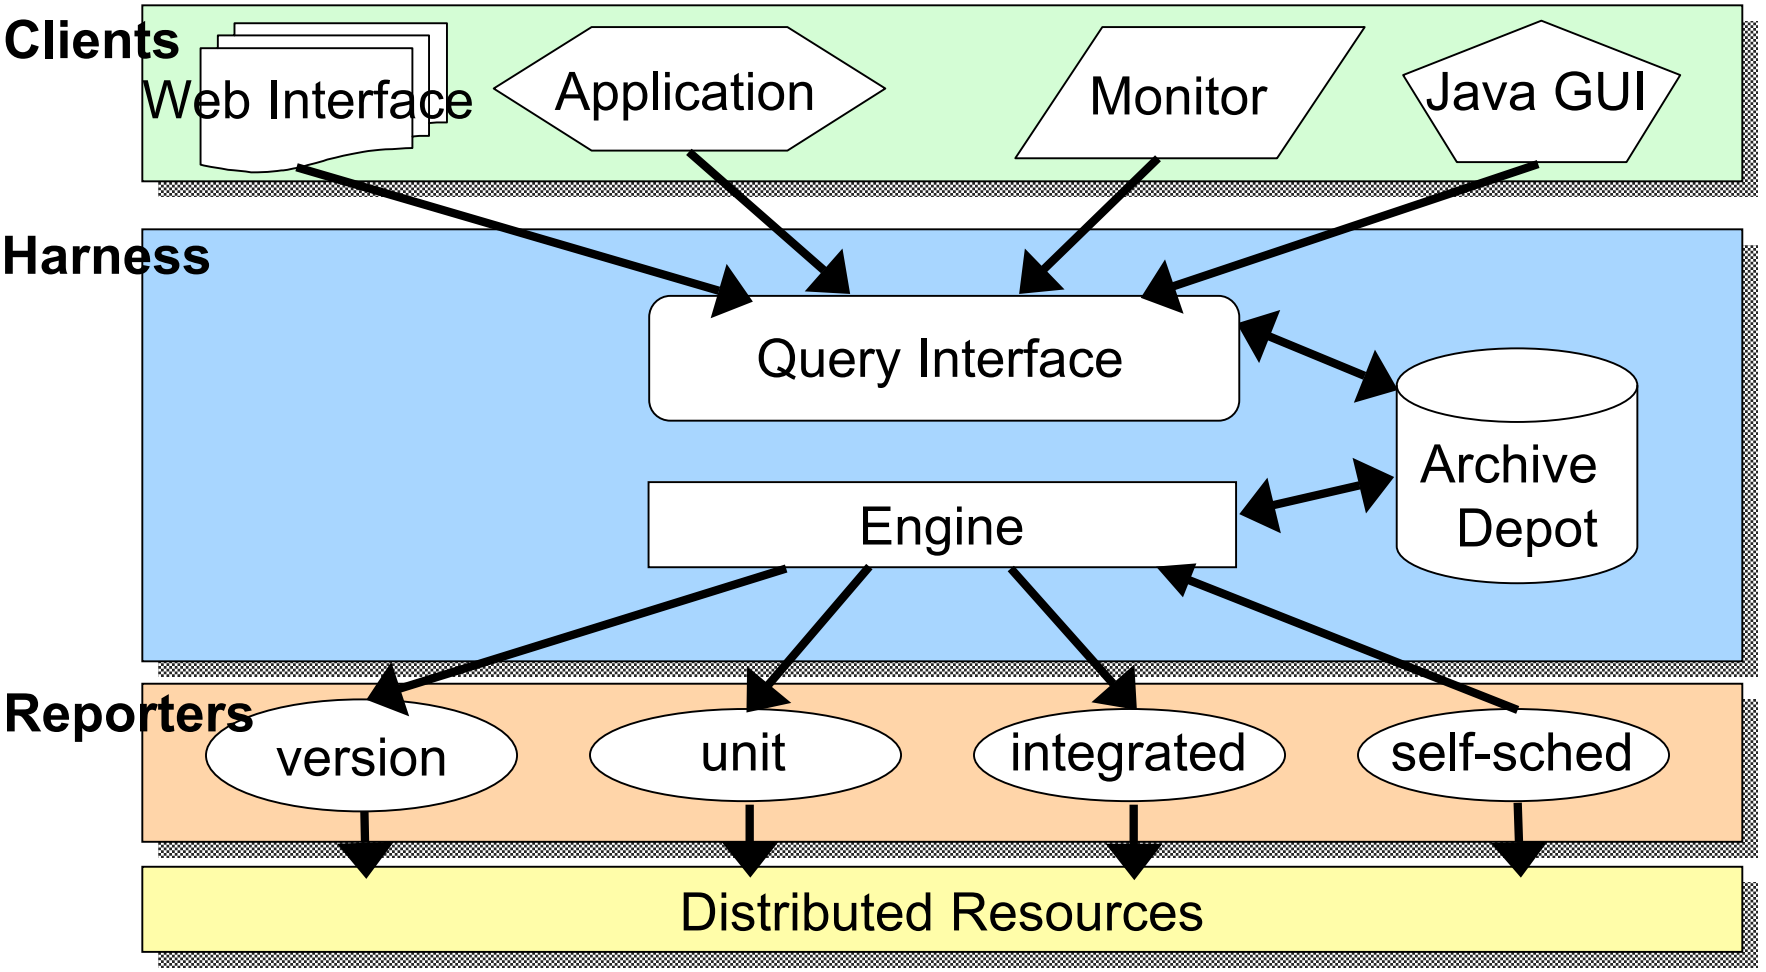

Across the Entire TeraGrid: A Single Language For Reporters

```
<?xml version="1.0" encoding="UTF-8" ?>  <!-- Generated by Turbo
      XML 2.3.1.100. Conforms to w3c
      http://www.w3.org/2001/XMLSchema  -->
  <xsd:schema xmlns:xsd="http://www.w3.org/2001/XMLSchema"
      elementFormDefault="qualified">
  <xsd:element name="INCA_Reporter">
  <xsd:complexType>
  <xsd:sequence>
  <xsd:element ref="INCA_Version" />
  <xsd:element ref="localtime" />
  <xsd:element ref="gmt" />
  <xsd:element ref="ipaddr" />
  <xsd:element ref="hostname" />
  <xsd:element ref="uname" />
  <xsd:element ref="url" />
  <xsd:element ref="name" />
  <xsd:element ref="description" />
  <xsd:element ref="version" />
  <xsd:element ref="INCA_Input" />
  <xsd:element ref="body" />
  <xsd:element ref="exit_status" />
  </xsd:sequence>
  </xsd:complexType>
  </xsd:element>
  <xsd:element name="gmt" type="xsd:string" />
  <xsd:element name="localtime" type="xsd:string" />
  <xsd:element name="hostname" type="xsd:string" />
  <xsd:element name="ipaddr" type="xsd:string" />
  <xsd:element name="uname" type="xsd:string" />
  <xsd:element name="url" type="xsd:string" />
  <xsd:element name="name" type="xsd:string" />
  <xsd:element name="version" type="xsd:string" />
  <xsd:element name="INCA_Version" type="xsd:string" />
```

```
<xsd:element name="description" type="xsd:string" />
<xsd:element name="exit_status" nillable="true" fixed="0">
<xsd:complexType mixed="true">
<xsd:choice>
<xsd:element ref="message" minOccurs="0" />
</xsd:choice>
</xsd:complexType>
</xsd:element>
<xsd:element name="message" type="xsd:string" />
<xsd:element name="body" abstract="true" />
<xsd:element name="ID" type="xsd:string" />
<xsd:element name="INCA_Input">
<xsd:complexType>
<xsd:sequence>
<xsd:element ref="input" minOccurs="0" maxOccurs="unbounded" />
</xsd:sequence>
</xsd:complexType>
</xsd:element>
<xsd:element name="input" />
<xsd:element name="verbose" type="xsd:integer"
    substitutionGroup="input" />
<xsd:element name="help" substitutionGroup="input">
<xsd:simpleType>
<xsd:restriction base="xsd:string">
<xsd:enumeration value="yes" />
<xsd:enumeration value="no" />
</xsd:restriction>
</xsd:simpleType>
</xsd:element>
</xsd:schema>
```

Current Status of V&V Program

- More than 900 components are tested
- Output is "engineer quality", and not prioritized or user readable
- Definition of "Up" underway
- Inca has been well received by other Grid communities, and we are beginning to receive requests for copies and presentations

In Practice…