

Web Service Composition Application Framework

Malik SAHEB
Arjuna Technologies Ltd

What is WS-CAF?



- Collection of 3 specifications designed to be used independently or together
 - WS-Context
 - Context service
 - WS-Coordination Framework
 - Framework for pluggable coordination protocols
 - WS-Transaction Management
 - Three transaction models for Web services



WS-Context

WS-Context



- Allows composite applications to share common information.
- Defines Notion of An “Activity”
 - An execution of a series of related interactions with a set of cooperating web services.
 - Operations correlated by a context associated with activity
- A Context
 - A way of doing correlation of messages
 - Context bound to one activity
 - Examples
 - Single-session sign-on
 - Transactions
 - Database session identifier

Context Structure



- An XML document containing a unique identifier and optional data specific to a related activity.
- Typically included in the SOAP header of messages to and from web services participating in an activity.
- Can be passed as a referenceable URI (by reference) or in its longer form (by value)

Context schema



```
<xs:complexType name="ContextType">
  <xs:sequence>
    <xs:any namespace="##other" processContents="lax" minOccurs="0"
      maxOccurs="unbounded"/>
    <xs:element name="context-identifier" type="tns:contextIdentifierType"/>
    <xs:element name="context-service" type="ref:ServiceRefType"
      minOccurs="0"/>
    <xs:element name="type" type="xs:anyURI"/>
    <xs:element name="context-manager" type="ref:ServiceRefType"
      minOccurs="0"/>
    <xs:element name="parent-context" type="tns:ContextType"
      minOccurs="0"/>
  </xs:sequence>
  <xs:attribute name="timeout" type="xs:int" use="optional"/>
  <xs:attribute ref="wsu:Id" use="optional"/>
</xs:complexType>
```

Contains



- *Context identifier* - URI, MUST be unique
 - with optional `wsu:Id` attribute.
- *Context service* – ServiceRefType (OPTIONAL)
 - locate the authority having generated the context
 - ServiceRefType = Generic structure to deal with addressing
- *Type* – URI
 - the type of the protocol supported by the context,
- *Context-manager* ServiceRefType (OPTIONAL)
 - to get data associated with a context-identifier
 - if present, the context has been passed by reference
- *parent-contexts* (OPTIONAL)
- *timeout* attribute (OPTIONAL)
 - indicates for how long the context information is valid;
- *wsu:Id* attribute (OPTIONAL)
 - used to support signing or encrypting the context structure.

WS-Context Services

arjuna

middleware for reliability



- Defines web services for maintaining contexts
- Ability to pass contexts by reference or by value.

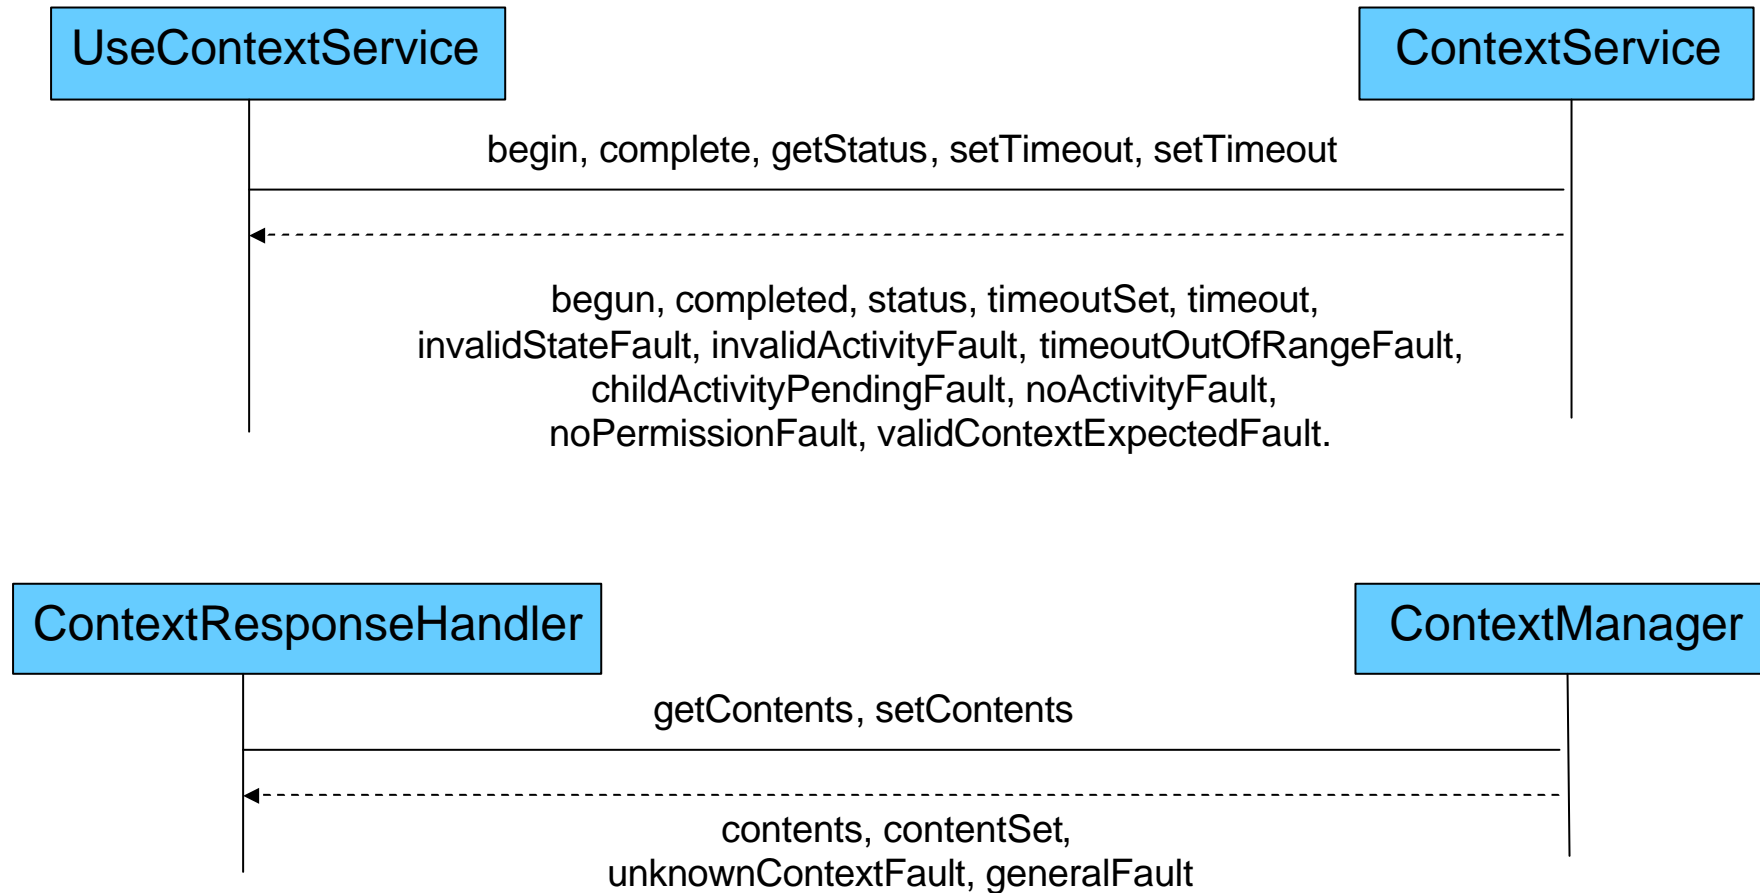
➤ Context Service

- creating - *begin*
- completing - *complete*
- getting status of a context – *getStatus*
 - *ACTIVE, COMPLETING, COMPLETED, NO_ACTIVITY, UNKOWN*
- Setting and getting timeout – *setTimeout, getTimeout*

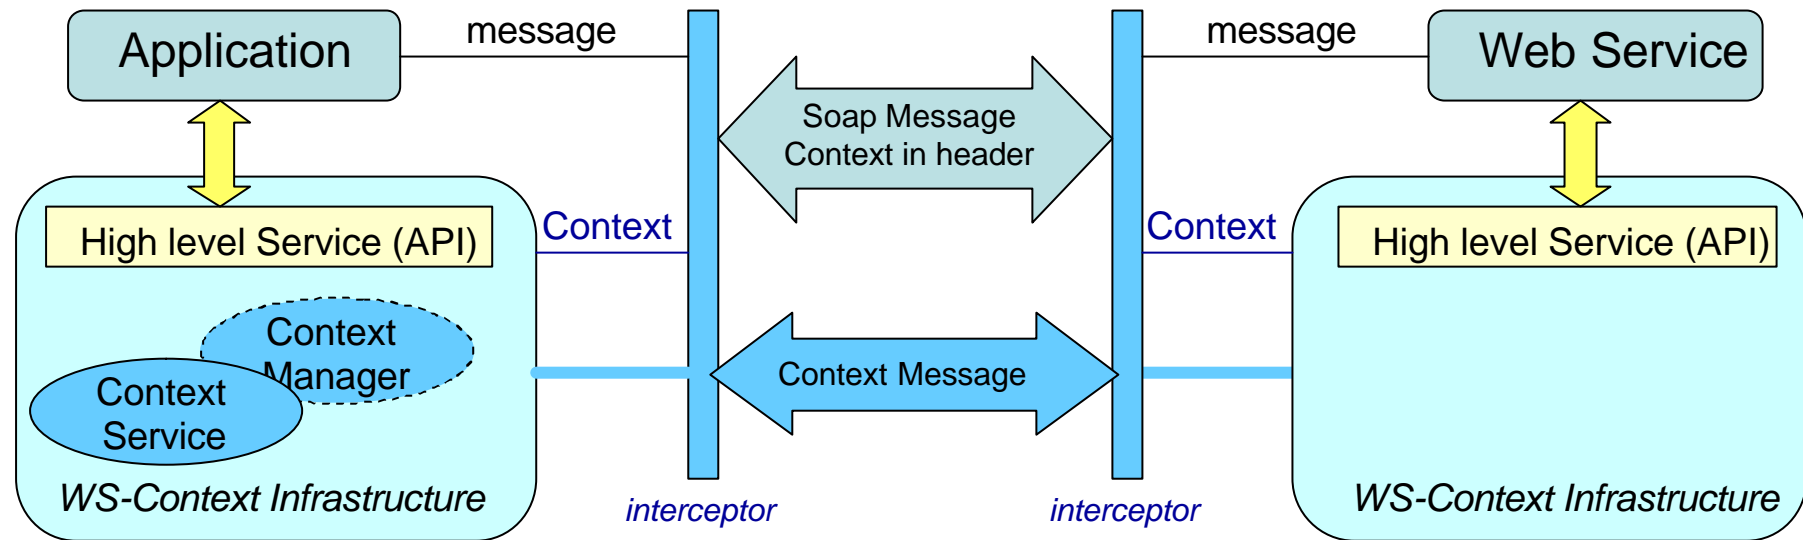
➤ Context Manager Service

- Obtaining/setting a content of a context got by reference – *getContents, setContents*

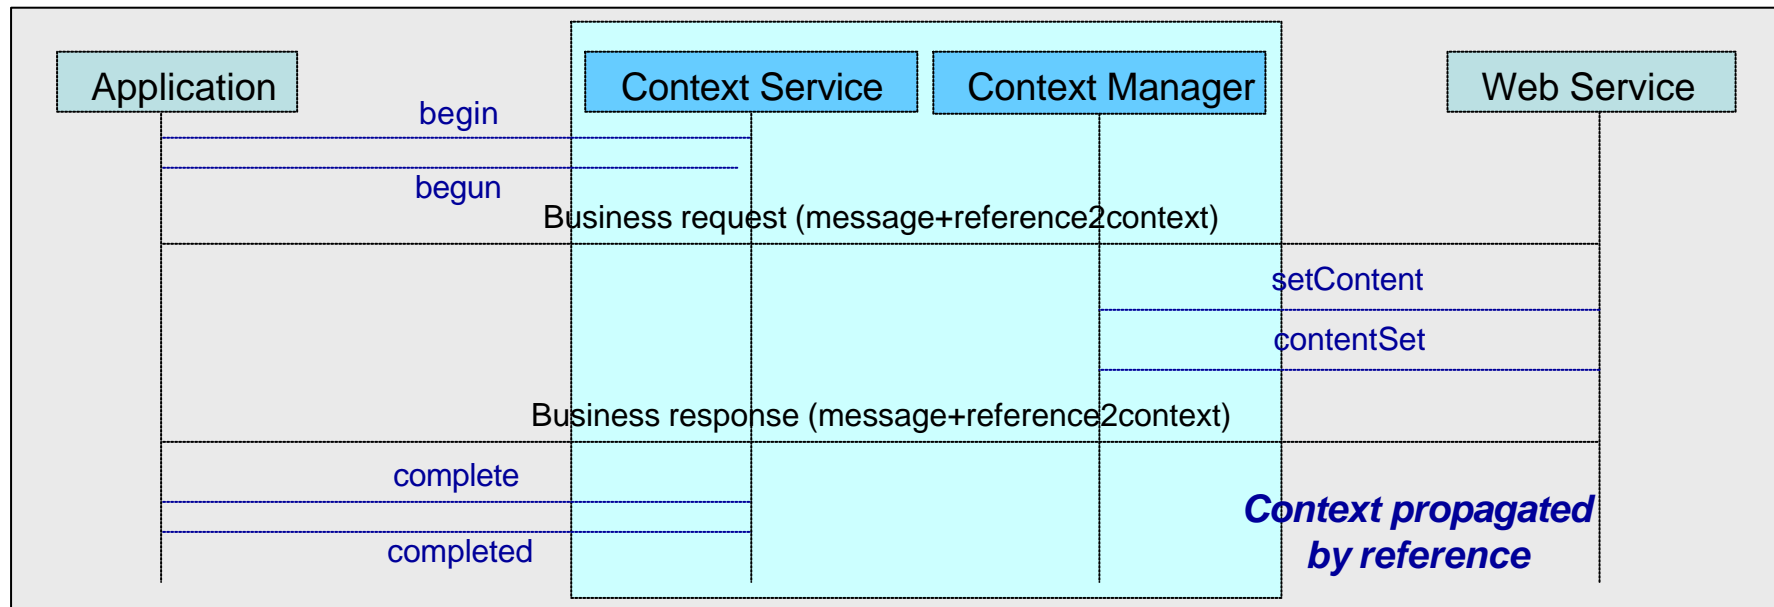
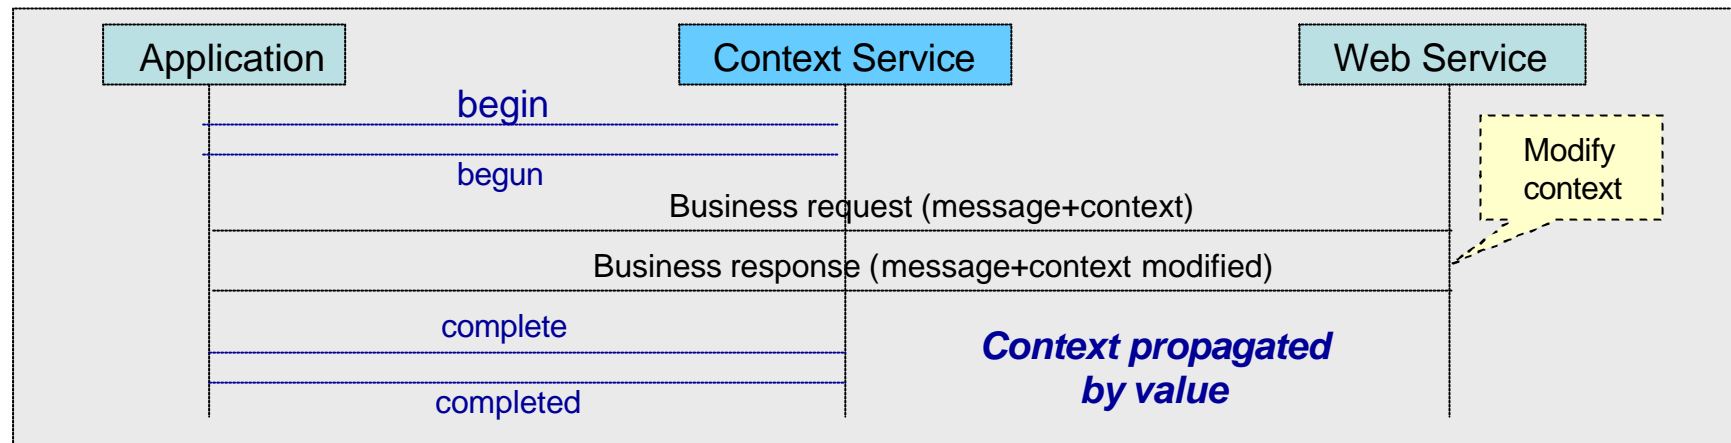
Interactions by Callback



Architectural Overview



Interactions Overview



WS-Coordination Framework

WS-CF

- Coordination is more fundamental than transactions
 - Security
 - Replication
 - Transactions
 - ...
- Coordination could be seen as “disseminating information by a *coordinator* to a number of *participants* to guarantee that all participants obtain a specific message”.

Goals



- Provide a general framework for coordination protocols
 - Existing implementations to be plugged in
 - New implementations can be supported
 - Defines coordinator and participant relationships
- Work with WS-Context
 - Augment context
 - Coordination Context
- Scope of activity becomes scope of coordination boundary

Context type



```
<xs:complexType name="ContextType">
  <xs:complexContent>
    <xs:extension base="wsctx:ContextType">
      <xs:sequence>
        <xs:element name="protocol-reference" type="tns:ProtocolReferenceType"/>
        <xs:element name="coordinator-reference" type="tns:CoordinatorReferenceType"
          maxOccurs="unbounded"/>
        <xs:any namespace="##any" processContents="lax" maxOccurs="unbounded"/>
      </xs:sequence>
    </xs:extension>
  </xs:complexContent>
</xs:complexType>
```

WS-CF

Main components



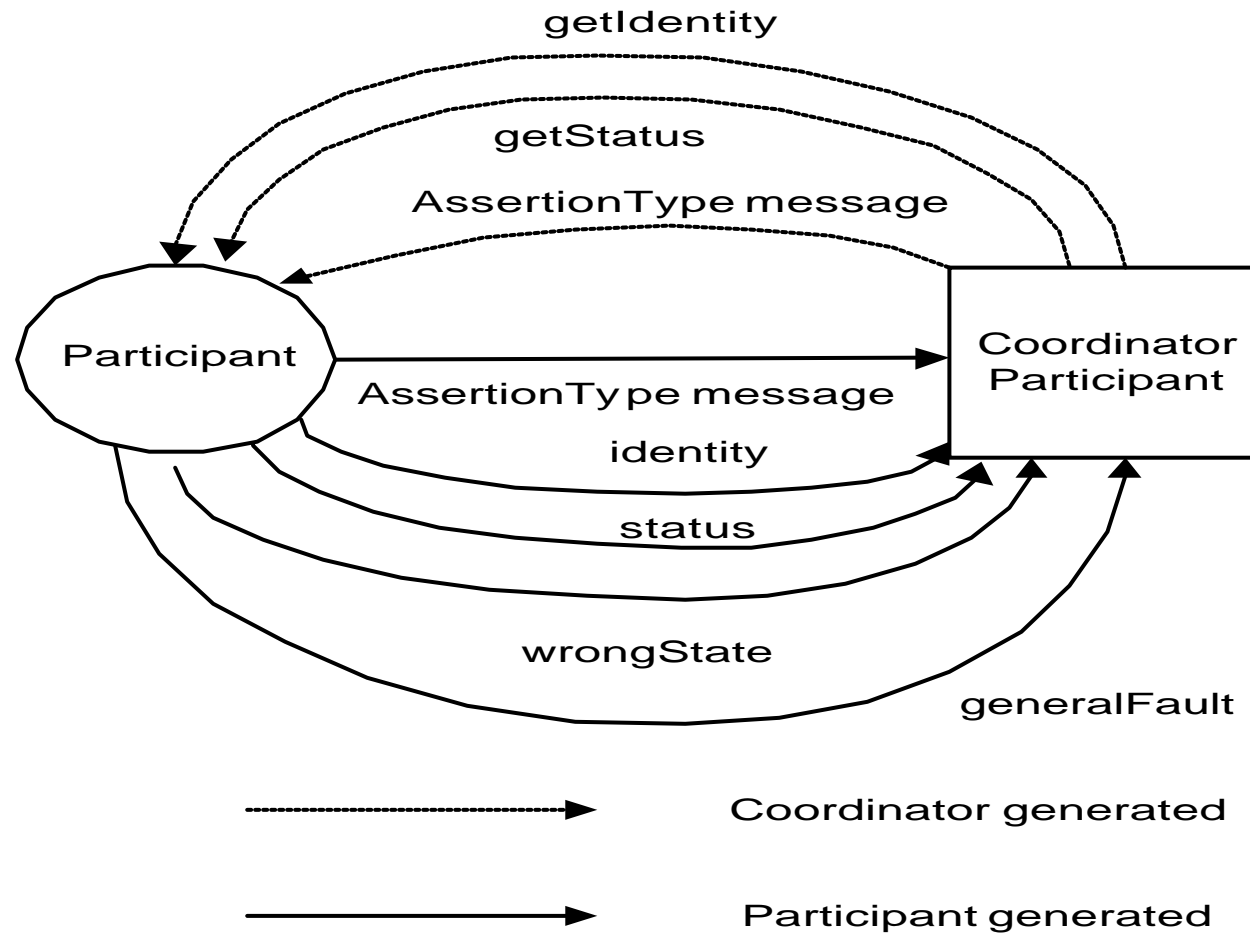
- *Coordinator:*
 - enable registration of participants triggered at *coordination points*.
- *Participant:*
 - The operation or operations that are performed as part of coordination sequence processing
 - A *Coordination Service*: Defines the behavior for a specific coordination model.
- *Coordination Service:*
 - provides a processing pattern that is used for outcome processing.
 - For example
 - ACID (prepare, commit, rollback)
 - Sagas
 - Real-time transactions
 - ...

Coordination protocol

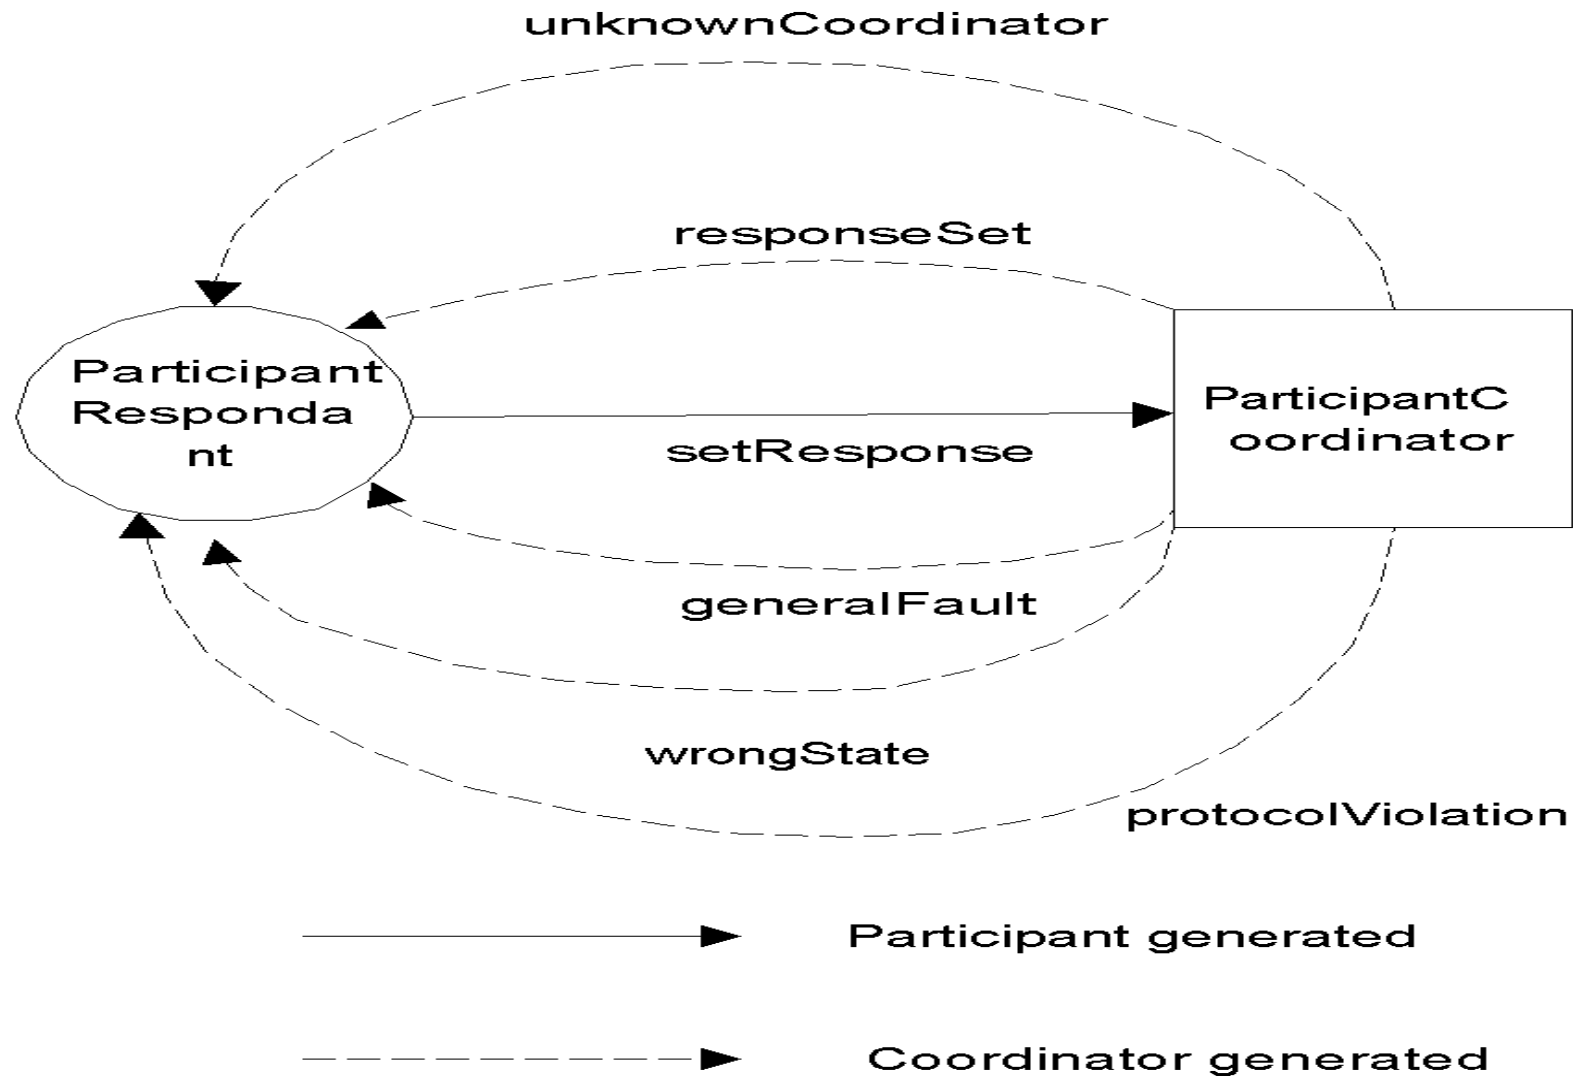


- Defined by message interactions between Coordinator and participants
 - Coordinator-to-participant
 - coordinator sends a protocol message to the participant and will eventually get a response.
 - Coordination status and identity
 - Participant-to-coordinator
 - a participant may autonomously communicate protocol messages to the coordinator.
 - Works in terms of AssertionTypes
- WC-CF protocol neutral
- Protocols Identified by URI

Coordinator-to-participant



Participant-to-coordinator



AssertionType



- “Base class” for all coordinator-to-participant message interactions
 - Requests and responses
- All protocol specific messages enhance this type
- One service (participant or coordinator) can accept multiple protocols

Qualifiers

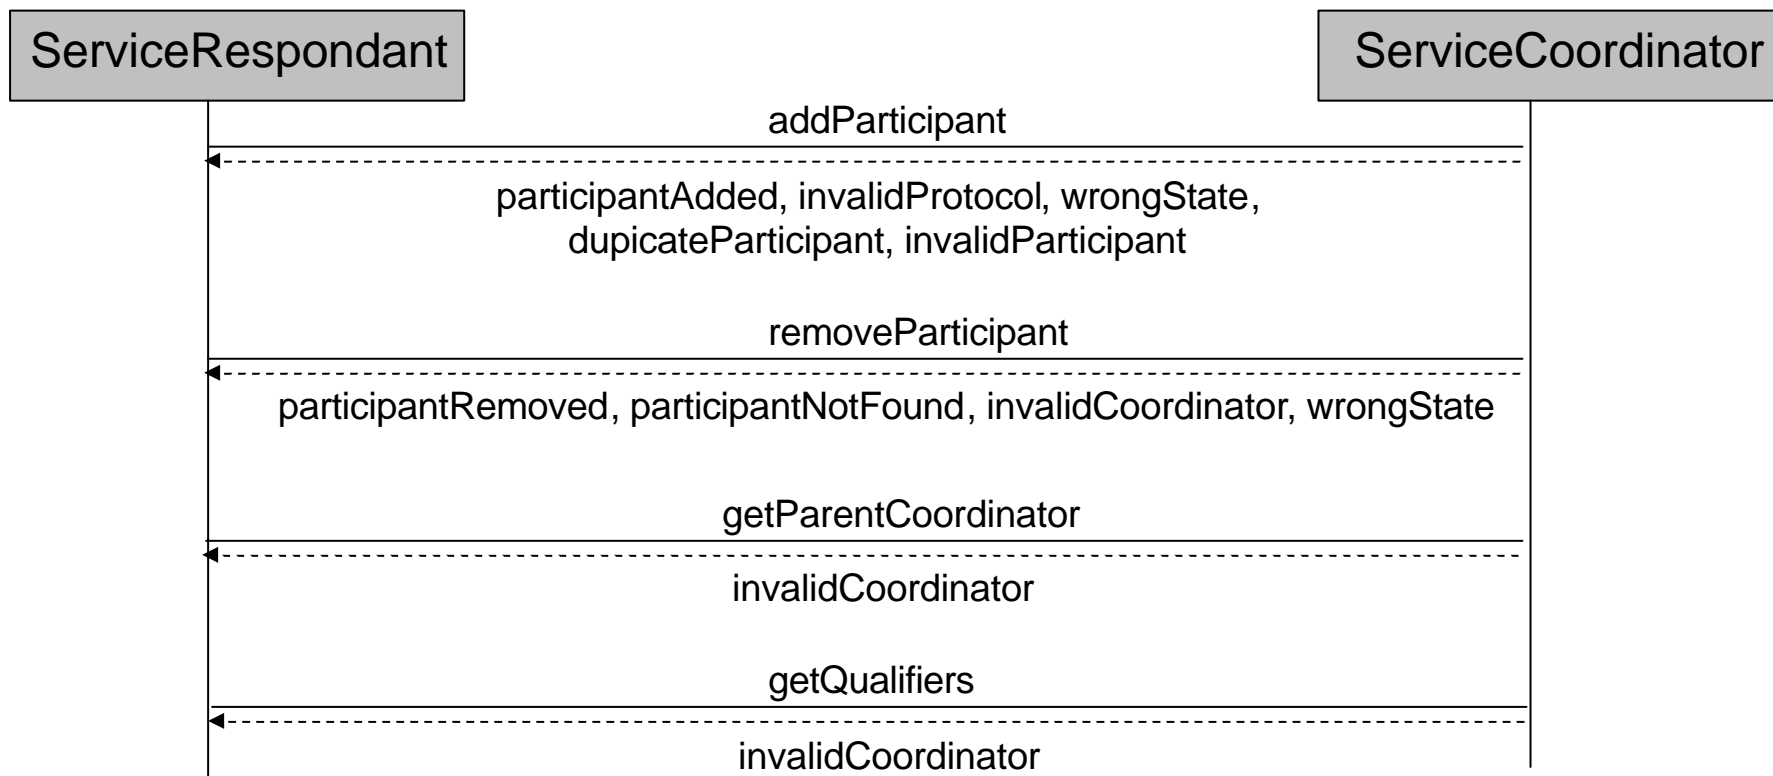


- Runtime protocol extensibility option
- Typically in enlist/delist
 - For coordinator/participant information
 - E.g., will cancel in 24 hours

Service-to-coordinator interactions



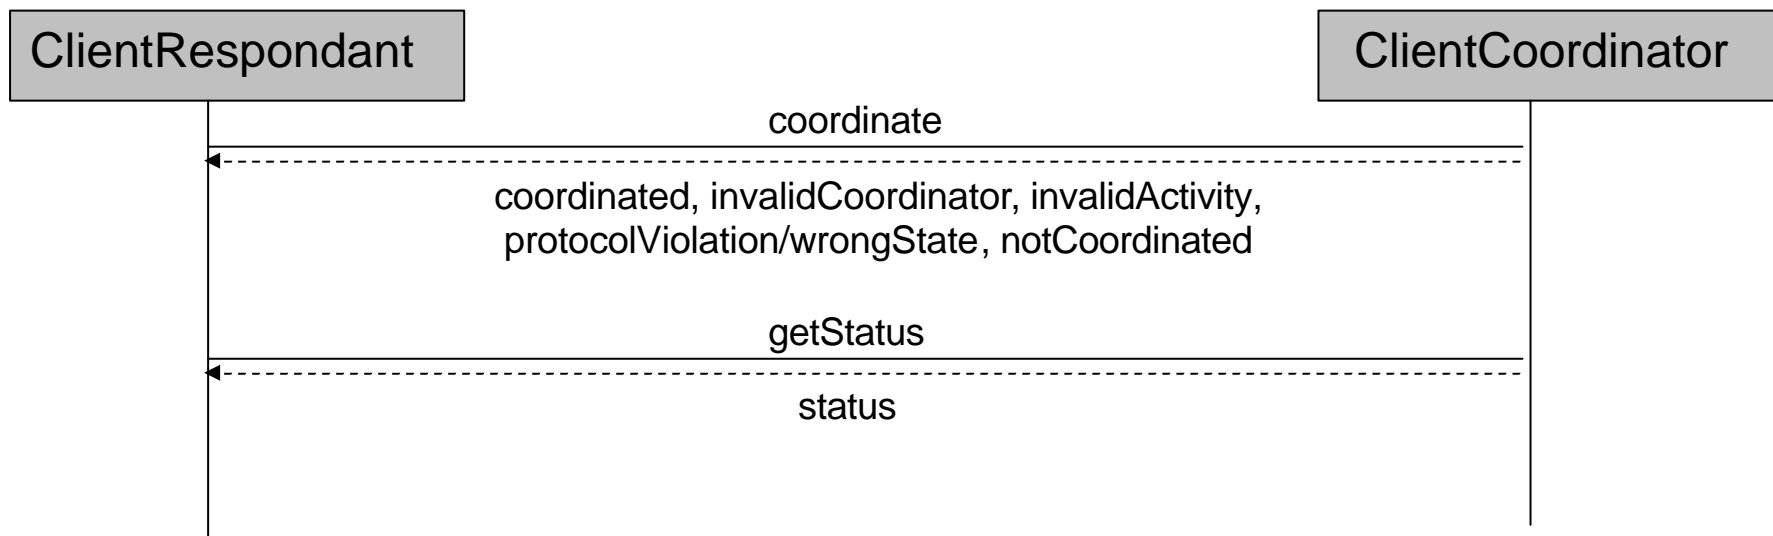
- define how a service may enlist or delist a participant with the coordinator



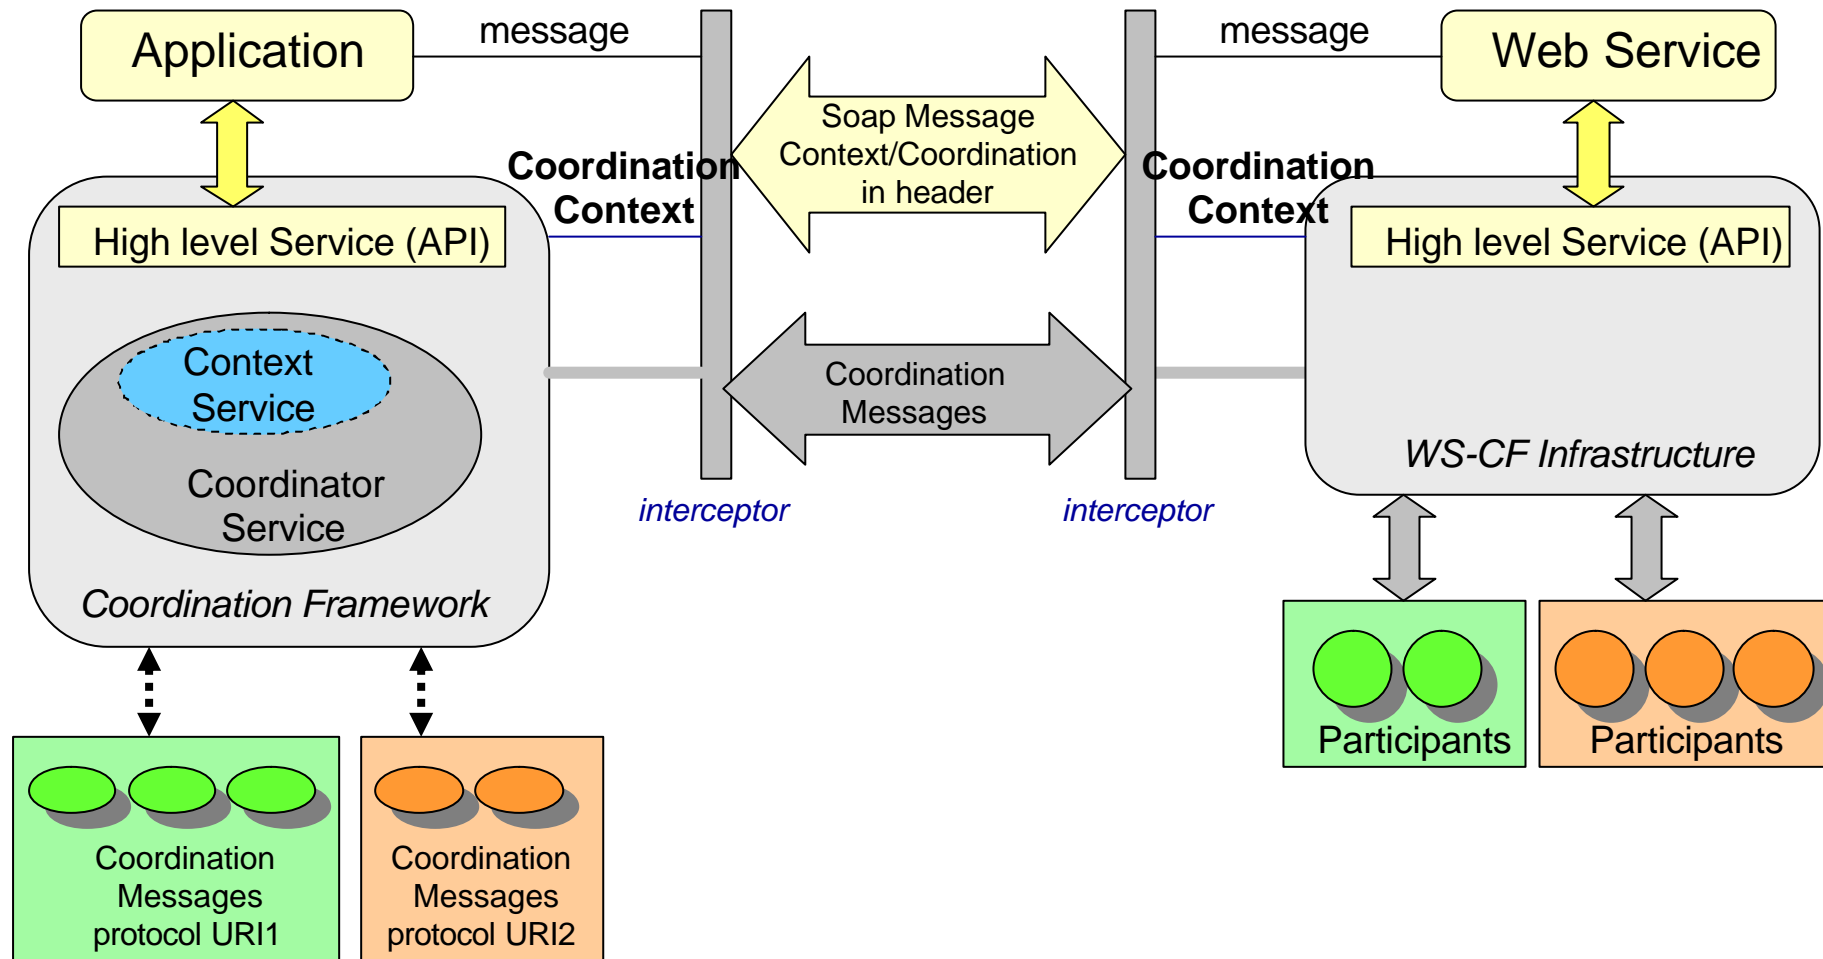
Client-to-coordinator interactions



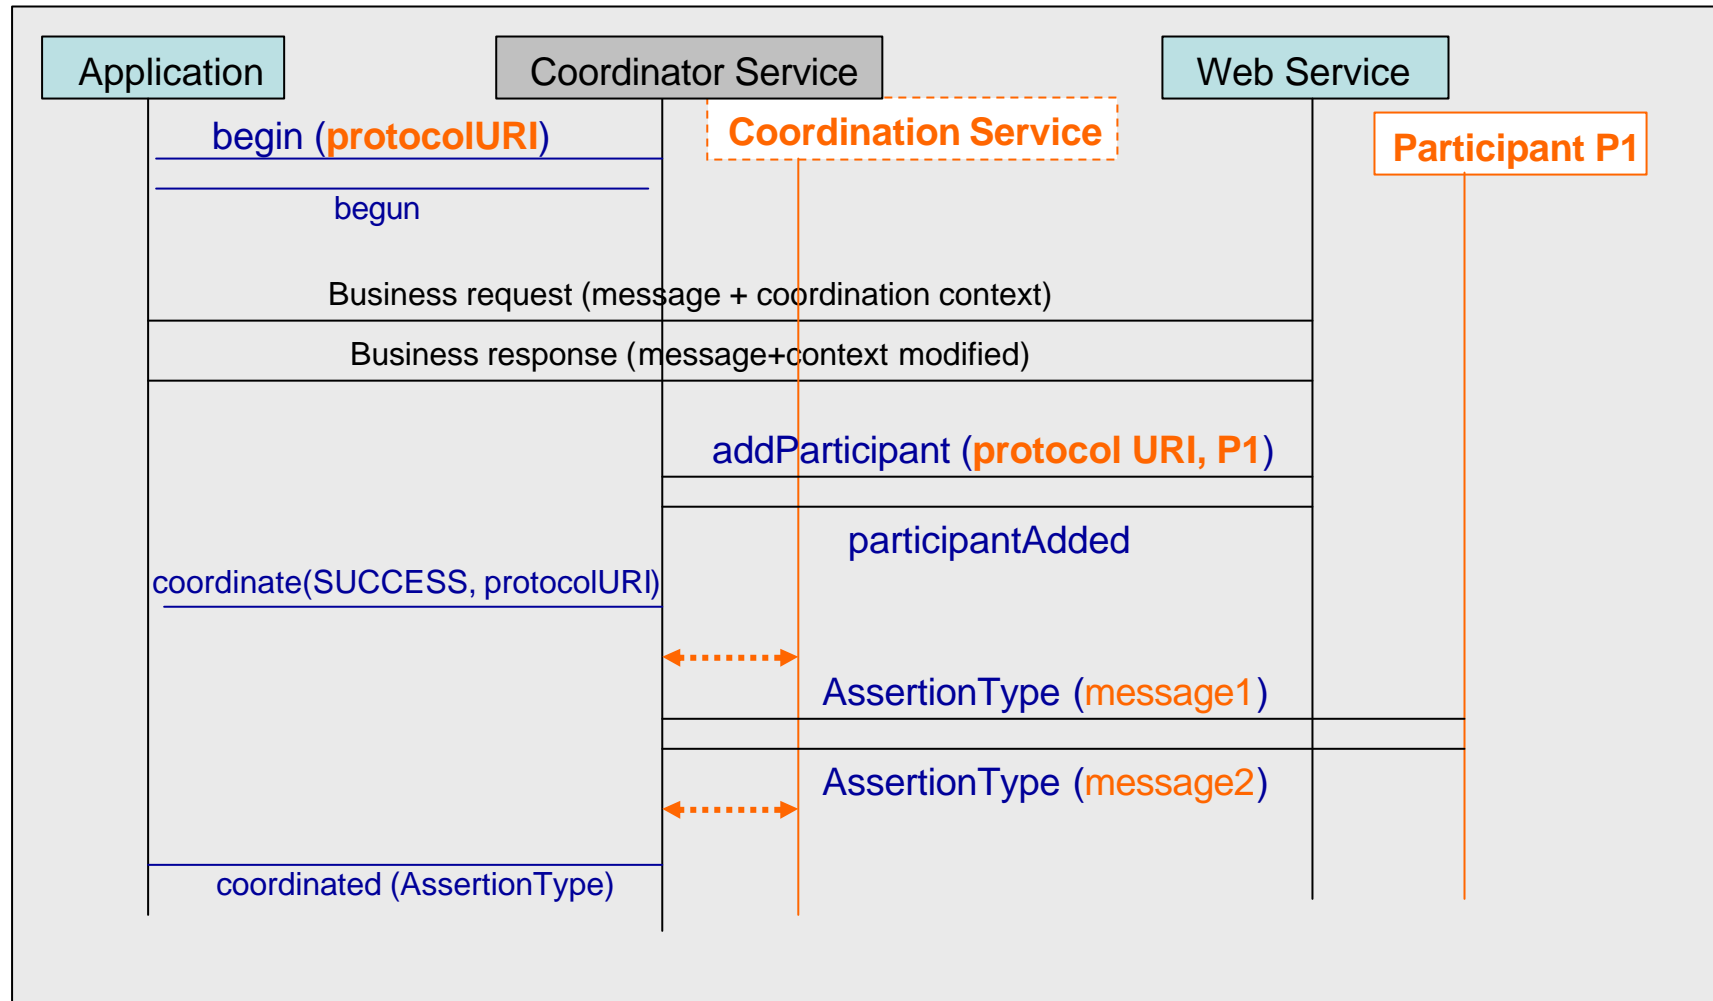
- Define how a client can obtain the status of the coordinator or ask it to perform coordination.



Architectural Overview



Interactions Overview




Interposition



- Important for security and performance reasons
 - Part of most distributed transaction protocols
- Subordinate coordinator
 - Participant as far as coordinator is concerned
 - Coordinator as far as participant is concerned
- Supported by WS-CF
 - Not mandated

Example



Recovery



- Distributed application federated into natural recovery (admin) domains
 - Can't mandate one specific recovery mechanism
 - Very application specific anyway
 - Have to allow administrative domains autonomy
- Therefore, support but not mandate

Recovery support



- RecoveryCoordinator
 - Drives recovery on behalf of participant
 - Participants may not be able to recover at same URI
 - Machine crash, domain migration, ...
- Coordinator can replace one endpoint with another to continue protocol

WS-Transaction Management

WS-TXM



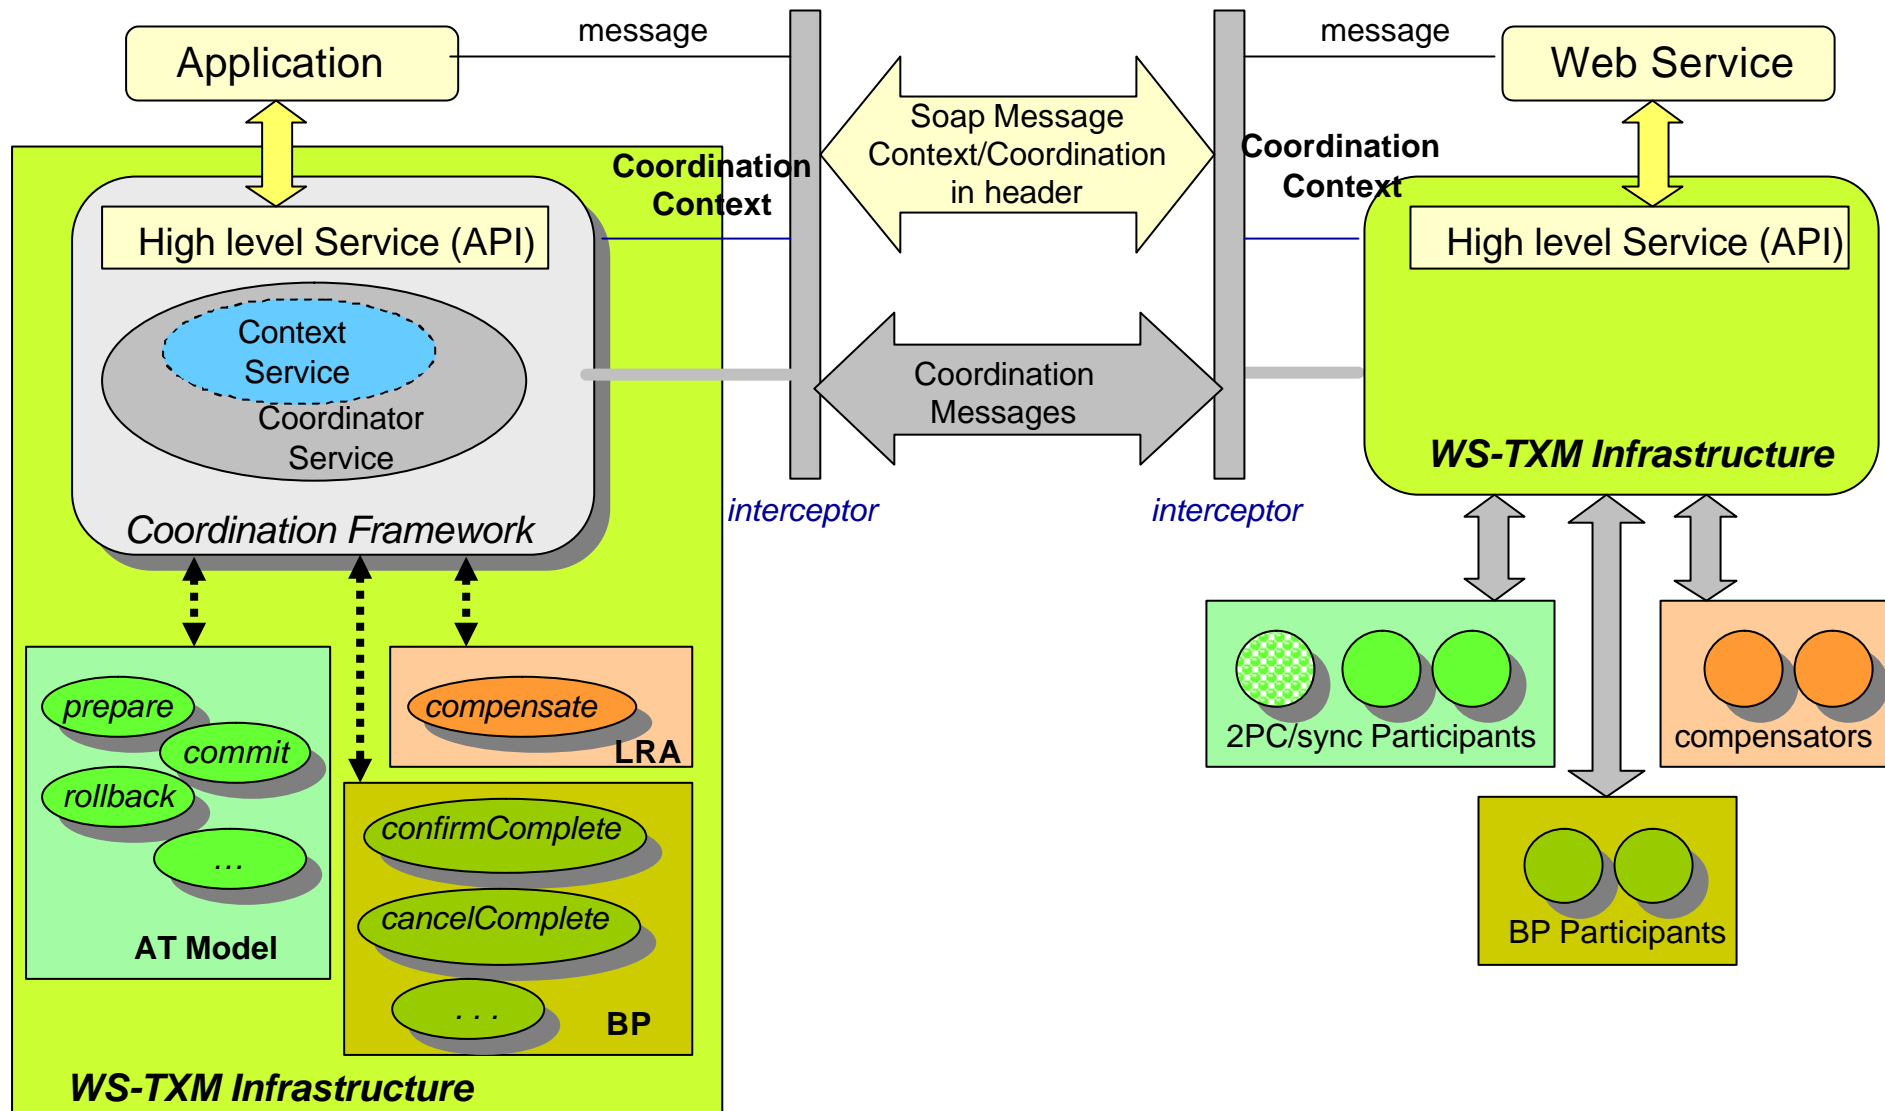
- Goals
 - Support range of use cases
 - “One-size does not fit all”
 - Therefore a single protocol cannot cope with all requirements
 - All requirements aren’t “two-phase”
- Builds on WS-CF and WS-Context
 - Define specific coordinators and participants
 - Augment context

Defines



- Three transaction models
 - ACID transaction
 - For interoperability and high-cost services where ACID transactions are a requirement
 - Long running action
 - Loosely coupled, long duration work that uses compensations
 - Business process
 - For treating all steps in an automated business process as part of a single logical transaction

Architectural Overview



ACID Transaction (Arjuna) middleware for reliability

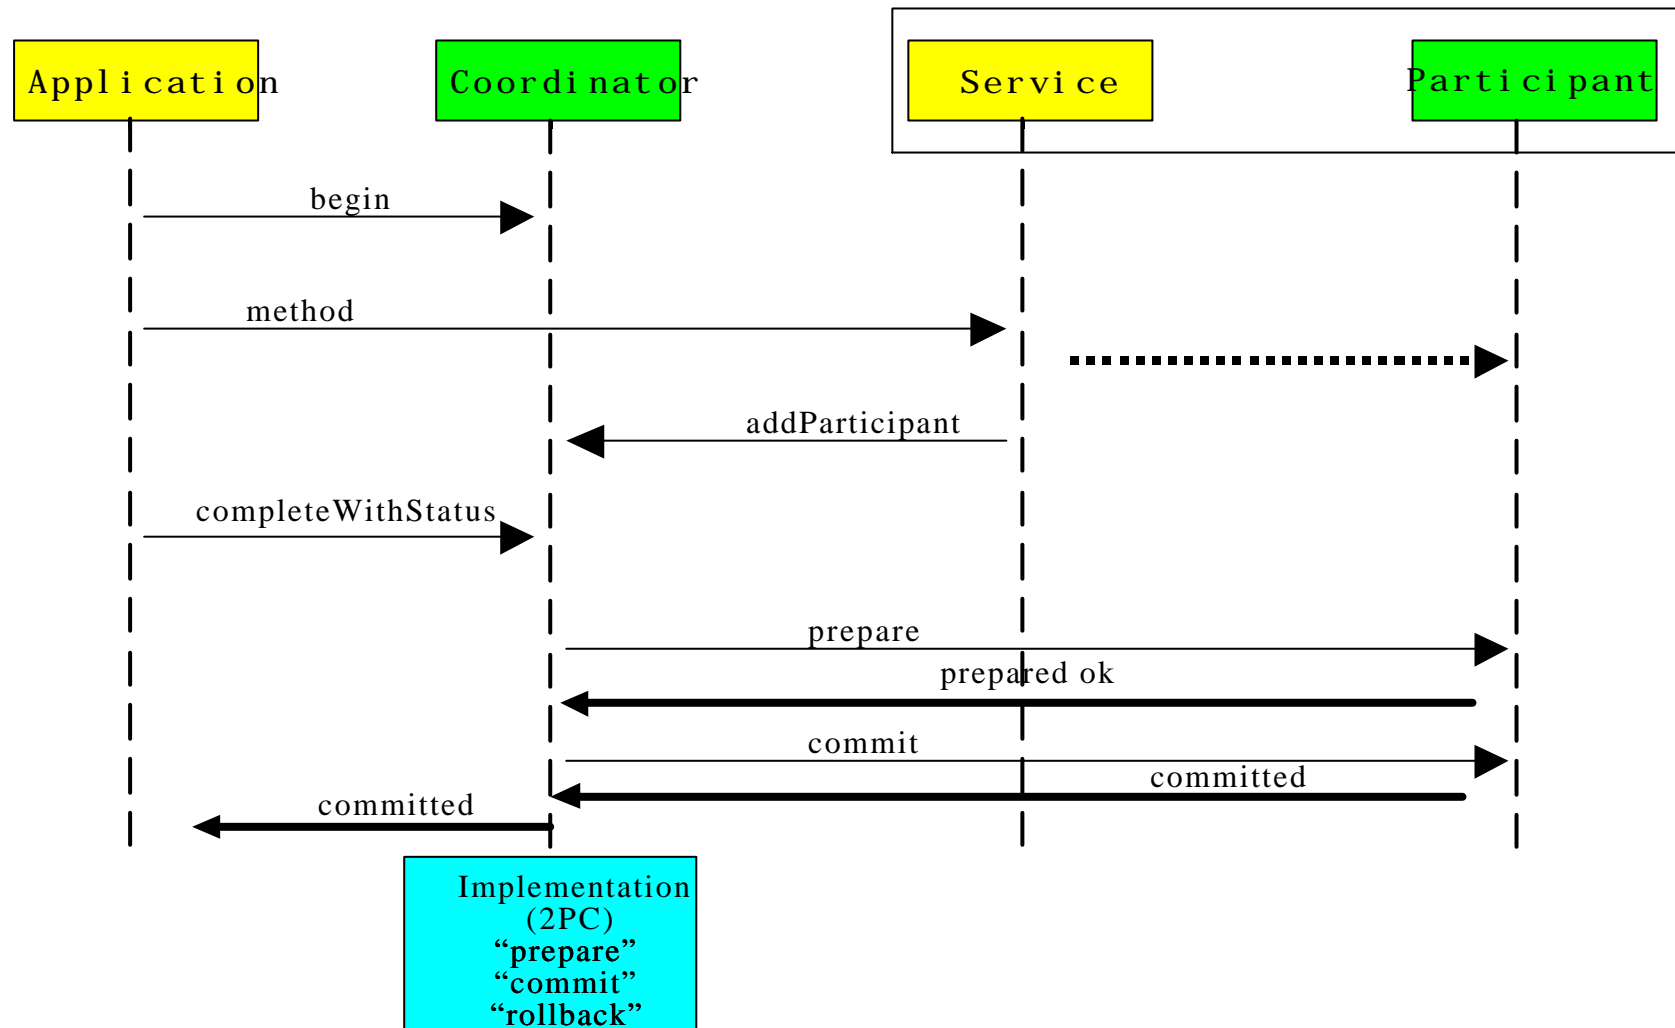
- Coordination-type URI
 - <http://www.webservicetransactions.org/wsdl/wstxm/tx-acid/2003/03>
- Traditional *ACID transaction* with two sub-protocols (different URIs)
 - Two-phase commit
 - <http://www.webservicetransactions.org/wsdl/wsTXM/tx-acid/2pc/2003/03>
 - Synchronizations
 - <http://www.webservicetransactions.org/wsdl/wsTXM/tx-acid/sync/2003/03>
- Interoperability across different vendor implementations
 - *removeParticipant* illegal
 - *wrongState* returned by coordinator
 - *coordinate* cannot be used
 - Bind the scope of activity to the scope of transaction

Model (AT)



- ACID semantics explicitly required
- Presumed rollback
- One-phase optimization
- Read-only optimization
- Heuristics

Example (AT)



2PC protocol messages (AT)



- Usual for two-phase commit
 - prepare
 - voteReadonly, voteCommit, voteRollback
 - And heuristics
 - commit
 - Heuristics
 - rollback
 - Heuristics

Synchronization messages (AT)



- beforeCompletion
 - Runs before two-phase commit begins
- afterCompletion
 - Runs after two-phase protocol

Long running action model (LRA)



- Protocol URI
 - <http://www.webservicetransactions.org/wstxm/tx-lra/2003/03>
- Specifically for long duration interactions
- ACID transactions are not appropriate
 - Can't lock resources for the duration
 - No assumed trust relationships
- Compensation actions used
 - Forward work to return the business state to consistency
 - E.g., credit your credit card and give you back interest payments

Relationship to context and coordination



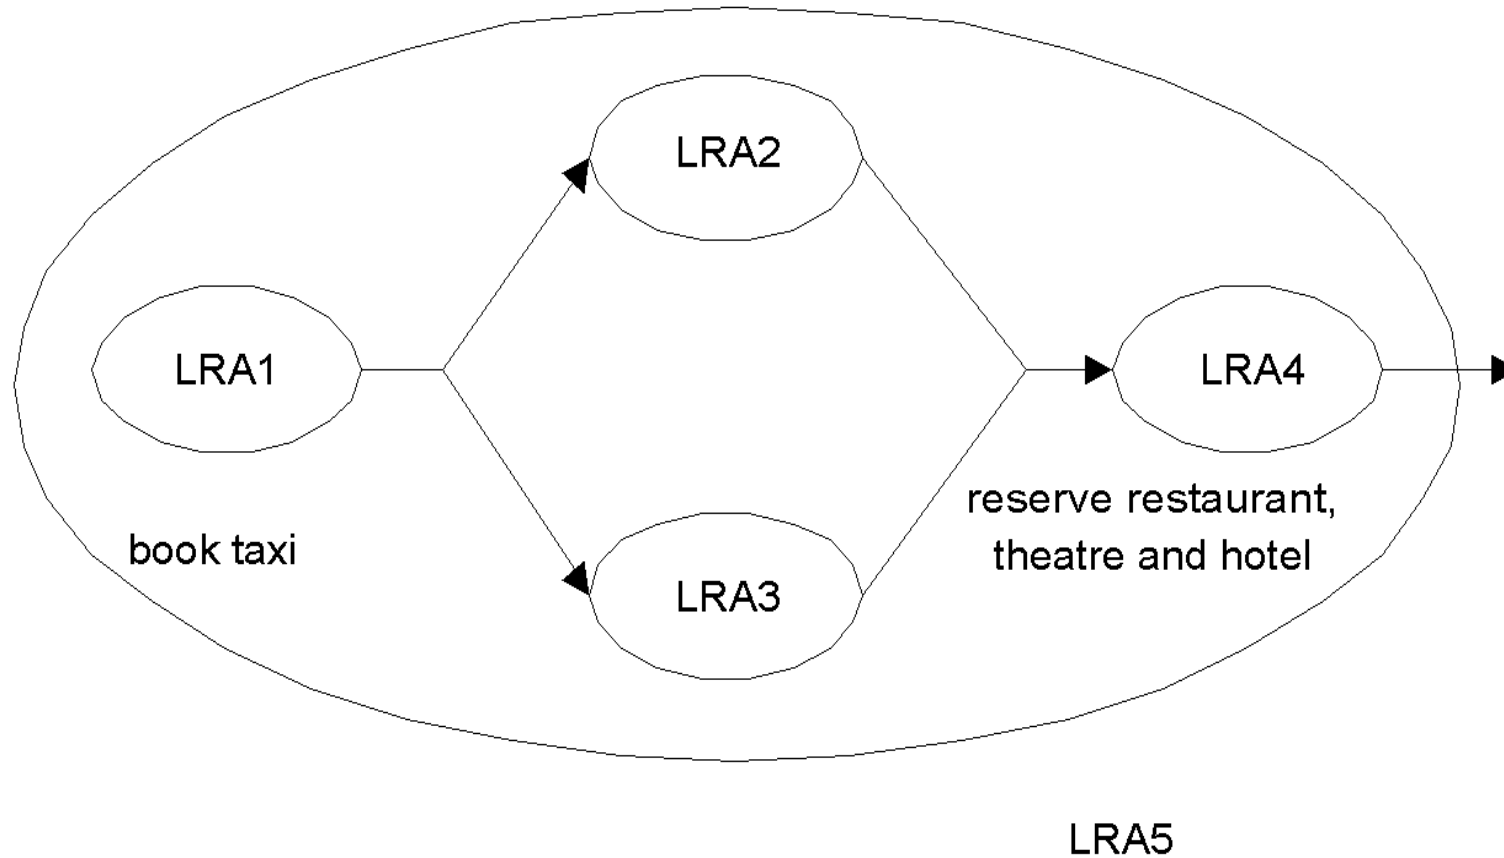
- Activity becomes the scope of business interactions
 - Travel agent, computer construction, etc.
- How services do work is not important
 - Back-end implementation choice
- If work can be compensated then compensation is bound to the activity
 - Non-atomic behaviour
- Activities can be nested

Activities and compensators



- Each activity is a unit of compensatable work
- Work performed must remain compensatable for duration of activity
 - TimeLimit qualifier
- Nested activities imply nested compensators
 - Could be different compensator from child to parent

Travel agent example



Context



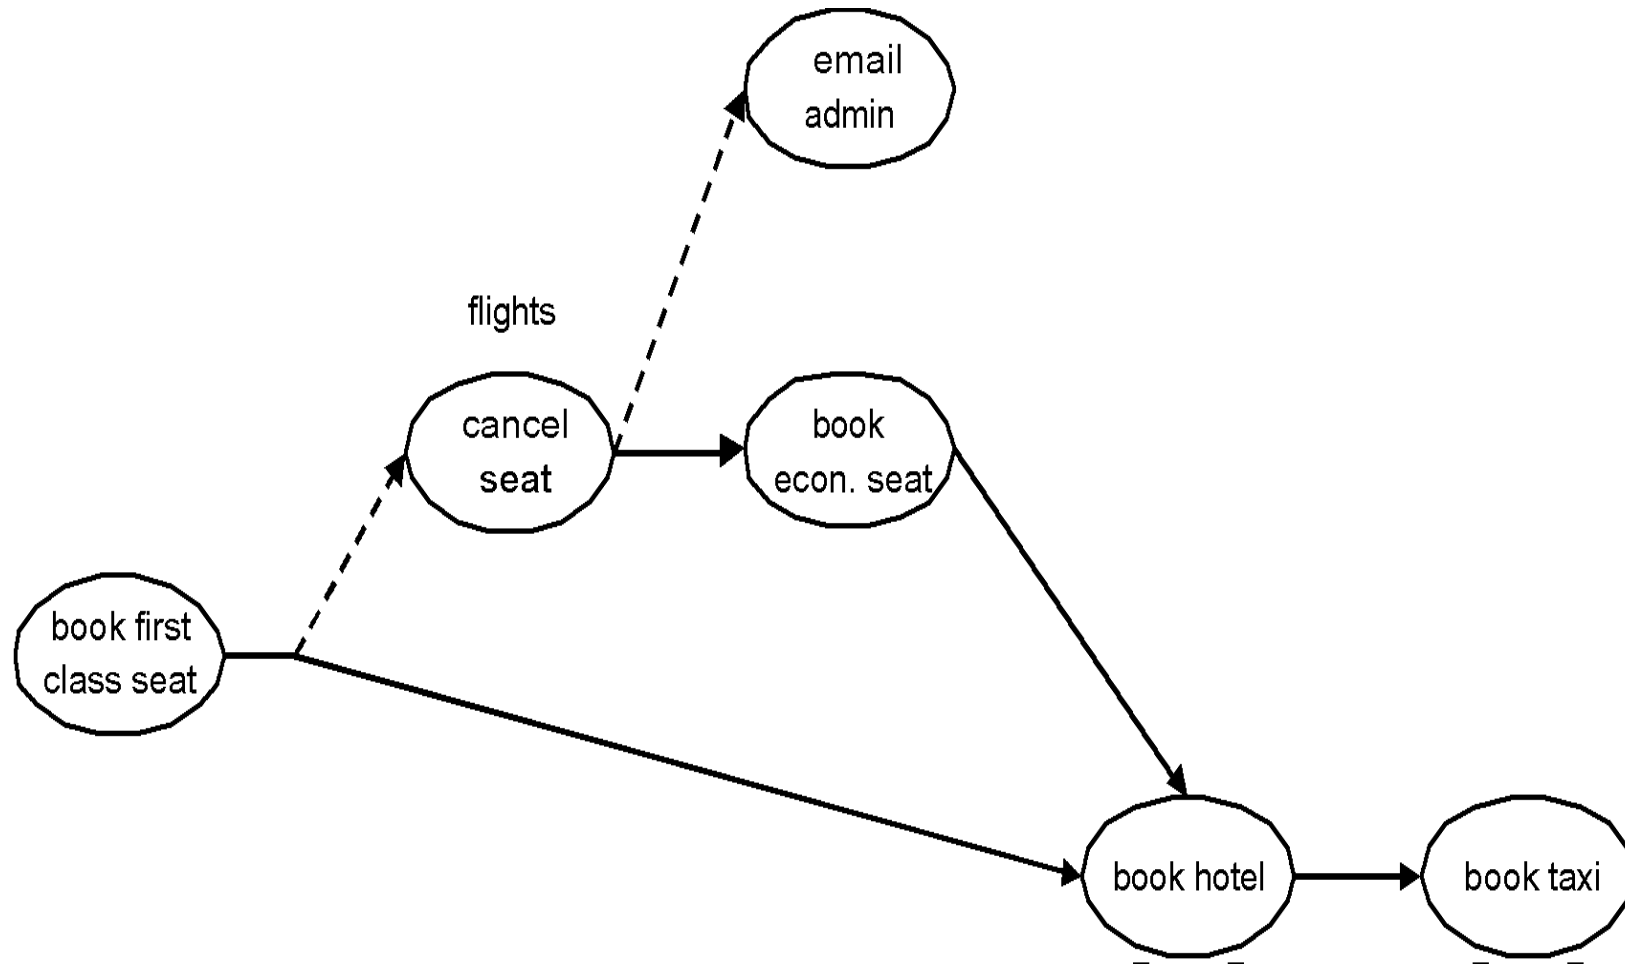
```
<xs:element name="context">
  <xs:complexType>
    <xs:complexContent>
      <xs:extension base="wstxm:ContextType">
        <xs:sequence>
          <xs:element name="lra-id" type="xs:anyURI"/>
          <xs:element name="coordinator-hierarchy">
            <xs:complexType>
              <xs:sequence>
                <xs:element name="coordinator-location" type="xs:anyURI" minOccurs="0"
                  maxOccurs="unbounded"/>
              </xs:sequence>
            </xs:complexType>
          </xs:element>
        </xs:sequence>
      </xs:extension>
    </xs:complexContent>
  </xs:complexType>
</xs:element>
```

To compensate or not?

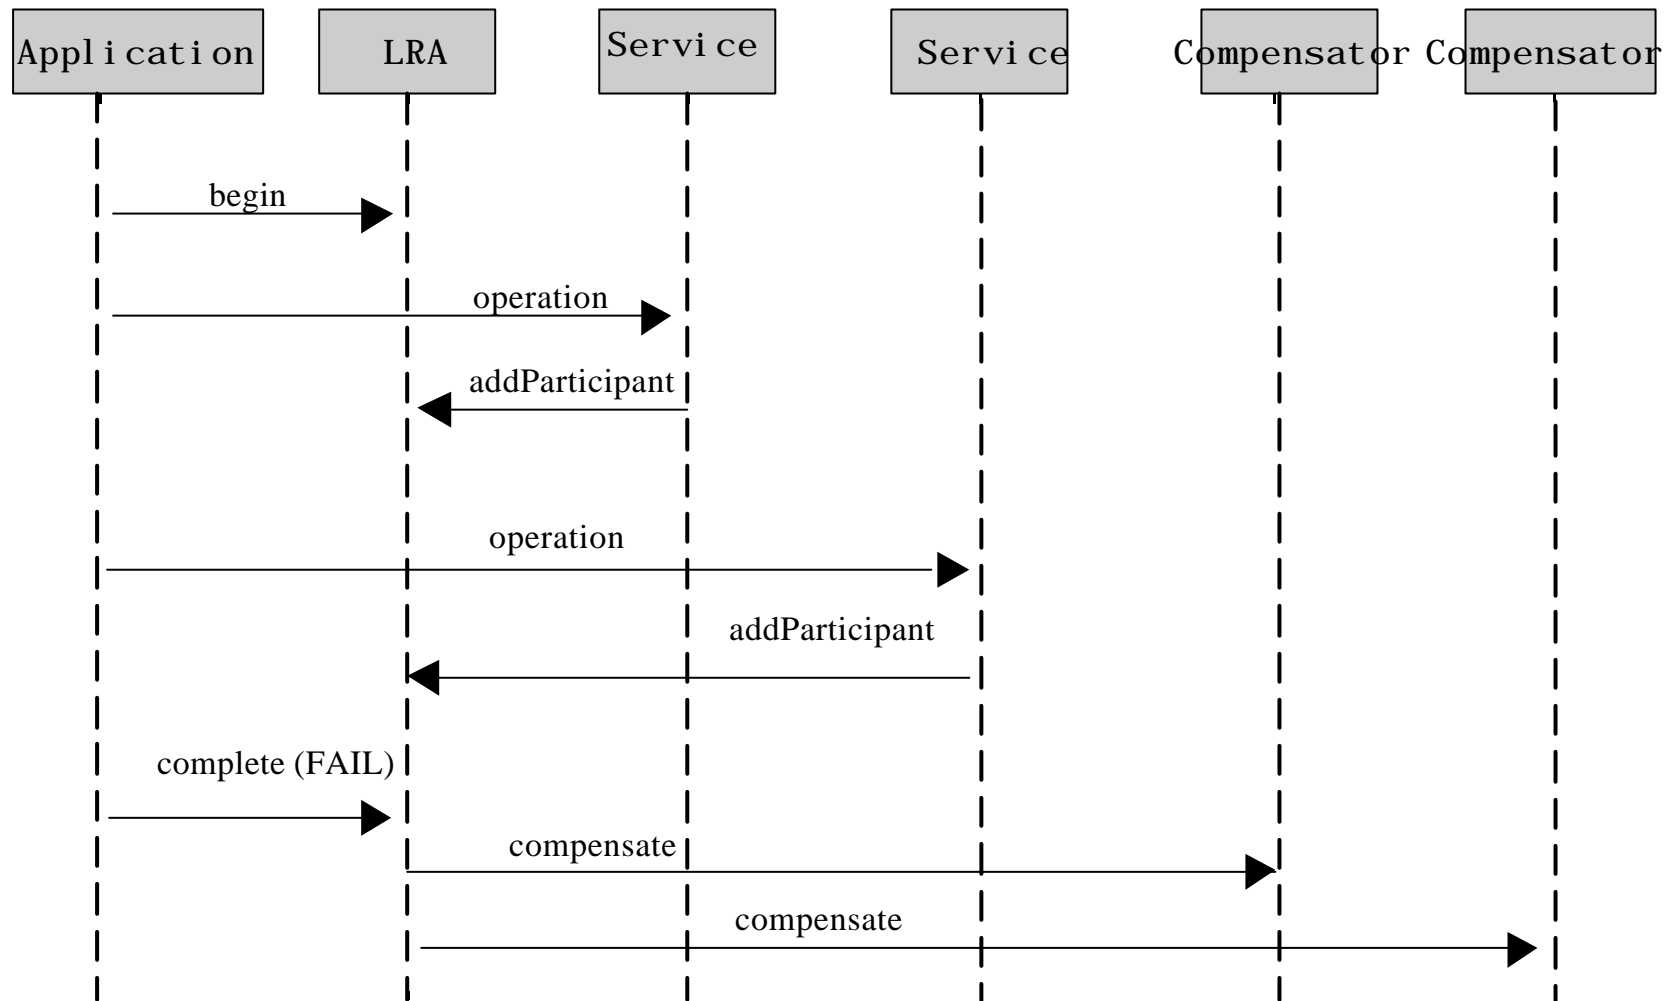


- Some services may not be able to compensate
- The user defines whether or not this is important
 - MustUnderstand
- Must be an explicit choice
 - Adverse consequences otherwise

Compensators



Example




Concurrent LRAs



- Activities can be concurrent
 - Therefore, LRAs can be too
- Allows selection of work within overall activity
 - E.g., choosing the cheapest flight

Selection of work



Business process model



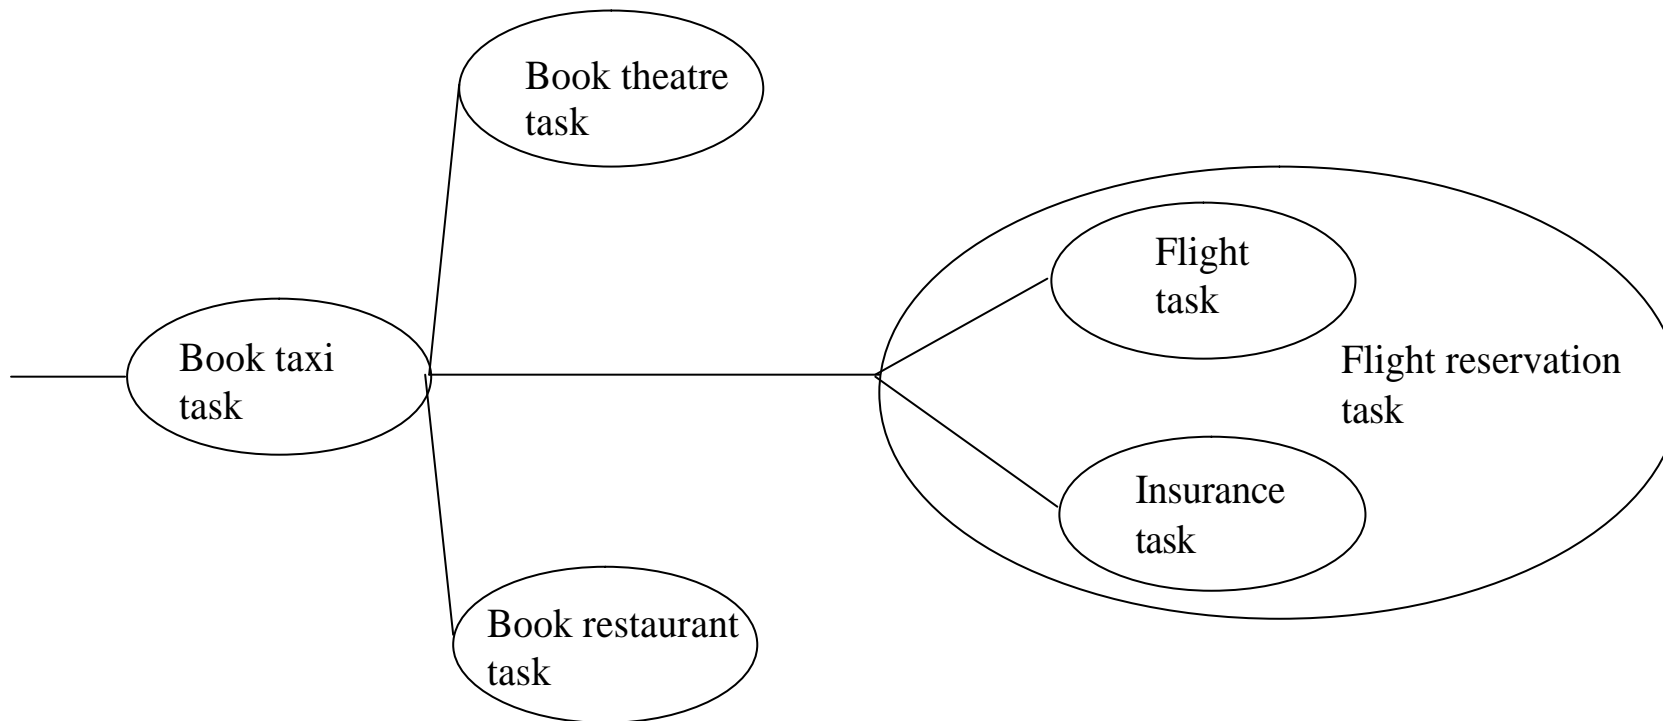
- Aimed at *long running* interactions that span different domains *and* models
 - Workflow
 - Messaging
 - Traditional ACID transactions
- Federated systems that can't/won't expose back-end implementations

Domains

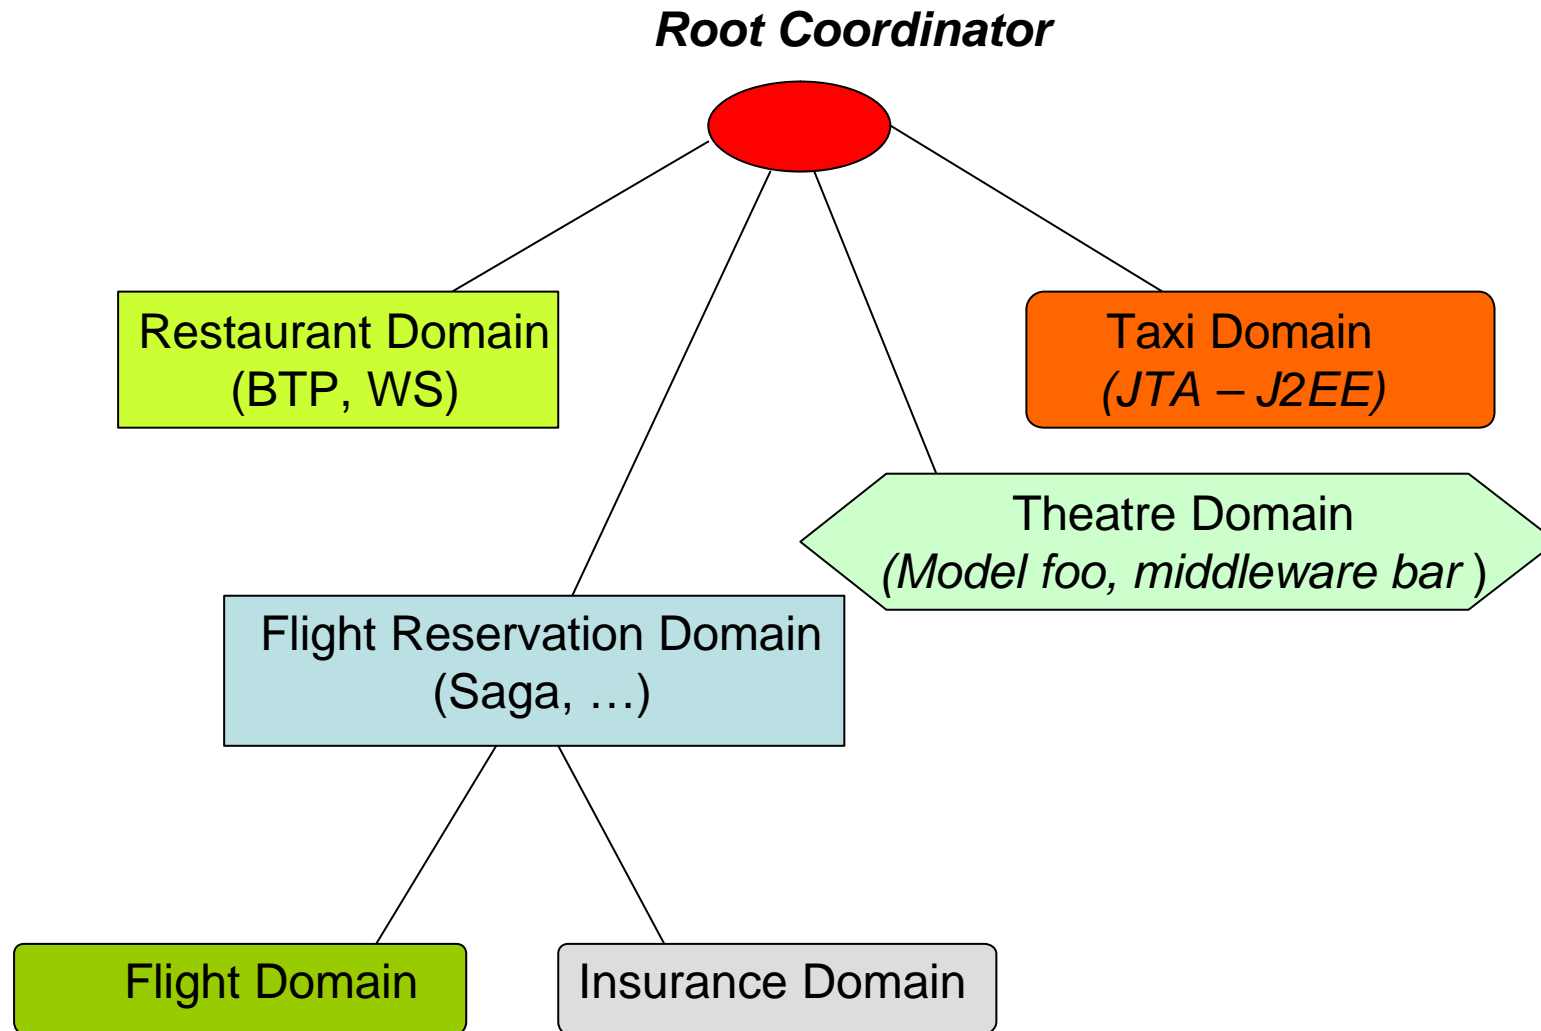


- All parties reside within *business domains*
 - Recursive structure is allowed
 - May represent a different transaction model
- Business process is split into *business tasks*
 - Execute within domains
 - Compensatable units of work
 - Forward compensation during activity is allowed
 - Keep business process making forward progress

Travel agent



Business Process interposition



Model



- Supports synchronous and asynchronous interactions
 - Users can submit work and call back later
 - Or interact synchronously (traditionally)
- Business Process manager
- Optimistic rather than pessimistic
 - Assumes failures are rare and can be handled offline if necessary

How does it work?



- Each domain is exposed as a subordinate coordinator
 - Responsible for mapping incoming BP requests to domain specific protocol
- Protocol messages
 - checkpoint, confirm, cancel, restart, workStatus

checkpoint/restart



- Application driven
 - E.g., via *coordinate* message
- Checkpoints are created across the domains
 - Uniquely identified
- Domains can then roll back to specific checkpoint

workstatus



- Domain calls back to coordinator to inform it of final status
- Or application can enquire
 - WorkCancelled
 - WorkCompleted
 - WorkProcessing

confirm/cancel



- BP termination protocol messages
- Map to success/failure of activity
- Because long-running, heuristics may occur
 - Mixed responses from domains
 - Sufficient information to allow administrative handling

Business Process Entities

